

# Problem Statement

- Ninjacart is India's largest fresh produce supply chain company. They are pioneers in solving one of the toughest supply chain problems of the world by leveraging innovative technology. They source fresh produce from farmers and deliver them to businesses within 12 hours. An integral component of their automation process is the development of robust classifiers which can distinguish between images of different types of vegetables, while also correctly labeling images that do not contain any one type of vegetable as noise.
- As a starting point, ninjacart has provided us with a dataset scraped from the web which contains train and test folders, each having 4 sub-folders with images of onions, potatoes, tomatoes and some market scenes. We have been tasked with preparing a multiclass classifier for identifying these vegetables. The dataset provided has all the required images to achieve the task.

## Ninjacart Business Case: Image Classification for Agricultural Supply Chain Optimization

### Why Ninjacart Needs Image Classification Modeling

Ninjacart, a leading agri-tech platform, connects farmers directly with retailers, restaurants, and other buyers, streamlining the agricultural supply chain in India. The provided dataset, consisting of images of **Onions**, **Potatoes**, **Tomatoes**, and **Indian market** scenes, indicates a focus on automating the identification and classification of agricultural produce and market environments. Image classification modeling is critical for Ninjacart's operations due to the following reasons:

- **Quality Control and Sorting:** Accurate classification of produce (e.g., Onions, Potatoes, Tomatoes) ensures consistent quality assessment. Automated identification can detect defects, categorize produce by type, and ensure only high-quality items reach customers, reducing returns and enhancing customer satisfaction.
- **Inventory Management:** Classifying produce and market-related images helps track inventory in real-time. For example, distinguishing between Onions, Potatoes, and Tomatoes allows Ninjacart to monitor stock levels, predict demand, and optimize procurement from farmers, minimizing waste and shortages.
- **Market Insights:** The "Indian market" class suggests a need to analyze market environments, possibly to identify vendor stalls, storage conditions, or market

trends. This can inform strategic decisions, such as targeting specific markets or optimizing logistics routes.

- **Operational Efficiency:** Manual sorting and identification are labor-intensive and error-prone. Automated image classification reduces human effort, speeds up processes, and lowers operational costs, enabling Ninjacart to scale efficiently.
- **Supply Chain Transparency:** Accurate classification enhances traceability by ensuring that produce is correctly identified from farm to buyer. This builds trust with customers and supports compliance with quality standards.

## How Image Classification Improves Ninjacart's Business

The development and deployment of advanced image classification models (both custom CNNs and pretrained models like ResNet50, InceptionV3, VGG16, and MobileNetV2) directly contribute to Ninjacart's business objectives in the following ways:

- **Enhanced Accuracy and Customer Trust:**
  - **Pretrained Models:** Fine-tuned ResNet50 (94.59%) and InceptionV3 (93.16%) achieve high accuracy, ensuring reliable classification of Onions, Potatoes, Tomatoes, and Indian market scenes. This reduces errors in produce identification, ensuring customers receive the correct items.
  - **Custom Models:** Model 5 (91.74%) offers competitive performance for resource-constrained environments, supporting cost-effective deployment.
  - **Impact:** Accurate classification minimizes customer complaints, enhances brand reputation, and fosters repeat business from retailers and restaurants.
- **Cost Reduction:**
  - **MobileNetV2:** With fewer parameters (3.2M–6.9M), it enables efficient deployment on edge devices (e.g., mobile apps for farmers or warehouse staff), reducing computational costs.
  - **Custom Models:** Models 2–7 (~462K parameters) are lightweight, making them suitable for low-resource settings like rural warehouses.
  - **Impact:** Lower operational costs through automation and efficient models allow Ninjacart to maintain competitive pricing and invest in scaling operations.
- **Improved Supply Chain Efficiency:**
  - Automated classification speeds up sorting and inventory tracking, reducing delays in the supply chain.
  - **Pretrained Models:** High accuracy (e.g., ResNet50 at 94.59%) ensures precise identification, enabling faster processing from farm to market.
  - **Custom Models:** Model 5's balance of accuracy and efficiency supports rapid deployment in warehouses.
  - **Impact:** Faster processing reduces spoilage of perishable goods like Tomatoes, improving profit margins and ensuring fresher produce for customers.
- **Data-Driven Decision Making:**

- Classifying **Indian market** images can provide insights into market dynamics, such as vendor density or produce quality in specific regions.
- **Pretrained Models:** Advanced architectures (e.g., InceptionV3) extract robust features, potentially identifying subtle market characteristics.
- **Impact:** Ninjacart can optimize sourcing strategies, prioritize high-demand regions, and tailor logistics to market conditions, enhancing operational agility.
- **Scalability and Adaptability:**
  - **Pretrained Models:** Leverage pretrained weights from large datasets (e.g., ImageNet), requiring fewer training epochs (5–9 for VGG16) and less labeled data.
  - **Custom Models:** Require more tuning but are adaptable to specific needs with techniques like Data Augmentation (Model 7) or L2 regularization (Model 5).
  - **Impact:** Scalable models support Ninjacart's growth across new regions and product categories, while adaptability ensures robustness to diverse produce types or market conditions.

## Challenges and Recommendations

- **Persistent Misclassifications:**
  - Both custom and pretrained models (e.g., VGG16, custom Model 5) struggle with **Tomato** ↔ **Indian market** (104–106 cases) and **Indian market** ↔ **Onion** misclassifications, likely due to visual similarities (e.g., color, texture).
  - **Recommendation:** Use feature visualization (e.g., Grad-CAM) to identify problematic features, collect more diverse **Indian market** data, and apply class-weighted loss to prioritize these classes.
- **Overfitting in Custom Models:**
  - Custom models (e.g., Models 3–6) show significant overfitting (training: 95–99%, validation: 90–94%), unlike pretrained models with minimal overfitting.
  - **Recommendation:** Enhance regularization (stronger L2, higher Dropout) or refine Data Augmentation in Model 7 to balance generalization and accuracy.
- **Model Selection for Ninjacart:**
  - **High-Accuracy Needs:** Deploy **fine-tuned ResNet50 (94.59%)** for warehouses or centralized processing units to maximize accuracy and reliability.
  - **Resource-Constrained Settings:** Use **MobileNetV2 (92.00%)** or custom **Model 5 (91.74%)** for edge devices or rural operations, balancing performance and efficiency.
  - **Hybrid Approach:** Combine ResNet50 for high-stakes tasks (e.g., quality control) and MobileNetV2 for field operations (e.g., farmer-facing apps).
- **Future Improvements:**
  - Explore ensemble models (e.g., ResNet50 + InceptionV3) to boost accuracy and robustness.

- Increase input resolution for custom models (e.g., from 128x128 to 224x224) to match pretrained models, if resources allow.
- Develop real-time classification apps for farmers to upload produce images, streamlining sourcing.

## Conclusion

Image classification modeling is pivotal for Ninjacart to enhance quality control, streamline inventory management, gain market insights, and improve operational efficiency. **Pretrained models**, particularly **fine-tuned ResNet50 (94.59%)**, offer superior accuracy and generalization, making them ideal for high-precision tasks.

**Custom Model 5 (91.74%)** and **MobileNetV2 (92.00%)** provide cost-effective alternatives for resource-limited environments. By addressing misclassifications and overfitting, Ninjacart can leverage these models to reduce costs, improve customer satisfaction, and scale its agri-tech platform, ultimately strengthening its position in India's agricultural supply chain.

## Dataset Link

[https://drive.google.com/file/d/1cIZX-IV\\_MLxKHSyeyTheX5OCQtNCUcqT/view?  
usp=sharing](https://drive.google.com/file/d/1cIZX-IV_MLxKHSyeyTheX5OCQtNCUcqT/view?usp=sharing)

## Context

- This dataset contains images of the following food items: noise-Indian market and images of vegetables- onion, potato and tomato.

## Data Collection

- The images in this dataset were scraped from Google.

## Content

- This dataset contains a folder train, which has a total of 3135 images, split into four folders as follows:
  - Tomato : 789
  - Potato : 898
  - Onion : 849
  - Indian market : 599

- This dataset contains another folder test which has a total of 351 images, split into four folders
  - Tomato : 106
  - potato : 83
  - onion : 81
  - Indian market : 81

## Inspiration

The objective is to develop a program that can recognize the vegetable item(s) in a photo and identify them for the user.

## Process:

- There are numerous libraries that can be used to work with images. Importing the libraries we feel most confident with is our first step: TensorFlow, keras, matplotlib, opencv, seaborn etc
- Download the Data (you can use gdown) with the provided Dataset Link and unzip it
- Visualize the data, use the dataset directory to create a list containing all the image paths in the training folder. You can use matplotlib or tensorflow to plot a grid sample of the images you fetched from the list of image paths.
- Plot a few of the images of each class to check their dimensions. [Note that the images are not all of uniform dimensions]
- Verify the count of images in each train and test folder by plotting histogram .
- Check each folder to see if the number of images matches the reported number.
- Split the dataset to a train and validation set.
- The provided data does not contain separate training and validation folders. For us to do hyperparameter tuning of our models, it is important to divide the dataset into an 80-20 split for training and validation respectively.
- Before fitting data to our model, we must make sure that each image is square-shaped so that we may resize it to the required dimensions and also perform rescaling which will rescale the inputs between 0-1 by dividing each value by 255.
- Use a model of your choice (could be vgg, resnet and mobilenet) and train it with an appropriate batch size.

- Using the pretrained weights of popular networks is a great way to do transfer learning, since the size of our original dataset is small.
- Obtain the testing accuracy to see how well your model generalizes.
- Compare and check the performance of multiple models using a confusion matrix.
- Implement a TensorBoard callback to log each of our model metrics for each model during the training process. [ recommended google colab or jupyter notebook]
- Plot the train/valid accuracy and loss graph for each model.
- If the model is overfitting , Try to tune your model by applying - :
- BatchNormalization and Dropout
- Callbacks : EarlyStopping, ModelCheckpoint and TensorBoard callback
- Data Augmentation

```
In [ ]: !pip install tensorflow
```

```
Requirement already satisfied: tensorflow in /usr/local/lib/python3.11/dist-packages (2.18.0)
Requirement already satisfied: absl-py>=1.0.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (1.4.0)
Requirement already satisfied: astunparse>=1.6.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (1.6.3)
Requirement already satisfied: flatbuffers>=24.3.25 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (25.2.10)
Requirement already satisfied: gast!=0.5.0,!0.5.1,!0.5.2,>0.2.1 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (0.6.0)
Requirement already satisfied: google-pasta>=0.1.1 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (0.2.0)
Requirement already satisfied: libclang>=13.0.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (18.1.1)
Requirement already satisfied: opt-einsum>=2.3.2 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (3.4.0)
Requirement already satisfied: packaging in /usr/local/lib/python3.11/dist-packages (from tensorflow) (25.0)
Requirement already satisfied: protobuf!=4.21.0,!4.21.1,!4.21.2,!4.21.3,!4.21.4,!4.21.5,<6.0.0dev,>=3.20.3 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (5.29.5)
Requirement already satisfied: requests<3,>=2.21.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (2.32.3)
Requirement already satisfied: setuptools in /usr/local/lib/python3.11/dist-packages (from tensorflow) (75.2.0)
Requirement already satisfied: six>=1.12.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (1.17.0)
Requirement already satisfied: termcolor>=1.1.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (3.1.0)
Requirement already satisfied: typing-extensions>=3.6.6 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (4.14.1)
Requirement already satisfied: wrapt>=1.11.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (1.17.2)
Requirement already satisfied: grpcio<2.0,>=1.24.3 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (1.73.1)
Requirement already satisfied: tensorboard<2.19,>=2.18 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (2.18.0)
Requirement already satisfied: keras>=3.5.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (3.8.0)
Requirement already satisfied: numpy<2.1.0,>=1.26.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (2.0.2)
Requirement already satisfied: h5py>=3.11.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (3.14.0)
Requirement already satisfied: ml-dtypes<0.5.0,>=0.4.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (0.4.1)
Requirement already satisfied: tensorflow-io-gcs-filesystem>=0.23.1 in /usr/local/lib/python3.11/dist-packages (from tensorflow) (0.37.1)
Requirement already satisfied: wheel<1.0,>=0.23.0 in /usr/local/lib/python3.11/dist-packages (from astunparse>=1.6.0->tensorflow) (0.45.1)
Requirement already satisfied: rich in /usr/local/lib/python3.11/dist-packages (from keras>=3.5.0->tensorflow) (13.9.4)
Requirement already satisfied: namex in /usr/local/lib/python3.11/dist-packages (from keras>=3.5.0->tensorflow) (0.1.0)
Requirement already satisfied: optree in /usr/local/lib/python3.11/dist-packages (from keras>=3.5.0->tensorflow) (0.16.0)
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.11/dist-packages (from requests<3,>=2.21.0->tensorflow) (3.4.2)
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.11/dist-packages (from requests<3,>=2.21.0->tensorflow) (3.10)
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.11/di
```

```
st-packages (from requests<3,>=2.21.0->tensorflow) (2.4.0)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.11/dist-packages (from requests<3,>=2.21.0->tensorflow) (2025.7.14)
Requirement already satisfied: markdown>=2.6.8 in /usr/local/lib/python3.11/dist-packages (from tensorflow<2.19,>=2.18->tensorflow) (3.8.2)
Requirement already satisfied: tensorflow-data-server<0.8.0,>=0.7.0 in /usr/local/lib/python3.11/dist-packages (from tensorflow<2.19,>=2.18->tensorflow) (0.7.2)
Requirement already satisfied: werkzeug>=1.0.1 in /usr/local/lib/python3.11/dist-packages (from tensorflow<2.19,>=2.18->tensorflow) (3.1.3)
Requirement already satisfied: MarkupSafe>=2.1.1 in /usr/local/lib/python3.11/dist-packages (from werkzeug>=1.0.1->tensorflow<2.19,>=2.18->tensorflow) (3.0.2)
Requirement already satisfied: markdown-it-py>=2.2.0 in /usr/local/lib/python3.11/dist-packages (from rich->keras>=3.5.0->tensorflow) (3.0.0)
Requirement already satisfied: pygments<3.0.0,>=2.13.0 in /usr/local/lib/python3.11/dist-packages (from rich->keras>=3.5.0->tensorflow) (2.19.2)
Requirement already satisfied: mdurl~0.1 in /usr/local/lib/python3.11/dist-packages (from markdown-it-py>=2.2.0->rich->keras>=3.5.0->tensorflow) (0.1.2)
```

```
In [ ]: # Import All Required Libraries
```

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from google.colab import drive
import random
import glob
import os
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers
from google.colab import files
from tensorflow.keras import Input
# from tensorflow.keras import metrics
from sklearn import metrics
from tensorflow.keras import layers, regularizers
import functools
from tensorflow.keras.optimizers import Adam, SGD
from tensorflow.keras.models import Model
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.applications import VGG16
from tensorflow.keras.applications import vgg16
from tensorflow.keras.applications import ResNet50
from tensorflow.keras.applications import resnet50
from tensorflow.keras.applications import InceptionV3
from tensorflow.keras.applications import inception_v3
from sklearn.metrics import confusion_matrix, classification_report
from tensorflow.keras.applications import MobileNetV2
from IPython.display import FileLink, display
from tensorflow.keras.applications import EfficientNetB0
```

```
In [ ]: !sudo apt install tree
```

```
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
The following NEW packages will be installed:
  tree
0 upgraded, 1 newly installed, 0 to remove and 35 not upgraded.
Need to get 47.9 kB of archives.
After this operation, 116 kB of additional disk space will be used.
Get:1 http://archive.ubuntu.com/ubuntu jammy/universe amd64 tree amd64 2.0.2-1 [4
7.9 kB]
Fetched 47.9 kB in 1s (66.5 kB/s)
debconf: unable to initialize frontend: Dialog
debconf: (No usable dialog-like program is installed, so the dialog based fronten
d cannot be used. at /usr/share/perl5/Debconf/FrontEnd/Dialog.pm line 78, <> line
1.)
debconf: falling back to frontend: Readline
debconf: unable to initialize frontend: Readline
debconf: (This frontend requires a controlling tty.)
debconf: falling back to frontend: Teletype
dpkg-preconfigure: unable to re-open stdin:
Selecting previously unselected package tree.
(Reading database ... 126284 files and directories currently installed.)
Preparing to unpack .../tree_2.0.2-1_amd64.deb ...
Unpacking tree (2.0.2-1) ...
Setting up tree (2.0.2-1) ...
Processing triggers for man-db (2.10.2-1) ...
```

```
In [ ]: !rm -rf /content/ninjacart_data
```

```
In [ ]: drive.mount('/content/drive')
```

```
Mounted at /content/drive
```

```
In [ ]: path = r"/content/drive/MyDrive/Colab Notebooks"
folder_name = "ninjacart_data"
folder_path = os.path.join(path, folder_name)
zip_path = os.path.join(path, "ninjacart_data.zip")

cwd = os.getcwd()
extract_path = os.path.join(cwd, folder_name)
print(extract_path)

if not os.path.isdir(extract_path):
    print("Folder doesn't exist. Unzipping...")
    !unzip -q "{zip_path}" -d "{cwd}"
    print("Unzipped successfully.")
else:
    print("Folder already exists.")
```

```
/content/ninjacart_data
Folder doesn't exist. Unzipping...
Unzipped successfully.
```

```
In [ ]: !cd /content/ninjacart_data ; tree -L 2
```

```

.
├── test
│   ├── indian market
│   ├── onion
│   ├── potato
│   └── tomato
└── train
    ├── indian market
    ├── onion
    ├── potato
    └── tomato

```

10 directories, 0 files

## Samples from each of these classes and Data Distribution across samples!

```
In [ ]: def get_image_filenames(base_path, target_paths=['train'], extensions=[".jpg", ".png"]):
    """
    Scans a base directory and collects sorted image filenames by category and target.

    Args:
        base_path (str): Root directory (e.g., "/content/ninjacart_data")
        target_paths (list): List of subfolders to search inside (e.g., ['train'])
        extensions (list): Valid image extensions

    Returns:
        dict: Structure like {'train': {'onion': [...], 'tomato': [...]}, ...}
    """
    result = {}

    for target in target_paths:
        target_dir = os.path.join(base_path, target)
        if not os.path.isdir(target_dir):
            print(f"Skipping missing directory: {target_dir}")
            continue

        category_dict = {}
        for category in os.listdir(target_dir):
            category_path = os.path.join(target_dir, category)
            if not os.path.isdir(category_path):
                continue

            image_files = sorted([
                f for f in os.listdir(category_path)
                if any(f.lower().endswith(ext) for ext in extensions)
            ])

            category_dict[category] = image_files

        result[target] = category_dict

    return result
```

```

def get_sub_folder_names(image_files_path):
    class_dirs = os.listdir(f'{image_files_path}')
    return class_dirs

# def get_categories(base_path):
#     class_dirs = os.listdir(base_path)
#     return class_dirs

def get_category_paths(base_path, traget_paths=['train']):
    get_category_paths = []
    categories = None
    sub_folder_paths = get_train_test_val_paths(base_path)
    for subfolder_path in sub_folder_paths:
        if subfolder_path.split('/')[-1] in traget_paths :
            categories = os.listdir(subfolder_path)
            for category in categories:
                category_path = os.path.join(subfolder_path, category)
                get_category_paths.append(category_path)
    return get_category_paths


def get_train_test_val_paths(base_path):
    sub_folder_paths = []
    for i in get_sub_folder_names(base_path):
        sub_folder_path = os.path.join(base_path, i)
        sub_folder_paths.append(sub_folder_path)
    return sub_folder_paths

def get_files_count(base_path, traget_paths=['train']):
    final_dict = {}
    for i in traget_paths:
        catergory_files_count_dict={}
        for category_path in get_category_paths(base_path, i):
            files = glob.glob(f'{category_path}/*')
            catergory_files_count_dict.update({category_path.split('/')[-1] : len(files)})
        final_dict.update({i:catergory_files_count_dict})

    return final_dict

def get_sample_images(base_path, traget_paths=['train']):
    image_dict = {}
    for category_path in get_category_paths(base_path, traget_paths):
        image_files = glob.glob(f'{category_path}/*')
        image_path = random.choice(image_files)
        category = image_path.split('/')[2]
        image_dict.update({category: tf.keras.utils.load_img(image_path)})
    return image_dict

def get_noise_images(noise_path, noise_images):
    image_dict = {}
    for noise_image in noise_images:
        image_path = os.path.join(noise_path, noise_image)
        image_name = image_path.split('/')[-1]
        image_dict.update({image_name: tf.keras.utils.load_img(image_path)})
    return image_dict


def plot_images(image_dict, title):
    plt.figure(figsize=(15, 12))

```

```

plt.suptitle(title)
for i, (cls, img) in enumerate(image_dict.items()):
    ax = plt.subplot(3, 4, i + 1)
    plt.imshow(img)
    plt.title(f'{cls}, {img.size}')
    plt.axis("off")

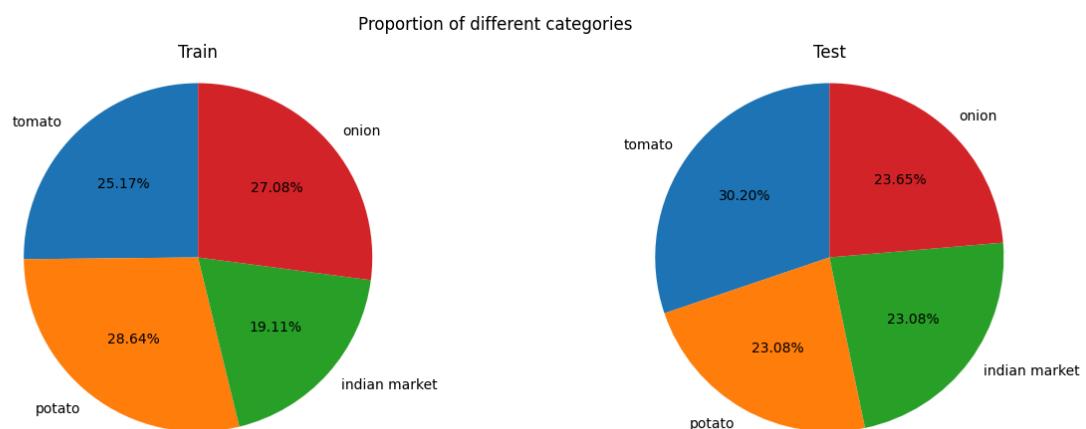
```

In [ ]: `base_data_path = 'nijacart_data'  
train_test_proportion = get_files_count(base_data_path, target_paths=['train', 'test'])  
pd.DataFrame(train_test_proportion)`

Out[ ]:

	train	test
<b>tomato</b>	789	106
<b>potato</b>	898	81
<b>indian market</b>	599	81
<b>onion</b>	849	83

In [ ]: `train_proportion = train_test_proportion['train']  
test_proportion = train_test_proportion['test']  
plt.figure(figsize=(15, 5))  
plt.suptitle('Proportion of different categories')  
plt.subplot(1, 2, 1)  
plt.title('Train')  
plt.pie(x=train_proportion.values(), labels=train_proportion.keys(), autopct='%0.0f')  
plt.axis('equal')  
plt.subplot(1, 2, 2)  
plt.title('Test')  
plt.pie(x=test_proportion.values(), labels=test_proportion.keys(), autopct='%0.0f')  
plt.axis('equal')  
plt.show()`



## Observation:

Train Data:

- Potato has the highest number of images at 898 (around 29% of the total). Onion is the second most frequent class with 849 images (around 27% of the total).
- Tomato is the third most frequent class with 789 images (around 25% of the total).

- Indian market has the lowest number of images at 599 (around 19% of the total).
- The training dataset is relatively balanced, with some slight variations in class distribution.

Test Data:

- Tomato has the highest number of images at 106 (around 30% of the total).
- Onion, Indian market, and Potato have a similar number of images (83, 81, and 81 respectively), each representing around 23-24% of the total.

The test dataset is also relatively balanced, with 'tomato' having a slightly higher representation compared to the other classes.

```
In [ ]: images = get_sample_images(base_data_path)
plot_images(images, title='One sample form each category from training data')
```

One sample form each category from training data



```
In [ ]: images = get_sample_images(base_data_path, traget_paths=['test'])
plot_images(images, title='One sample form each category from Testing data')
```

One sample form each category from Testing data



## Noise Data Handling

```
In [ ]: noise_images = ['market11407.png', 'market11408.png', 'market11403.png', 'market
train_indian_market_path = os.path.join(base_data_path, 'train', 'indian market'

noise_img_dict = get_noise_images(train_indian_market_path, noise_images)
plot_images(noise_img_dict, "Noise images from the 'indian market' folder from t
```

Noise images from the 'indian market' folder from the train dataset



```
In [ ]: def remove_noise_images(base_data_path, noise_images_list):
    """Removes specified noise images from the 'indian market' folder in the train
    train_indian_market_path = os.path.join(base_data_path, 'train', 'indian market')
    removed_count = 0
    for image_name in noise_images_list:
        image_path = os.path.join(train_indian_market_path, image_name)
        if os.path.exists(image_path):
            os.remove(image_path)
            print(f"Removed: {image_path}")
            removed_count += 1
        else:
            print(f"Image not found: {image_path}")
    print(f"Finished removing {removed_count} noise images.")

# Define the base data path and the list of noise images
base_data_path = 'ninjacart_data'
noise_images = ['market11407.png', 'market11408.png', 'market11403.png', 'market11404.png',
                'market11405.png', 'market11406.png']

# Call the function to remove the noise images
remove_noise_images(base_data_path, noise_images)
```

Removed: ninjacart\_data/train/indian market/market11407.png  
Removed: ninjacart\_data/train/indian market/market11408.png  
Removed: ninjacart\_data/train/indian market/market11403.png  
Removed: ninjacart\_data/train/indian market/market11404.png  
Removed: ninjacart\_data/train/indian market/market11405.png  
Removed: ninjacart\_data/train/indian market/market11406.png  
Finished removing 6 noise images.

```
In [ ]: remove_noise_images(base_data_path, noise_images)
```

Image not found: ninjacart\_data/train/indian market/market11407.png  
Image not found: ninjacart\_data/train/indian market/market11408.png  
Image not found: ninjacart\_data/train/indian market/market11403.png  
Image not found: ninjacart\_data/train/indian market/market11404.png  
Image not found: ninjacart\_data/train/indian market/market11405.png  
Image not found: ninjacart\_data/train/indian market/market11406.png  
Finished removing 0 noise images.

# Splitting the Train data into Train(80%) and Validation(20%) data sets

In [ ]:

```
import os
import glob
import shutil
from sklearn.model_selection import train_test_split

def create_directory_if_not_exists(directory_path):
    """Creates a directory if it doesn't exist."""
    if not os.path.exists(directory_path):
        os.makedirs(directory_path)
        print(f"Created directory: {directory_path}")
    else:
        print(f"Directory already exists: {directory_path}")

def split_data_for_category(category_path, validation_category_path, test_size=0
    """Splits image files for a given category into train and validation sets."""
    image_files = glob.glob(os.path.join(category_path, '*'))
    train_files, val_files = train_test_split(image_files, test_size=test_size,
                                              random_state=42)
    return train_files, val_files

def move_files(file_list, destination_path):
    """Moves a list of files to a specified destination."""
    for file_path in file_list:
        file_name = os.path.basename(file_path)
        destination = os.path.join(destination_path, file_name)
        shutil.move(file_path, destination)

# Define the base path for the data
base_data_path = 'nijacart_data'
train_path = os.path.join(base_data_path, 'train')
validation_path = os.path.join(base_data_path, 'validation')

# Create the validation directory if it doesn't exist
create_directory_if_not_exists(validation_path)

# Get the list of categories (subfolders) in the training data
categories = os.listdir(train_path)

# Iterate through each category and split and move the data
for category in categories:
    category_train_path = os.path.join(train_path, category)
    category_validation_path = os.path.join(validation_path, category)

    # Create the category subfolder in the validation directory
    create_directory_if_not_exists(category_validation_path)

    # Split the image files for the current category
    train_files, val_files = split_data_for_category(category_train_path, category_validation_path)

    # Move the validation files to the new validation directory
    print(f"Moving {len(val_files)} images from {category} to validation folder.")
    move_files(val_files, category_validation_path)

print("Data splitting and validation folder creation complete.")
```

```

Created directory: ninjacart_data/validation
Created directory: ninjacart_data/validation/tomato
Moving 158 images from tomato to validation folder...
Created directory: ninjacart_data/validation/potato
Moving 180 images from potato to validation folder...
Created directory: ninjacart_data/validation/indian market
Moving 119 images from indian market to validation folder...
Created directory: ninjacart_data/validation/onion
Moving 170 images from onion to validation folder...
Data splitting and validation folder creation complete.

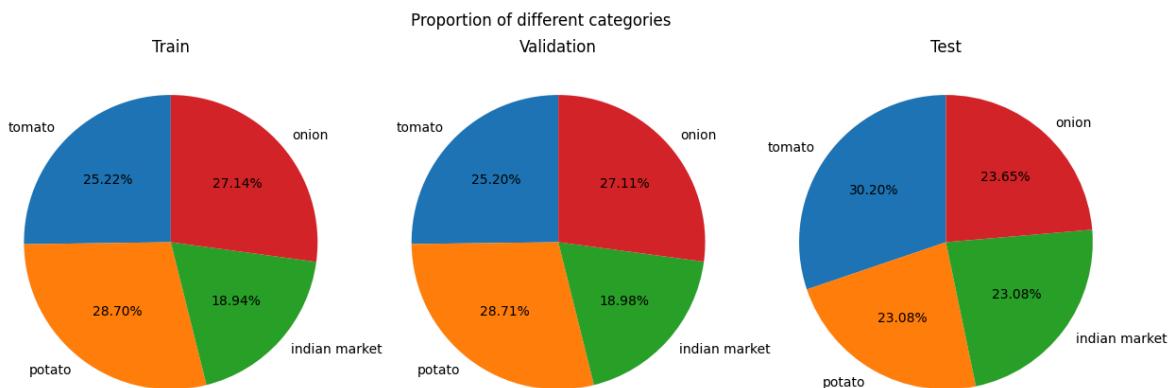
```

```
In [ ]: train_test_val_proportions = get_files_count(base_data_path, target_paths=['train', 'validation', 'test'])
pd.DataFrame(train_test_val_proportions)
```

```
Out[ ]:
```

	train	validation	test
<b>tomato</b>	631	158	106
<b>potato</b>	718	180	81
<b>indian market</b>	474	119	81
<b>onion</b>	679	170	83

```
In [ ]: train_proportion = train_test_val_proportions['train']
test_proportion = train_test_val_proportions['test']
validation_proportion = train_test_val_proportions['validation']
plt.figure(figsize=(15, 5))
plt.suptitle('Proportion of different categories')
plt.subplot(1, 3, 1)
plt.title('Train')
plt.pie(x=train_proportion.values(), labels=train_proportion.keys(), autopct='%.0f%%')
plt.axis('equal')
plt.subplot(1, 3, 2)
plt.title('Validation')
plt.pie(x=validation_proportion.values(), labels=validation_proportion.keys(), autopct='%.0f%%')
plt.axis('equal')
plt.subplot(1, 3, 3)
plt.title('Test')
plt.pie(x=test_proportion.values(), labels=test_proportion.keys(), autopct='%.0f%%')
plt.axis('equal')
plt.show()
```



Train Data (after splitting for validation):

- Potato now has 718 images, making it the most frequent class.

- Onion has 679 images, the second most frequent.
- Tomato has 631 images, the third most frequent.
- Indian market has 479 images, the least frequent. The training dataset is still relatively balanced, with 'potato' and 'onion' having slightly more images than 'tomato' and 'indian market'.

Validation Data:

- Potato has the highest number of images at 180.
- Onion has 170 images.
- Tomato has 158 images.
- Indian market has the lowest number of images at 120. The validation dataset maintains a similar distribution to the new training data, which is expected after the split.

Test Data:

- Tomato has the highest number of images at 106.
- Onion, Indian market, and Potato have a similar number of images (83, 81, and 81 respectively). The test dataset is also relatively balanced, with 'tomato' having a slightly higher representation compared to the other classes.

Overall, the data distribution across the classes is fairly consistent among the updated train, validation, and test sets.

```
In [ ]: images = get_sample_images(base_data_path, ['train'])
plot_images(images, title='One sample form each category from training data')
```

One sample form each category from training data



```
In [ ]: images = get_sample_images(base_data_path, ['validation'])
plot_images(images, title='One sample form each category from validation data')
```

One sample from each category from validation data



```
In [ ]: images = get_sample_images(base_data_path, ['test'])
plot_images(images, title='One sample from each category from test data')
```

One sample from each category from test data



```
In [ ]: image_names_dict = get_image_filenames('nиняcart_data', target_paths=['train',
```

## Preprocessing

- Shape Preprocessing
- Value Preprocessing

## Shape Preprocessing

```
In [ ]: def keras_load_data(base_path, target_paths=['train'], _label_mode='int', _batch_size=32):
    keras_data_dict = {}
    for target_path in target_paths:
        _shuffle = True if target_path == 'train' else False

        print(f'Loading {target_path} data...')
        path = os.path.join(base_path, target_path)
        data = tf.keras.utils.image_dataset_from_directory(path, shuffle=_shuffle, label_mode=_label_mode)
        keras_data_dict.update({target_path: data})

    return keras_data_dict

keras_loaded_data = keras_load_data(base_data_path, target_paths=['train', 'validation'])
```

```
Loading train data...
Found 2502 files belonging to 4 classes.
Loading validation data...
Found 627 files belonging to 4 classes.
Loading test data...
Found 351 files belonging to 4 classes.
```

```
In [ ]: keras_loaded_data
```

```
Out[ ]: {'train': <_PrefetchDataset element_spec=(TensorSpec(shape=(128, 128, 3), dtype=tf.float32, name=None), TensorSpec(shape=(), dtype=tf.int32, name=None))>, 'validation': <_PrefetchDataset element_spec=(TensorSpec(shape=(128, 128, 3), dtype=tf.float32, name=None), TensorSpec(shape=(), dtype=tf.int32, name=None))>, 'test': <_PrefetchDataset element_spec=(TensorSpec(shape=(128, 128, 3), dtype=tf.float32, name=None), TensorSpec(shape=(), dtype=tf.int32, name=None))>}
```

## Value Preprocessing

```
In [ ]: data_rescale = keras.Sequential( name = "data_rescale", layers = [ layers.Rescaling(1./255) ] ) # Perform Data Processing on the train, val, test dataset train_ds = keras_loaded_data['train'].map(lambda x, y: (data_rescale(x), y)) validation_ds = keras_loaded_data['validation'].map(lambda x, y: (data_rescale(x), y)) test_ds = keras_loaded_data['test'].map(lambda x, y: (data_rescale(x), y))
```

## Sanity check on Preprocessed data

```
In [ ]: def store_processed_images(dest_dir_name, keras_loaded_data, target_paths=['train']): """Stores processed images from a keras_loaded_data dictionary into a directory""" for target_path in target_paths: print(f'Storing {target_path} data...') source_path = os.path.join(base_data_path, target_path) dest_path = os.path.join(dest_dir_name, target_path) # Get class names from the source directory class_names = sorted(os.listdir(source_path)) create_directory_if_not_exists(dest_path) # Create category folders and counters category_counters = {} for category in class_names: category_path = os.path.join(dest_path, category) create_directory_if_not_exists(category_path) category_counters[category] = 1 # Initialize counter # Process each image individually # Iterate over the dataset which contains rescaled images for image, label in keras_loaded_data[target_path]: label = int(label.numpy()) category = class_names[label] img_path = os.path.join(dest_path, category) # Create the file name: e.g., onion1.jpg filename = f'{category}{category_counters[category]}.jpg' category_counters[category] += 1 # Increment counter
```

```
# Encode and save
encoded_image = tf.cast(image, tf.uint8)
jpeg_encoded = tf.io.encode_jpeg(encoded_image)
tf.io.write_file(os.path.join(img_path, filename), jpeg_encoded)
```

In [ ]: `store_processed_images('nijacart_data_processed', keras_loaded_data, target_pat`

```
Storing train data...
Created directory: ninjacart_data_processed/train
Created directory: ninjacart_data_processed/train/indian market
Created directory: ninjacart_data_processed/train/onion
Created directory: ninjacart_data_processed/train/potato
Created directory: ninjacart_data_processed/train/tomato
Storing validation data...
Created directory: ninjacart_data_processed/validation
Created directory: ninjacart_data_processed/validation/indian market
Created directory: ninjacart_data_processed/validation/onion
Created directory: ninjacart_data_processed/validation/potato
Created directory: ninjacart_data_processed/validation/tomato
Storing test data...
Created directory: ninjacart_data_processed/test
Created directory: ninjacart_data_processed/test/indian market
Created directory: ninjacart_data_processed/test/onion
Created directory: ninjacart_data_processed/test/potato
Created directory: ninjacart_data_processed/test/tomato
```

In [ ]: `def make_archive_zip(scr_path, dest_path, target_paths=['train']):
 create_directory_if_not_exists(dest_path)
 for target_path in target_paths:
 source_path = os.path.join(scr_path, target_path)
 # destination_path = os.path.join(dest_dir, target_)
 shutil.make_archive(os.path.join(dest_path, target_path), 'zip', source_path
 # files.download(f'{os.path.join(dest_path, target_path)}.zip')`

In [ ]: `make_archive_zip('nijacart_data_processed', 'nijacart_data_processed_zip_files`

```
Created directory: ninjacart_data_processed_zip_files
```

In [ ]: `make_archive_zip('nijacart_data', 'nijacart_zip_files', target_paths=['train',`

```
Created directory: ninjacart_zip_files
```

In [ ]: `base_data_path = 'nijacart_data_processed'
images = get_sample_images(base_data_path, ['train'])
plot_images(images, title='One sample form each category from training data afte`

One sample form each category from training data after preprocessing



In [ ]: `images = get_sample_images(base_data_path, ['test'])
plot_images(images, title='One sample form each category from test data after pr`

One sample form each category from test data after preprocessing



```
In [ ]: images = get_sample_images(base_data_path, ['validation'])
plot_images(images, title='One sample form each category from validation data af
```

One sample form each category from validation data after preprocessing



## BaseLine Model

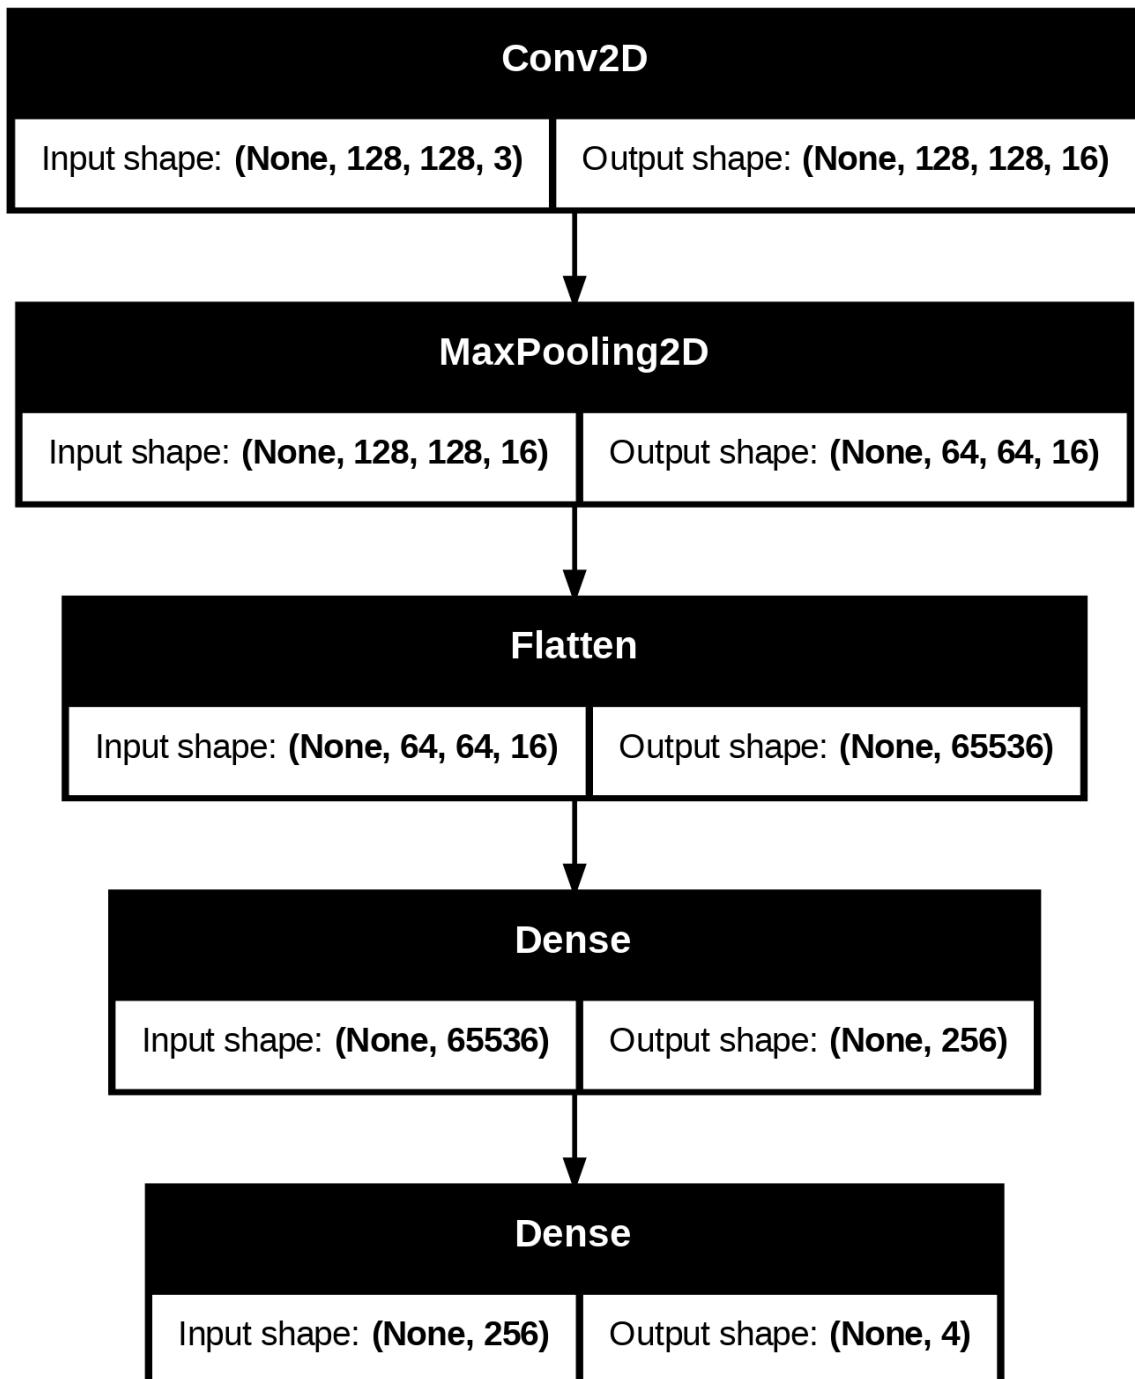
```
In [ ]: num_classes = 4
hidden_size = 256
height = 128
width = 128

model = keras.Sequential(
    name="model_cnn_without_hyperparameter_tuning",
    layers=[
        Input(shape=(height, width, 3)),
        layers.Conv2D(filters=16, kernel_size=3, strides=1, padding="same", activation="relu"),
        layers.MaxPooling2D(),
        layers.Flatten(),
        layers.Dense(units=hidden_size, activation='relu'),
        layers.Dense(units=num_classes, activation='softmax')
    ]
)
```

```
In [ ]: create_directory_if_not_exists('nunjacart_models')
create_directory_if_not_exists('nunjacart_models/model1_BaseLineModel')
tf.keras.utils.plot_model(model, to_file="/content/nunjacart_models/model1_Basel
```

Created directory: nunjacart\_models  
Created directory: nunjacart\_models/model1\_BaseLineModel

```
Out[ ]:
```



```
In [ ]: model.summary()
```

Model: "model\_cnn\_without\_hyperparameter\_tuning"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 128, 128, 16)	448
max_pooling2d (MaxPooling2D)	(None, 64, 64, 16)	0
flatten (Flatten)	(None, 65536)	0
dense (Dense)	(None, 256)	16,777,472
dense_1 (Dense)	(None, 4)	1,028

```
Total params: 16,778,948 (64.01 MB)
Trainable params: 16,778,948 (64.01 MB)
Non-trainable params: 0 (0.00 B)
```

## b. Compile the model with cross-entropy loss and adam optimizer:

```
In [ ]: model.compile(optimizer='adam',
                      loss='sparse_categorical_crossentropy',
                      metrics=['accuracy'])
```

## Training and Evaluation with CNNs

```
In [ ]: epochs = 10
batch_size = 30
train_ds_batched = keras_loaded_data['train'].batch(batch_size)
validation_ds_batched = keras_loaded_data['validation'].batch(batch_size)
test_ds_batched = keras_loaded_data['test'].batch(batch_size)

model_fit = model.fit(train_ds_batched, validation_data=validation_ds_batched, e
```

Epoch 1/10  
84/84 13s 111ms/step - accuracy: 0.3656 - loss: 2111.0425 - val\_accuracy: 0.7193 - val\_loss: 18.4912  
Epoch 2/10  
84/84 8s 97ms/step - accuracy: 0.7775 - loss: 11.4944 - val\_accuracy: 0.7927 - val\_loss: 8.3506  
Epoch 3/10  
84/84 6s 77ms/step - accuracy: 0.8469 - loss: 3.2771 - val\_accuracy: 0.7943 - val\_loss: 6.3301  
Epoch 4/10  
84/84 8s 96ms/step - accuracy: 0.8398 - loss: 4.1331 - val\_accuracy: 0.7289 - val\_loss: 10.8417  
Epoch 5/10  
84/84 11s 106ms/step - accuracy: 0.8747 - loss: 2.6937 - val\_accuracy: 0.8070 - val\_loss: 6.0605  
Epoch 6/10  
84/84 6s 74ms/step - accuracy: 0.9499 - loss: 0.6330 - val\_accuracy: 0.8006 - val\_loss: 6.1888  
Epoch 7/10  
84/84 9s 110ms/step - accuracy: 0.9457 - loss: 0.6710 - val\_accuracy: 0.8150 - val\_loss: 6.0428  
Epoch 8/10  
84/84 7s 79ms/step - accuracy: 0.9711 - loss: 0.2189 - val\_accuracy: 0.8182 - val\_loss: 5.8333  
Epoch 9/10  
84/84 10s 77ms/step - accuracy: 0.9857 - loss: 0.1303 - val\_accuracy: 0.7799 - val\_loss: 6.6017  
Epoch 10/10  
84/84 9s 109ms/step - accuracy: 0.9783 - loss: 0.1390 - val\_accuracy: 0.7512 - val\_loss: 6.1099

```
In [ ]: fig, axes = plt.subplots(nrows=1, ncols=2, figsize=(15, 5))
ax = axes.ravel()

#accuracy graph
```

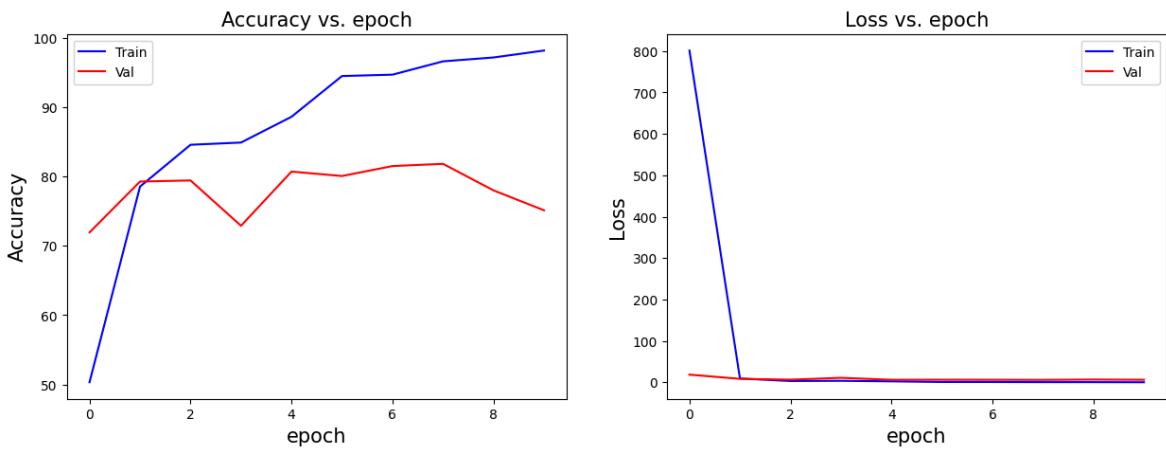
```

ax[0].plot(range(0,model_fit.params['epochs']), [acc * 100 for acc in model_fit.history['accuracy']])
ax[0].plot(range(0,model_fit.params['epochs']), [acc * 100 for acc in model_fit.history['val_accuracy']])
ax[0].set_title('Accuracy vs. epoch', fontsize=15)
ax[0].set_ylabel('Accuracy', fontsize=15)
ax[0].set_xlabel('epoch', fontsize=15)
ax[0].legend()

#Loss graph
ax[1].plot(range(0,model_fit.params['epochs']), model_fit.history['loss'], label='Train')
ax[1].plot(range(0,model_fit.params['epochs']), model_fit.history['val_loss'], label='Val')
ax[1].set_title('Loss vs. epoch', fontsize=15)
ax[1].set_ylabel('Loss', fontsize=15)
ax[1].set_xlabel('epoch', fontsize=15)
ax[1].legend()

#display the graph
plt.show()

```



```
In [ ]: model.save_weights("/content/ninjacart_models/model1_BaseLineModel/BaseLineModel")
```

```

In [ ]: # Load model from pretrained checkpoints (optional)
model.load_weights("/content/ninjacart_models/model1_BaseLineModel/BaseLineModel")

true_categories = tf.concat([y for x, y in test_ds_batched], axis=0)
images = tf.concat([x for x, y in test_ds_batched], axis=0)
y_pred = model.predict(test_ds_batched)
class_names = keras_loaded_data['test'].class_names
predicted_categories = tf.argmax(y_pred, axis=1)

test_acc = metrics.accuracy_score(true_categories, predicted_categories) * 100
print(f'\nTest Accuracy: {test_acc:.2f}\n')

```

12/12 ————— 1s 111ms/step

Test Accuracy: 72.08%

```

In [ ]: def ConfusionMatrix(model, ds, label_list):
    # Note: This logic doesn't work with shuffled datasets
    # run model prediction and obtain probabilities
    y_pred = model.predict(ds)
    # get list of predicted classes by taking argmax of the probabilities(y_pred)
    predicted_categories = tf.argmax(y_pred, axis=1)
    # create list of all "y's Labels, by iterating over test dataset
    true_categories = tf.concat([y for x, y in ds], axis=0)
    # generate confusion matrix and plot it

```

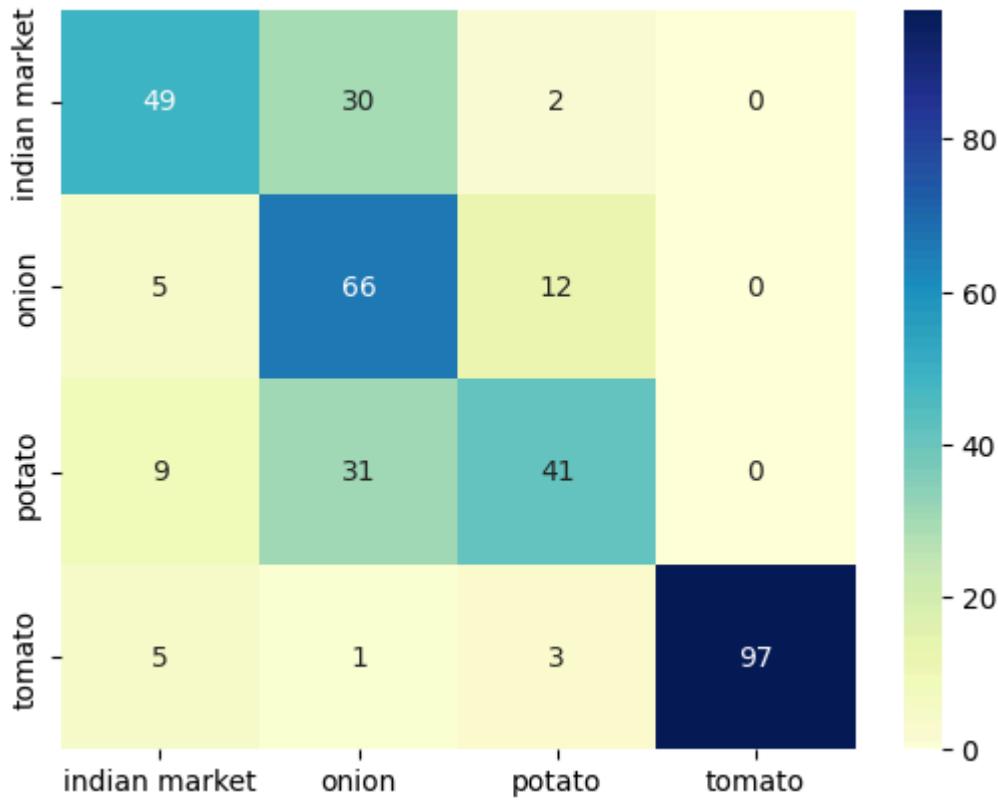
```

cm = metrics.confusion_matrix(true_categories,predicted_categories) # last b
sns.heatmap(cm, annot=True, xticklabels=label_list, yticklabels=label_list,
plt.show()

ConfusionMatrix(model, test_ds_batched, keras_loaded_data['test'].class_names)

```

12/12 ————— 1s 65ms/step



```

In [ ]: def plot_image(pred_array, true_label, img):
    plt.grid(False)
    plt.xticks([])
    plt.yticks([])

    # Explicitly cast the image data to uint8 before plotting
    img_to_plot = tf.cast(img, tf.uint8)
    plt.imshow(img_to_plot, cmap=plt.cm.binary)

    predicted_label = np.argmax(pred_array)
    if predicted_label == true_label:
        color = 'blue'
    else:
        color = 'red'

    plt.xlabel("{} {:.2f}% ({})".format(class_names[predicted_label],
                                         100*np.max(pred_array),
                                         class_names[true_label]),
                                         color=color)

# function to plot barplot of class probabilities (pred_array)
def plot_value_array(pred_array, true_label):
    plt.grid(False)
    plt.xticks(range(4))
    plt.yticks([])
    thisplot = plt.bar(range(4), pred_array, color="#777777")
    plt.ylim([0, 1])
    predicted_label = np.argmax(pred_array)

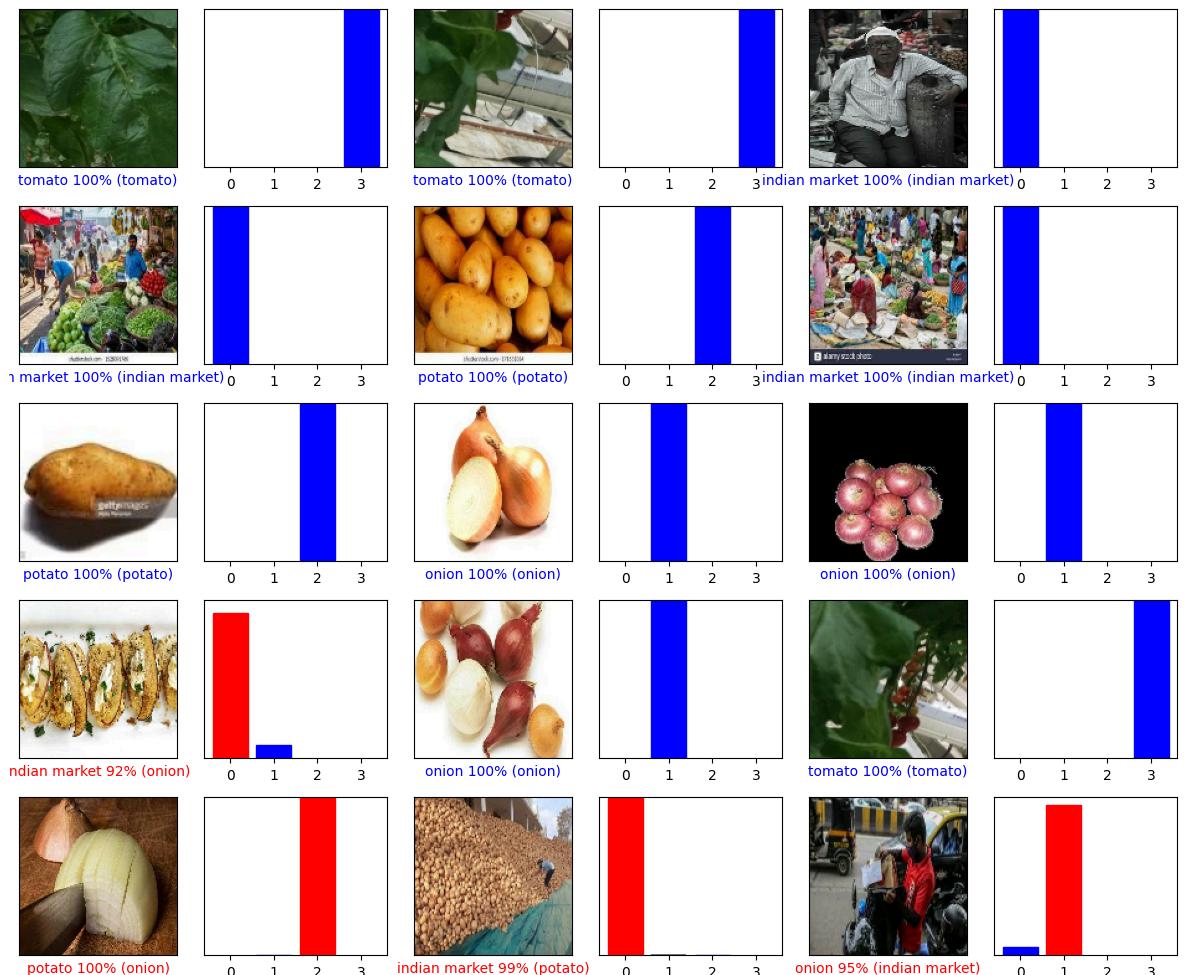
```

```
thisplot[predicted_label].set_color('red')
thisplot[true_label].set_color('blue')
```

```
In [ ]: true_categories = tf.concat([y for x, y in test_ds_batched], axis=0)
images = tf.concat([x for x, y in test_ds_batched], axis=0)
y_pred = model.predict(test_ds_batched)
class_names = keras_loaded_data['test'].class_names

# print(y_pred)
# print(true_categories)
# Randomly sample 15 test images and plot it with their predicted Labels, and the
indices = random.sample(range(len(images)), 15)
# Color correct predictions in blue and incorrect predictions in red.
num_rows = 5
num_cols = 3
num_images = num_rows*num_cols
plt.figure(figsize=(2*2*num_cols, 2*num_rows))
for i,index in enumerate(indices):
    plt.subplot(num_rows, 2*num_cols, 2*i+1)
    plot_image(y_pred[index], true_categories[index], images[index])
    plt.subplot(num_rows, 2*num_cols, 2*i+2)
    plot_value_array(y_pred[index], true_categories[index])
plt.tight_layout()
plt.show()
```

12/12 ━━━━━━ 1s 115ms/step



## Model2 - Using Global Average pooling

```
In [ ]: def arch_1(height=128, width=128):
    num_classes = 4
    hidden_size = 256

    model = keras.Sequential(
        name="model_cnn_1",
        layers=[
            Input(shape=(height, width, 3)),
            layers.Conv2D(filters=16, kernel_size=3, padding="same", activation=
            layers.MaxPooling2D(),
            layers.Conv2D(filters=32, kernel_size=3, padding="same", activation=
            layers.MaxPooling2D(),
            layers.Conv2D(filters=64, kernel_size=3, padding="same", activation=
            layers.MaxPooling2D(),
            layers.Conv2D(filters=128, kernel_size=3, padding="same", activation=
            layers.MaxPooling2D(),
            layers.Conv2D(filters=256, kernel_size=3, padding="same", activation=
            # layers.MaxPooling2D(),
            # layers.Flatten(),
            layers.GlobalAveragePooling2D(),
            layers.Dense(units=hidden_size, activation='relu'),
            layers.Dense(units=num_classes, activation='softmax')
        ]
    )
    return model
```

```
In [ ]: model = arch_1()
model.summary()
```

Model: "model\_cnn\_1"

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 128, 128, 16)	448
max_pooling2d_1 (MaxPooling2D)	(None, 64, 64, 16)	0
conv2d_2 (Conv2D)	(None, 64, 64, 32)	4,640
max_pooling2d_2 (MaxPooling2D)	(None, 32, 32, 32)	0
conv2d_3 (Conv2D)	(None, 32, 32, 64)	18,496
max_pooling2d_3 (MaxPooling2D)	(None, 16, 16, 64)	0
conv2d_4 (Conv2D)	(None, 16, 16, 128)	73,856
max_pooling2d_4 (MaxPooling2D)	(None, 8, 8, 128)	0
conv2d_5 (Conv2D)	(None, 8, 8, 256)	295,168
global_average_pooling2d (GlobalAveragePooling2D)	(None, 256)	0
dense_2 (Dense)	(None, 256)	65,792
dense_3 (Dense)	(None, 4)	1,028

Total params: 459,428 (1.75 MB)

Trainable params: 459,428 (1.75 MB)

Non-trainable params: 0 (0.00 B)

```
In [ ]: create_directory_if_not_exists('/content/ninjacart_models/model2_with_GAP')
```

Created directory: /content/ninjacart\_models/model2\_with\_GAP

```
In [ ]: # model_fit = compile_train_v1(model, train_ds, val_ds)
base_data_path = '/content/ninjacart_data'
keras_loaded_data = keras_load_data(base_data_path, target_paths=['train', 'vali
```

Loading train data...  
 Found 2502 files belonging to 4 classes.  
 Loading validation data...  
 Found 627 files belonging to 4 classes.  
 Loading test data...  
 Found 351 files belonging to 4 classes.

```
In [ ]: train_ds = keras_loaded_data['train'].map(lambda x, y: (data_rescale(x), y))
validation_ds = keras_loaded_data['validation'].map(lambda x, y: (data_rescale(x), y))
test_ds = keras_loaded_data['test'].map(lambda x, y: (data_rescale(x), y))
```

```
In [ ]: def compile_train_v1(model, train_ds, val_ds, _callback, epochs=10):
    model.compile(optimizer='adam',
                  loss='categorical_crossentropy',
                  metrics=['accuracy'])

    model_fit = model.fit(train_ds, validation_data=val_ds, epochs=epochs, callbacks=_callback)
    return model_fit
```

```
In [ ]: ckpt_path="/content/ninjacart_models/model2_with_GAP/gap_model_10_epochs.weights
callbacks=[keras.callbacks.ModelCheckpoint(ckpt_path, save_weights_only=True, monitor='val_accuracy')]
```

```
In [ ]: epochs = 10
batch_size = 30
train_ds_batched = keras_loaded_data['train'].batch(batch_size)
validation_ds_batched = keras_loaded_data['validation'].batch(batch_size)
test_ds_batched = keras_loaded_data['test'].batch(batch_size)
model_fit = compile_train_v1(model, train_ds_batched, validation_ds_batched, callbacks)
```

Epoch 1/10  
**84/84** 16s 131ms/step - accuracy: 0.3064 - loss: 7.1938 - val\_accuracy: 0.5614 - val\_loss: 0.9352  
Epoch 2/10  
**84/84** 7s 78ms/step - accuracy: 0.5533 - loss: 0.9396 - val\_accuracy: 0.6651 - val\_loss: 0.8803  
Epoch 3/10  
**84/84** 8s 97ms/step - accuracy: 0.6997 - loss: 0.7457 - val\_accuracy: 0.7895 - val\_loss: 0.5413  
Epoch 4/10  
**84/84** 7s 86ms/step - accuracy: 0.8185 - loss: 0.4770 - val\_accuracy: 0.8166 - val\_loss: 0.4605  
Epoch 5/10  
**84/84** 10s 77ms/step - accuracy: 0.8206 - loss: 0.4283 - val\_accuracy: 0.7911 - val\_loss: 0.5262  
Epoch 6/10  
**84/84** 8s 97ms/step - accuracy: 0.8077 - loss: 0.4919 - val\_accuracy: 0.8724 - val\_loss: 0.3469  
Epoch 7/10  
**84/84** 7s 78ms/step - accuracy: 0.8450 - loss: 0.3964 - val\_accuracy: 0.8389 - val\_loss: 0.4023  
Epoch 8/10  
**84/84** 8s 96ms/step - accuracy: 0.8337 - loss: 0.4038 - val\_accuracy: 0.8628 - val\_loss: 0.3343  
Epoch 9/10  
**84/84** 7s 79ms/step - accuracy: 0.8675 - loss: 0.3098 - val\_accuracy: 0.7799 - val\_loss: 0.5083  
Epoch 10/10  
**84/84** 9s 110ms/step - accuracy: 0.8445 - loss: 0.3846 - val\_accuracy: 0.8565 - val\_loss: 0.3463

```
In [ ]: # helper function to annotate maximum values in the plots
def annot_max(x,y, xytext=(0.94,0.96), ax=None, only_y=True):
    xmax = x[np.argmax(y)]
    ymax = max(y)
    if only_y:
        text = "{:.2f}%".format(ymax)
    else:
        text= "x={:.2f}, y={:.2f}%".format(xmax, ymax)
```

```

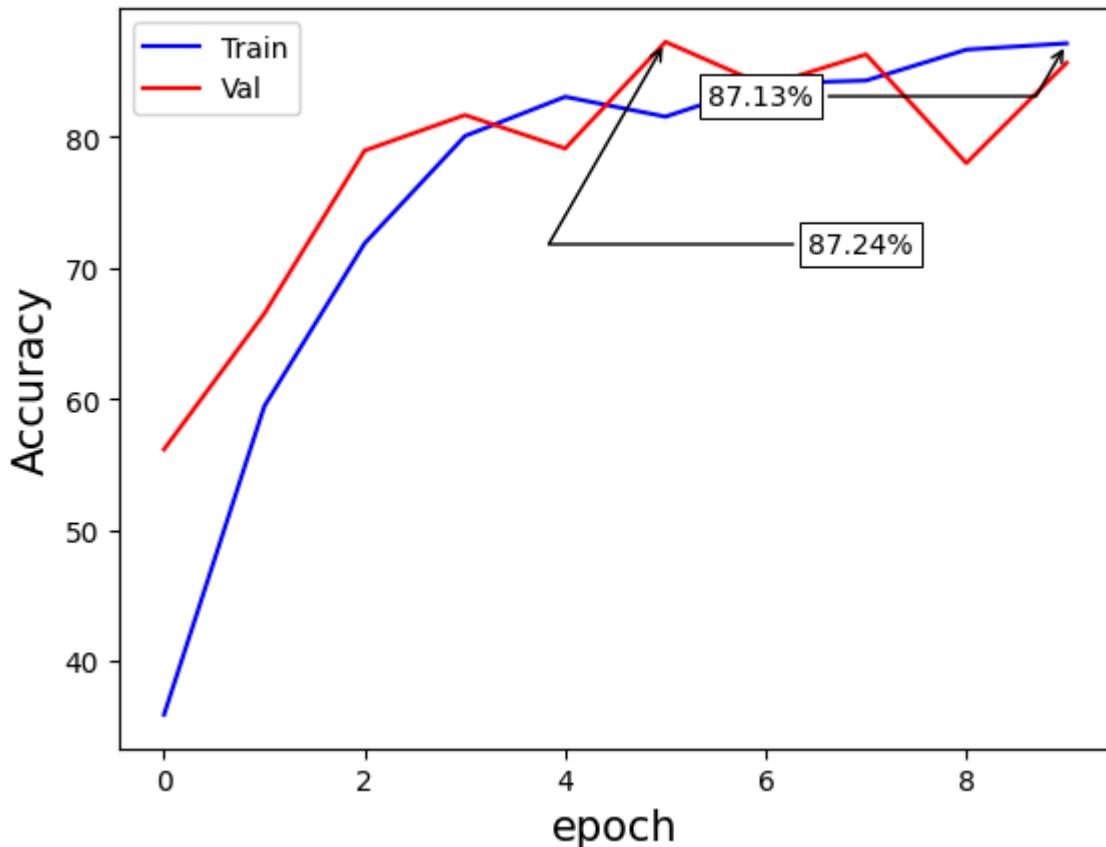
if not ax:
    ax=plt.gca()
bbox_props = dict(boxstyle="square,pad=0.3", fc="w", ec="k", lw=0.72)
arrowprops=dict(arrowstyle="->",connectionstyle="angle,angleA=0,angleB=60")
kw = dict(xycoords='data',textcoords="axes fraction",
          arrowprops=arrowprops, bbox=bbox_props, ha="right", va="top")
ax.annotate(text, xy=(xmax, ymax), xytext=xytext, **kw)

def plot_accuracy(model_fit):
    #accuracy graph
    x = range(0,len(model_fit.history['accuracy']))
    y_train = [acc * 100 for acc in model_fit.history['accuracy']]
    y_val = [acc * 100 for acc in model_fit.history['val_accuracy']]

    plt.plot(x, y_train, label='Train', color='b')
    annot_max(x, y_train, xytext=(0.7,0.9))
    plt.plot(x, y_val, label='Val', color='r')
    annot_max(x, y_val, xytext=(0.8,0.7))
    plt.ylabel('Accuracy', fontsize=15)
    plt.xlabel('epoch', fontsize=15)
    plt.legend()
    plt.show()

```

In [ ]: `plot_accuracy(model_fit)`



In [ ]: `def print_accuracy_stats(model, ds, class_names, weights_path):`  
`model.load_weights(weights_path)`

`# Collect true labels by iterating through the dataset`  
`true_categories_list = []`  
`for _, labels in ds:`  
`# Check rank using tf.rank and compare as integers`  
`if tf.rank(labels).numpy() > 1:`

```

        true_categories_list.append(tf.argmax(labels, axis=1).numpy())
    else:
        true_categories_list.append(labels.numpy())

true_categories = np.concatenate(true_categories_list, axis=0)

y_pred = model.predict(ds)
predicted_categories = tf.argmax(y_pred, axis=1)

test_acc = metrics.accuracy_score(true_categories, predicted_categories) * 1
print(f'\nTest Accuracy: {test_acc:.2f}%\n')

# Note: This doesn't work with shuffled datasets - This note is still relevant f
def plot_confusion_matrix(model, ds, class_names, weights_path):
    model.load_weights(weights_path)

    # Collect true Labels by iterating through the dataset
    true_categories_list = []
    for _, labels in ds:
        # Check rank using tf.rank and compare as integers
        if tf.rank(labels).numpy() > 1:
            true_categories_list.append(tf.argmax(labels, axis=1).numpy())
        else:
            true_categories_list.append(labels.numpy())

    true_categories = np.concatenate(true_categories_list, axis=0)

    y_pred = model.predict(ds)
    predicted_categories = tf.argmax(y_pred, axis=1)
    cm = metrics.confusion_matrix(true_categories, predicted_categories)
    sns.heatmap(cm, annot=True, xticklabels=class_names, yticklabels=class_names)
    plt.show()

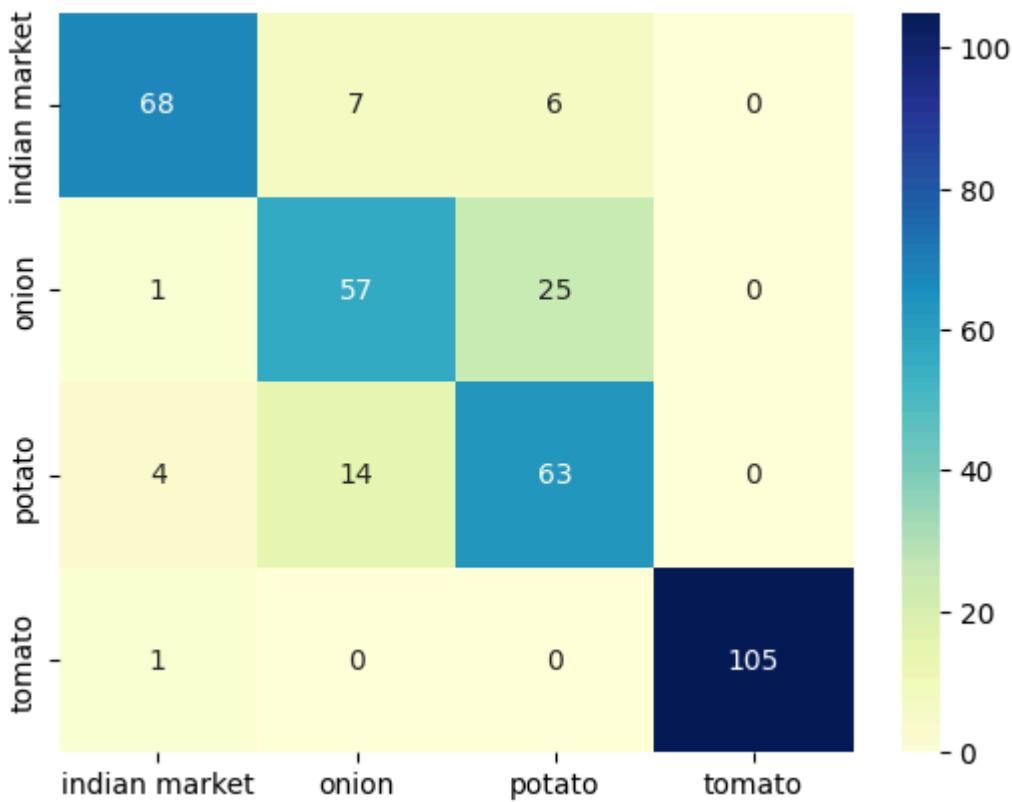
```

In [ ]: `print_accuracy_stats(model, test_ds_batched, class_names, ckpt_path)`  
`plot_confusion_matrix(model, test_ds_batched, class_names, ckpt_path)`

**12/12** ————— **2s** 130ms/step

Test Accuracy: 83.48%

**12/12** ————— **1s** 91ms/step



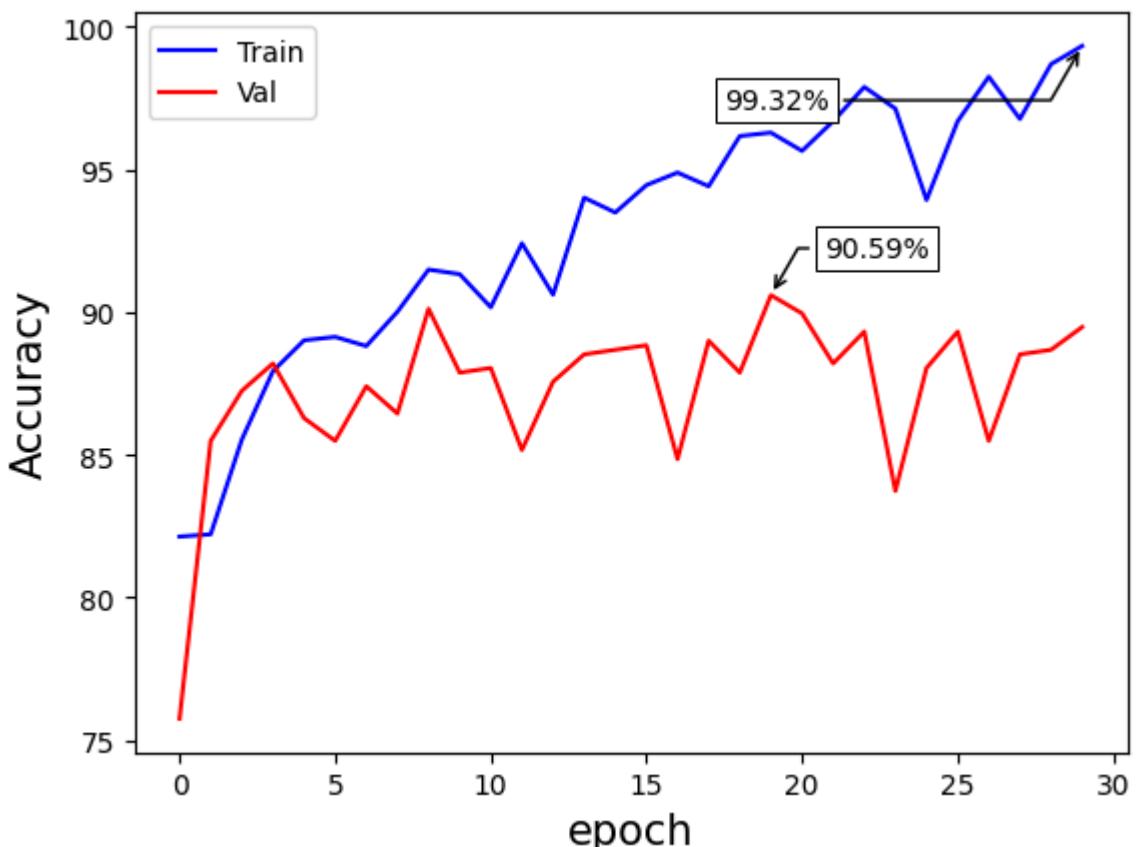
```
In [ ]: ckpt_path="/content/ninjacart_models/model2_with_GAP/gap_model_30_epochs.weights  
callbacks=[  
    keras.callbacks.ModelCheckpoint(ckpt_path, save_weights_only=True, monit  
]
```

```
In [ ]: # model_fit = compile_train_v1(model, train_ds_batched, validation_ds_batched, e  
model_fit = compile_train_v1(model, train_ds_batched, validation_ds_batched, ca
```

Epoch 1/30  
**84/84** 12s 101ms/step - accuracy: 0.7902 - loss: 0.5468 - val\_accuracy: 0.7576 - val\_loss: 0.5714  
Epoch 2/30  
**84/84** 8s 91ms/step - accuracy: 0.8280 - loss: 0.4315 - val\_accuracy: 0.8549 - val\_loss: 0.3476  
Epoch 3/30  
**84/84** 8s 101ms/step - accuracy: 0.8555 - loss: 0.3695 - val\_accuracy: 0.8724 - val\_loss: 0.3577  
Epoch 4/30  
**84/84** 7s 79ms/step - accuracy: 0.8828 - loss: 0.3003 - val\_accuracy: 0.8820 - val\_loss: 0.3140  
Epoch 5/30  
**84/84** 8s 96ms/step - accuracy: 0.8790 - loss: 0.2836 - val\_accuracy: 0.8628 - val\_loss: 0.3941  
Epoch 6/30  
**84/84** 7s 78ms/step - accuracy: 0.8956 - loss: 0.2710 - val\_accuracy: 0.8549 - val\_loss: 0.3926  
Epoch 7/30  
**84/84** 8s 96ms/step - accuracy: 0.8827 - loss: 0.2906 - val\_accuracy: 0.8740 - val\_loss: 0.3982  
Epoch 8/30  
**84/84** 7s 78ms/step - accuracy: 0.9042 - loss: 0.2550 - val\_accuracy: 0.8644 - val\_loss: 0.3169  
Epoch 9/30  
**84/84** 8s 97ms/step - accuracy: 0.9094 - loss: 0.2134 - val\_accuracy: 0.9011 - val\_loss: 0.3030  
Epoch 10/30  
**84/84** 7s 77ms/step - accuracy: 0.9192 - loss: 0.2035 - val\_accuracy: 0.8788 - val\_loss: 0.3984  
Epoch 11/30  
**84/84** 8s 96ms/step - accuracy: 0.8906 - loss: 0.2831 - val\_accuracy: 0.8804 - val\_loss: 0.3886  
Epoch 12/30  
**84/84** 8s 94ms/step - accuracy: 0.9310 - loss: 0.1905 - val\_accuracy: 0.8517 - val\_loss: 0.3802  
Epoch 13/30  
**84/84** 7s 83ms/step - accuracy: 0.9038 - loss: 0.2302 - val\_accuracy: 0.8756 - val\_loss: 0.3763  
Epoch 14/30  
**84/84** 9s 109ms/step - accuracy: 0.9268 - loss: 0.1815 - val\_accuracy: 0.8852 - val\_loss: 0.3459  
Epoch 15/30  
**84/84** 7s 79ms/step - accuracy: 0.9419 - loss: 0.1509 - val\_accuracy: 0.8868 - val\_loss: 0.3082  
Epoch 16/30  
**84/84** 8s 97ms/step - accuracy: 0.9462 - loss: 0.1356 - val\_accuracy: 0.8884 - val\_loss: 0.3958  
Epoch 17/30  
**84/84** 7s 78ms/step - accuracy: 0.9519 - loss: 0.1133 - val\_accuracy: 0.8485 - val\_loss: 0.4256  
Epoch 18/30  
**84/84** 8s 97ms/step - accuracy: 0.9367 - loss: 0.1665 - val\_accuracy: 0.8900 - val\_loss: 0.3541  
Epoch 19/30  
**84/84** 8s 95ms/step - accuracy: 0.9567 - loss: 0.1024 - val\_accuracy: 0.8788 - val\_loss: 0.4244  
Epoch 20/30  
**84/84** 9s 81ms/step - accuracy: 0.9664 - loss: 0.0905 - val\_accuracy: 0.9059 - val\_loss: 0.3596

```
Epoch 21/30
84/84 8s 98ms/step - accuracy: 0.9515 - loss: 0.1391 - val_accuracy: 0.8995 - val_loss: 0.4170
Epoch 22/30
84/84 7s 79ms/step - accuracy: 0.9771 - loss: 0.0608 - val_accuracy: 0.8820 - val_loss: 0.3782
Epoch 23/30
84/84 8s 98ms/step - accuracy: 0.9723 - loss: 0.0752 - val_accuracy: 0.8931 - val_loss: 0.4028
Epoch 24/30
84/84 8s 95ms/step - accuracy: 0.9808 - loss: 0.0562 - val_accuracy: 0.8373 - val_loss: 0.7279
Epoch 25/30
84/84 9s 79ms/step - accuracy: 0.9191 - loss: 0.2369 - val_accuracy: 0.8804 - val_loss: 0.3938
Epoch 26/30
84/84 8s 97ms/step - accuracy: 0.9757 - loss: 0.0714 - val_accuracy: 0.8931 - val_loss: 0.4045
Epoch 27/30
84/84 7s 79ms/step - accuracy: 0.9809 - loss: 0.0687 - val_accuracy: 0.8549 - val_loss: 0.7350
Epoch 28/30
84/84 8s 98ms/step - accuracy: 0.9682 - loss: 0.0872 - val_accuracy: 0.8852 - val_loss: 0.4622
Epoch 29/30
84/84 11s 108ms/step - accuracy: 0.9798 - loss: 0.0626 - val_accuracy: 0.8868 - val_loss: 0.6122
Epoch 30/30
84/84 7s 80ms/step - accuracy: 0.9904 - loss: 0.0247 - val_accuracy: 0.8947 - val_loss: 0.5931
```

```
In [ ]: plot_accuracy(model_fit)
```

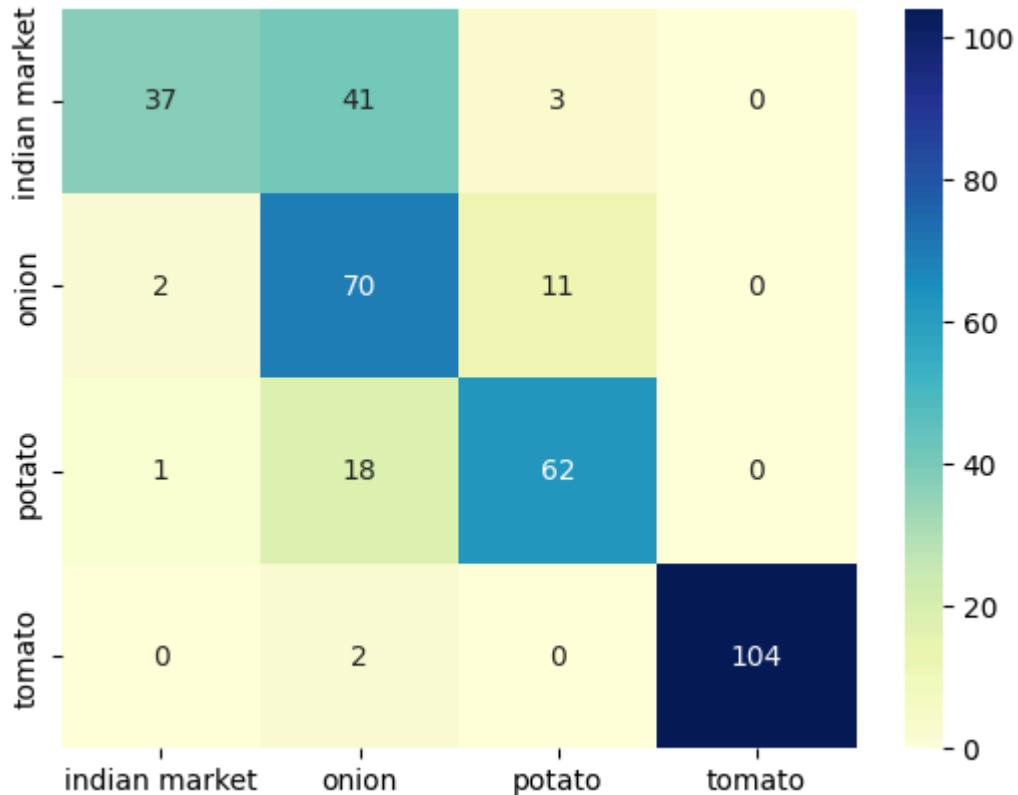


```
In [ ]: print_accuracy_stats(model, test_ds_batched, class_names, ckpt_path)
plot_confusion_matrix(model, test_ds_batched, class_names, ckpt_path)
```

12/12 ━━━━━━━━ 2s 109ms/step

Test Accuracy: 77.78%

12/12 ━━━━━━━━ 1s 109ms/step



## Model3 - Using Dropout and Batch Normalization Layers

```
In [ ]: def arch_2(height=128, width=128):
    num_classes = 4
    hidden_size = 256

    model = keras.Sequential(
        name="model_cnn_2",
        layers=[
            Input(shape=(height, width, 3)),
            layers.Conv2D(filters=16, kernel_size=3, padding="same"),
            layers.Activation("relu"),
            layers.BatchNormalization(),
            layers.MaxPooling2D(),
            layers.Conv2D(filters=32, kernel_size=3, padding="same"),
            layers.Activation("relu"),
            layers.BatchNormalization(),
            layers.MaxPooling2D(),
            layers.Conv2D(filters=64, kernel_size=3, padding="same"),
            layers.Activation("relu"),
            layers.BatchNormalization(),
            layers.MaxPooling2D(),
```

```
        layers.Conv2D(filters=128, kernel_size=3, padding="same"),
        layers.Activation("relu"),
        layers.BatchNormalization(),
        layers.MaxPooling2D(),
        layers.Conv2D(filters=256, kernel_size=3, padding="same"),
        layers.Activation("relu"),
        layers.BatchNormalization(),
        # layers.MaxPooling2D(),
        # layers.Flatten(),
        layers.GlobalAveragePooling2D(),
        layers.Dense(units=hidden_size),
        layers.Activation("relu"),
        layers.BatchNormalization(),
        layers.Dropout(0.5),
        layers.Dense(units=num_classes, activation='softmax')
    ]
)
return model
```

```
In [ ]: model = arch_2()
model.summary()
```

```
Model: "model_cnn_2"
```

Layer (type)	Output Shape	Param #
conv2d_6 (Conv2D)	(None, 128, 128, 16)	448
activation (Activation)	(None, 128, 128, 16)	0
batch_normalization (BatchNormalization)	(None, 128, 128, 16)	64
max_pooling2d_5 (MaxPooling2D)	(None, 64, 64, 16)	0
conv2d_7 (Conv2D)	(None, 64, 64, 32)	4,640
activation_1 (Activation)	(None, 64, 64, 32)	0
batch_normalization_1 (BatchNormalization)	(None, 64, 64, 32)	128
max_pooling2d_6 (MaxPooling2D)	(None, 32, 32, 32)	0
conv2d_8 (Conv2D)	(None, 32, 32, 64)	18,496
activation_2 (Activation)	(None, 32, 32, 64)	0
batch_normalization_2 (BatchNormalization)	(None, 32, 32, 64)	256
max_pooling2d_7 (MaxPooling2D)	(None, 16, 16, 64)	0
conv2d_9 (Conv2D)	(None, 16, 16, 128)	73,856
activation_3 (Activation)	(None, 16, 16, 128)	0
batch_normalization_3 (BatchNormalization)	(None, 16, 16, 128)	512
max_pooling2d_8 (MaxPooling2D)	(None, 8, 8, 128)	0
conv2d_10 (Conv2D)	(None, 8, 8, 256)	295,168
activation_4 (Activation)	(None, 8, 8, 256)	0
batch_normalization_4 (BatchNormalization)	(None, 8, 8, 256)	1,024
global_average_pooling2d_1 (GlobalAveragePooling2D)	(None, 256)	0
dense_4 (Dense)	(None, 256)	65,792
activation_5 (Activation)	(None, 256)	0
batch_normalization_5 (BatchNormalization)	(None, 256)	1,024
dropout (Dropout)	(None, 256)	0

dense_5 (Dense)	(None, 4)	1,028
-----------------	-----------	-------

Total params: 462,436 (1.76 MB)

Trainable params: 460,932 (1.76 MB)

Non-trainable params: 1,504 (5.88 KB)

```
In [ ]: create_directory_if_not_exists('/content/ninjacart_models/model3_with_GAP_BN_DO'
    ckpt_path="/content/ninjacart_models/model3_with_GAP_BN_DO/model3_with_GAP_BN_DO"
    callbacks=[
        keras.callbacks.ModelCheckpoint(ckpt_path, save_weights_only=True, monitor='val_accuracy')
    ]
```

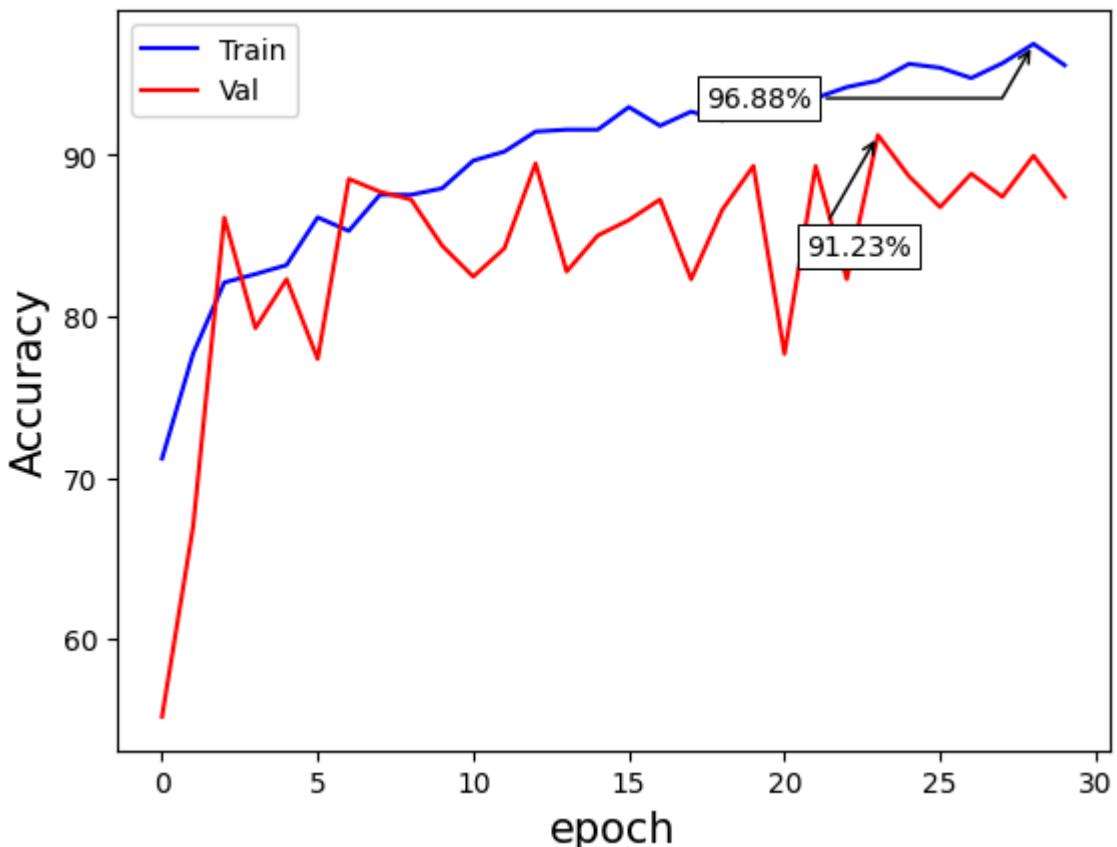
Created directory: /content/ninjacart\_models/model3\_with\_GAP\_BN\_DO

```
In [ ]: model_fit = compile_train_v1(model, train_ds_batched, validation_ds_batched, callbacks)
```

Epoch 1/30  
**84/84** 20s 150ms/step - accuracy: 0.6507 - loss: 1.0929 - val\_accuracy: 0.5518 - val\_loss: 1.2198  
Epoch 2/30  
**84/84** 8s 98ms/step - accuracy: 0.7661 - loss: 0.6533 - val\_accuracy: 0.6699 - val\_loss: 1.0389  
Epoch 3/30  
**84/84** 7s 81ms/step - accuracy: 0.8057 - loss: 0.5186 - val\_accuracy: 0.8612 - val\_loss: 0.4023  
Epoch 4/30  
**84/84** 8s 99ms/step - accuracy: 0.8229 - loss: 0.4633 - val\_accuracy: 0.7927 - val\_loss: 0.5806  
Epoch 5/30  
**84/84** 7s 81ms/step - accuracy: 0.8094 - loss: 0.4932 - val\_accuracy: 0.8230 - val\_loss: 0.4313  
Epoch 6/30  
**84/84** 8s 96ms/step - accuracy: 0.8460 - loss: 0.3985 - val\_accuracy: 0.7735 - val\_loss: 0.7846  
Epoch 7/30  
**84/84** 8s 95ms/step - accuracy: 0.8386 - loss: 0.3958 - val\_accuracy: 0.8852 - val\_loss: 0.2918  
Epoch 8/30  
**84/84** 7s 83ms/step - accuracy: 0.8824 - loss: 0.3383 - val\_accuracy: 0.8772 - val\_loss: 0.3710  
Epoch 9/30  
**84/84** 8s 97ms/step - accuracy: 0.8642 - loss: 0.3824 - val\_accuracy: 0.8724 - val\_loss: 0.3984  
Epoch 10/30  
**84/84** 7s 78ms/step - accuracy: 0.8717 - loss: 0.3053 - val\_accuracy: 0.8437 - val\_loss: 0.3744  
Epoch 11/30  
**84/84** 8s 97ms/step - accuracy: 0.8925 - loss: 0.2877 - val\_accuracy: 0.8246 - val\_loss: 0.5320  
Epoch 12/30  
**84/84** 7s 81ms/step - accuracy: 0.8938 - loss: 0.2657 - val\_accuracy: 0.8421 - val\_loss: 0.4545  
Epoch 13/30  
**84/84** 8s 99ms/step - accuracy: 0.9018 - loss: 0.2628 - val\_accuracy: 0.8947 - val\_loss: 0.3347  
Epoch 14/30  
**84/84** 7s 80ms/step - accuracy: 0.9276 - loss: 0.2057 - val\_accuracy: 0.8278 - val\_loss: 0.4184  
Epoch 15/30  
**84/84** 8s 98ms/step - accuracy: 0.9081 - loss: 0.2518 - val\_accuracy: 0.8501 - val\_loss: 0.3844  
Epoch 16/30  
**84/84** 8s 90ms/step - accuracy: 0.9260 - loss: 0.2272 - val\_accuracy: 0.8596 - val\_loss: 0.4256  
Epoch 17/30  
**84/84** 9s 80ms/step - accuracy: 0.9094 - loss: 0.2412 - val\_accuracy: 0.8724 - val\_loss: 0.3552  
Epoch 18/30  
**84/84** 8s 99ms/step - accuracy: 0.9356 - loss: 0.1895 - val\_accuracy: 0.8230 - val\_loss: 0.6499  
Epoch 19/30  
**84/84** 7s 80ms/step - accuracy: 0.9206 - loss: 0.2080 - val\_accuracy: 0.8660 - val\_loss: 0.4578  
Epoch 20/30  
**84/84** 8s 99ms/step - accuracy: 0.9264 - loss: 0.1801 - val\_accuracy: 0.8931 - val\_loss: 0.3578

```
Epoch 21/30
84/84 8s 92ms/step - accuracy: 0.9322 - loss: 0.1737 - val_accuracy: 0.7767 - val_loss: 0.5710
Epoch 22/30
84/84 15s 153ms/step - accuracy: 0.9345 - loss: 0.1684 - val_accuracy: 0.8931 - val_loss: 0.2970
Epoch 23/30
84/84 16s 98ms/step - accuracy: 0.9503 - loss: 0.1410 - val_accuracy: 0.8230 - val_loss: 0.6571
Epoch 24/30
84/84 7s 81ms/step - accuracy: 0.9480 - loss: 0.1369 - val_accuracy: 0.9123 - val_loss: 0.2605
Epoch 25/30
84/84 8s 99ms/step - accuracy: 0.9502 - loss: 0.1227 - val_accuracy: 0.8868 - val_loss: 0.3576
Epoch 26/30
84/84 8s 95ms/step - accuracy: 0.9598 - loss: 0.1099 - val_accuracy: 0.8676 - val_loss: 0.4969
Epoch 27/30
84/84 9s 79ms/step - accuracy: 0.9427 - loss: 0.1796 - val_accuracy: 0.8884 - val_loss: 0.5453
Epoch 28/30
84/84 10s 113ms/step - accuracy: 0.9491 - loss: 0.1484 - val_accuracy: 0.8740 - val_loss: 0.4570
Epoch 29/30
84/84 7s 88ms/step - accuracy: 0.9685 - loss: 0.1061 - val_accuracy: 0.8995 - val_loss: 0.4009
Epoch 30/30
84/84 10s 80ms/step - accuracy: 0.9573 - loss: 0.1091 - val_accuracy: 0.8740 - val_loss: 0.3829
```

```
In [ ]: plot_accuracy(model_fit)
```



```
In [ ]: create_directory_if_not_exists('/content/ninjacart_models/model4_with_GAP_BN_DO_ckpt_path = '/content/ninjacart_models/model4_with_GAP_BN_DO_CB/model4_with_GAP_callbacks = [  
    keras.callbacks.ReduceLROnPlateau(  
        monitor="val_loss", factor=0.3, patience=5, min_lr=0.00001  
    ),  
    keras.callbacks.ModelCheckpoint(ckpt_path, save_weights_only=True, monitor="val_loss", mode='min')  
    keras.callbacks.EarlyStopping(  
        monitor="val_loss", patience=10, min_delta=0.001, mode='min'  
    )]
```

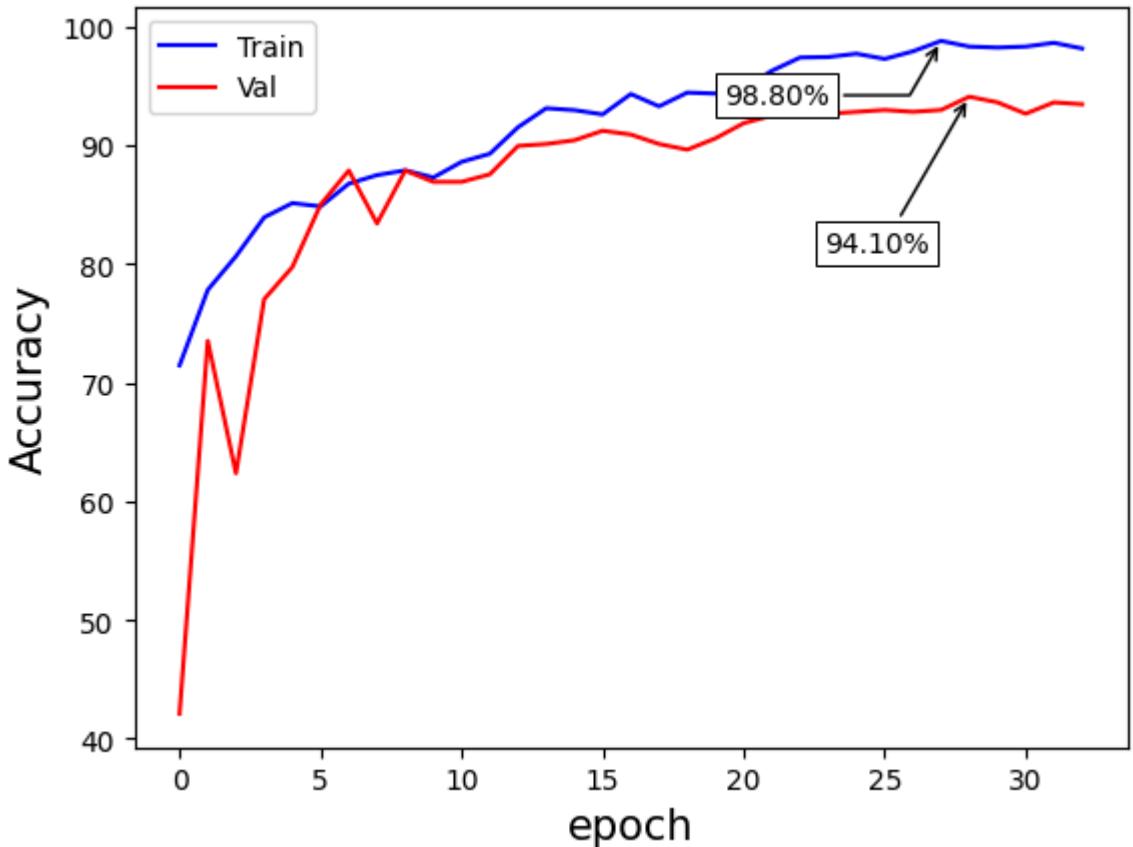
Created directory: /content/ninjacart\_models/model4\_with\_GAP\_BN\_DO\_CB

```
In [ ]: model = arch_2()  
model_fit = compile_train_v1(model, train_ds_batched, validation_ds_batched, callbacks)
```

Epoch 1/100  
**84/84** 20s 148ms/step - accuracy: 0.6546 - loss: 1.0235 - val\_accuracy: 0.4211 - val\_loss: 1.6457 - learning\_rate: 0.0010  
Epoch 2/100  
**84/84** 12s 84ms/step - accuracy: 0.7689 - loss: 0.7054 - val\_accuracy: 0.7352 - val\_loss: 0.6958 - learning\_rate: 0.0010  
Epoch 3/100  
**84/84** 11s 97ms/step - accuracy: 0.7913 - loss: 0.5327 - val\_accuracy: 0.6236 - val\_loss: 1.0406 - learning\_rate: 0.0010  
Epoch 4/100  
**84/84** 8s 97ms/step - accuracy: 0.8357 - loss: 0.4497 - val\_accuracy: 0.7703 - val\_loss: 0.8803 - learning\_rate: 0.0010  
Epoch 5/100  
**84/84** 7s 85ms/step - accuracy: 0.8478 - loss: 0.3982 - val\_accuracy: 0.7974 - val\_loss: 0.6586 - learning\_rate: 0.0010  
Epoch 6/100  
**84/84** 9s 112ms/step - accuracy: 0.8367 - loss: 0.4439 - val\_accuracy: 0.8501 - val\_loss: 0.3897 - learning\_rate: 0.0010  
Epoch 7/100  
**84/84** 7s 79ms/step - accuracy: 0.8771 - loss: 0.3185 - val\_accuracy: 0.8788 - val\_loss: 0.2904 - learning\_rate: 0.0010  
Epoch 8/100  
**84/84** 8s 96ms/step - accuracy: 0.8783 - loss: 0.3194 - val\_accuracy: 0.8341 - val\_loss: 0.4680 - learning\_rate: 0.0010  
Epoch 9/100  
**84/84** 7s 79ms/step - accuracy: 0.8830 - loss: 0.3105 - val\_accuracy: 0.8788 - val\_loss: 0.2962 - learning\_rate: 0.0010  
Epoch 10/100  
**84/84** 8s 96ms/step - accuracy: 0.8654 - loss: 0.3260 - val\_accuracy: 0.8692 - val\_loss: 0.4599 - learning\_rate: 0.0010  
Epoch 11/100  
**84/84** 8s 91ms/step - accuracy: 0.8762 - loss: 0.3073 - val\_accuracy: 0.8692 - val\_loss: 0.3927 - learning\_rate: 0.0010  
Epoch 12/100  
**84/84** 9s 80ms/step - accuracy: 0.8945 - loss: 0.2942 - val\_accuracy: 0.8756 - val\_loss: 0.3991 - learning\_rate: 0.0010  
Epoch 13/100  
**84/84** 8s 98ms/step - accuracy: 0.9093 - loss: 0.2433 - val\_accuracy: 0.8995 - val\_loss: 0.2802 - learning\_rate: 3.0000e-04  
Epoch 14/100  
**84/84** 7s 81ms/step - accuracy: 0.9249 - loss: 0.2029 - val\_accuracy: 0.9011 - val\_loss: 0.2565 - learning\_rate: 3.0000e-04  
Epoch 15/100  
**84/84** 10s 80ms/step - accuracy: 0.9221 - loss: 0.2035 - val\_accuracy: 0.9043 - val\_loss: 0.2977 - learning\_rate: 3.0000e-04  
Epoch 16/100  
**84/84** 8s 99ms/step - accuracy: 0.9175 - loss: 0.2046 - val\_accuracy: 0.9123 - val\_loss: 0.2378 - learning\_rate: 3.0000e-04  
Epoch 17/100  
**84/84** 7s 79ms/step - accuracy: 0.9517 - loss: 0.1522 - val\_accuracy: 0.9091 - val\_loss: 0.2763 - learning\_rate: 3.0000e-04  
Epoch 18/100  
**84/84** 8s 97ms/step - accuracy: 0.9307 - loss: 0.1715 - val\_accuracy: 0.9011 - val\_loss: 0.2528 - learning\_rate: 3.0000e-04  
Epoch 19/100  
**84/84** 8s 95ms/step - accuracy: 0.9490 - loss: 0.1337 - val\_accuracy: 0.8963 - val\_loss: 0.3023 - learning\_rate: 3.0000e-04  
Epoch 20/100  
**84/84** 7s 83ms/step - accuracy: 0.9389 - loss: 0.1866 - val\_accuracy: 0.9059 - val\_loss: 0.3227 - learning\_rate: 3.0000e-04

```
Epoch 21/100
84/84 11s 99ms/step - accuracy: 0.9457 - loss: 0.1436 - val_accuracy: 0.9187 - val_loss: 0.2626 - learning_rate: 3.0000e-04
Epoch 22/100
84/84 7s 85ms/step - accuracy: 0.9604 - loss: 0.1075 - val_accuracy: 0.9250 - val_loss: 0.2174 - learning_rate: 9.0000e-05
Epoch 23/100
84/84 8s 94ms/step - accuracy: 0.9692 - loss: 0.0894 - val_accuracy: 0.9314 - val_loss: 0.1972 - learning_rate: 9.0000e-05
Epoch 24/100
84/84 9s 103ms/step - accuracy: 0.9780 - loss: 0.0664 - val_accuracy: 0.9266 - val_loss: 0.2125 - learning_rate: 9.0000e-05
Epoch 25/100
84/84 7s 79ms/step - accuracy: 0.9762 - loss: 0.0690 - val_accuracy: 0.9282 - val_loss: 0.2080 - learning_rate: 9.0000e-05
Epoch 26/100
84/84 8s 96ms/step - accuracy: 0.9727 - loss: 0.0838 - val_accuracy: 0.9298 - val_loss: 0.2088 - learning_rate: 9.0000e-05
Epoch 27/100
84/84 7s 79ms/step - accuracy: 0.9799 - loss: 0.0587 - val_accuracy: 0.9282 - val_loss: 0.2380 - learning_rate: 9.0000e-05
Epoch 28/100
84/84 8s 97ms/step - accuracy: 0.9881 - loss: 0.0433 - val_accuracy: 0.9298 - val_loss: 0.2251 - learning_rate: 9.0000e-05
Epoch 29/100
84/84 7s 80ms/step - accuracy: 0.9830 - loss: 0.0521 - val_accuracy: 0.9410 - val_loss: 0.1993 - learning_rate: 2.7000e-05
Epoch 30/100
84/84 8s 98ms/step - accuracy: 0.9808 - loss: 0.0529 - val_accuracy: 0.9362 - val_loss: 0.1995 - learning_rate: 2.7000e-05
Epoch 31/100
84/84 7s 87ms/step - accuracy: 0.9810 - loss: 0.0450 - val_accuracy: 0.9266 - val_loss: 0.2125 - learning_rate: 2.7000e-05
Epoch 32/100
84/84 10s 79ms/step - accuracy: 0.9854 - loss: 0.0440 - val_accuracy: 0.9362 - val_loss: 0.2238 - learning_rate: 2.7000e-05
Epoch 33/100
84/84 8s 97ms/step - accuracy: 0.9819 - loss: 0.0450 - val_accuracy: 0.9346 - val_loss: 0.2128 - learning_rate: 2.7000e-05
```

```
In [ ]: plot_accuracy(model_fit)
```

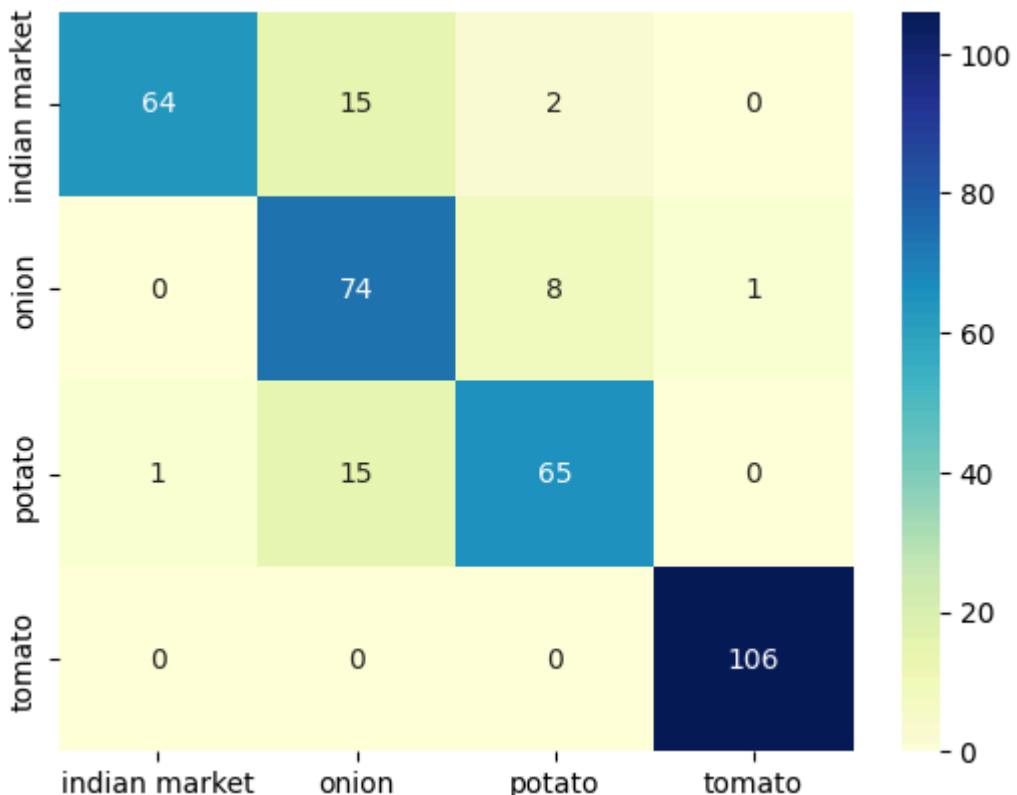


```
In [ ]: print_accuracy_stats(model, test_ds_batched, class_names, ckpt_path)
plot_confusion_matrix(model, test_ds_batched, class_names, ckpt_path)
```

12/12 ━━━━━━━━ 2s 113ms/step

Test Accuracy: 88.03%

12/12 ━━━━━━━━ 1s 67ms/step



```
In [ ]: def arch_3(height=128, width=128):
    num_classes = 4
    hidden_size = 256

    model = keras.Sequential(
        name="model_cnn_3",
        layers=[
            layers.Conv2D(filters=16, kernel_size=3, padding="same", input_shape=(height, width, 3),
                         kernel_regularizer=regularizers.l2(1e-3)),
            layers.Activation("relu"),
            layers.BatchNormalization(),
            layers.MaxPooling2D(),
            layers.Conv2D(filters=32, kernel_size=3, padding="same",
                         kernel_regularizer=regularizers.l2(1e-3)),
            layers.Activation("relu"),
            layers.BatchNormalization(),
            layers.MaxPooling2D(),
            layers.Conv2D(filters=64, kernel_size=3, padding="same",
                         kernel_regularizer=regularizers.l2(1e-3)),
            layers.Activation("relu"),
            layers.BatchNormalization(),
            layers.MaxPooling2D(),
            layers.Conv2D(filters=128, kernel_size=3, padding="same",
                         kernel_regularizer=regularizers.l2(1e-3)),
            layers.Activation("relu"),
            layers.BatchNormalization(),
            layers.MaxPooling2D(),
            layers.Conv2D(filters=256, kernel_size=3, padding="same",
                         kernel_regularizer=regularizers.l2(1e-3)),
            layers.Activation("relu"),
            layers.BatchNormalization(),
            # layers.MaxPooling2D(),
            # layers.Flatten(),
            layers.GlobalAveragePooling2D(),
            layers.Dense(units=hidden_size, kernel_regularizer=regularizers.l2(1e-3)),
            layers.Activation("relu"),
            layers.BatchNormalization(),
            layers.Dropout(0.5),
            layers.Dense(units=num_classes, activation='softmax')
        ]
    )
    return model
```

```
In [ ]: model = arch_3()
model.summary()
```

```
/usr/local/lib/python3.11/dist-packages/keras/src/layers/convolutional/base_conv.py:107: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.
    super().__init__(activity_regularizer=activity_regularizer, **kwargs)
Model: "model_cnn_3"
```

Layer (type)	Output Shape	Param #
conv2d_16 (Conv2D)	(None, 128, 128, 16)	448
activation_12 (Activation)	(None, 128, 128, 16)	0
batch_normalization_12 (BatchNormalization)	(None, 128, 128, 16)	64
max_pooling2d_13 (MaxPooling2D)	(None, 64, 64, 16)	0
conv2d_17 (Conv2D)	(None, 64, 64, 32)	4,640
activation_13 (Activation)	(None, 64, 64, 32)	0
batch_normalization_13 (BatchNormalization)	(None, 64, 64, 32)	128
max_pooling2d_14 (MaxPooling2D)	(None, 32, 32, 32)	0
conv2d_18 (Conv2D)	(None, 32, 32, 64)	18,496
activation_14 (Activation)	(None, 32, 32, 64)	0
batch_normalization_14 (BatchNormalization)	(None, 32, 32, 64)	256
max_pooling2d_15 (MaxPooling2D)	(None, 16, 16, 64)	0
conv2d_19 (Conv2D)	(None, 16, 16, 128)	73,856
activation_15 (Activation)	(None, 16, 16, 128)	0
batch_normalization_15 (BatchNormalization)	(None, 16, 16, 128)	512
max_pooling2d_16 (MaxPooling2D)	(None, 8, 8, 128)	0
conv2d_20 (Conv2D)	(None, 8, 8, 256)	295,168
activation_16 (Activation)	(None, 8, 8, 256)	0
batch_normalization_16 (BatchNormalization)	(None, 8, 8, 256)	1,024
global_average_pooling2d_3 (GlobalAveragePooling2D)	(None, 256)	0
dense_8 (Dense)	(None, 256)	65,792
activation_17 (Activation)	(None, 256)	0
batch_normalization_17 (BatchNormalization)	(None, 256)	1,024
dropout_2 (Dropout)	(None, 256)	0

dense_9 (Dense)	(None, 4)	1,028
-----------------	-----------	-------

Total params: 462,436 (1.76 MB)

Trainable params: 460,932 (1.76 MB)

Non-trainable params: 1,504 (5.88 KB)

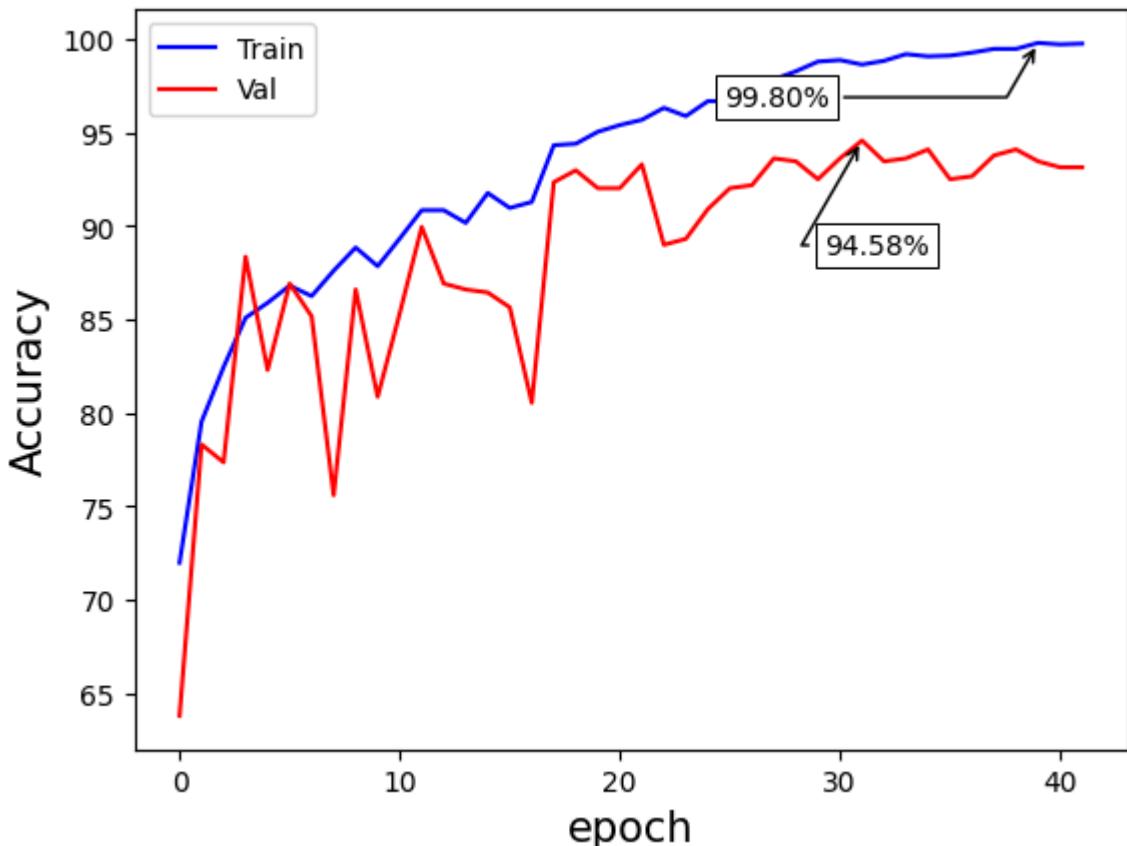
```
In [ ]: create_directory_if_not_exists('/content/ninjacart_models/model5_with_GAP_BN_DO_ckpt_path = '/content/ninjacart_models/model5_with_GAP_BN_DO_CB_LR/model5_with_G callbacks = [
    keras.callbacks.ReduceLROnPlateau(
        monitor="val_loss", factor=0.3, patience=5, min_lr=0.00001
    ),
    keras.callbacks.ModelCheckpoint(ckpt_path, save_weights_only=True, monit
    keras.callbacks.EarlyStopping(
        monitor="val_loss", patience=10, min_delta=0.001, mode='min'
    )
model_fit = compile_train_v1(model, train_ds_batched, validation_ds_batched, cal
```

```
Created directory: /content/ninjacart_models/model5_with_GAP_BN_DO_CB_LR
Epoch 1/100
84/84 _____ 23s 161ms/step - accuracy: 0.6537 - loss: 1.6525 - val_accuracy: 0.6380 - val_loss: 1.6854 - learning_rate: 0.0010
Epoch 2/100
84/84 _____ 8s 98ms/step - accuracy: 0.7872 - loss: 1.1510 - val_accuracy: 0.7831 - val_loss: 1.3219 - learning_rate: 0.0010
Epoch 3/100
84/84 _____ 7s 80ms/step - accuracy: 0.8129 - loss: 1.0824 - val_accuracy: 0.7735 - val_loss: 1.2467 - learning_rate: 0.0010
Epoch 4/100
84/84 _____ 8s 97ms/step - accuracy: 0.8313 - loss: 1.0211 - val_accuracy: 0.8836 - val_loss: 0.8421 - learning_rate: 0.0010
Epoch 5/100
84/84 _____ 7s 79ms/step - accuracy: 0.8448 - loss: 0.9219 - val_accuracy: 0.8230 - val_loss: 1.1713 - learning_rate: 0.0010
Epoch 6/100
84/84 _____ 8s 96ms/step - accuracy: 0.8664 - loss: 0.8048 - val_accuracy: 0.8692 - val_loss: 0.8172 - learning_rate: 0.0010
Epoch 7/100
84/84 _____ 7s 79ms/step - accuracy: 0.8644 - loss: 0.8291 - val_accuracy: 0.8517 - val_loss: 0.9049 - learning_rate: 0.0010
Epoch 8/100
84/84 _____ 8s 98ms/step - accuracy: 0.8750 - loss: 0.7582 - val_accuracy: 0.7560 - val_loss: 1.2067 - learning_rate: 0.0010
Epoch 9/100
84/84 _____ 8s 95ms/step - accuracy: 0.8830 - loss: 0.7068 - val_accuracy: 0.8660 - val_loss: 0.7985 - learning_rate: 0.0010
Epoch 10/100
84/84 _____ 7s 83ms/step - accuracy: 0.8693 - loss: 0.7028 - val_accuracy: 0.8086 - val_loss: 0.8129 - learning_rate: 0.0010
Epoch 11/100
84/84 _____ 8s 98ms/step - accuracy: 0.8928 - loss: 0.6310 - val_accuracy: 0.8533 - val_loss: 0.7042 - learning_rate: 0.0010
Epoch 12/100
84/84 _____ 7s 80ms/step - accuracy: 0.9026 - loss: 0.5855 - val_accuracy: 0.8995 - val_loss: 0.5608 - learning_rate: 0.0010
Epoch 13/100
84/84 _____ 8s 98ms/step - accuracy: 0.9110 - loss: 0.5537 - val_accuracy: 0.8692 - val_loss: 0.7774 - learning_rate: 0.0010
Epoch 14/100
84/84 _____ 7s 79ms/step - accuracy: 0.9010 - loss: 0.5425 - val_accuracy: 0.8660 - val_loss: 0.6432 - learning_rate: 0.0010
Epoch 15/100
84/84 _____ 8s 99ms/step - accuracy: 0.9100 - loss: 0.5163 - val_accuracy: 0.8644 - val_loss: 0.6661 - learning_rate: 0.0010
Epoch 16/100
84/84 _____ 8s 95ms/step - accuracy: 0.9210 - loss: 0.4655 - val_accuracy: 0.8565 - val_loss: 0.5637 - learning_rate: 0.0010
Epoch 17/100
84/84 _____ 7s 84ms/step - accuracy: 0.9051 - loss: 0.5011 - val_accuracy: 0.8054 - val_loss: 0.8058 - learning_rate: 0.0010
Epoch 18/100
84/84 _____ 9s 109ms/step - accuracy: 0.9369 - loss: 0.4201 - val_accuracy: 0.9234 - val_loss: 0.4437 - learning_rate: 3.0000e-04
Epoch 19/100
84/84 _____ 7s 80ms/step - accuracy: 0.9350 - loss: 0.3728 - val_accuracy: 0.9298 - val_loss: 0.4168 - learning_rate: 3.0000e-04
Epoch 20/100
84/84 _____ 8s 98ms/step - accuracy: 0.9502 - loss: 0.3490 - val_accuracy: 0.9502 - val_loss: 0.3490 - learning_rate: 3.0000e-04
```

```
accuracy: 0.9203 - val_loss: 0.4485 - learning_rate: 3.0000e-04
Epoch 21/100
84/84 _____ 7s 79ms/step - accuracy: 0.9515 - loss: 0.3416 - val_a
accuracy: 0.9203 - val_loss: 0.4509 - learning_rate: 3.0000e-04
Epoch 22/100
84/84 _____ 8s 98ms/step - accuracy: 0.9605 - loss: 0.3122 - val_a
accuracy: 0.9330 - val_loss: 0.3934 - learning_rate: 3.0000e-04
Epoch 23/100
84/84 _____ 7s 86ms/step - accuracy: 0.9635 - loss: 0.3083 - val_a
accuracy: 0.8900 - val_loss: 0.5436 - learning_rate: 3.0000e-04
Epoch 24/100
84/84 _____ 10s 79ms/step - accuracy: 0.9601 - loss: 0.2977 - val_a
accuracy: 0.8931 - val_loss: 0.4722 - learning_rate: 3.0000e-04
Epoch 25/100
84/84 _____ 8s 97ms/step - accuracy: 0.9714 - loss: 0.2747 - val_a
accuracy: 0.9091 - val_loss: 0.5256 - learning_rate: 3.0000e-04
Epoch 26/100
84/84 _____ 7s 80ms/step - accuracy: 0.9654 - loss: 0.2719 - val_a
accuracy: 0.9203 - val_loss: 0.4288 - learning_rate: 3.0000e-04
Epoch 27/100
84/84 _____ 8s 96ms/step - accuracy: 0.9797 - loss: 0.2400 - val_a
accuracy: 0.9219 - val_loss: 0.4566 - learning_rate: 3.0000e-04
Epoch 28/100
84/84 _____ 7s 85ms/step - accuracy: 0.9779 - loss: 0.2311 - val_a
accuracy: 0.9362 - val_loss: 0.4018 - learning_rate: 9.0000e-05
Epoch 29/100
84/84 _____ 8s 91ms/step - accuracy: 0.9861 - loss: 0.2193 - val_a
accuracy: 0.9346 - val_loss: 0.3686 - learning_rate: 9.0000e-05
Epoch 30/100
84/84 _____ 8s 95ms/step - accuracy: 0.9876 - loss: 0.2118 - val_a
accuracy: 0.9250 - val_loss: 0.4008 - learning_rate: 9.0000e-05
Epoch 31/100
84/84 _____ 9s 80ms/step - accuracy: 0.9908 - loss: 0.2017 - val_a
accuracy: 0.9362 - val_loss: 0.3642 - learning_rate: 9.0000e-05
Epoch 32/100
84/84 _____ 8s 99ms/step - accuracy: 0.9910 - loss: 0.1995 - val_a
accuracy: 0.9458 - val_loss: 0.3479 - learning_rate: 9.0000e-05
Epoch 33/100
84/84 _____ 7s 79ms/step - accuracy: 0.9926 - loss: 0.1952 - val_a
accuracy: 0.9346 - val_loss: 0.3743 - learning_rate: 9.0000e-05
Epoch 34/100
84/84 _____ 10s 80ms/step - accuracy: 0.9902 - loss: 0.1939 - val_a
accuracy: 0.9362 - val_loss: 0.3631 - learning_rate: 9.0000e-05
Epoch 35/100
84/84 _____ 10s 82ms/step - accuracy: 0.9895 - loss: 0.1943 - val_a
accuracy: 0.9410 - val_loss: 0.3774 - learning_rate: 9.0000e-05
Epoch 36/100
84/84 _____ 8s 99ms/step - accuracy: 0.9935 - loss: 0.1818 - val_a
accuracy: 0.9250 - val_loss: 0.4271 - learning_rate: 9.0000e-05
Epoch 37/100
84/84 _____ 7s 81ms/step - accuracy: 0.9926 - loss: 0.1842 - val_a
accuracy: 0.9266 - val_loss: 0.3950 - learning_rate: 9.0000e-05
Epoch 38/100
84/84 _____ 8s 99ms/step - accuracy: 0.9970 - loss: 0.1730 - val_a
accuracy: 0.9378 - val_loss: 0.3673 - learning_rate: 2.7000e-05
Epoch 39/100
84/84 _____ 7s 79ms/step - accuracy: 0.9942 - loss: 0.1788 - val_a
accuracy: 0.9410 - val_loss: 0.3755 - learning_rate: 2.7000e-05
Epoch 40/100
84/84 _____ 8s 97ms/step - accuracy: 0.9986 - loss: 0.1673 - val_a
```

```
accuracy: 0.9346 - val_loss: 0.3809 - learning_rate: 2.7000e-05
Epoch 41/100
84/84 8s 96ms/step - accuracy: 0.9970 - loss: 0.1662 - val_a
ccuracy: 0.9314 - val_loss: 0.3748 - learning_rate: 2.7000e-05
Epoch 42/100
84/84 9s 80ms/step - accuracy: 0.9986 - loss: 0.1646 - val_a
ccuracy: 0.9314 - val_loss: 0.3848 - learning_rate: 2.7000e-05
```

```
In [ ]: plot_accuracy(model_fit)
```

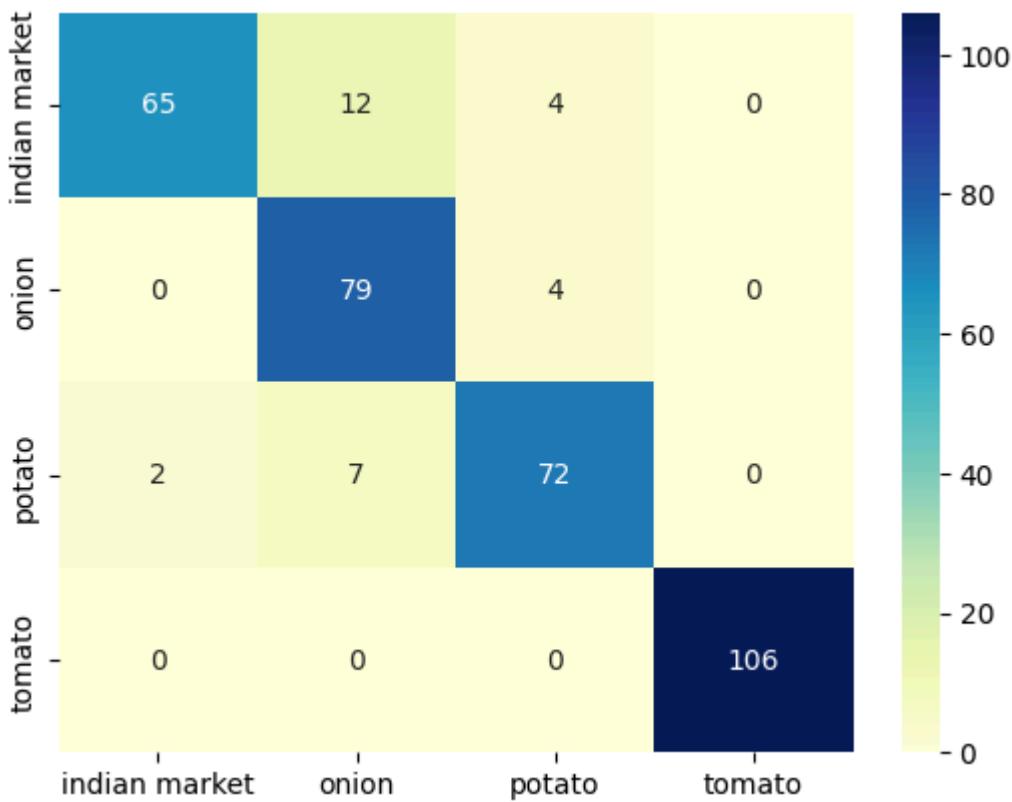


```
In [ ]: print_accuracy_stats(model, test_ds_batched, class_names, ckpt_path)
plot_confusion_matrix(model, test_ds_batched, class_names, ckpt_path)
```

```
12/12 2s 165ms/step
```

```
Test Accuracy: 91.74%
```

```
12/12 1s 84ms/step
```



```
In [ ]: def arch_3(height=128, width=128):
    num_classes = 4
    hidden_size = 256

    model = keras.Sequential(
        name="model_cnn_3",
        layers=[
            layers.Conv2D(filters=16, kernel_size=3, padding="same", input_shape=(height, width, 3),
                         kernel_regularizer=regularizers.l2(1e-3)),
            layers.Activation("relu"),
            layers.BatchNormalization(),
            layers.MaxPooling2D(),
            layers.Conv2D(filters=32, kernel_size=3, padding="same",
                         kernel_regularizer=regularizers.l2(1e-3)),
            layers.Activation("relu"),
            layers.BatchNormalization(),
            layers.MaxPooling2D(),
            layers.Conv2D(filters=64, kernel_size=3, padding="same",
                         kernel_regularizer=regularizers.l2(1e-3)),
            layers.Activation("relu"),
            layers.BatchNormalization(),
            layers.MaxPooling2D(),
            layers.Conv2D(filters=128, kernel_size=3, padding="same",
                         kernel_regularizer=regularizers.l2(1e-3)),
            layers.Activation("relu"),
            layers.BatchNormalization(),
            layers.MaxPooling2D(),
            layers.Conv2D(filters=256, kernel_size=3, padding="same",
                         kernel_regularizer=regularizers.l2(1e-3)),
            layers.Activation("relu"),
            layers.BatchNormalization(),
            # layers.MaxPooling2D(),
            # layers.Flatten(),
            layers.GlobalAveragePooling2D(),
            layers.Dense(hidden_size),
            layers.Activation("relu"),
            layers.Dense(num_classes)
        ]
    )
```

```
        layers.Dense(units=hidden_size, kernel_regularizer=regularizers.l2(1),
          layers.Activation("relu"),
          layers.BatchNormalization(),
          layers.Dropout(0.5),
          layers.Dense(units=num_classes, activation='softmax')
      ]
)
return model
```

```
In [ ]: model = arch_3()
model.summary()
```

```
/usr/local/lib/python3.11/dist-packages/keras/src/layers/convolutional/base_conv.
py:107: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a laye
r. When using Sequential models, prefer using an `Input(shape)` object as the fir
st layer in the model instead.
    super().__init__(activity_regularizer=activity_regularizer, **kwargs)
Model: "model_cnn_3"
```

Layer (type)	Output Shape	Param #
conv2d_21 (Conv2D)	(None, 128, 128, 16)	448
activation_18 (Activation)	(None, 128, 128, 16)	0
batch_normalization_18 (BatchNormalization)	(None, 128, 128, 16)	64
max_pooling2d_17 (MaxPooling2D)	(None, 64, 64, 16)	0
conv2d_22 (Conv2D)	(None, 64, 64, 32)	4,640
activation_19 (Activation)	(None, 64, 64, 32)	0
batch_normalization_19 (BatchNormalization)	(None, 64, 64, 32)	128
max_pooling2d_18 (MaxPooling2D)	(None, 32, 32, 32)	0
conv2d_23 (Conv2D)	(None, 32, 32, 64)	18,496
activation_20 (Activation)	(None, 32, 32, 64)	0
batch_normalization_20 (BatchNormalization)	(None, 32, 32, 64)	256
max_pooling2d_19 (MaxPooling2D)	(None, 16, 16, 64)	0
conv2d_24 (Conv2D)	(None, 16, 16, 128)	73,856
activation_21 (Activation)	(None, 16, 16, 128)	0
batch_normalization_21 (BatchNormalization)	(None, 16, 16, 128)	512
max_pooling2d_20 (MaxPooling2D)	(None, 8, 8, 128)	0
conv2d_25 (Conv2D)	(None, 8, 8, 256)	295,168
activation_22 (Activation)	(None, 8, 8, 256)	0
batch_normalization_22 (BatchNormalization)	(None, 8, 8, 256)	1,024
global_average_pooling2d_4 (GlobalAveragePooling2D)	(None, 256)	0
dense_10 (Dense)	(None, 256)	65,792
activation_23 (Activation)	(None, 256)	0
batch_normalization_23 (BatchNormalization)	(None, 256)	1,024
dropout_3 (Dropout)	(None, 256)	0

dense_11 (Dense)	(None, 4)	1,028
------------------	-----------	-------

Total params: 462,436 (1.76 MB)

Trainable params: 460,932 (1.76 MB)

Non-trainable params: 1,504 (5.88 KB)

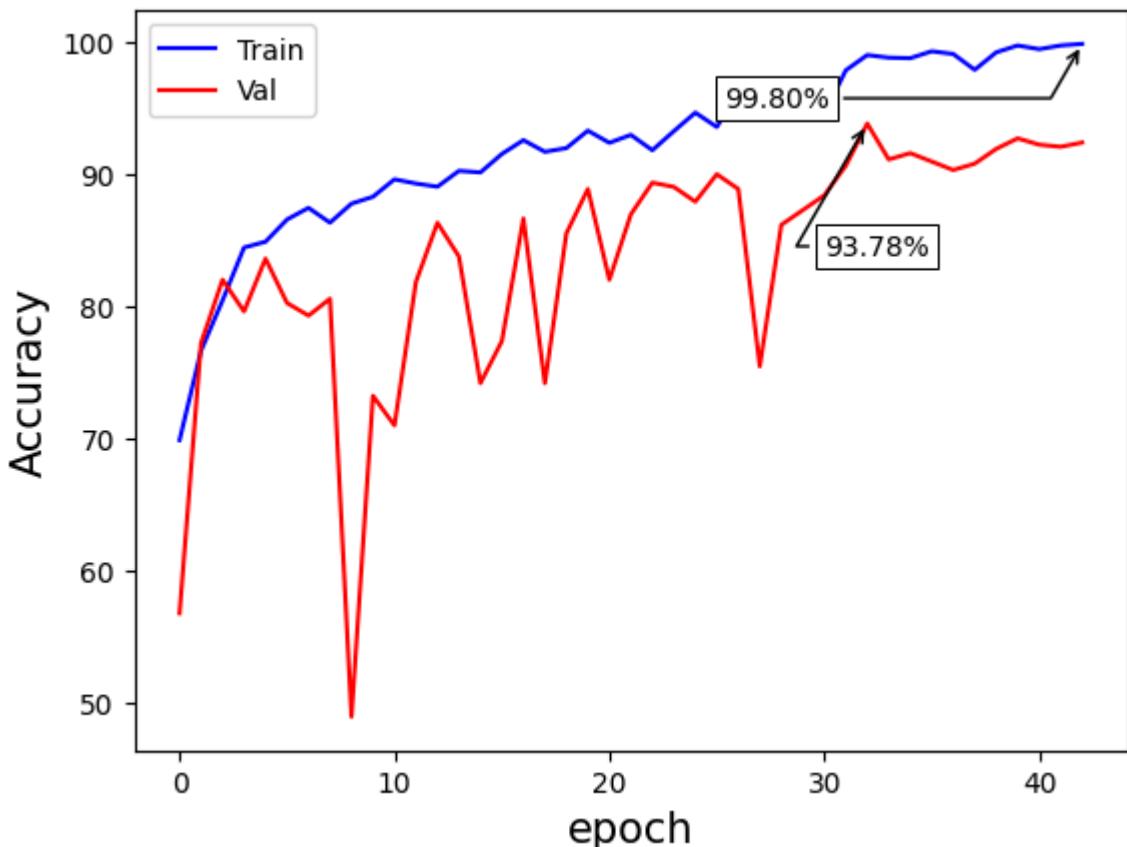
```
In [ ]: create_directory_if_not_exists('/content/ninjacart_models/model6_with_GAP_BN_DO_ckpt_path = '/content/ninjacart_models/model6_with_GAP_BN_DO_CB_LR_L2/model6_wit
callbacks = [
    keras.callbacks.ReduceLROnPlateau(
        monitor="val_loss", factor=0.3, patience=5, min_lr=0.00001
    ),
    keras.callbacks.ModelCheckpoint(ckpt_path, save_weights_only=True, monit
    keras.callbacks.EarlyStopping(
        monitor="val_loss", patience=10, min_delta=0.001, mode='min'
    )
model_fit = compile_train_v1(model, train_ds_batched, validation_ds_batched, cal
```

```
Created directory: /content/ninjacart_models/model6_with_GAP_BN_DO_CB_LR_L2
Epoch 1/100
84/84 _____ 20s 151ms/step - accuracy: 0.6297 - loss: 1.6929 - val_
accuracy: 0.5678 - val_loss: 1.7282 - learning_rate: 0.0010
Epoch 2/100
84/84 _____ 8s 100ms/step - accuracy: 0.7579 - loss: 1.2804 - val_
accuracy: 0.7719 - val_loss: 1.1630 - learning_rate: 0.0010
Epoch 3/100
84/84 _____ 7s 81ms/step - accuracy: 0.8041 - loss: 1.1072 - val_a
ccuracy: 0.8198 - val_loss: 1.0329 - learning_rate: 0.0010
Epoch 4/100
84/84 _____ 8s 99ms/step - accuracy: 0.8433 - loss: 0.9580 - val_a
ccuracy: 0.7959 - val_loss: 1.1170 - learning_rate: 0.0010
Epoch 5/100
84/84 _____ 7s 85ms/step - accuracy: 0.8463 - loss: 0.9204 - val_a
ccuracy: 0.8357 - val_loss: 1.0488 - learning_rate: 0.0010
Epoch 6/100
84/84 _____ 10s 81ms/step - accuracy: 0.8711 - loss: 0.8317 - val_
accuracy: 0.8022 - val_loss: 1.0198 - learning_rate: 0.0010
Epoch 7/100
84/84 _____ 8s 99ms/step - accuracy: 0.8689 - loss: 0.8218 - val_a
ccuracy: 0.7927 - val_loss: 1.1416 - learning_rate: 0.0010
Epoch 8/100
84/84 _____ 7s 81ms/step - accuracy: 0.8559 - loss: 0.8323 - val_a
ccuracy: 0.8054 - val_loss: 0.9580 - learning_rate: 0.0010
Epoch 9/100
84/84 _____ 8s 101ms/step - accuracy: 0.8792 - loss: 0.7382 - val_
accuracy: 0.4896 - val_loss: 2.1495 - learning_rate: 0.0010
Epoch 10/100
84/84 _____ 8s 93ms/step - accuracy: 0.8806 - loss: 0.6793 - val_a
ccuracy: 0.7321 - val_loss: 1.2706 - learning_rate: 0.0010
Epoch 11/100
84/84 _____ 11s 97ms/step - accuracy: 0.8966 - loss: 0.6414 - val_
accuracy: 0.7097 - val_loss: 1.4492 - learning_rate: 0.0010
Epoch 12/100
84/84 _____ 8s 99ms/step - accuracy: 0.8671 - loss: 0.6652 - val_a
ccuracy: 0.8182 - val_loss: 0.8225 - learning_rate: 0.0010
Epoch 13/100
84/84 _____ 8s 98ms/step - accuracy: 0.8891 - loss: 0.6094 - val_a
ccuracy: 0.8628 - val_loss: 0.7391 - learning_rate: 0.0010
Epoch 14/100
84/84 _____ 9s 81ms/step - accuracy: 0.9029 - loss: 0.5711 - val_a
ccuracy: 0.8373 - val_loss: 0.7184 - learning_rate: 0.0010
Epoch 15/100
84/84 _____ 8s 99ms/step - accuracy: 0.9057 - loss: 0.5347 - val_a
ccuracy: 0.7416 - val_loss: 0.8825 - learning_rate: 0.0010
Epoch 16/100
84/84 _____ 7s 81ms/step - accuracy: 0.9065 - loss: 0.5151 - val_a
ccuracy: 0.7735 - val_loss: 0.8146 - learning_rate: 0.0010
Epoch 17/100
84/84 _____ 8s 99ms/step - accuracy: 0.9224 - loss: 0.4532 - val_a
ccuracy: 0.8660 - val_loss: 0.6297 - learning_rate: 0.0010
Epoch 18/100
84/84 _____ 8s 96ms/step - accuracy: 0.9084 - loss: 0.4649 - val_a
ccuracy: 0.7416 - val_loss: 1.1390 - learning_rate: 0.0010
Epoch 19/100
84/84 _____ 7s 82ms/step - accuracy: 0.9109 - loss: 0.4821 - val_a
ccuracy: 0.8549 - val_loss: 0.7070 - learning_rate: 0.0010
Epoch 20/100
84/84 _____ 9s 109ms/step - accuracy: 0.9409 - loss: 0.4057 - val_
```

```
accuracy: 0.8884 - val_loss: 0.5481 - learning_rate: 0.0010
Epoch 21/100
84/84 _____ 7s 79ms/step - accuracy: 0.9170 - loss: 0.4361 - val_a
ccuracy: 0.8198 - val_loss: 0.7245 - learning_rate: 0.0010
Epoch 22/100
84/84 _____ 8s 98ms/step - accuracy: 0.9331 - loss: 0.4084 - val_a
ccuracy: 0.8692 - val_loss: 0.5348 - learning_rate: 0.0010
Epoch 23/100
84/84 _____ 7s 79ms/step - accuracy: 0.9058 - loss: 0.4778 - val_a
ccuracy: 0.8931 - val_loss: 0.4944 - learning_rate: 0.0010
Epoch 24/100
84/84 _____ 8s 96ms/step - accuracy: 0.9408 - loss: 0.3577 - val_a
ccuracy: 0.8900 - val_loss: 0.4988 - learning_rate: 0.0010
Epoch 25/100
84/84 _____ 7s 84ms/step - accuracy: 0.9487 - loss: 0.3365 - val_a
ccuracy: 0.8788 - val_loss: 0.5545 - learning_rate: 0.0010
Epoch 26/100
84/84 _____ 8s 92ms/step - accuracy: 0.9294 - loss: 0.3667 - val_a
ccuracy: 0.8995 - val_loss: 0.4603 - learning_rate: 0.0010
Epoch 27/100
84/84 _____ 8s 95ms/step - accuracy: 0.9684 - loss: 0.2783 - val_a
ccuracy: 0.8884 - val_loss: 0.5072 - learning_rate: 0.0010
Epoch 28/100
84/84 _____ 7s 80ms/step - accuracy: 0.9457 - loss: 0.3128 - val_a
ccuracy: 0.7544 - val_loss: 0.8014 - learning_rate: 0.0010
Epoch 29/100
84/84 _____ 9s 110ms/step - accuracy: 0.9562 - loss: 0.2888 - val_a
ccuracy: 0.8612 - val_loss: 0.6769 - learning_rate: 0.0010
Epoch 30/100
84/84 _____ 7s 79ms/step - accuracy: 0.9525 - loss: 0.2969 - val_a
ccuracy: 0.8724 - val_loss: 0.5400 - learning_rate: 0.0010
Epoch 31/100
84/84 _____ 8s 97ms/step - accuracy: 0.9509 - loss: 0.3005 - val_a
ccuracy: 0.8836 - val_loss: 0.5809 - learning_rate: 0.0010
Epoch 32/100
84/84 _____ 7s 80ms/step - accuracy: 0.9712 - loss: 0.2377 - val_a
ccuracy: 0.9059 - val_loss: 0.4371 - learning_rate: 3.0000e-04
Epoch 33/100
84/84 _____ 10s 81ms/step - accuracy: 0.9929 - loss: 0.1937 - val_a
ccuracy: 0.9378 - val_loss: 0.3583 - learning_rate: 3.0000e-04
Epoch 34/100
84/84 _____ 8s 96ms/step - accuracy: 0.9916 - loss: 0.1830 - val_a
ccuracy: 0.9107 - val_loss: 0.4522 - learning_rate: 3.0000e-04
Epoch 35/100
84/84 _____ 10s 96ms/step - accuracy: 0.9880 - loss: 0.1873 - val_a
ccuracy: 0.9155 - val_loss: 0.4408 - learning_rate: 3.0000e-04
Epoch 36/100
84/84 _____ 7s 81ms/step - accuracy: 0.9926 - loss: 0.1735 - val_a
ccuracy: 0.9091 - val_loss: 0.4858 - learning_rate: 3.0000e-04
Epoch 37/100
84/84 _____ 8s 97ms/step - accuracy: 0.9928 - loss: 0.1713 - val_a
ccuracy: 0.9027 - val_loss: 0.5327 - learning_rate: 3.0000e-04
Epoch 38/100
84/84 _____ 7s 81ms/step - accuracy: 0.9697 - loss: 0.2258 - val_a
ccuracy: 0.9075 - val_loss: 0.4925 - learning_rate: 3.0000e-04
Epoch 39/100
84/84 _____ 8s 96ms/step - accuracy: 0.9800 - loss: 0.1942 - val_a
ccuracy: 0.9187 - val_loss: 0.4286 - learning_rate: 9.0000e-05
Epoch 40/100
84/84 _____ 7s 80ms/step - accuracy: 0.9966 - loss: 0.1477 - val_a
```

```
accuracy: 0.9266 - val_loss: 0.4120 - learning_rate: 9.0000e-05
Epoch 41/100
84/84 ----- 8s 98ms/step - accuracy: 0.9944 - loss: 0.1541 - val_a
ccuracy: 0.9219 - val_loss: 0.4128 - learning_rate: 9.0000e-05
Epoch 42/100
84/84 ----- 7s 80ms/step - accuracy: 0.9958 - loss: 0.1462 - val_a
ccuracy: 0.9203 - val_loss: 0.4367 - learning_rate: 9.0000e-05
Epoch 43/100
84/84 ----- 10s 79ms/step - accuracy: 0.9975 - loss: 0.1564 - val_a
accuracy: 0.9234 - val_loss: 0.4461 - learning_rate: 9.0000e-05
```

```
In [ ]: plot_accuracy(model_fit)
```

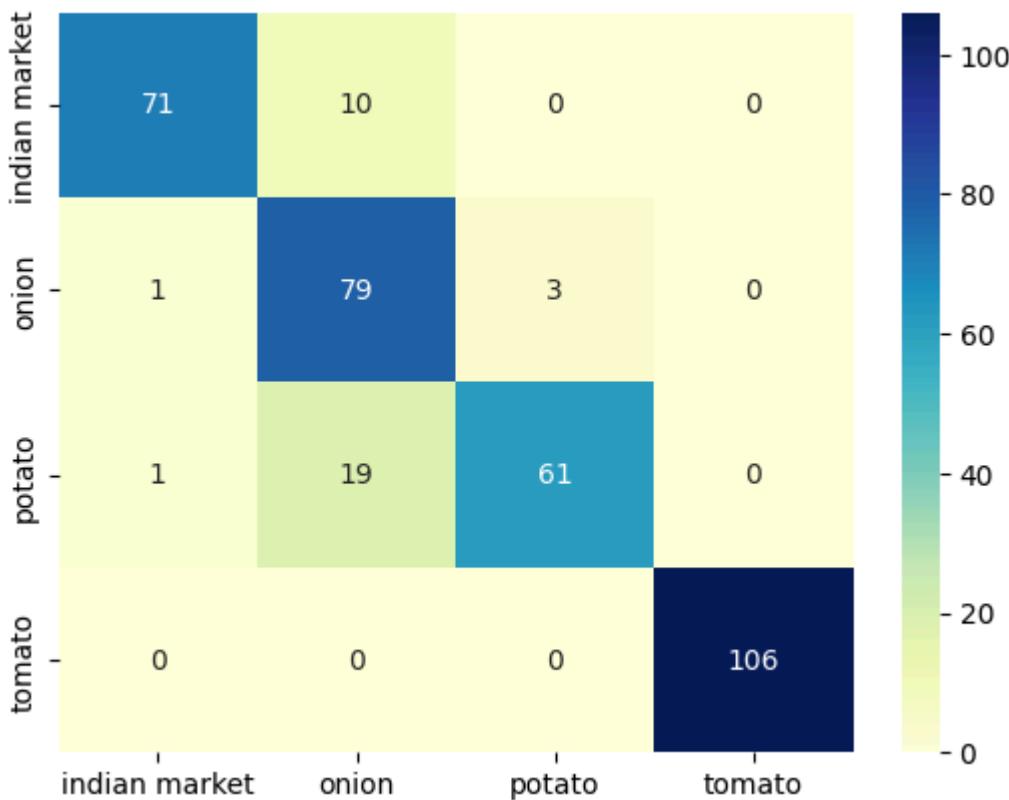


```
In [ ]: print_accuracy_stats(model, test_ds_batched, class_names, ckpt_path)
plot_confusion_matrix(model, test_ds_batched, class_names, ckpt_path)
```

```
12/12 ----- 3s 208ms/step
```

Test Accuracy: 90.31%

```
12/12 ----- 1s 63ms/step
```



```
In [ ]: def preprocess_v2(train_data, val_data, test_data, target_height=128, target_width=156):
    # Data Processing Stage with resizing and rescaling operations #same as before
    data_preprocess = keras.Sequential(
        name="data_preprocess",
        layers=[
            layers.Resizing(target_height, target_width),
            layers.Rescaling(1.0/255),
        ]
    )

    # Data Processing Stage with resizing and rescaling operations
    data_augmentation = keras.Sequential(
        name="data_augmentation",
        layers=[
            layers.Resizing(156, 156), # First resize to 156,156
            layers.RandomCrop(target_height, target_width), # Then randomly crop
            # layers.RandomBrightness(0.2), # Modify brightness by 0.2 factor
            layers.Rescaling(1.0/255), # Finally rescale
        ]
    )

    # Perform Data Processing on the train, val, test dataset
    train_ds = train_data.map(
        lambda x, y: (data_augmentation(x), y), num_parallel_calls=tf.data.AUTOTUNE
    ).prefetch(tf.data.AUTOTUNE)
    val_ds = val_data.map(
        lambda x, y: (data_preprocess(x), y), num_parallel_calls=tf.data.AUTOTUNE
    ).prefetch(tf.data.AUTOTUNE)
    test_ds = test_data.map(
        lambda x, y: (data_preprocess(x), y), num_parallel_calls=tf.data.AUTOTUNE
    ).prefetch(tf.data.AUTOTUNE)

    return train_ds, val_ds, test_ds
```

```
In [ ]: def keras_load_data_v2(base_path, target_paths=['train'], _label_mode='int', _ba
      keras_data_dict = {}
      for target_path in target_paths:
          _shuffle = True if target_path == 'train' else False

          print(f'Loading {target_path} data...')
          path = os.path.join(base_path, target_path)
          data = tf.keras.utils.image_dataset_from_directory(path, shuffle=_shuffle, i
          keras_data_dict.update({target_path: data})

      return keras_data_dict

# keras_loaded_data = keras_Load_data(base_data_path, target_paths=['train', 'va
```

```
In [ ]: base_data_path = 'nijacart_data'
keras_loaded_data = keras_load_data(base_data_path, target_paths=['train', 'vali

Loading train data...
Found 2502 files belonging to 4 classes.
Loading validation data...
Found 627 files belonging to 4 classes.
Loading test data...
Found 351 files belonging to 4 classes.
```

```
In [ ]: train_ds, val_ds, test_ds = preprocess_v2(keras_loaded_data['train'], keras_load
```

```
In [ ]: model = arch_3()
model.summary()
```

```
/usr/local/lib/python3.11/dist-packages/keras/src/layers/convolutional/base_conv.
py:107: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a laye
r. When using Sequential models, prefer using an `Input(shape)` object as the fir
st layer in the model instead.
    super().__init__(activity_regularizer=activity_regularizer, **kwargs)
Model: "model_cnn_3"
```

Layer (type)	Output Shape	Param #
conv2d_26 (Conv2D)	(None, 128, 128, 16)	448
activation_24 (Activation)	(None, 128, 128, 16)	0
batch_normalization_24 (BatchNormalization)	(None, 128, 128, 16)	64
max_pooling2d_21 (MaxPooling2D)	(None, 64, 64, 16)	0
conv2d_27 (Conv2D)	(None, 64, 64, 32)	4,640
activation_25 (Activation)	(None, 64, 64, 32)	0
batch_normalization_25 (BatchNormalization)	(None, 64, 64, 32)	128
max_pooling2d_22 (MaxPooling2D)	(None, 32, 32, 32)	0
conv2d_28 (Conv2D)	(None, 32, 32, 64)	18,496
activation_26 (Activation)	(None, 32, 32, 64)	0
batch_normalization_26 (BatchNormalization)	(None, 32, 32, 64)	256
max_pooling2d_23 (MaxPooling2D)	(None, 16, 16, 64)	0
conv2d_29 (Conv2D)	(None, 16, 16, 128)	73,856
activation_27 (Activation)	(None, 16, 16, 128)	0
batch_normalization_27 (BatchNormalization)	(None, 16, 16, 128)	512
max_pooling2d_24 (MaxPooling2D)	(None, 8, 8, 128)	0
conv2d_30 (Conv2D)	(None, 8, 8, 256)	295,168
activation_28 (Activation)	(None, 8, 8, 256)	0
batch_normalization_28 (BatchNormalization)	(None, 8, 8, 256)	1,024
global_average_pooling2d_5 (GlobalAveragePooling2D)	(None, 256)	0
dense_12 (Dense)	(None, 256)	65,792
activation_29 (Activation)	(None, 256)	0
batch_normalization_29 (BatchNormalization)	(None, 256)	1,024
dropout_4 (Dropout)	(None, 256)	0

dense_13 (Dense)	(None, 4)	1,028
------------------	-----------	-------

Total params: 462,436 (1.76 MB)

Trainable params: 460,932 (1.76 MB)

Non-trainable params: 1,504 (5.88 KB)

```
In [ ]: create_directory_if_not_exists('/content/ninjacart_models/model7_agumentation')
ckpt_path = '/content/ninjacart_models/model7_agumentation/model7_agumentation.w
callbacks = [
    keras.callbacks.ReduceLROnPlateau(
        monitor="val_loss", factor=0.3, patience=5, min_lr=0.00001
    ),
    keras.callbacks.ModelCheckpoint(ckpt_path, save_weights_only=True, monit
    keras.callbacks.EarlyStopping(
        monitor="val_loss", patience=10, min_delta=0.001, mode='min'
    )
model_fit = compile_train_v1(model, train_ds, val_ds, callbacks, epochs=100)
```

```
Created directory: /content/ninjacart_models/model7_agumentation
Epoch 1/100
84/84 _____ 20s 129ms/step - accuracy: 0.6591 - loss: 1.6233 - val_
accuracy: 0.2871 - val_loss: 2.6854 - learning_rate: 0.0010
Epoch 2/100
84/84 _____ 12s 81ms/step - accuracy: 0.7841 - loss: 1.2075 - val_
accuracy: 0.4019 - val_loss: 2.3617 - learning_rate: 0.0010
Epoch 3/100
84/84 _____ 8s 99ms/step - accuracy: 0.8022 - loss: 1.0842 - val_a
ccuracy: 0.4530 - val_loss: 2.5060 - learning_rate: 0.0010
Epoch 4/100
84/84 _____ 8s 94ms/step - accuracy: 0.8167 - loss: 0.9574 - val_a
ccuracy: 0.4721 - val_loss: 2.2748 - learning_rate: 0.0010
Epoch 5/100
84/84 _____ 7s 82ms/step - accuracy: 0.8382 - loss: 0.9240 - val_a
ccuracy: 0.6890 - val_loss: 1.3911 - learning_rate: 0.0010
Epoch 6/100
84/84 _____ 12s 100ms/step - accuracy: 0.8240 - loss: 0.9358 - val_
accuracy: 0.7687 - val_loss: 0.9811 - learning_rate: 0.0010
Epoch 7/100
84/84 _____ 8s 96ms/step - accuracy: 0.8616 - loss: 0.8356 - val_a
ccuracy: 0.8278 - val_loss: 0.8526 - learning_rate: 0.0010
Epoch 8/100
84/84 _____ 9s 80ms/step - accuracy: 0.8783 - loss: 0.7607 - val_a
ccuracy: 0.8373 - val_loss: 0.9362 - learning_rate: 0.0010
Epoch 9/100
84/84 _____ 10s 80ms/step - accuracy: 0.8569 - loss: 0.7717 - val_
accuracy: 0.7943 - val_loss: 0.8395 - learning_rate: 0.0010
Epoch 10/100
84/84 _____ 12s 97ms/step - accuracy: 0.8591 - loss: 0.7577 - val_
accuracy: 0.7687 - val_loss: 0.8842 - learning_rate: 0.0010
Epoch 11/100
84/84 _____ 10s 98ms/step - accuracy: 0.8836 - loss: 0.6901 - val_
accuracy: 0.7257 - val_loss: 1.2509 - learning_rate: 0.0010
Epoch 12/100
84/84 _____ 9s 78ms/step - accuracy: 0.8667 - loss: 0.6769 - val_a
ccuracy: 0.4242 - val_loss: 2.2943 - learning_rate: 0.0010
Epoch 13/100
84/84 _____ 8s 101ms/step - accuracy: 0.8785 - loss: 0.6280 - val_
accuracy: 0.8389 - val_loss: 0.7043 - learning_rate: 0.0010
Epoch 14/100
84/84 _____ 8s 95ms/step - accuracy: 0.8703 - loss: 0.6118 - val_a
ccuracy: 0.8118 - val_loss: 0.7136 - learning_rate: 0.0010
Epoch 15/100
84/84 _____ 7s 85ms/step - accuracy: 0.8799 - loss: 0.5811 - val_a
ccuracy: 0.7145 - val_loss: 1.1678 - learning_rate: 0.0010
Epoch 16/100
84/84 _____ 11s 100ms/step - accuracy: 0.8851 - loss: 0.5513 - val_
accuracy: 0.7624 - val_loss: 0.8755 - learning_rate: 0.0010
Epoch 17/100
84/84 _____ 11s 104ms/step - accuracy: 0.8838 - loss: 0.5497 - val_
accuracy: 0.8070 - val_loss: 0.7161 - learning_rate: 0.0010
Epoch 18/100
84/84 _____ 7s 80ms/step - accuracy: 0.9018 - loss: 0.4901 - val_a
ccuracy: 0.7097 - val_loss: 0.9437 - learning_rate: 0.0010
Epoch 19/100
84/84 _____ 8s 100ms/step - accuracy: 0.9053 - loss: 0.4646 - val_
accuracy: 0.7719 - val_loss: 0.7965 - learning_rate: 3.0000e-04
Epoch 20/100
84/84 _____ 10s 97ms/step - accuracy: 0.9241 - loss: 0.4063 - val_
```

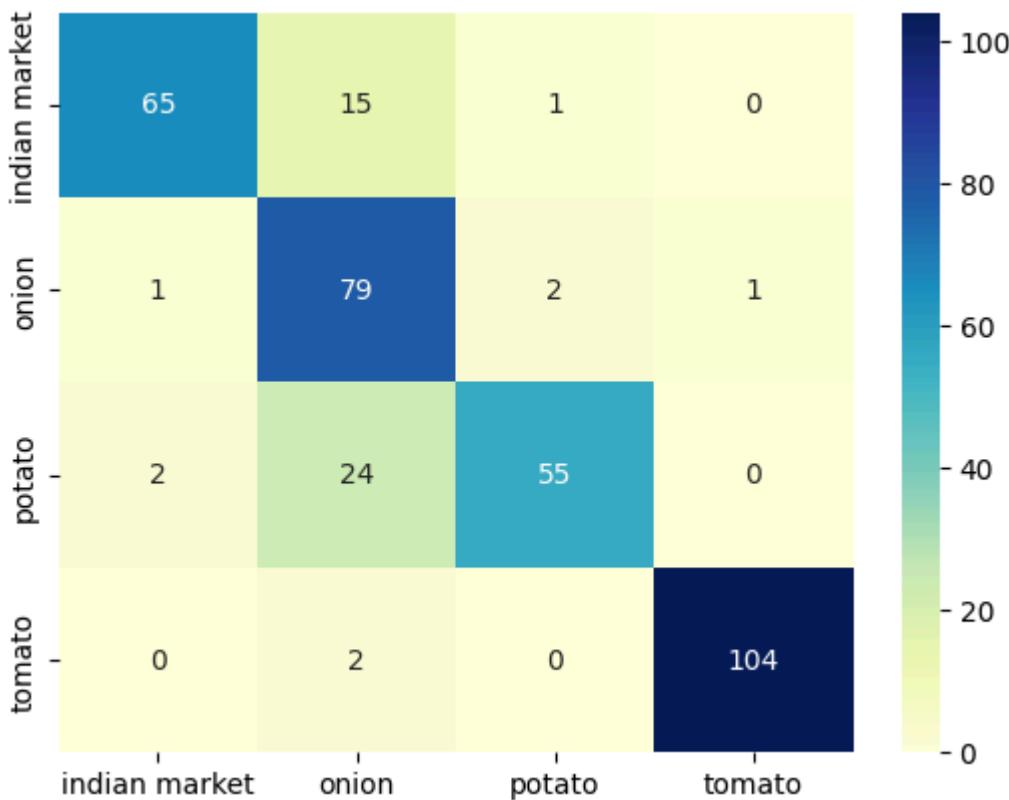
```
accuracy: 0.8660 - val_loss: 0.5157 - learning_rate: 3.0000e-04
Epoch 21/100
84/84 ----- 7s 83ms/step - accuracy: 0.9291 - loss: 0.3889 - val_a
ccuracy: 0.8102 - val_loss: 0.6698 - learning_rate: 3.0000e-04
Epoch 22/100
84/84 ----- 8s 101ms/step - accuracy: 0.9221 - loss: 0.3971 - val_
accuracy: 0.9027 - val_loss: 0.4391 - learning_rate: 3.0000e-04
Epoch 23/100
84/84 ----- 9s 80ms/step - accuracy: 0.9313 - loss: 0.3724 - val_a
ccuracy: 0.8278 - val_loss: 0.6514 - learning_rate: 3.0000e-04
Epoch 24/100
84/84 ----- 10s 80ms/step - accuracy: 0.9205 - loss: 0.3787 - val_
accuracy: 0.7528 - val_loss: 0.8363 - learning_rate: 3.0000e-04
Epoch 25/100
84/84 ----- 11s 91ms/step - accuracy: 0.9321 - loss: 0.3690 - val_
accuracy: 0.8533 - val_loss: 0.5563 - learning_rate: 3.0000e-04
Epoch 26/100
84/84 ----- 9s 103ms/step - accuracy: 0.9354 - loss: 0.3479 - val_
accuracy: 0.8708 - val_loss: 0.4998 - learning_rate: 3.0000e-04
Epoch 27/100
84/84 ----- 7s 79ms/step - accuracy: 0.9386 - loss: 0.3260 - val_a
ccuracy: 0.8724 - val_loss: 0.5075 - learning_rate: 3.0000e-04
Epoch 28/100
84/84 ----- 11s 94ms/step - accuracy: 0.9442 - loss: 0.3286 - val_
accuracy: 0.8900 - val_loss: 0.4901 - learning_rate: 9.0000e-05
Epoch 29/100
84/84 ----- 11s 100ms/step - accuracy: 0.9443 - loss: 0.2921 - val_
_accuracy: 0.8884 - val_loss: 0.4681 - learning_rate: 9.0000e-05
Epoch 30/100
84/84 ----- 7s 80ms/step - accuracy: 0.9624 - loss: 0.2640 - val_a
ccuracy: 0.8309 - val_loss: 0.5919 - learning_rate: 9.0000e-05
Epoch 31/100
84/84 ----- 8s 100ms/step - accuracy: 0.9628 - loss: 0.2685 - val_
accuracy: 0.8293 - val_loss: 0.6415 - learning_rate: 9.0000e-05
Epoch 32/100
84/84 ----- 10s 94ms/step - accuracy: 0.9563 - loss: 0.2707 - val_
accuracy: 0.8740 - val_loss: 0.5283 - learning_rate: 9.0000e-05
```

```
In [ ]: print_accuracy_stats(model, test_ds, class_names, ckpt_path)
plot_confusion_matrix(model, test_ds, class_names, ckpt_path)
```

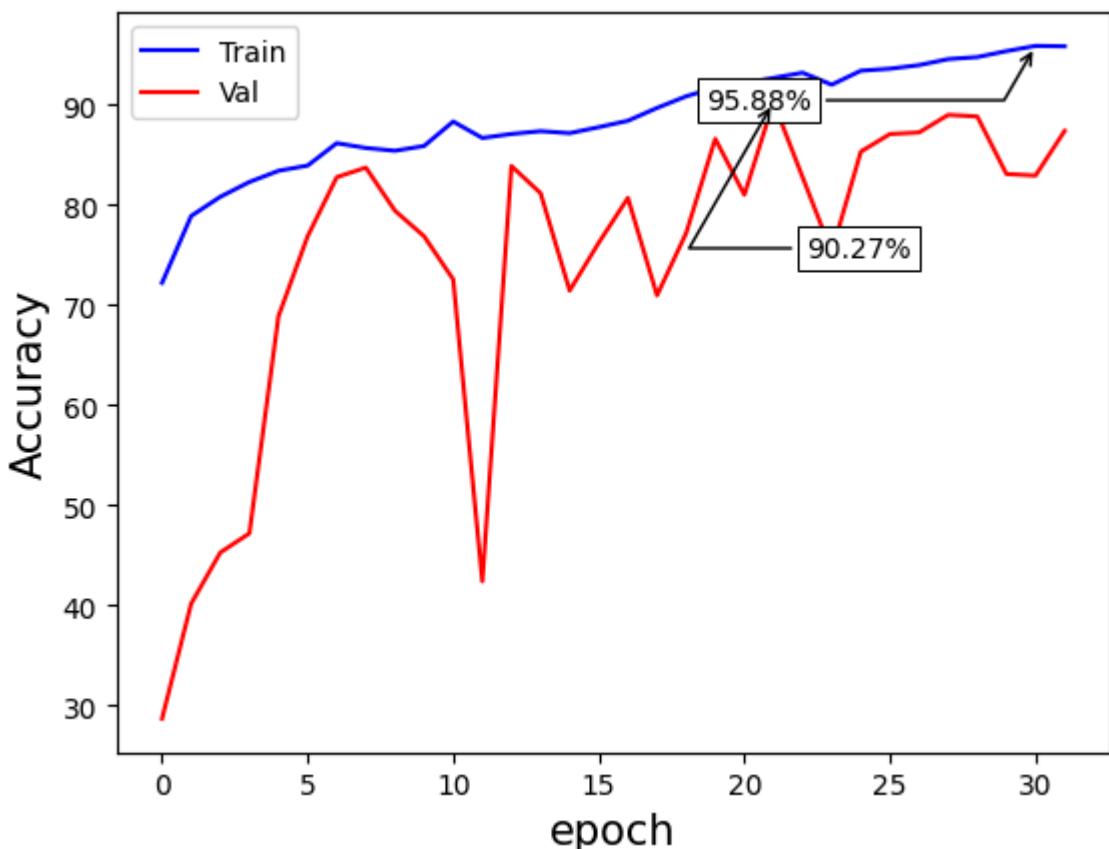
```
12/12 ----- 2s 94ms/step
```

Test Accuracy: 86.32%

```
12/12 ----- 1s 66ms/step
```



```
In [ ]: plot_accuracy(model_fit)
```



## VGG16

```
In [ ]: def plot_predictions_with_bars(images, true_labels, pred_probs, class_names):
    pred_classes = np.argmax(pred_probs, axis=1)
```

```

true_classes = np.argmax(true_labels, axis=1)

# Filter for correctly predicted images
correct_indices = [i for i, (true, pred) in enumerate(zip(true_classes, pred))

# If there are fewer than 5 correct predictions, plot all of them, otherwise
num_images_to_plot = min(5, len(correct_indices))
if num_images_to_plot == 0:
    print("No correctly predicted images in this batch to plot.")
    return

# Randomly sample indices from the correctly predicted ones
indices_to_plot = random.sample(correct_indices, num_images_to_plot)

fig, axs = plt.subplots(num_images_to_plot, 2, figsize=(8, num_images_to_plot))
if num_images_to_plot == 1:
    axs = np.expand_dims(axs, axis=0)

for i, original_index in enumerate(indices_to_plot):
    img = images[original_index]
    true_idx = true_classes[original_index]
    pred_idx = pred_classes[original_index]
    confidence = pred_probs[original_index][pred_idx] * 100
    pred_label = class_names[pred_idx]
    true_label = class_names[true_idx]
    color = 'blue' # Since we are only plotting correct predictions

    axs[i, 0].imshow(img)
    axs[i, 0].axis('off')
    axs[i, 0].set_title(f'{pred_label} {confidence:.0f}% ({true_label})', color=color)

    axs[i, 1].barh(range(len(class_names)), pred_probs[original_index],
                   color=['blue' if j == pred_idx else 'gray' for j in range(len(class_names))])
    axs[i, 1].set_xlim([0, 1])
    axs[i, 1].set_yticks(range(len(class_names)))
    axs[i, 1].set_yticklabels(class_names, rotation=45)
    axs[i, 1].invert_yaxis()

plt.subplots_adjust(hspace=0.5, wspace=0.4)
plt.show()

```

```

In [ ]: def plot_training_history(history):
    # Debug print to inspect keys
    print("History keys:", history.history.keys())

    acc = history.history.get('accuracy', [])
    val_acc = history.history.get('val_accuracy', [])
    loss = history.history.get('loss', [])
    val_loss = history.history.get('val_loss', [])
    epochs = list(range(len(acc)))

    plt.figure(figsize=(14, 5))

    # Accuracy Plot
    plt.subplot(1, 2, 1)
    plt.plot(epochs, acc, 'b-', label='Training Accuracy')
    plt.plot(epochs, val_acc, 'r-', label='Validation Accuracy')
    plt.title('Training vs Validation Accuracy')
    plt.xlabel('Epochs')
    plt.ylabel('Accuracy')

```

```

plt.xticks(epochs)
plt.legend()

# Loss Plot
plt.subplot(1, 2, 2)
plt.plot(epochs, loss, 'b-', label='Training Loss')
plt.plot(epochs, val_loss, 'r-', label='Validation Loss')
plt.title('Training vs Validation Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.xticks(epochs)
plt.legend()

plt.tight_layout()
plt.show()

```

```

In [ ]: def plot_confusion_matrix(model, data_generator, class_names):
    y_true = data_generator.classes
    y_pred_probs = model.predict(data_generator)
    y_pred = np.argmax(y_pred_probs, axis=1)

    cm = confusion_matrix(y_true, y_pred)

    plt.figure(figsize=(10, 8))
    sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
                xticklabels=class_names,
                yticklabels=class_names)
    plt.title('Confusion Matrix')
    plt.xlabel('Predicted Labels')
    plt.ylabel('True Labels')
    plt.show()

    print("\nClassification Report:\n")
    print(classification_report(y_true, y_pred, target_names=class_names))

```

## Before fine-tuning

```

In [ ]: # 1. Define directories
train_dir = "/content/ninjacart_data/train"
val_dir = "/content/ninjacart_data/validation"
test_dir = "/content/ninjacart_data/test"

# 2. Image Data Generators
train_datagen = ImageDataGenerator(rescale=1./255, horizontal_flip=True)
val_datagen = ImageDataGenerator(rescale=1./255)
test_datagen = ImageDataGenerator(rescale=1./255)

train_generator = train_datagen.flow_from_directory(
    train_dir, target_size=(224, 224), batch_size=32, class_mode='categorical'
)

val_generator = val_datagen.flow_from_directory(
    val_dir, target_size=(224, 224), batch_size=32, class_mode='categorical'
)

test_generator = test_datagen.flow_from_directory(
    test_dir, target_size=(224, 224), batch_size=32, class_mode='categorical', s
)

```

```

# 3. Load base VGG16 model
base_model = VGG16(weights='imagenet', include_top=False, input_shape=(224, 224, 3))

# Freeze all layers first
for layer in base_model.layers:
    layer.trainable = False

# 4. Add custom classification head
x = base_model.output
x = tf.keras.layers.GlobalAveragePooling2D()(x)
x = tf.keras.layers.Dense(512, activation='relu')(x)
x = tf.keras.layers.Dropout(0.5)(x)
predictions = tf.keras.layers.Dense(train_generator.num_classes, activation='softmax')(x)

model = Model(inputs=base_model.input, outputs=predictions)

# 5. Compile the model
model.compile(optimizer=Adam(learning_rate=1e-4),
              loss='categorical_crossentropy',
              metrics=['accuracy'])

# 6. Train head only
history = model.fit(train_generator, validation_data=val_generator, epochs=10)

```

Found 2502 images belonging to 4 classes.  
 Found 627 images belonging to 4 classes.  
 Found 351 images belonging to 4 classes.  
 Downloading data from [https://storage.googleapis.com/tensorflow/keras-applications/vgg16/vgg16\\_weights\\_tf\\_dim\\_ordering\\_tf\\_kernels\\_notop.h5](https://storage.googleapis.com/tensorflow/keras-applications/vgg16/vgg16_weights_tf_dim_ordering_tf_kernels_notop.h5)  
 58889256/58889256 ————— 0s 0us/step

/usr/local/lib/python3.11/dist-packages/keras/src/trainers/data\_adapters/py\_dataset\_adapter.py:121: UserWarning: Your `PyDataset` class should call `super().\_\_init\_\_(\*\*kwargs)` in its constructor. `\*\*kwargs` can include `workers`, `use\_multiprocessing`, `max\_queue\_size`. Do not pass these arguments to `fit()`, as they will be ignored.

self.\_warn\_if\_super\_not\_called()

```
Epoch 1/10
79/79 43s 390ms/step - accuracy: 0.3548 - loss: 1.3781 - val
_accuracy: 0.7767 - val_loss: 0.9356
Epoch 2/10
79/79 17s 219ms/step - accuracy: 0.6406 - loss: 0.9452 - val
_accuracy: 0.8182 - val_loss: 0.7171
Epoch 3/10
79/79 22s 238ms/step - accuracy: 0.7335 - loss: 0.7616 - val
_accuracy: 0.8246 - val_loss: 0.5903
Epoch 4/10
79/79 19s 242ms/step - accuracy: 0.7947 - loss: 0.6030 - val
_accuracy: 0.8517 - val_loss: 0.5124
Epoch 5/10
79/79 17s 215ms/step - accuracy: 0.8182 - loss: 0.5376 - val
_accuracy: 0.8581 - val_loss: 0.4674
Epoch 6/10
79/79 18s 226ms/step - accuracy: 0.8215 - loss: 0.4927 - val
_accuracy: 0.8581 - val_loss: 0.4282
Epoch 7/10
79/79 18s 220ms/step - accuracy: 0.8407 - loss: 0.4504 - val
_accuracy: 0.8628 - val_loss: 0.4037
Epoch 8/10
79/79 18s 223ms/step - accuracy: 0.8666 - loss: 0.4109 - val
_accuracy: 0.8756 - val_loss: 0.3832
Epoch 9/10
79/79 18s 224ms/step - accuracy: 0.8739 - loss: 0.3872 - val
_accuracy: 0.8852 - val_loss: 0.3657
Epoch 10/10
79/79 17s 217ms/step - accuracy: 0.8672 - loss: 0.3582 - val
_accuracy: 0.8852 - val_loss: 0.3490
```

```
In [ ]: model.summary()
```

```
Model: "functional_10"
```

Layer (type)	Output Shape	Param #
input_layer_10 (InputLayer)	(None, 224, 224, 3)	0
block1_conv1 (Conv2D)	(None, 224, 224, 64)	1,792
block1_conv2 (Conv2D)	(None, 224, 224, 64)	36,928
block1_pool (MaxPooling2D)	(None, 112, 112, 64)	0
block2_conv1 (Conv2D)	(None, 112, 112, 128)	73,856
block2_conv2 (Conv2D)	(None, 112, 112, 128)	147,584
block2_pool (MaxPooling2D)	(None, 56, 56, 128)	0
block3_conv1 (Conv2D)	(None, 56, 56, 256)	295,168
block3_conv2 (Conv2D)	(None, 56, 56, 256)	590,080
block3_conv3 (Conv2D)	(None, 56, 56, 256)	590,080
block3_pool (MaxPooling2D)	(None, 28, 28, 256)	0
block4_conv1 (Conv2D)	(None, 28, 28, 512)	1,180,160
block4_conv2 (Conv2D)	(None, 28, 28, 512)	2,359,808
block4_conv3 (Conv2D)	(None, 28, 28, 512)	2,359,808
block4_pool (MaxPooling2D)	(None, 14, 14, 512)	0
block5_conv1 (Conv2D)	(None, 14, 14, 512)	2,359,808
block5_conv2 (Conv2D)	(None, 14, 14, 512)	2,359,808
block5_conv3 (Conv2D)	(None, 14, 14, 512)	2,359,808
block5_pool (MaxPooling2D)	(None, 7, 7, 512)	0
global_average_pooling2d_6 (GlobalAveragePooling2D)	(None, 512)	0
dense_14 (Dense)	(None, 512)	262,656
dropout_5 (Dropout)	(None, 512)	0
dense_15 (Dense)	(None, 4)	2,052

Total params: 15,508,814 (59.16 MB)

Trainable params: 264,708 (1.01 MB)

Non-trainable params: 14,714,688 (56.13 MB)

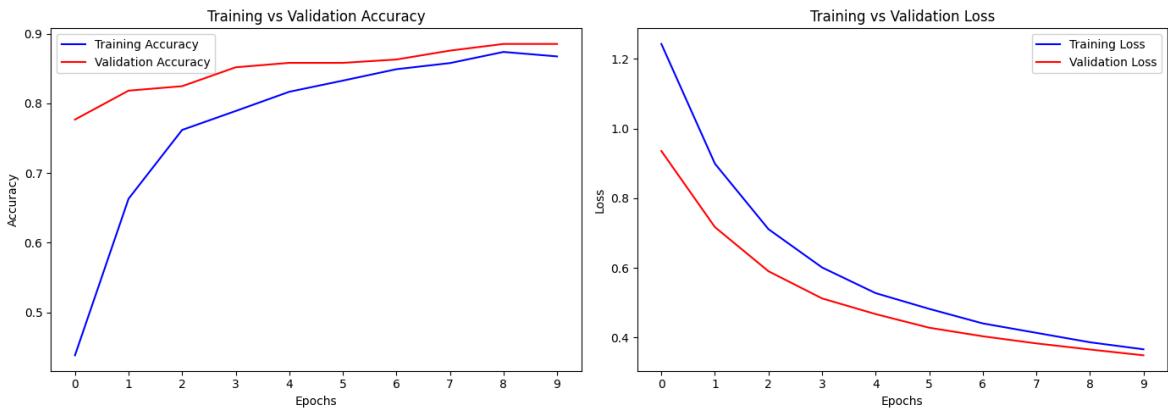
Optimizer params: 529,418 (2.02 MB)

```
In [ ]: create_directory_if_not_exists('/content/ninjacart_models/model8_VGG_GAP_DP_Fine
model.save("/content/ninjacart_models/model8_VGG_GAP_DP_FineTune/model8_VGG_GAP_
```

```
Created directory: /content/ninjacart_models/model8_VGG_GAP_DP_FineTune
```

```
In [ ]: plot_training_history(history)
```

```
History keys: dict_keys(['accuracy', 'loss', 'val_accuracy', 'val_loss'])
```



```
In [ ]: inception_model_finetune = tf.keras.models.load_model("/content/ninjacart_models")
loss, acc = inception_model_finetune.evaluate(test_generator, verbose=2)
print('Before FineTuning')
print("Restored model, accuracy: {:.2f}%".format(100 * acc))
```

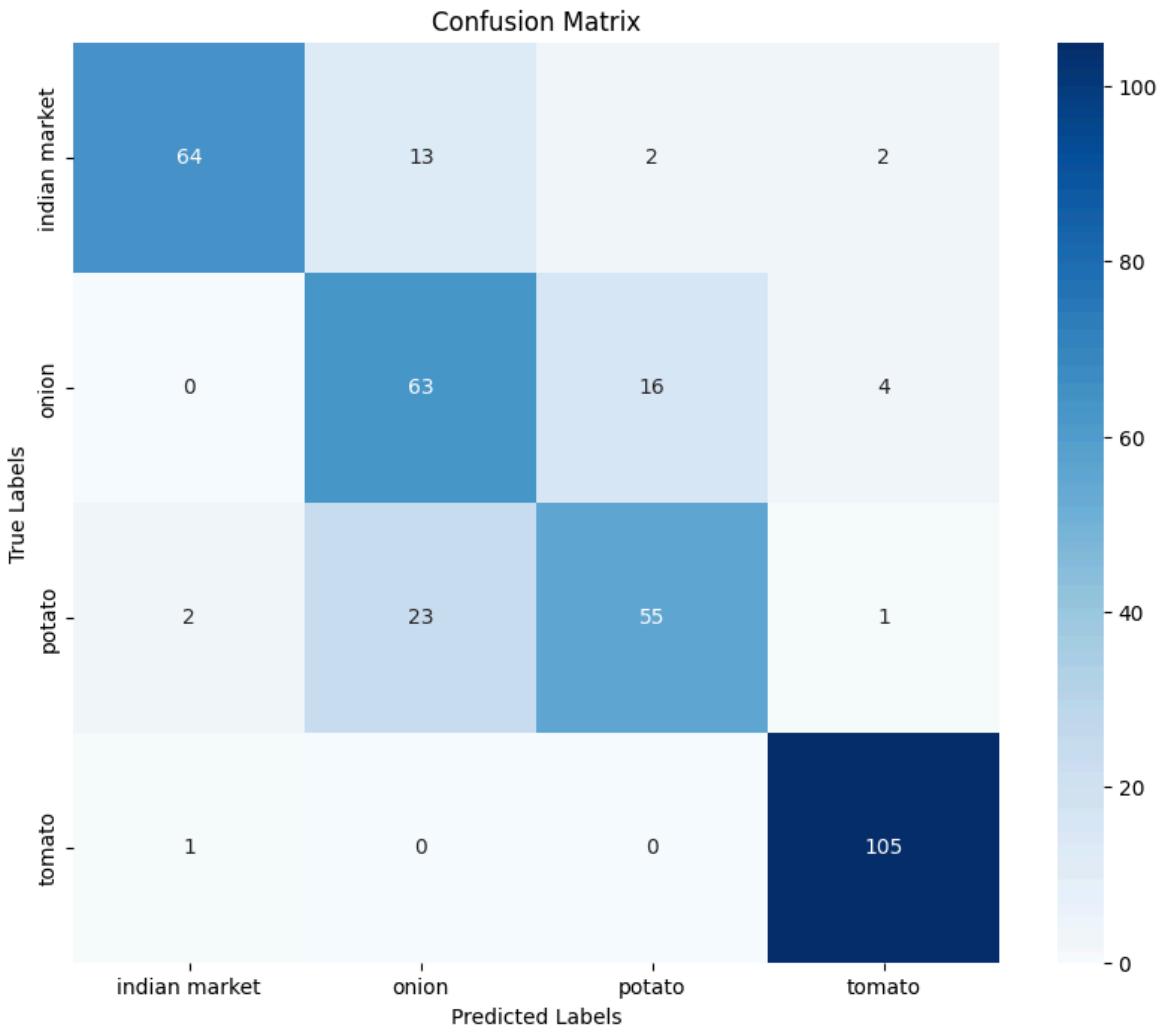
```
11/11 - 13s - 1s/step - accuracy: 0.8177 - loss: 0.4606
```

```
Before FineTuning
```

```
Restored model, accuracy: 81.77%
```

```
In [ ]: class_names = list(test_generator.class_indices.keys())
plot_confusion_matrix(model, test_generator, class_names)
```

```
11/11 ━━━━━━━━ 4s 342ms/step
```



Classification Report:

	precision	recall	f1-score	support
indian market	0.96	0.79	0.86	81
onion	0.64	0.76	0.69	83
potato	0.75	0.68	0.71	81
tomato	0.94	0.99	0.96	106
accuracy			0.82	351
macro avg	0.82	0.80	0.81	351
weighted avg	0.83	0.82	0.82	351

## After fine-tuning

```
In [ ]: # 7. Unfreeze last few layers of base model for fine-tuning
for layer in base_model.layers[-4:]:
    layer.trainable = True

# 8. Recompile with Lower LR
model.compile(optimizer=Adam(learning_rate=1e-5),
              loss='categorical_crossentropy',
              metrics=['accuracy'])

# 9. Continue training (fine-tune)
history_finetune = model.fit(train_generator, validation_data=val_generator, epochs=10)
```

```
Epoch 1/5
79/79 28s 293ms/step - accuracy: 0.8762 - loss: 0.3222 - val
_accuracy: 0.9059 - val_loss: 0.2376
Epoch 2/5
79/79 19s 239ms/step - accuracy: 0.9167 - loss: 0.2213 - val
_accuracy: 0.9266 - val_loss: 0.1941
Epoch 3/5
79/79 19s 239ms/step - accuracy: 0.9344 - loss: 0.1673 - val
_accuracy: 0.9250 - val_loss: 0.1914
Epoch 4/5
79/79 20s 237ms/step - accuracy: 0.9582 - loss: 0.1365 - val
_accuracy: 0.9458 - val_loss: 0.1654
Epoch 5/5
79/79 21s 270ms/step - accuracy: 0.9658 - loss: 0.0977 - val
_accuracy: 0.9458 - val_loss: 0.1527
```

```
In [ ]: model.summary()
```

```
Model: "functional_10"
```

Layer (type)	Output Shape	Param #
input_layer_10 (InputLayer)	(None, 224, 224, 3)	0
block1_conv1 (Conv2D)	(None, 224, 224, 64)	1,792
block1_conv2 (Conv2D)	(None, 224, 224, 64)	36,928
block1_pool (MaxPooling2D)	(None, 112, 112, 64)	0
block2_conv1 (Conv2D)	(None, 112, 112, 128)	73,856
block2_conv2 (Conv2D)	(None, 112, 112, 128)	147,584
block2_pool (MaxPooling2D)	(None, 56, 56, 128)	0
block3_conv1 (Conv2D)	(None, 56, 56, 256)	295,168
block3_conv2 (Conv2D)	(None, 56, 56, 256)	590,080
block3_conv3 (Conv2D)	(None, 56, 56, 256)	590,080
block3_pool (MaxPooling2D)	(None, 28, 28, 256)	0
block4_conv1 (Conv2D)	(None, 28, 28, 512)	1,180,160
block4_conv2 (Conv2D)	(None, 28, 28, 512)	2,359,808
block4_conv3 (Conv2D)	(None, 28, 28, 512)	2,359,808
block4_pool (MaxPooling2D)	(None, 14, 14, 512)	0
block5_conv1 (Conv2D)	(None, 14, 14, 512)	2,359,808
block5_conv2 (Conv2D)	(None, 14, 14, 512)	2,359,808
block5_conv3 (Conv2D)	(None, 14, 14, 512)	2,359,808
block5_pool (MaxPooling2D)	(None, 7, 7, 512)	0
global_average_pooling2d_6 (GlobalAveragePooling2D)	(None, 512)	0
dense_14 (Dense)	(None, 512)	262,656
dropout_5 (Dropout)	(None, 512)	0
dense_15 (Dense)	(None, 4)	2,052

Total params: 29,667,662 (113.17 MB)

Trainable params: 7,344,132 (28.02 MB)

Non-trainable params: 7,635,264 (29.13 MB)

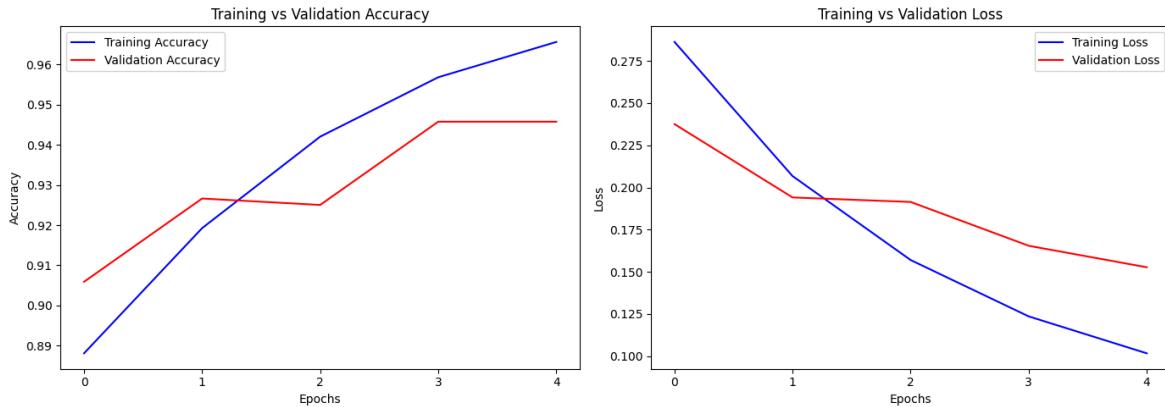
Optimizer params: 14,688,266 (56.03 MB)

```
In [ ]: create_directory_if_not_exists('/content/ninjacart_models/model8_VGG_GAP_DP_Fine
model.save("/content/ninjacart_models/model8_VGG_GAP_DP_FineTune/model8_VGG_GAP_
```

```
Directory already exists: /content/ninjacart_models/model8_VGG_GAP_DP_FineTune
```

```
In [ ]: plot_training_history(history_finetune)
```

```
History keys: dict_keys(['accuracy', 'loss', 'val_accuracy', 'val_loss'])
```



```
In [ ]: inception_model_finetune = tf.keras.models.load_model("/content/ninjacart_models/  
loss, acc = inception_model_finetune.evaluate(test_generator, verbose=2)  
print('After FineTuning')  
print("Restored model, accuracy: {:.5.2f}%".format(100 * acc))
```

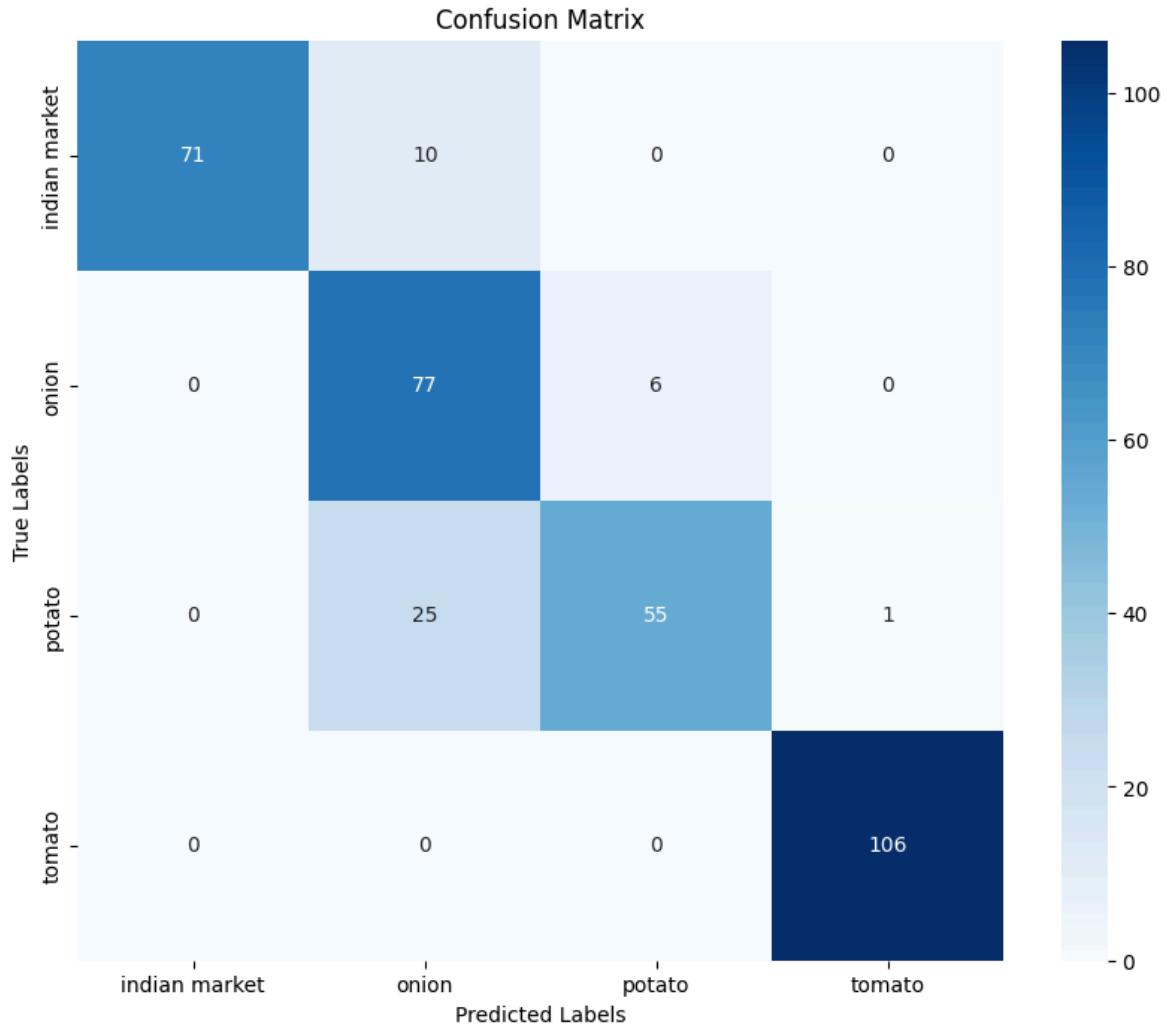
```
11/11 - 5s - 479ms/step - accuracy: 0.8803 - loss: 0.3373
```

```
After FineTuning
```

```
Restored model, accuracy: 88.03%
```

```
In [ ]: class_names = list(test_generator.class_indices.keys())  
plot_confusion_matrix(model, test_generator, class_names)
```

```
11/11 ━━━━━━━━ 5s 389ms/step
```



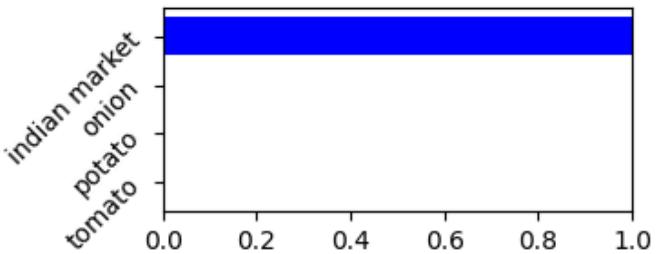
Classification Report:

	precision	recall	f1-score	support
indian market	1.00	0.88	0.93	81
onion	0.69	0.93	0.79	83
potato	0.90	0.68	0.77	81
tomato	0.99	1.00	1.00	106
accuracy			0.88	351
macro avg	0.89	0.87	0.87	351
weighted avg	0.90	0.88	0.88	351

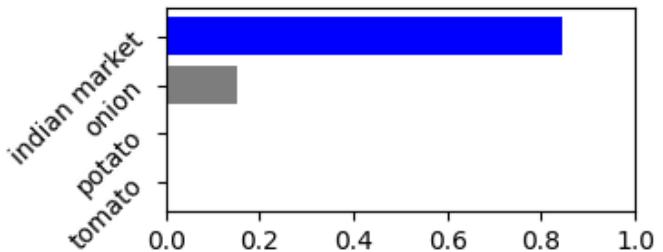
```
In [ ]: test_generator.reset()
images, labels = next(test_generator)
pred_probs = model.predict(images)
plot_predictions_with_bars(images[:5], labels[:5], pred_probs[:5], class_names)
```

1/1 ━━━━━━ 1s 1s/step

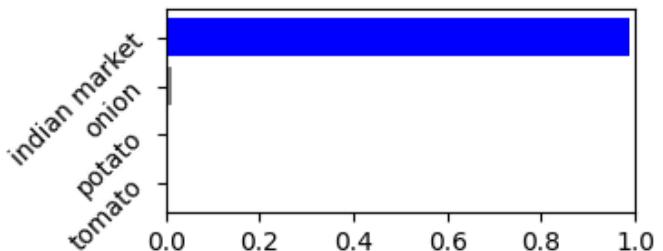
indian market 100% (indian market)



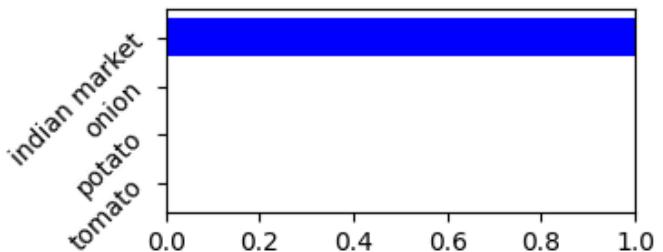
indian market 85% (indian market)



indian market 99% (indian market)



indian market 100% (indian market)



## InceptionV3 model

```
In [ ]: datagen = ImageDataGenerator(preprocessing_function=inception_v3.preprocess_input)

train_data = datagen.flow_from_directory(
    '/content/ninjacart_data/train',
    target_size=(299, 299),
    batch_size=32,
    class_mode='categorical'
)

val_data = datagen.flow_from_directory(
    '/content/ninjacart_data/validation',
    target_size=(299, 299),
    batch_size=32,
    class_mode='categorical',
    shuffle=False
)

test_data = datagen.flow_from_directory(
```

```

    '/content/ninjacart_data/test',
    target_size=(299, 299),
    batch_size=32,
    class_mode='categorical',
    shuffle=False
)

```

Found 2502 images belonging to 4 classes.  
 Found 627 images belonging to 4 classes.  
 Found 351 images belonging to 4 classes.

## Without FineTuning

```
In [ ]: num_classes = 4
# Load base InceptionV3 model without top classification Layer
base_model = InceptionV3(weights='imagenet', include_top=False, input_shape=(299, 299, 3))

# Freeze all layers first
for layer in base_model.layers:
    layer.trainable = False

# Add custom top layers
x = base_model.output
x = tf.keras.layers.GlobalAveragePooling2D()(x)
x = tf.keras.layers.Dropout(0.5)(x)          # Optional dropout
x = tf.keras.layers.Dense(1024, activation='relu')(x) # Fully connected layer
predictions = tf.keras.layers.Dense(num_classes, activation='softmax')(x)

# Final model
inception_model = Model(inputs=base_model.input, outputs=predictions)

# Compile the model
inception_model.compile(optimizer=Adam(learning_rate=0.001), loss='categorical_crossentropy')

# Train only the top layers
before_ft_inc_hist = inception_model.fit(train_data, epochs=10, validation_data=val_data)
```

Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/inception\_v3/inception\_v3\_weights\_tf\_dim\_ordering\_tf\_kernels\_notop.h5  
 87910968/87910968 ————— 0s 0us/step

```
/usr/local/lib/python3.11/dist-packages/keras/src/trainers/data_adapters/py_dataset_adapter.py:121: UserWarning: Your `PyDataset` class should call `super().__init__(**kwargs)` in its constructor. `**kwargs` can include `workers`, `use_multiprocessing`, `max_queue_size`. Do not pass these arguments to `fit()`, as they will be ignored.
    self._warn_if_super_not_called()
```

```
Epoch 1/10
79/79 72s 542ms/step - accuracy: 0.8215 - loss: 0.5547 - val
_accuracy: 0.9458 - val_loss: 0.1409
Epoch 2/10
79/79 22s 274ms/step - accuracy: 0.9487 - loss: 0.1236 - val
_accuracy: 0.9681 - val_loss: 0.0904
Epoch 3/10
79/79 41s 271ms/step - accuracy: 0.9622 - loss: 0.0980 - val
_accuracy: 0.9585 - val_loss: 0.1406
Epoch 4/10
79/79 20s 259ms/step - accuracy: 0.9701 - loss: 0.0806 - val
_accuracy: 0.9793 - val_loss: 0.0686
Epoch 5/10
79/79 21s 261ms/step - accuracy: 0.9763 - loss: 0.0622 - val
_accuracy: 0.9649 - val_loss: 0.1100
Epoch 6/10
79/79 20s 259ms/step - accuracy: 0.9663 - loss: 0.1068 - val
_accuracy: 0.9793 - val_loss: 0.0650
Epoch 7/10
79/79 21s 269ms/step - accuracy: 0.9752 - loss: 0.0593 - val
_accuracy: 0.9793 - val_loss: 0.0754
Epoch 8/10
79/79 20s 257ms/step - accuracy: 0.9726 - loss: 0.0816 - val
_accuracy: 0.9681 - val_loss: 0.0930
Epoch 9/10
79/79 22s 278ms/step - accuracy: 0.9736 - loss: 0.0870 - val
_accuracy: 0.9745 - val_loss: 0.0675
Epoch 10/10
79/79 21s 260ms/step - accuracy: 0.9813 - loss: 0.0550 - val
_accuracy: 0.9713 - val_loss: 0.0905
```

```
In [ ]: inception_model.summary()
```

```
Model: "functional_11"
```

Layer (type)	Output Shape	Param #	Connected to
input_layer_11 (InputLayer)	(None, 299, 299, 3)	0	-
conv2d_31 (Conv2D)	(None, 149, 149, 32)	864	input_layer_11[0]...
batch_normalizatio... (BatchNormalizatio...)	(None, 149, 149, 32)	96	conv2d_31[0][0]
activation_30 (Activation)	(None, 149, 149, 32)	0	batch_normalizat...
conv2d_32 (Conv2D)	(None, 147, 147, 32)	9,216	activation_30[0]...
batch_normalizatio... (BatchNormalizatio...)	(None, 147, 147, 32)	96	conv2d_32[0][0]
activation_31 (Activation)	(None, 147, 147, 32)	0	batch_normalizat...
conv2d_33 (Conv2D)	(None, 147, 147, 64)	18,432	activation_31[0]...
batch_normalizatio... (BatchNormalizatio...)	(None, 147, 147, 64)	192	conv2d_33[0][0]
activation_32 (Activation)	(None, 147, 147, 64)	0	batch_normalizat...
max_pooling2d_25 (MaxPooling2D)	(None, 73, 73, 64)	0	activation_32[0]...
conv2d_34 (Conv2D)	(None, 73, 73, 80)	5,120	max_pooling2d_25...
batch_normalizatio... (BatchNormalizatio...)	(None, 73, 73, 80)	240	conv2d_34[0][0]
activation_33 (Activation)	(None, 73, 73, 80)	0	batch_normalizat...
conv2d_35 (Conv2D)	(None, 71, 71, 192)	138,240	activation_33[0]...
batch_normalizatio... (BatchNormalizatio...)	(None, 71, 71, 192)	576	conv2d_35[0][0]
activation_34 (Activation)	(None, 71, 71, 192)	0	batch_normalizat...
max_pooling2d_26 (MaxPooling2D)	(None, 35, 35, 192)	0	activation_34[0]...
conv2d_39 (Conv2D)	(None, 35, 35, 64)	12,288	max_pooling2d_26...

batch_normalization_39 (BatchNormalization)	(None, 35, 35, 64)	192	conv2d_39[0][0]
activation_38 (Activation)	(None, 35, 35, 64)	0	batch_normalizat...
conv2d_37 (Conv2D)	(None, 35, 35, 48)	9,216	max_pooling2d_26...
conv2d_40 (Conv2D)	(None, 35, 35, 96)	55,296	activation_38[0]...
batch_normalization_37 (BatchNormalization)	(None, 35, 35, 48)	144	conv2d_37[0][0]
batch_normalization_38 (BatchNormalization)	(None, 35, 35, 96)	288	conv2d_40[0][0]
activation_36 (Activation)	(None, 35, 35, 48)	0	batch_normalizat...
activation_39 (Activation)	(None, 35, 35, 96)	0	batch_normalizat...
average_pooling2d (AveragePooling2D)	(None, 35, 35, 192)	0	max_pooling2d_26...
conv2d_36 (Conv2D)	(None, 35, 35, 64)	12,288	max_pooling2d_26...
conv2d_38 (Conv2D)	(None, 35, 35, 64)	76,800	activation_36[0]...
conv2d_41 (Conv2D)	(None, 35, 35, 96)	82,944	activation_39[0]...
conv2d_42 (Conv2D)	(None, 35, 35, 32)	6,144	average_pooling2...
batch_normalization_36 (BatchNormalization)	(None, 35, 35, 64)	192	conv2d_36[0][0]
batch_normalization_37 (BatchNormalization)	(None, 35, 35, 64)	192	conv2d_38[0][0]
batch_normalization_38 (BatchNormalization)	(None, 35, 35, 96)	288	conv2d_41[0][0]
batch_normalization_39 (BatchNormalization)	(None, 35, 35, 32)	96	conv2d_42[0][0]
activation_35 (Activation)	(None, 35, 35, 64)	0	batch_normalizat...
activation_37 (Activation)	(None, 35, 35, 64)	0	batch_normalizat...
activation_40 (Activation)	(None, 35, 35, 96)	0	batch_normalizat...

activation_41 (Activation)	(None, 35, 35, 32)	0	batch_normalizat...
mixed0 (Concatenate)	(None, 35, 35, 256)	0	activation_35[0]... activation_37[0]... activation_40[0]... activation_41[0]...
conv2d_46 (Conv2D)	(None, 35, 35, 64)	16,384	mixed0[0][0]
batch_normalizatio... (BatchNormalizatio...)	(None, 35, 35, 64)	192	conv2d_46[0][0]
activation_45 (Activation)	(None, 35, 35, 64)	0	batch_normalizat...
conv2d_44 (Conv2D)	(None, 35, 35, 48)	12,288	mixed0[0][0]
conv2d_47 (Conv2D)	(None, 35, 35, 96)	55,296	activation_45[0]...
batch_normalizatio... (BatchNormalizatio...)	(None, 35, 35, 48)	144	conv2d_44[0][0]
batch_normalizatio... (BatchNormalizatio...)	(None, 35, 35, 96)	288	conv2d_47[0][0]
activation_43 (Activation)	(None, 35, 35, 48)	0	batch_normalizat...
activation_46 (Activation)	(None, 35, 35, 96)	0	batch_normalizat...
average_pooling2d_1 (AveragePooling2D)	(None, 35, 35, 256)	0	mixed0[0][0]
conv2d_43 (Conv2D)	(None, 35, 35, 64)	16,384	mixed0[0][0]
conv2d_45 (Conv2D)	(None, 35, 35, 64)	76,800	activation_43[0]...
conv2d_48 (Conv2D)	(None, 35, 35, 96)	82,944	activation_46[0]...
conv2d_49 (Conv2D)	(None, 35, 35, 64)	16,384	average_pooling2...
batch_normalizatio... (BatchNormalizatio...)	(None, 35, 35, 64)	192	conv2d_43[0][0]
batch_normalizatio... (BatchNormalizatio...)	(None, 35, 35, 64)	192	conv2d_45[0][0]
batch_normalizatio... (BatchNormalizatio...)	(None, 35, 35, 96)	288	conv2d_48[0][0]
batch_normalizatio... (BatchNormalizatio...)	(None, 35, 35,	192	conv2d_49[0][0]

(BatchNormalizatio...	64)		
activation_42 (Activation)	(None, 35, 35, 64)	0	batch_normalizat...
activation_44 (Activation)	(None, 35, 35, 64)	0	batch_normalizat...
activation_47 (Activation)	(None, 35, 35, 96)	0	batch_normalizat...
activation_48 (Activation)	(None, 35, 35, 64)	0	batch_normalizat...
mixed1 (Concatenate)	(None, 35, 35, 288)	0	activation_42[0]... activation_44[0]... activation_47[0]... activation_48[0]...
conv2d_53 (Conv2D)	(None, 35, 35, 64)	18,432	mixed1[0][0]
batch_normalizatio... (BatchNormalizatio...	(None, 35, 35, 64)	192	conv2d_53[0][0]
activation_52 (Activation)	(None, 35, 35, 64)	0	batch_normalizat...
conv2d_51 (Conv2D)	(None, 35, 35, 48)	13,824	mixed1[0][0]
conv2d_54 (Conv2D)	(None, 35, 35, 96)	55,296	activation_52[0]...
batch_normalizatio... (BatchNormalizatio...	(None, 35, 35, 48)	144	conv2d_51[0][0]
batch_normalizatio... (BatchNormalizatio...	(None, 35, 35, 96)	288	conv2d_54[0][0]
activation_50 (Activation)	(None, 35, 35, 48)	0	batch_normalizat...
activation_53 (Activation)	(None, 35, 35, 96)	0	batch_normalizat...
average_pooling2d_2 (AveragePooling2D)	(None, 35, 35, 288)	0	mixed1[0][0]
conv2d_50 (Conv2D)	(None, 35, 35, 64)	18,432	mixed1[0][0]
conv2d_52 (Conv2D)	(None, 35, 35, 64)	76,800	activation_50[0]...
conv2d_55 (Conv2D)	(None, 35, 35, 96)	82,944	activation_53[0]...
conv2d_56 (Conv2D)	(None, 35, 35, 64)	18,432	average_pooling2...

batch_normalizatio... (BatchNormalizatio...)	(None, 35, 35, 64)	192	conv2d_50[0][0]
batch_normalizatio... (BatchNormalizatio...)	(None, 35, 35, 64)	192	conv2d_52[0][0]
batch_normalizatio... (BatchNormalizatio...)	(None, 35, 35, 96)	288	conv2d_55[0][0]
batch_normalizatio... (BatchNormalizatio...)	(None, 35, 35, 64)	192	conv2d_56[0][0]
activation_49 (Activation)	(None, 35, 35, 64)	0	batch_normalizat...
activation_51 (Activation)	(None, 35, 35, 64)	0	batch_normalizat...
activation_54 (Activation)	(None, 35, 35, 96)	0	batch_normalizat...
activation_55 (Activation)	(None, 35, 35, 64)	0	batch_normalizat...
mixed2 (Concatenate)	(None, 35, 35, 288)	0	activation_49[0]... activation_51[0]... activation_54[0]... activation_55[0]...
conv2d_58 (Conv2D)	(None, 35, 35, 64)	18,432	mixed2[0][0]
batch_normalizatio... (BatchNormalizatio...)	(None, 35, 35, 64)	192	conv2d_58[0][0]
activation_57 (Activation)	(None, 35, 35, 64)	0	batch_normalizat...
conv2d_59 (Conv2D)	(None, 35, 35, 96)	55,296	activation_57[0]...
batch_normalizatio... (BatchNormalizatio...)	(None, 35, 35, 96)	288	conv2d_59[0][0]
activation_58 (Activation)	(None, 35, 35, 96)	0	batch_normalizat...
conv2d_57 (Conv2D)	(None, 17, 17, 384)	995,328	mixed2[0][0]
conv2d_60 (Conv2D)	(None, 17, 17, 96)	82,944	activation_58[0]...
batch_normalizatio... (BatchNormalizatio...)	(None, 17, 17, 384)	1,152	conv2d_57[0][0]
batch_normalizatio... (BatchNormalizatio...)	(None, 17, 17, 96)	288	conv2d_60[0][0]

activation_56 (Activation)	(None, 17, 17, 384)	0	batch_normalizat...
activation_59 (Activation)	(None, 17, 17, 96)	0	batch_normalizat...
max_pooling2d_27 (MaxPooling2D)	(None, 17, 17, 288)	0	mixed2[0][0]
mixed3 (Concatenate)	(None, 17, 17, 768)	0	activation_56[0]... activation_59[0]... max_pooling2d_27...
conv2d_65 (Conv2D)	(None, 17, 17, 128)	98,304	mixed3[0][0]
batch_normalizatio... (BatchNormalizatio...)	(None, 17, 17, 128)	384	conv2d_65[0][0]
activation_64 (Activation)	(None, 17, 17, 128)	0	batch_normalizat...
conv2d_66 (Conv2D)	(None, 17, 17, 128)	114,688	activation_64[0]...
batch_normalizatio... (BatchNormalizatio...)	(None, 17, 17, 128)	384	conv2d_66[0][0]
activation_65 (Activation)	(None, 17, 17, 128)	0	batch_normalizat...
conv2d_62 (Conv2D)	(None, 17, 17, 128)	98,304	mixed3[0][0]
conv2d_67 (Conv2D)	(None, 17, 17, 128)	114,688	activation_65[0]...
batch_normalizatio... (BatchNormalizatio...)	(None, 17, 17, 128)	384	conv2d_62[0][0]
batch_normalizatio... (BatchNormalizatio...)	(None, 17, 17, 128)	384	conv2d_67[0][0]
activation_61 (Activation)	(None, 17, 17, 128)	0	batch_normalizat...
activation_66 (Activation)	(None, 17, 17, 128)	0	batch_normalizat...
conv2d_63 (Conv2D)	(None, 17, 17, 128)	114,688	activation_61[0]...
conv2d_68 (Conv2D)	(None, 17, 17, 128)	114,688	activation_66[0]...
batch_normalizatio... (BatchNormalizatio...)	(None, 17, 17, 128)	384	conv2d_63[0][0]
batch_normalizatio... (BatchNormalizatio...)	(None, 17, 17, 128)	384	conv2d_68[0][0]

activation_62 (Activation)	(None, 17, 17, 128)	0	batch_normalizat...
activation_67 (Activation)	(None, 17, 17, 128)	0	batch_normalizat...
average_pooling2d_3 (AveragePooling2D)	(None, 17, 17, 768)	0	mixed3[0][0]
conv2d_61 (Conv2D)	(None, 17, 17, 192)	147,456	mixed3[0][0]
conv2d_64 (Conv2D)	(None, 17, 17, 192)	172,032	activation_62[0]...
conv2d_69 (Conv2D)	(None, 17, 17, 192)	172,032	activation_67[0]...
conv2d_70 (Conv2D)	(None, 17, 17, 192)	147,456	average_pooling2...
batch_normalizatio... (BatchNormalizatio...)	(None, 17, 17, 192)	576	conv2d_61[0][0]
batch_normalizatio... (BatchNormalizatio...)	(None, 17, 17, 192)	576	conv2d_64[0][0]
batch_normalizatio... (BatchNormalizatio...)	(None, 17, 17, 192)	576	conv2d_69[0][0]
batch_normalizatio... (BatchNormalizatio...)	(None, 17, 17, 192)	576	conv2d_70[0][0]
activation_60 (Activation)	(None, 17, 17, 192)	0	batch_normalizat...
activation_63 (Activation)	(None, 17, 17, 192)	0	batch_normalizat...
activation_68 (Activation)	(None, 17, 17, 192)	0	batch_normalizat...
activation_69 (Activation)	(None, 17, 17, 192)	0	batch_normalizat...
mixed4 (Concatenate)	(None, 17, 17, 768)	0	activation_60[0]... activation_63[0]... activation_68[0]... activation_69[0]...
conv2d_75 (Conv2D)	(None, 17, 17, 160)	122,880	mixed4[0][0]
batch_normalizatio... (BatchNormalizatio...)	(None, 17, 17, 160)	480	conv2d_75[0][0]
activation_74 (Activation)	(None, 17, 17, 160)	0	batch_normalizat...

conv2d_76 (Conv2D)	(None, 17, 17, 160)	179,200	activation_74[0]...
batch_normalizatio... (BatchNormalizatio...)	(None, 17, 17, 160)	480	conv2d_76[0][0]
activation_75 (Activation)	(None, 17, 17, 160)	0	batch_normalizat...
conv2d_72 (Conv2D)	(None, 17, 17, 160)	122,880	mixed4[0][0]
conv2d_77 (Conv2D)	(None, 17, 17, 160)	179,200	activation_75[0]...
batch_normalizatio... (BatchNormalizatio...)	(None, 17, 17, 160)	480	conv2d_72[0][0]
batch_normalizatio... (BatchNormalizatio...)	(None, 17, 17, 160)	480	conv2d_77[0][0]
activation_71 (Activation)	(None, 17, 17, 160)	0	batch_normalizat...
activation_76 (Activation)	(None, 17, 17, 160)	0	batch_normalizat...
conv2d_73 (Conv2D)	(None, 17, 17, 160)	179,200	activation_71[0]...
conv2d_78 (Conv2D)	(None, 17, 17, 160)	179,200	activation_76[0]...
batch_normalizatio... (BatchNormalizatio...)	(None, 17, 17, 160)	480	conv2d_73[0][0]
batch_normalizatio... (BatchNormalizatio...)	(None, 17, 17, 160)	480	conv2d_78[0][0]
activation_72 (Activation)	(None, 17, 17, 160)	0	batch_normalizat...
activation_77 (Activation)	(None, 17, 17, 160)	0	batch_normalizat...
average_pooling2d_4 (AveragePooling2D)	(None, 17, 17, 768)	0	mixed4[0][0]
conv2d_71 (Conv2D)	(None, 17, 17, 192)	147,456	mixed4[0][0]
conv2d_74 (Conv2D)	(None, 17, 17, 192)	215,040	activation_72[0]...
conv2d_79 (Conv2D)	(None, 17, 17, 192)	215,040	activation_77[0]...
conv2d_80 (Conv2D)	(None, 17, 17, 192)	147,456	average_pooling2...

batch_normalization_71 (BatchNormalization)	(None, 17, 17, 192)	576	conv2d_71[0][0]
batch_normalization_74 (BatchNormalization)	(None, 17, 17, 192)	576	conv2d_74[0][0]
batch_normalization_79 (BatchNormalization)	(None, 17, 17, 192)	576	conv2d_79[0][0]
batch_normalization_80 (BatchNormalization)	(None, 17, 17, 192)	576	conv2d_80[0][0]
activation_70 (Activation)	(None, 17, 17, 192)	0	batch_normalizat...
activation_73 (Activation)	(None, 17, 17, 192)	0	batch_normalizat...
activation_78 (Activation)	(None, 17, 17, 192)	0	batch_normalizat...
activation_79 (Activation)	(None, 17, 17, 192)	0	batch_normalizat...
mixed5 (Concatenate)	(None, 17, 17, 768)	0	activation_70[0]... activation_73[0]... activation_78[0]... activation_79[0]...
conv2d_85 (Conv2D)	(None, 17, 17, 160)	122,880	mixed5[0][0]
batch_normalization_85 (BatchNormalization)	(None, 17, 17, 160)	480	conv2d_85[0][0]
activation_84 (Activation)	(None, 17, 17, 160)	0	batch_normalizat...
conv2d_86 (Conv2D)	(None, 17, 17, 160)	179,200	activation_84[0]...
batch_normalization_86 (BatchNormalization)	(None, 17, 17, 160)	480	conv2d_86[0][0]
activation_85 (Activation)	(None, 17, 17, 160)	0	batch_normalizat...
conv2d_82 (Conv2D)	(None, 17, 17, 160)	122,880	mixed5[0][0]
conv2d_87 (Conv2D)	(None, 17, 17, 160)	179,200	activation_85[0]...
batch_normalization_82 (BatchNormalization)	(None, 17, 17, 160)	480	conv2d_82[0][0]
batch_normalization_87 (BatchNormalization)	(None, 17, 17, 160)	480	conv2d_87[0][0]
activation_81	(None, 17, 17,	0	batch_normalizat...

(Activation)	160)		
activation_86 (Activation)	(None, 17, 17, 160)	0	batch_normalizat...
conv2d_83 (Conv2D)	(None, 17, 17, 160)	179,200	activation_81[0]...
conv2d_88 (Conv2D)	(None, 17, 17, 160)	179,200	activation_86[0]...
batch_normalizatio... (BatchNormalizatio...)	(None, 17, 17, 160)	480	conv2d_83[0][0]
batch_normalizatio... (BatchNormalizatio...)	(None, 17, 17, 160)	480	conv2d_88[0][0]
activation_82 (Activation)	(None, 17, 17, 160)	0	batch_normalizat...
activation_87 (Activation)	(None, 17, 17, 160)	0	batch_normalizat...
average_pooling2d_5 (AveragePooling2D)	(None, 17, 17, 768)	0	mixed5[0][0]
conv2d_81 (Conv2D)	(None, 17, 17, 192)	147,456	mixed5[0][0]
conv2d_84 (Conv2D)	(None, 17, 17, 192)	215,040	activation_82[0]...
conv2d_89 (Conv2D)	(None, 17, 17, 192)	215,040	activation_87[0]...
conv2d_90 (Conv2D)	(None, 17, 17, 192)	147,456	average_pooling2...
batch_normalizatio... (BatchNormalizatio...)	(None, 17, 17, 192)	576	conv2d_81[0][0]
batch_normalizatio... (BatchNormalizatio...)	(None, 17, 17, 192)	576	conv2d_84[0][0]
batch_normalizatio... (BatchNormalizatio...)	(None, 17, 17, 192)	576	conv2d_89[0][0]
batch_normalizatio... (BatchNormalizatio...)	(None, 17, 17, 192)	576	conv2d_90[0][0]
activation_80 (Activation)	(None, 17, 17, 192)	0	batch_normalizat...
activation_83 (Activation)	(None, 17, 17, 192)	0	batch_normalizat...
activation_88 (Activation)	(None, 17, 17, 192)	0	batch_normalizat...
activation_89	(None, 17, 17,	0	batch_normalizat...

(Activation)	192)		
mixed6 (Concatenate)	(None, 17, 17, 768)	0	activation_80[0]... activation_83[0]... activation_88[0]... activation_89[0]...
conv2d_95 (Conv2D)	(None, 17, 17, 192)	147,456	mixed6[0][0]
batch_normalizatio... (BatchNormalizatio...)	(None, 17, 17, 192)	576	conv2d_95[0][0]
activation_94 (Activation)	(None, 17, 17, 192)	0	batch_normalizat...
conv2d_96 (Conv2D)	(None, 17, 17, 192)	258,048	activation_94[0]...
batch_normalizatio... (BatchNormalizatio...)	(None, 17, 17, 192)	576	conv2d_96[0][0]
activation_95 (Activation)	(None, 17, 17, 192)	0	batch_normalizat...
conv2d_92 (Conv2D)	(None, 17, 17, 192)	147,456	mixed6[0][0]
conv2d_97 (Conv2D)	(None, 17, 17, 192)	258,048	activation_95[0]...
batch_normalizatio... (BatchNormalizatio...)	(None, 17, 17, 192)	576	conv2d_92[0][0]
batch_normalizatio... (BatchNormalizatio...)	(None, 17, 17, 192)	576	conv2d_97[0][0]
activation_91 (Activation)	(None, 17, 17, 192)	0	batch_normalizat...
activation_96 (Activation)	(None, 17, 17, 192)	0	batch_normalizat...
conv2d_93 (Conv2D)	(None, 17, 17, 192)	258,048	activation_91[0]...
conv2d_98 (Conv2D)	(None, 17, 17, 192)	258,048	activation_96[0]...
batch_normalizatio... (BatchNormalizatio...)	(None, 17, 17, 192)	576	conv2d_93[0][0]
batch_normalizatio... (BatchNormalizatio...)	(None, 17, 17, 192)	576	conv2d_98[0][0]
activation_92 (Activation)	(None, 17, 17, 192)	0	batch_normalizat...
activation_97 (Activation)	(None, 17, 17, 192)	0	batch_normalizat...

average_pooling2d_6 (AveragePooling2D)	(None, 17, 17, 768)	0	mixed6[0][0]
conv2d_91 (Conv2D)	(None, 17, 17, 192)	147,456	mixed6[0][0]
conv2d_94 (Conv2D)	(None, 17, 17, 192)	258,048	activation_92[0]...
conv2d_99 (Conv2D)	(None, 17, 17, 192)	258,048	activation_97[0]...
conv2d_100 (Conv2D)	(None, 17, 17, 192)	147,456	average_pooling2...
batch_normalizatio... (BatchNormalizatio...)	(None, 17, 17, 192)	576	conv2d_91[0][0]
batch_normalizatio... (BatchNormalizatio...)	(None, 17, 17, 192)	576	conv2d_94[0][0]
batch_normalizatio... (BatchNormalizatio...)	(None, 17, 17, 192)	576	conv2d_99[0][0]
batch_normalizatio... (BatchNormalizatio...)	(None, 17, 17, 192)	576	conv2d_100[0][0]
activation_90 (Activation)	(None, 17, 17, 192)	0	batch_normalizat...
activation_93 (Activation)	(None, 17, 17, 192)	0	batch_normalizat...
activation_98 (Activation)	(None, 17, 17, 192)	0	batch_normalizat...
activation_99 (Activation)	(None, 17, 17, 192)	0	batch_normalizat...
mixed7 (Concatenate)	(None, 17, 17, 768)	0	activation_90[0]... activation_93[0]... activation_98[0]... activation_99[0]...
conv2d_103 (Conv2D)	(None, 17, 17, 192)	147,456	mixed7[0][0]
batch_normalizatio... (BatchNormalizatio...)	(None, 17, 17, 192)	576	conv2d_103[0][0]
activation_102 (Activation)	(None, 17, 17, 192)	0	batch_normalizat...
conv2d_104 (Conv2D)	(None, 17, 17, 192)	258,048	activation_102[0]...
batch_normalizatio... (BatchNormalizatio...)	(None, 17, 17, 192)	576	conv2d_104[0][0]

activation_103 (Activation)	(None, 17, 17, 192)	0	batch_normalizat...
conv2d_101 (Conv2D)	(None, 17, 17, 192)	147,456	mixed7[0][0]
conv2d_105 (Conv2D)	(None, 17, 17, 192)	258,048	activation_103[0...]
batch_normalizatio... (BatchNormalizatio...)	(None, 17, 17, 192)	576	conv2d_101[0][0]
batch_normalizatio... (BatchNormalizatio...)	(None, 17, 17, 192)	576	conv2d_105[0][0]
activation_100 (Activation)	(None, 17, 17, 192)	0	batch_normalizat...
activation_104 (Activation)	(None, 17, 17, 192)	0	batch_normalizat...
conv2d_102 (Conv2D)	(None, 8, 8, 320)	552,960	activation_100[0...]
conv2d_106 (Conv2D)	(None, 8, 8, 192)	331,776	activation_104[0...]
batch_normalizatio... (BatchNormalizatio...)	(None, 8, 8, 320)	960	conv2d_102[0][0]
batch_normalizatio... (BatchNormalizatio...)	(None, 8, 8, 192)	576	conv2d_106[0][0]
activation_101 (Activation)	(None, 8, 8, 320)	0	batch_normalizat...
activation_105 (Activation)	(None, 8, 8, 192)	0	batch_normalizat...
max_pooling2d_28 (MaxPooling2D)	(None, 8, 8, 768)	0	mixed7[0][0]
mixed8 (Concatenate)	(None, 8, 8, 1280)	0	activation_101[0...] activation_105[0...] max_pooling2d_28[0]
conv2d_111 (Conv2D)	(None, 8, 8, 448)	573,440	mixed8[0][0]
batch_normalizatio... (BatchNormalizatio...)	(None, 8, 8, 448)	1,344	conv2d_111[0][0]
activation_110 (Activation)	(None, 8, 8, 448)	0	batch_normalizat...
conv2d_108 (Conv2D)	(None, 8, 8, 384)	491,520	mixed8[0][0]
conv2d_112 (Conv2D)	(None, 8, 8, 384)	1,548,288	activation_110[0...]
batch_normalizatio... (BatchNormalizatio...)	(None, 8, 8, 384)	1,152	conv2d_108[0][0]
batch_normalizatio... (BatchNormalizatio...)	(None, 8, 8, 384)	1,152	conv2d_112[0][0]

(BatchNormalizatio...			
activation_107 (Activation)	(None, 8, 8, 384)	0	batch_normalizat...
activation_111 (Activation)	(None, 8, 8, 384)	0	batch_normalizat...
conv2d_109 (Conv2D)	(None, 8, 8, 384)	442,368	activation_107[0...]
conv2d_110 (Conv2D)	(None, 8, 8, 384)	442,368	activation_107[0...]
conv2d_113 (Conv2D)	(None, 8, 8, 384)	442,368	activation_111[0...]
conv2d_114 (Conv2D)	(None, 8, 8, 384)	442,368	activation_111[0...]
average_pooling2d_7 (AveragePooling2D)	(None, 8, 8, 1280)	0	mixed8[0][0]
conv2d_107 (Conv2D)	(None, 8, 8, 320)	409,600	mixed8[0][0]
batch_normalizatio... (BatchNormalizatio...	(None, 8, 8, 384)	1,152	conv2d_109[0][0]
batch_normalizatio... (BatchNormalizatio...	(None, 8, 8, 384)	1,152	conv2d_110[0][0]
batch_normalizatio... (BatchNormalizatio...	(None, 8, 8, 384)	1,152	conv2d_113[0][0]
batch_normalizatio... (BatchNormalizatio...	(None, 8, 8, 384)	1,152	conv2d_114[0][0]
conv2d_115 (Conv2D)	(None, 8, 8, 192)	245,760	average_pooling2...
batch_normalizatio... (BatchNormalizatio...	(None, 8, 8, 320)	960	conv2d_107[0][0]
activation_108 (Activation)	(None, 8, 8, 384)	0	batch_normalizat...
activation_109 (Activation)	(None, 8, 8, 384)	0	batch_normalizat...
activation_112 (Activation)	(None, 8, 8, 384)	0	batch_normalizat...
activation_113 (Activation)	(None, 8, 8, 384)	0	batch_normalizat...
batch_normalizatio... (BatchNormalizatio...	(None, 8, 8, 192)	576	conv2d_115[0][0]
activation_106 (Activation)	(None, 8, 8, 320)	0	batch_normalizat...
mixed9_0 (Concatenate)	(None, 8, 8, 768)	0	activation_108[0...] activation_109[0...]
concatenate	(None, 8, 8, 768)	0	activation_112[0...]

(Concatenate)			activation_113[0...]
activation_114 (Activation)	(None, 8, 8, 192)	0	batch_normalizat...
mixed9 (Concatenate)	(None, 8, 8, 2048)	0	activation_106[0...] mixed9_0[0][0], concatenate[0][0...] activation_114[0...]
conv2d_120 (Conv2D)	(None, 8, 8, 448)	917,504	mixed9[0][0]
batch_normalizatio... (BatchNormalizatio...)	(None, 8, 8, 448)	1,344	conv2d_120[0][0]
activation_119 (Activation)	(None, 8, 8, 448)	0	batch_normalizat...
conv2d_117 (Conv2D)	(None, 8, 8, 384)	786,432	mixed9[0][0]
conv2d_121 (Conv2D)	(None, 8, 8, 384)	1,548,288	activation_119[0...]
batch_normalizatio... (BatchNormalizatio...)	(None, 8, 8, 384)	1,152	conv2d_117[0][0]
batch_normalizatio... (BatchNormalizatio...)	(None, 8, 8, 384)	1,152	conv2d_121[0][0]
activation_116 (Activation)	(None, 8, 8, 384)	0	batch_normalizat...
activation_120 (Activation)	(None, 8, 8, 384)	0	batch_normalizat...
conv2d_118 (Conv2D)	(None, 8, 8, 384)	442,368	activation_116[0...]
conv2d_119 (Conv2D)	(None, 8, 8, 384)	442,368	activation_116[0...]
conv2d_122 (Conv2D)	(None, 8, 8, 384)	442,368	activation_120[0...]
conv2d_123 (Conv2D)	(None, 8, 8, 384)	442,368	activation_120[0...]
average_pooling2d_8 (AveragePooling2D)	(None, 8, 8, 2048)	0	mixed9[0][0]
conv2d_116 (Conv2D)	(None, 8, 8, 320)	655,360	mixed9[0][0]
batch_normalizatio... (BatchNormalizatio...)	(None, 8, 8, 384)	1,152	conv2d_118[0][0]
batch_normalizatio... (BatchNormalizatio...)	(None, 8, 8, 384)	1,152	conv2d_119[0][0]
batch_normalizatio... (BatchNormalizatio...)	(None, 8, 8, 384)	1,152	conv2d_122[0][0]
batch_normalizatio... (BatchNormalizatio...)	(None, 8, 8, 384)	1,152	conv2d_123[0][0]
conv2d_124 (Conv2D)	(None, 8, 8, 192)	393,216	average_pooling2...

batch_normalizatio... (BatchNormalizatio...)	(None, 8, 8, 320)	960	conv2d_116[0][0]
activation_117 (Activation)	(None, 8, 8, 384)	0	batch_normalizat...
activation_118 (Activation)	(None, 8, 8, 384)	0	batch_normalizat...
activation_121 (Activation)	(None, 8, 8, 384)	0	batch_normalizat...
activation_122 (Activation)	(None, 8, 8, 384)	0	batch_normalizat...
batch_normalizatio... (BatchNormalizatio...)	(None, 8, 8, 192)	576	conv2d_124[0][0]
activation_115 (Activation)	(None, 8, 8, 320)	0	batch_normalizat...
mixed9_1 (Concatenate)	(None, 8, 8, 768)	0	activation_117[0]... activation_118[0]...
concatenate_1 (Concatenate)	(None, 8, 8, 768)	0	activation_121[0]... activation_122[0]...
activation_123 (Activation)	(None, 8, 8, 192)	0	batch_normalizat...
mixed10 (Concatenate)	(None, 8, 8, 2048)	0	activation_115[0]... mixed9_1[0][0], concatenate_1[0]... activation_123[0]...
global_average_poo... (GlobalAveragePool...)	(None, 2048)	0	mixed10[0][0]
dropout_6 (Dropout)	(None, 2048)	0	global_average_p...
dense_16 (Dense)	(None, 1024)	2,098,176	dropout_6[0][0]
dense_17 (Dense)	(None, 4)	4,100	dense_16[0][0]

Total params: 28,109,614 (107.23 MB)

Trainable params: 2,102,276 (8.02 MB)

Non-trainable params: 21,802,784 (83.17 MB)

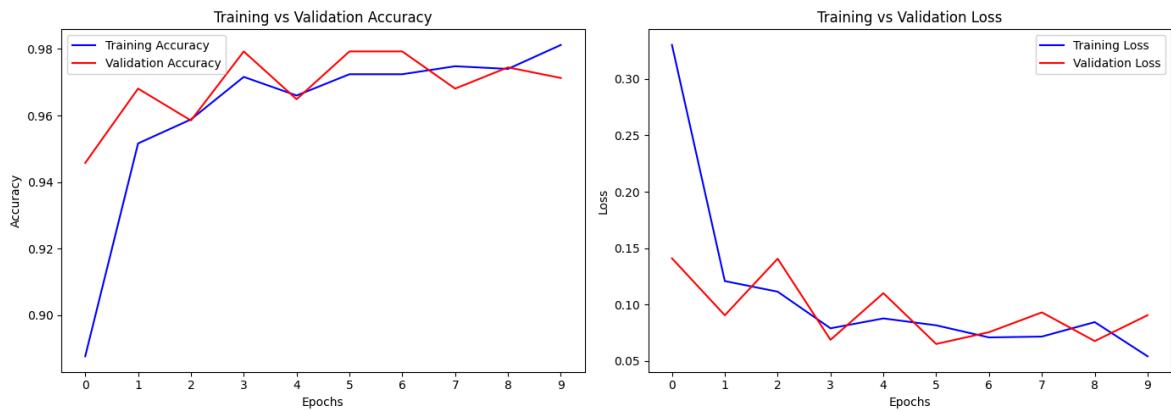
Optimizer params: 4,204,554 (16.04 MB)

```
In [ ]: create_directory_if_not_exists('/content/ninjacart_models/model9_Inception_GAP_D_inception_model.save("/content/ninjacart_models/model9_Inception_GAP_DP_FineTune
```

Created directory: /content/ninjacart\_models/model9\_Inception\_GAP\_DP\_FineTune

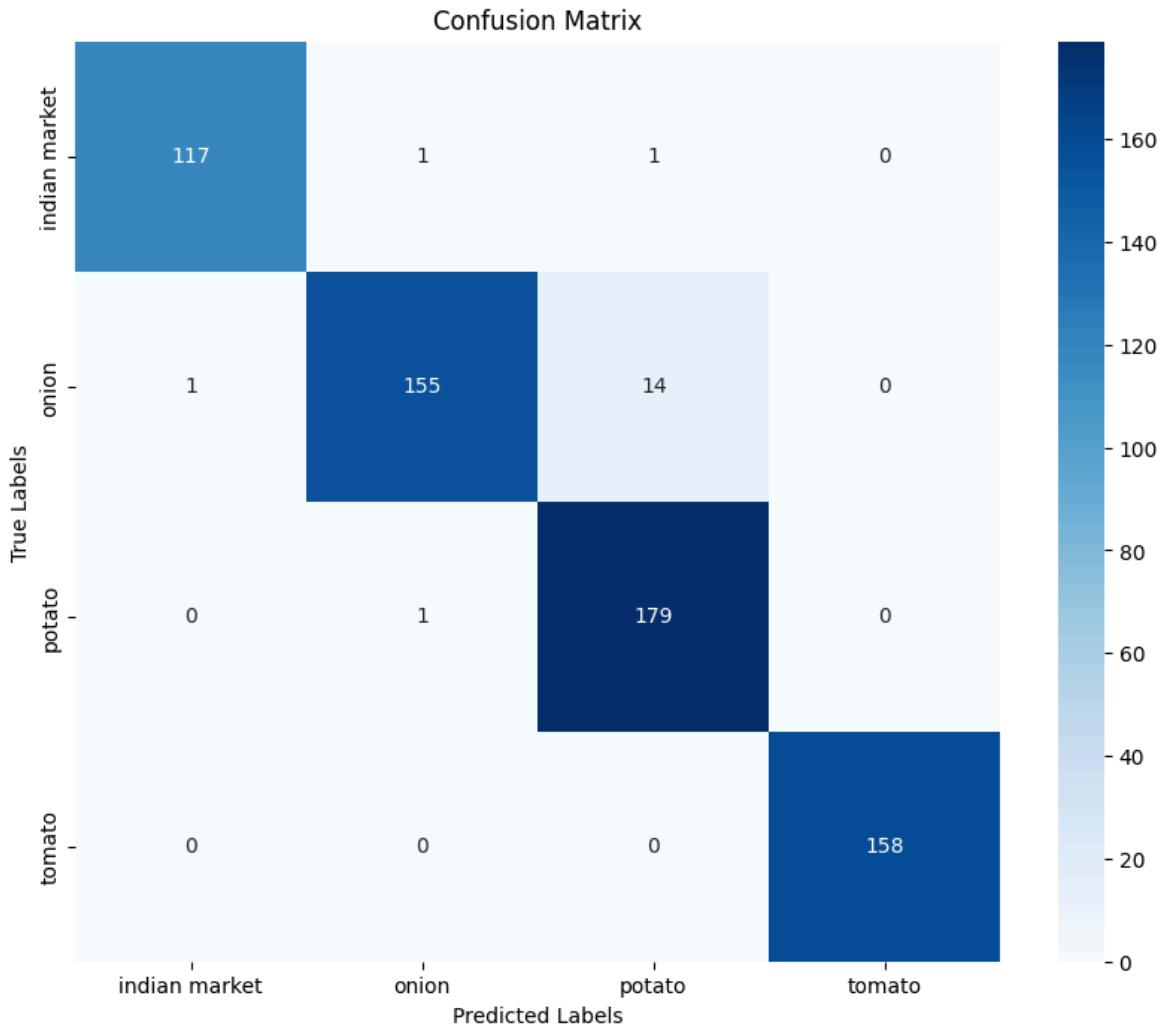
```
In [ ]: plot_training_history(before_ft_inc_hist)
```

History keys: dict\_keys(['accuracy', 'loss', 'val\_accuracy', 'val\_loss'])



```
In [ ]: class_names = list(val_data.class_indices.keys())
plot_confusion_matrix(inception_model, val_data, class_names)
```

20/20 ━━━━━━ 17s 551ms/step



### Classification Report:

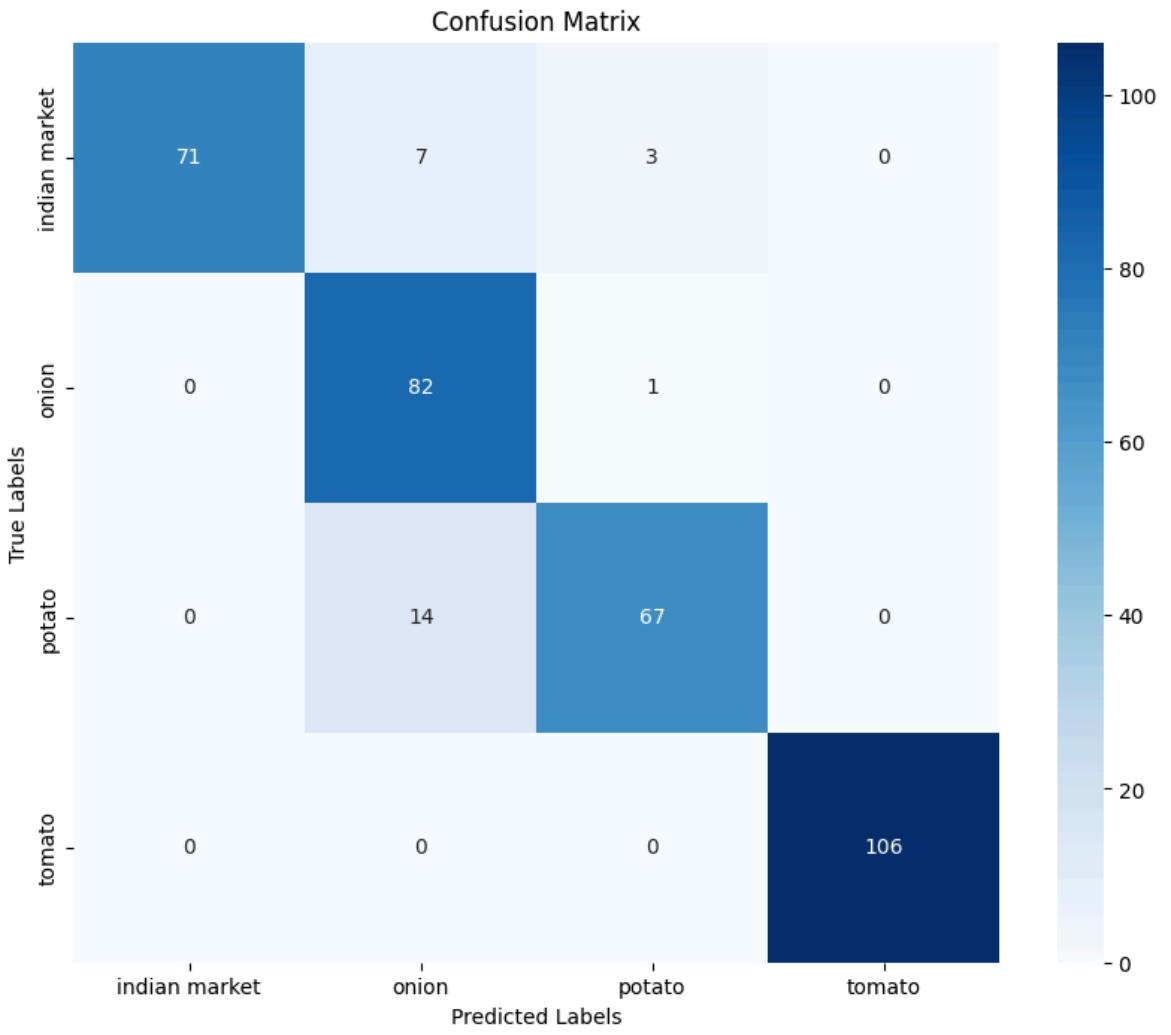
	precision	recall	f1-score	support
indian market	0.99	0.98	0.99	119
onion	0.99	0.91	0.95	170
potato	0.92	0.99	0.96	180
tomato	1.00	1.00	1.00	158
accuracy			0.97	627
macro avg	0.98	0.97	0.97	627
weighted avg	0.97	0.97	0.97	627

```
In [ ]: inception_model_finetune = tf.keras.models.load_model("/content/ninjacart_models")
loss, acc = inception_model_finetune.evaluate(test_data, verbose=2)
print('Before FineTuning')
print("Restored model, accuracy: {:.5.2f}%".format(100 * acc))
```

```
/usr/local/lib/python3.11/dist-packages/keras/src/trainers/data_adapters/py_dataset_adapter.py:121: UserWarning: Your `PyDataset` class should call `super().__init__(**kwargs)` in its constructor. `**kwargs` can include `workers`, `use_multiprocessing`, `max_queue_size`. Do not pass these arguments to `fit()`, as they will be ignored.
    self._warn_if_super_not_called()
11/11 - 20s - 2s/step - accuracy: 0.9288 - loss: 0.2621
Before FineTuning
Restored model, accuracy: 92.88%
```

```
In [ ]: class_names = list(test_data.class_indices.keys())
plot_confusion_matrix(inception_model, test_data, class_names)
```

```
11/11 ━━━━━━━━ 5s 516ms/step
```



Classification Report:

	precision	recall	f1-score	support
indian market	1.00	0.88	0.93	81
onion	0.80	0.99	0.88	83
potato	0.94	0.83	0.88	81
tomato	1.00	1.00	1.00	106
accuracy			0.93	351
macro avg	0.93	0.92	0.92	351
weighted avg	0.94	0.93	0.93	351

## FineTuning

```
In [ ]: # Unfreeze last few layers for fine-tuning
for layer in base_model.layers[-50:]: # You can change -50 to -30 or so
    layer.trainable = True

# Compile again with a lower learning rate
inception_model.compile(optimizer=Adam(learning_rate=1e-5), loss='categorical_crossentropy')

# Continue training (fine-tuning)
after_ft_inc_hist = inception_model.fit(train_data, epochs=10, validation_data=val_data)
```

```
Epoch 1/10
79/79 63s 507ms/step - accuracy: 0.9282 - loss: 0.1738 - val_accuracy: 0.9713 - val_loss: 0.0729
Epoch 2/10
79/79 21s 269ms/step - accuracy: 0.9714 - loss: 0.0755 - val_accuracy: 0.9697 - val_loss: 0.0667
Epoch 3/10
79/79 21s 259ms/step - accuracy: 0.9757 - loss: 0.0686 - val_accuracy: 0.9729 - val_loss: 0.0579
Epoch 4/10
79/79 21s 271ms/step - accuracy: 0.9720 - loss: 0.0711 - val_accuracy: 0.9793 - val_loss: 0.0550
Epoch 5/10
79/79 21s 263ms/step - accuracy: 0.9881 - loss: 0.0325 - val_accuracy: 0.9809 - val_loss: 0.0548
Epoch 6/10
79/79 22s 282ms/step - accuracy: 0.9899 - loss: 0.0286 - val_accuracy: 0.9777 - val_loss: 0.0536
Epoch 7/10
79/79 19s 245ms/step - accuracy: 0.9900 - loss: 0.0297 - val_accuracy: 0.9777 - val_loss: 0.0542
Epoch 8/10
79/79 22s 273ms/step - accuracy: 0.9940 - loss: 0.0176 - val_accuracy: 0.9793 - val_loss: 0.0516
Epoch 9/10
79/79 20s 255ms/step - accuracy: 0.9937 - loss: 0.0199 - val_accuracy: 0.9761 - val_loss: 0.0538
Epoch 10/10
79/79 22s 280ms/step - accuracy: 0.9914 - loss: 0.0216 - val_accuracy: 0.9777 - val_loss: 0.0551
```

```
In [ ]: inception_model.summary()
```

```
Model: "functional_11"
```

Layer (type)	Output Shape	Param #	Connected to
input_layer_11 (InputLayer)	(None, 299, 299, 3)	0	-
conv2d_31 (Conv2D)	(None, 149, 149, 32)	864	input_layer_11[0]...
batch_normalizatio... (BatchNormalizatio...)	(None, 149, 149, 32)	96	conv2d_31[0][0]
activation_30 (Activation)	(None, 149, 149, 32)	0	batch_normalizat...
conv2d_32 (Conv2D)	(None, 147, 147, 32)	9,216	activation_30[0]...
batch_normalizatio... (BatchNormalizatio...)	(None, 147, 147, 32)	96	conv2d_32[0][0]
activation_31 (Activation)	(None, 147, 147, 32)	0	batch_normalizat...
conv2d_33 (Conv2D)	(None, 147, 147, 64)	18,432	activation_31[0]...
batch_normalizatio... (BatchNormalizatio...)	(None, 147, 147, 64)	192	conv2d_33[0][0]
activation_32 (Activation)	(None, 147, 147, 64)	0	batch_normalizat...
max_pooling2d_25 (MaxPooling2D)	(None, 73, 73, 64)	0	activation_32[0]...
conv2d_34 (Conv2D)	(None, 73, 73, 80)	5,120	max_pooling2d_25...
batch_normalizatio... (BatchNormalizatio...)	(None, 73, 73, 80)	240	conv2d_34[0][0]
activation_33 (Activation)	(None, 73, 73, 80)	0	batch_normalizat...
conv2d_35 (Conv2D)	(None, 71, 71, 192)	138,240	activation_33[0]...
batch_normalizatio... (BatchNormalizatio...)	(None, 71, 71, 192)	576	conv2d_35[0][0]
activation_34 (Activation)	(None, 71, 71, 192)	0	batch_normalizat...
max_pooling2d_26 (MaxPooling2D)	(None, 35, 35, 192)	0	activation_34[0]...
conv2d_39 (Conv2D)	(None, 35, 35, 64)	12,288	max_pooling2d_26...

batch_normalization_39 (BatchNormalization)	(None, 35, 35, 64)	192	conv2d_39[0][0]
activation_38 (Activation)	(None, 35, 35, 64)	0	batch_normalizat...
conv2d_37 (Conv2D)	(None, 35, 35, 48)	9,216	max_pooling2d_26...
conv2d_40 (Conv2D)	(None, 35, 35, 96)	55,296	activation_38[0]...
batch_normalization_37 (BatchNormalization)	(None, 35, 35, 48)	144	conv2d_37[0][0]
batch_normalization_38 (BatchNormalization)	(None, 35, 35, 96)	288	conv2d_40[0][0]
activation_36 (Activation)	(None, 35, 35, 48)	0	batch_normalizat...
activation_39 (Activation)	(None, 35, 35, 96)	0	batch_normalizat...
average_pooling2d (AveragePooling2D)	(None, 35, 35, 192)	0	max_pooling2d_26...
conv2d_36 (Conv2D)	(None, 35, 35, 64)	12,288	max_pooling2d_26...
conv2d_38 (Conv2D)	(None, 35, 35, 64)	76,800	activation_36[0]...
conv2d_41 (Conv2D)	(None, 35, 35, 96)	82,944	activation_39[0]...
conv2d_42 (Conv2D)	(None, 35, 35, 32)	6,144	average_pooling2...
batch_normalization_36 (BatchNormalization)	(None, 35, 35, 64)	192	conv2d_36[0][0]
batch_normalization_37 (BatchNormalization)	(None, 35, 35, 64)	192	conv2d_38[0][0]
batch_normalization_38 (BatchNormalization)	(None, 35, 35, 96)	288	conv2d_41[0][0]
batch_normalization_39 (BatchNormalization)	(None, 35, 35, 32)	96	conv2d_42[0][0]
activation_35 (Activation)	(None, 35, 35, 64)	0	batch_normalizat...
activation_37 (Activation)	(None, 35, 35, 64)	0	batch_normalizat...
activation_40 (Activation)	(None, 35, 35, 96)	0	batch_normalizat...

activation_41 (Activation)	(None, 35, 35, 32)	0	batch_normalizat...
mixed0 (Concatenate)	(None, 35, 35, 256)	0	activation_35[0]... activation_37[0]... activation_40[0]... activation_41[0]...
conv2d_46 (Conv2D)	(None, 35, 35, 64)	16,384	mixed0[0][0]
batch_normalizatio... (BatchNormalizatio...)	(None, 35, 35, 64)	192	conv2d_46[0][0]
activation_45 (Activation)	(None, 35, 35, 64)	0	batch_normalizat...
conv2d_44 (Conv2D)	(None, 35, 35, 48)	12,288	mixed0[0][0]
conv2d_47 (Conv2D)	(None, 35, 35, 96)	55,296	activation_45[0]...
batch_normalizatio... (BatchNormalizatio...)	(None, 35, 35, 48)	144	conv2d_44[0][0]
batch_normalizatio... (BatchNormalizatio...)	(None, 35, 35, 96)	288	conv2d_47[0][0]
activation_43 (Activation)	(None, 35, 35, 48)	0	batch_normalizat...
activation_46 (Activation)	(None, 35, 35, 96)	0	batch_normalizat...
average_pooling2d_1 (AveragePooling2D)	(None, 35, 35, 256)	0	mixed0[0][0]
conv2d_43 (Conv2D)	(None, 35, 35, 64)	16,384	mixed0[0][0]
conv2d_45 (Conv2D)	(None, 35, 35, 64)	76,800	activation_43[0]...
conv2d_48 (Conv2D)	(None, 35, 35, 96)	82,944	activation_46[0]...
conv2d_49 (Conv2D)	(None, 35, 35, 64)	16,384	average_pooling2...
batch_normalizatio... (BatchNormalizatio...)	(None, 35, 35, 64)	192	conv2d_43[0][0]
batch_normalizatio... (BatchNormalizatio...)	(None, 35, 35, 64)	192	conv2d_45[0][0]
batch_normalizatio... (BatchNormalizatio...)	(None, 35, 35, 96)	288	conv2d_48[0][0]
batch_normalizatio... (BatchNormalizatio...)	(None, 35, 35,	192	conv2d_49[0][0]

(BatchNormalizatio...	64)		
activation_42 (Activation)	(None, 35, 35, 64)	0	batch_normalizat...
activation_44 (Activation)	(None, 35, 35, 64)	0	batch_normalizat...
activation_47 (Activation)	(None, 35, 35, 96)	0	batch_normalizat...
activation_48 (Activation)	(None, 35, 35, 64)	0	batch_normalizat...
mixed1 (Concatenate)	(None, 35, 35, 288)	0	activation_42[0]... activation_44[0]... activation_47[0]... activation_48[0]...
conv2d_53 (Conv2D)	(None, 35, 35, 64)	18,432	mixed1[0][0]
batch_normalizatio... (BatchNormalizatio...	(None, 35, 35, 64)	192	conv2d_53[0][0]
activation_52 (Activation)	(None, 35, 35, 64)	0	batch_normalizat...
conv2d_51 (Conv2D)	(None, 35, 35, 48)	13,824	mixed1[0][0]
conv2d_54 (Conv2D)	(None, 35, 35, 96)	55,296	activation_52[0]...
batch_normalizatio... (BatchNormalizatio...	(None, 35, 35, 48)	144	conv2d_51[0][0]
batch_normalizatio... (BatchNormalizatio...	(None, 35, 35, 96)	288	conv2d_54[0][0]
activation_50 (Activation)	(None, 35, 35, 48)	0	batch_normalizat...
activation_53 (Activation)	(None, 35, 35, 96)	0	batch_normalizat...
average_pooling2d_2 (AveragePooling2D)	(None, 35, 35, 288)	0	mixed1[0][0]
conv2d_50 (Conv2D)	(None, 35, 35, 64)	18,432	mixed1[0][0]
conv2d_52 (Conv2D)	(None, 35, 35, 64)	76,800	activation_50[0]...
conv2d_55 (Conv2D)	(None, 35, 35, 96)	82,944	activation_53[0]...
conv2d_56 (Conv2D)	(None, 35, 35, 64)	18,432	average_pooling2...

batch_normalizatio... (BatchNormalizatio...)	(None, 35, 35, 64)	192	conv2d_50[0][0]
batch_normalizatio... (BatchNormalizatio...)	(None, 35, 35, 64)	192	conv2d_52[0][0]
batch_normalizatio... (BatchNormalizatio...)	(None, 35, 35, 96)	288	conv2d_55[0][0]
batch_normalizatio... (BatchNormalizatio...)	(None, 35, 35, 64)	192	conv2d_56[0][0]
activation_49 (Activation)	(None, 35, 35, 64)	0	batch_normalizat...
activation_51 (Activation)	(None, 35, 35, 64)	0	batch_normalizat...
activation_54 (Activation)	(None, 35, 35, 96)	0	batch_normalizat...
activation_55 (Activation)	(None, 35, 35, 64)	0	batch_normalizat...
mixed2 (Concatenate)	(None, 35, 35, 288)	0	activation_49[0]... activation_51[0]... activation_54[0]... activation_55[0]...
conv2d_58 (Conv2D)	(None, 35, 35, 64)	18,432	mixed2[0][0]
batch_normalizatio... (BatchNormalizatio...)	(None, 35, 35, 64)	192	conv2d_58[0][0]
activation_57 (Activation)	(None, 35, 35, 64)	0	batch_normalizat...
conv2d_59 (Conv2D)	(None, 35, 35, 96)	55,296	activation_57[0]...
batch_normalizatio... (BatchNormalizatio...)	(None, 35, 35, 96)	288	conv2d_59[0][0]
activation_58 (Activation)	(None, 35, 35, 96)	0	batch_normalizat...
conv2d_57 (Conv2D)	(None, 17, 17, 384)	995,328	mixed2[0][0]
conv2d_60 (Conv2D)	(None, 17, 17, 96)	82,944	activation_58[0]...
batch_normalizatio... (BatchNormalizatio...)	(None, 17, 17, 384)	1,152	conv2d_57[0][0]
batch_normalizatio... (BatchNormalizatio...)	(None, 17, 17, 96)	288	conv2d_60[0][0]

activation_56 (Activation)	(None, 17, 17, 384)	0	batch_normalizat...
activation_59 (Activation)	(None, 17, 17, 96)	0	batch_normalizat...
max_pooling2d_27 (MaxPooling2D)	(None, 17, 17, 288)	0	mixed2[0][0]
mixed3 (Concatenate)	(None, 17, 17, 768)	0	activation_56[0]... activation_59[0]... max_pooling2d_27...
conv2d_65 (Conv2D)	(None, 17, 17, 128)	98,304	mixed3[0][0]
batch_normalizatio... (BatchNormalizatio...)	(None, 17, 17, 128)	384	conv2d_65[0][0]
activation_64 (Activation)	(None, 17, 17, 128)	0	batch_normalizat...
conv2d_66 (Conv2D)	(None, 17, 17, 128)	114,688	activation_64[0]...
batch_normalizatio... (BatchNormalizatio...)	(None, 17, 17, 128)	384	conv2d_66[0][0]
activation_65 (Activation)	(None, 17, 17, 128)	0	batch_normalizat...
conv2d_62 (Conv2D)	(None, 17, 17, 128)	98,304	mixed3[0][0]
conv2d_67 (Conv2D)	(None, 17, 17, 128)	114,688	activation_65[0]...
batch_normalizatio... (BatchNormalizatio...)	(None, 17, 17, 128)	384	conv2d_62[0][0]
batch_normalizatio... (BatchNormalizatio...)	(None, 17, 17, 128)	384	conv2d_67[0][0]
activation_61 (Activation)	(None, 17, 17, 128)	0	batch_normalizat...
activation_66 (Activation)	(None, 17, 17, 128)	0	batch_normalizat...
conv2d_63 (Conv2D)	(None, 17, 17, 128)	114,688	activation_61[0]...
conv2d_68 (Conv2D)	(None, 17, 17, 128)	114,688	activation_66[0]...
batch_normalizatio... (BatchNormalizatio...)	(None, 17, 17, 128)	384	conv2d_63[0][0]
batch_normalizatio... (BatchNormalizatio...)	(None, 17, 17, 128)	384	conv2d_68[0][0]

activation_62 (Activation)	(None, 17, 17, 128)	0	batch_normalizat...
activation_67 (Activation)	(None, 17, 17, 128)	0	batch_normalizat...
average_pooling2d_3 (AveragePooling2D)	(None, 17, 17, 768)	0	mixed3[0][0]
conv2d_61 (Conv2D)	(None, 17, 17, 192)	147,456	mixed3[0][0]
conv2d_64 (Conv2D)	(None, 17, 17, 192)	172,032	activation_62[0]...
conv2d_69 (Conv2D)	(None, 17, 17, 192)	172,032	activation_67[0]...
conv2d_70 (Conv2D)	(None, 17, 17, 192)	147,456	average_pooling2...
batch_normalizatio... (BatchNormalizatio...)	(None, 17, 17, 192)	576	conv2d_61[0][0]
batch_normalizatio... (BatchNormalizatio...)	(None, 17, 17, 192)	576	conv2d_64[0][0]
batch_normalizatio... (BatchNormalizatio...)	(None, 17, 17, 192)	576	conv2d_69[0][0]
batch_normalizatio... (BatchNormalizatio...)	(None, 17, 17, 192)	576	conv2d_70[0][0]
activation_60 (Activation)	(None, 17, 17, 192)	0	batch_normalizat...
activation_63 (Activation)	(None, 17, 17, 192)	0	batch_normalizat...
activation_68 (Activation)	(None, 17, 17, 192)	0	batch_normalizat...
activation_69 (Activation)	(None, 17, 17, 192)	0	batch_normalizat...
mixed4 (Concatenate)	(None, 17, 17, 768)	0	activation_60[0]... activation_63[0]... activation_68[0]... activation_69[0]...
conv2d_75 (Conv2D)	(None, 17, 17, 160)	122,880	mixed4[0][0]
batch_normalizatio... (BatchNormalizatio...)	(None, 17, 17, 160)	480	conv2d_75[0][0]
activation_74 (Activation)	(None, 17, 17, 160)	0	batch_normalizat...

conv2d_76 (Conv2D)	(None, 17, 17, 160)	179,200	activation_74[0]...
batch_normalizatio... (BatchNormalizatio...)	(None, 17, 17, 160)	480	conv2d_76[0][0]
activation_75 (Activation)	(None, 17, 17, 160)	0	batch_normalizat...
conv2d_72 (Conv2D)	(None, 17, 17, 160)	122,880	mixed4[0][0]
conv2d_77 (Conv2D)	(None, 17, 17, 160)	179,200	activation_75[0]...
batch_normalizatio... (BatchNormalizatio...)	(None, 17, 17, 160)	480	conv2d_72[0][0]
batch_normalizatio... (BatchNormalizatio...)	(None, 17, 17, 160)	480	conv2d_77[0][0]
activation_71 (Activation)	(None, 17, 17, 160)	0	batch_normalizat...
activation_76 (Activation)	(None, 17, 17, 160)	0	batch_normalizat...
conv2d_73 (Conv2D)	(None, 17, 17, 160)	179,200	activation_71[0]...
conv2d_78 (Conv2D)	(None, 17, 17, 160)	179,200	activation_76[0]...
batch_normalizatio... (BatchNormalizatio...)	(None, 17, 17, 160)	480	conv2d_73[0][0]
batch_normalizatio... (BatchNormalizatio...)	(None, 17, 17, 160)	480	conv2d_78[0][0]
activation_72 (Activation)	(None, 17, 17, 160)	0	batch_normalizat...
activation_77 (Activation)	(None, 17, 17, 160)	0	batch_normalizat...
average_pooling2d_4 (AveragePooling2D)	(None, 17, 17, 768)	0	mixed4[0][0]
conv2d_71 (Conv2D)	(None, 17, 17, 192)	147,456	mixed4[0][0]
conv2d_74 (Conv2D)	(None, 17, 17, 192)	215,040	activation_72[0]...
conv2d_79 (Conv2D)	(None, 17, 17, 192)	215,040	activation_77[0]...
conv2d_80 (Conv2D)	(None, 17, 17, 192)	147,456	average_pooling2...

batch_normalization_71 (BatchNormalization)	(None, 17, 17, 192)	576	conv2d_71[0][0]
batch_normalization_74 (BatchNormalization)	(None, 17, 17, 192)	576	conv2d_74[0][0]
batch_normalization_79 (BatchNormalization)	(None, 17, 17, 192)	576	conv2d_79[0][0]
batch_normalization_80 (BatchNormalization)	(None, 17, 17, 192)	576	conv2d_80[0][0]
activation_70 (Activation)	(None, 17, 17, 192)	0	batch_normalizat...
activation_73 (Activation)	(None, 17, 17, 192)	0	batch_normalizat...
activation_78 (Activation)	(None, 17, 17, 192)	0	batch_normalizat...
activation_79 (Activation)	(None, 17, 17, 192)	0	batch_normalizat...
mixed5 (Concatenate)	(None, 17, 17, 768)	0	activation_70[0]... activation_73[0]... activation_78[0]... activation_79[0]...
conv2d_85 (Conv2D)	(None, 17, 17, 160)	122,880	mixed5[0][0]
batch_normalization_85 (BatchNormalization)	(None, 17, 17, 160)	480	conv2d_85[0][0]
activation_84 (Activation)	(None, 17, 17, 160)	0	batch_normalizat...
conv2d_86 (Conv2D)	(None, 17, 17, 160)	179,200	activation_84[0]...
batch_normalization_86 (BatchNormalization)	(None, 17, 17, 160)	480	conv2d_86[0][0]
activation_85 (Activation)	(None, 17, 17, 160)	0	batch_normalizat...
conv2d_82 (Conv2D)	(None, 17, 17, 160)	122,880	mixed5[0][0]
conv2d_87 (Conv2D)	(None, 17, 17, 160)	179,200	activation_85[0]...
batch_normalization_82 (BatchNormalization)	(None, 17, 17, 160)	480	conv2d_82[0][0]
batch_normalization_87 (BatchNormalization)	(None, 17, 17, 160)	480	conv2d_87[0][0]
activation_81	(None, 17, 17,	0	batch_normalizat...

(Activation)	160)		
activation_86 (Activation)	(None, 17, 17, 160)	0	batch_normalizat...
conv2d_83 (Conv2D)	(None, 17, 17, 160)	179,200	activation_81[0]...
conv2d_88 (Conv2D)	(None, 17, 17, 160)	179,200	activation_86[0]...
batch_normalizatio... (BatchNormalizatio...)	(None, 17, 17, 160)	480	conv2d_83[0][0]
batch_normalizatio... (BatchNormalizatio...)	(None, 17, 17, 160)	480	conv2d_88[0][0]
activation_82 (Activation)	(None, 17, 17, 160)	0	batch_normalizat...
activation_87 (Activation)	(None, 17, 17, 160)	0	batch_normalizat...
average_pooling2d_5 (AveragePooling2D)	(None, 17, 17, 768)	0	mixed5[0][0]
conv2d_81 (Conv2D)	(None, 17, 17, 192)	147,456	mixed5[0][0]
conv2d_84 (Conv2D)	(None, 17, 17, 192)	215,040	activation_82[0]...
conv2d_89 (Conv2D)	(None, 17, 17, 192)	215,040	activation_87[0]...
conv2d_90 (Conv2D)	(None, 17, 17, 192)	147,456	average_pooling2...
batch_normalizatio... (BatchNormalizatio...)	(None, 17, 17, 192)	576	conv2d_81[0][0]
batch_normalizatio... (BatchNormalizatio...)	(None, 17, 17, 192)	576	conv2d_84[0][0]
batch_normalizatio... (BatchNormalizatio...)	(None, 17, 17, 192)	576	conv2d_89[0][0]
batch_normalizatio... (BatchNormalizatio...)	(None, 17, 17, 192)	576	conv2d_90[0][0]
activation_80 (Activation)	(None, 17, 17, 192)	0	batch_normalizat...
activation_83 (Activation)	(None, 17, 17, 192)	0	batch_normalizat...
activation_88 (Activation)	(None, 17, 17, 192)	0	batch_normalizat...
activation_89	(None, 17, 17,	0	batch_normalizat...

(Activation)	192)		
mixed6 (Concatenate)	(None, 17, 17, 768)	0	activation_80[0]... activation_83[0]... activation_88[0]... activation_89[0]...
conv2d_95 (Conv2D)	(None, 17, 17, 192)	147,456	mixed6[0][0]
batch_normalizatio... (BatchNormalizatio...)	(None, 17, 17, 192)	576	conv2d_95[0][0]
activation_94 (Activation)	(None, 17, 17, 192)	0	batch_normalizat...
conv2d_96 (Conv2D)	(None, 17, 17, 192)	258,048	activation_94[0]...
batch_normalizatio... (BatchNormalizatio...)	(None, 17, 17, 192)	576	conv2d_96[0][0]
activation_95 (Activation)	(None, 17, 17, 192)	0	batch_normalizat...
conv2d_92 (Conv2D)	(None, 17, 17, 192)	147,456	mixed6[0][0]
conv2d_97 (Conv2D)	(None, 17, 17, 192)	258,048	activation_95[0]...
batch_normalizatio... (BatchNormalizatio...)	(None, 17, 17, 192)	576	conv2d_92[0][0]
batch_normalizatio... (BatchNormalizatio...)	(None, 17, 17, 192)	576	conv2d_97[0][0]
activation_91 (Activation)	(None, 17, 17, 192)	0	batch_normalizat...
activation_96 (Activation)	(None, 17, 17, 192)	0	batch_normalizat...
conv2d_93 (Conv2D)	(None, 17, 17, 192)	258,048	activation_91[0]...
conv2d_98 (Conv2D)	(None, 17, 17, 192)	258,048	activation_96[0]...
batch_normalizatio... (BatchNormalizatio...)	(None, 17, 17, 192)	576	conv2d_93[0][0]
batch_normalizatio... (BatchNormalizatio...)	(None, 17, 17, 192)	576	conv2d_98[0][0]
activation_92 (Activation)	(None, 17, 17, 192)	0	batch_normalizat...
activation_97 (Activation)	(None, 17, 17, 192)	0	batch_normalizat...

average_pooling2d_6 (AveragePooling2D)	(None, 17, 17, 768)	0	mixed6[0][0]
conv2d_91 (Conv2D)	(None, 17, 17, 192)	147,456	mixed6[0][0]
conv2d_94 (Conv2D)	(None, 17, 17, 192)	258,048	activation_92[0]...
conv2d_99 (Conv2D)	(None, 17, 17, 192)	258,048	activation_97[0]...
conv2d_100 (Conv2D)	(None, 17, 17, 192)	147,456	average_pooling2...
batch_normalizatio... (BatchNormalizatio...)	(None, 17, 17, 192)	576	conv2d_91[0][0]
batch_normalizatio... (BatchNormalizatio...)	(None, 17, 17, 192)	576	conv2d_94[0][0]
batch_normalizatio... (BatchNormalizatio...)	(None, 17, 17, 192)	576	conv2d_99[0][0]
batch_normalizatio... (BatchNormalizatio...)	(None, 17, 17, 192)	576	conv2d_100[0][0]
activation_90 (Activation)	(None, 17, 17, 192)	0	batch_normalizat...
activation_93 (Activation)	(None, 17, 17, 192)	0	batch_normalizat...
activation_98 (Activation)	(None, 17, 17, 192)	0	batch_normalizat...
activation_99 (Activation)	(None, 17, 17, 192)	0	batch_normalizat...
mixed7 (Concatenate)	(None, 17, 17, 768)	0	activation_90[0]... activation_93[0]... activation_98[0]... activation_99[0]...
conv2d_103 (Conv2D)	(None, 17, 17, 192)	147,456	mixed7[0][0]
batch_normalizatio... (BatchNormalizatio...)	(None, 17, 17, 192)	576	conv2d_103[0][0]
activation_102 (Activation)	(None, 17, 17, 192)	0	batch_normalizat...
conv2d_104 (Conv2D)	(None, 17, 17, 192)	258,048	activation_102[0]...
batch_normalizatio... (BatchNormalizatio...)	(None, 17, 17, 192)	576	conv2d_104[0][0]

activation_103 (Activation)	(None, 17, 17, 192)	0	batch_normalizat...
conv2d_101 (Conv2D)	(None, 17, 17, 192)	147,456	mixed7[0][0]
conv2d_105 (Conv2D)	(None, 17, 17, 192)	258,048	activation_103[0...]
batch_normalizatio... (BatchNormalizatio...)	(None, 17, 17, 192)	576	conv2d_101[0][0]
batch_normalizatio... (BatchNormalizatio...)	(None, 17, 17, 192)	576	conv2d_105[0][0]
activation_100 (Activation)	(None, 17, 17, 192)	0	batch_normalizat...
activation_104 (Activation)	(None, 17, 17, 192)	0	batch_normalizat...
conv2d_102 (Conv2D)	(None, 8, 8, 320)	552,960	activation_100[0...]
conv2d_106 (Conv2D)	(None, 8, 8, 192)	331,776	activation_104[0...]
batch_normalizatio... (BatchNormalizatio...)	(None, 8, 8, 320)	960	conv2d_102[0][0]
batch_normalizatio... (BatchNormalizatio...)	(None, 8, 8, 192)	576	conv2d_106[0][0]
activation_101 (Activation)	(None, 8, 8, 320)	0	batch_normalizat...
activation_105 (Activation)	(None, 8, 8, 192)	0	batch_normalizat...
max_pooling2d_28 (MaxPooling2D)	(None, 8, 8, 768)	0	mixed7[0][0]
mixed8 (Concatenate)	(None, 8, 8, 1280)	0	activation_101[0...] activation_105[0...] max_pooling2d_28[0]
conv2d_111 (Conv2D)	(None, 8, 8, 448)	573,440	mixed8[0][0]
batch_normalizatio... (BatchNormalizatio...)	(None, 8, 8, 448)	1,344	conv2d_111[0][0]
activation_110 (Activation)	(None, 8, 8, 448)	0	batch_normalizat...
conv2d_108 (Conv2D)	(None, 8, 8, 384)	491,520	mixed8[0][0]
conv2d_112 (Conv2D)	(None, 8, 8, 384)	1,548,288	activation_110[0...]
batch_normalizatio... (BatchNormalizatio...)	(None, 8, 8, 384)	1,152	conv2d_108[0][0]
batch_normalizatio... (BatchNormalizatio...)	(None, 8, 8, 384)	1,152	conv2d_112[0][0]

(BatchNormalizatio...			
activation_107 (Activation)	(None, 8, 8, 384)	0	batch_normalizat...
activation_111 (Activation)	(None, 8, 8, 384)	0	batch_normalizat...
conv2d_109 (Conv2D)	(None, 8, 8, 384)	442,368	activation_107[0...]
conv2d_110 (Conv2D)	(None, 8, 8, 384)	442,368	activation_107[0...]
conv2d_113 (Conv2D)	(None, 8, 8, 384)	442,368	activation_111[0...]
conv2d_114 (Conv2D)	(None, 8, 8, 384)	442,368	activation_111[0...]
average_pooling2d_7 (AveragePooling2D)	(None, 8, 8, 1280)	0	mixed8[0][0]
conv2d_107 (Conv2D)	(None, 8, 8, 320)	409,600	mixed8[0][0]
batch_normalizatio... (BatchNormalizatio...	(None, 8, 8, 384)	1,152	conv2d_109[0][0]
batch_normalizatio... (BatchNormalizatio...	(None, 8, 8, 384)	1,152	conv2d_110[0][0]
batch_normalizatio... (BatchNormalizatio...	(None, 8, 8, 384)	1,152	conv2d_113[0][0]
batch_normalizatio... (BatchNormalizatio...	(None, 8, 8, 384)	1,152	conv2d_114[0][0]
conv2d_115 (Conv2D)	(None, 8, 8, 192)	245,760	average_pooling2...
batch_normalizatio... (BatchNormalizatio...	(None, 8, 8, 320)	960	conv2d_107[0][0]
activation_108 (Activation)	(None, 8, 8, 384)	0	batch_normalizat...
activation_109 (Activation)	(None, 8, 8, 384)	0	batch_normalizat...
activation_112 (Activation)	(None, 8, 8, 384)	0	batch_normalizat...
activation_113 (Activation)	(None, 8, 8, 384)	0	batch_normalizat...
batch_normalizatio... (BatchNormalizatio...	(None, 8, 8, 192)	576	conv2d_115[0][0]
activation_106 (Activation)	(None, 8, 8, 320)	0	batch_normalizat...
mixed9_0 (Concatenate)	(None, 8, 8, 768)	0	activation_108[0...] activation_109[0...]
concatenate	(None, 8, 8, 768)	0	activation_112[0...]

(Concatenate)			activation_113[0...]
activation_114 (Activation)	(None, 8, 8, 192)	0	batch_normalizat...
mixed9 (Concatenate)	(None, 8, 8, 2048)	0	activation_106[0...] mixed9_0[0][0], concatenate[0][0...] activation_114[0...]
conv2d_120 (Conv2D)	(None, 8, 8, 448)	917,504	mixed9[0][0]
batch_normalizatio... (BatchNormalizatio...)	(None, 8, 8, 448)	1,344	conv2d_120[0][0]
activation_119 (Activation)	(None, 8, 8, 448)	0	batch_normalizat...
conv2d_117 (Conv2D)	(None, 8, 8, 384)	786,432	mixed9[0][0]
conv2d_121 (Conv2D)	(None, 8, 8, 384)	1,548,288	activation_119[0...]
batch_normalizatio... (BatchNormalizatio...)	(None, 8, 8, 384)	1,152	conv2d_117[0][0]
batch_normalizatio... (BatchNormalizatio...)	(None, 8, 8, 384)	1,152	conv2d_121[0][0]
activation_116 (Activation)	(None, 8, 8, 384)	0	batch_normalizat...
activation_120 (Activation)	(None, 8, 8, 384)	0	batch_normalizat...
conv2d_118 (Conv2D)	(None, 8, 8, 384)	442,368	activation_116[0...]
conv2d_119 (Conv2D)	(None, 8, 8, 384)	442,368	activation_116[0...]
conv2d_122 (Conv2D)	(None, 8, 8, 384)	442,368	activation_120[0...]
conv2d_123 (Conv2D)	(None, 8, 8, 384)	442,368	activation_120[0...]
average_pooling2d_8 (AveragePooling2D)	(None, 8, 8, 2048)	0	mixed9[0][0]
conv2d_116 (Conv2D)	(None, 8, 8, 320)	655,360	mixed9[0][0]
batch_normalizatio... (BatchNormalizatio...)	(None, 8, 8, 384)	1,152	conv2d_118[0][0]
batch_normalizatio... (BatchNormalizatio...)	(None, 8, 8, 384)	1,152	conv2d_119[0][0]
batch_normalizatio... (BatchNormalizatio...)	(None, 8, 8, 384)	1,152	conv2d_122[0][0]
batch_normalizatio... (BatchNormalizatio...)	(None, 8, 8, 384)	1,152	conv2d_123[0][0]
conv2d_124 (Conv2D)	(None, 8, 8, 192)	393,216	average_pooling2...

batch_normalizatio... (BatchNormalizatio...)	(None, 8, 8, 320)	960	conv2d_116[0][0]
activation_117 (Activation)	(None, 8, 8, 384)	0	batch_normalizat...
activation_118 (Activation)	(None, 8, 8, 384)	0	batch_normalizat...
activation_121 (Activation)	(None, 8, 8, 384)	0	batch_normalizat...
activation_122 (Activation)	(None, 8, 8, 384)	0	batch_normalizat...
batch_normalizatio... (BatchNormalizatio...)	(None, 8, 8, 192)	576	conv2d_124[0][0]
activation_115 (Activation)	(None, 8, 8, 320)	0	batch_normalizat...
mixed9_1 (Concatenate)	(None, 8, 8, 768)	0	activation_117[0]... activation_118[0]...
concatenate_1 (Concatenate)	(None, 8, 8, 768)	0	activation_121[0]... activation_122[0]...
activation_123 (Activation)	(None, 8, 8, 192)	0	batch_normalizat...
mixed10 (Concatenate)	(None, 8, 8, 2048)	0	activation_115[0]... mixed9_1[0][0], concatenate_1[0]... activation_123[0]...
global_average_poo... (GlobalAveragePool...)	(None, 2048)	0	mixed10[0][0]
dropout_6 (Dropout)	(None, 2048)	0	global_average_p...
dense_16 (Dense)	(None, 1024)	2,098,176	dropout_6[0][0]
dense_17 (Dense)	(None, 4)	4,100	dense_16[0][0]

Total params: 42,456,238 (161.96 MB)

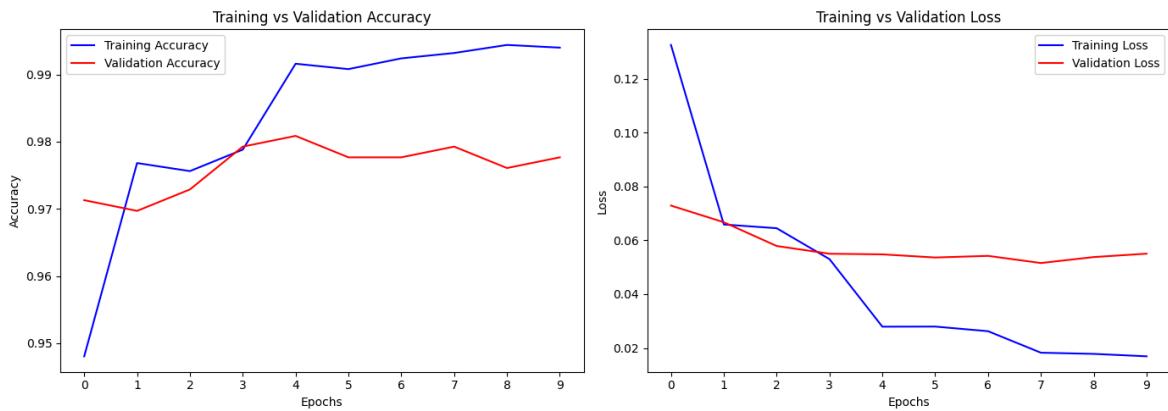
Trainable params: 9,275,588 (35.38 MB)

Non-trainable params: 14,629,472 (55.81 MB)

Optimizer params: 18,551,178 (70.77 MB)

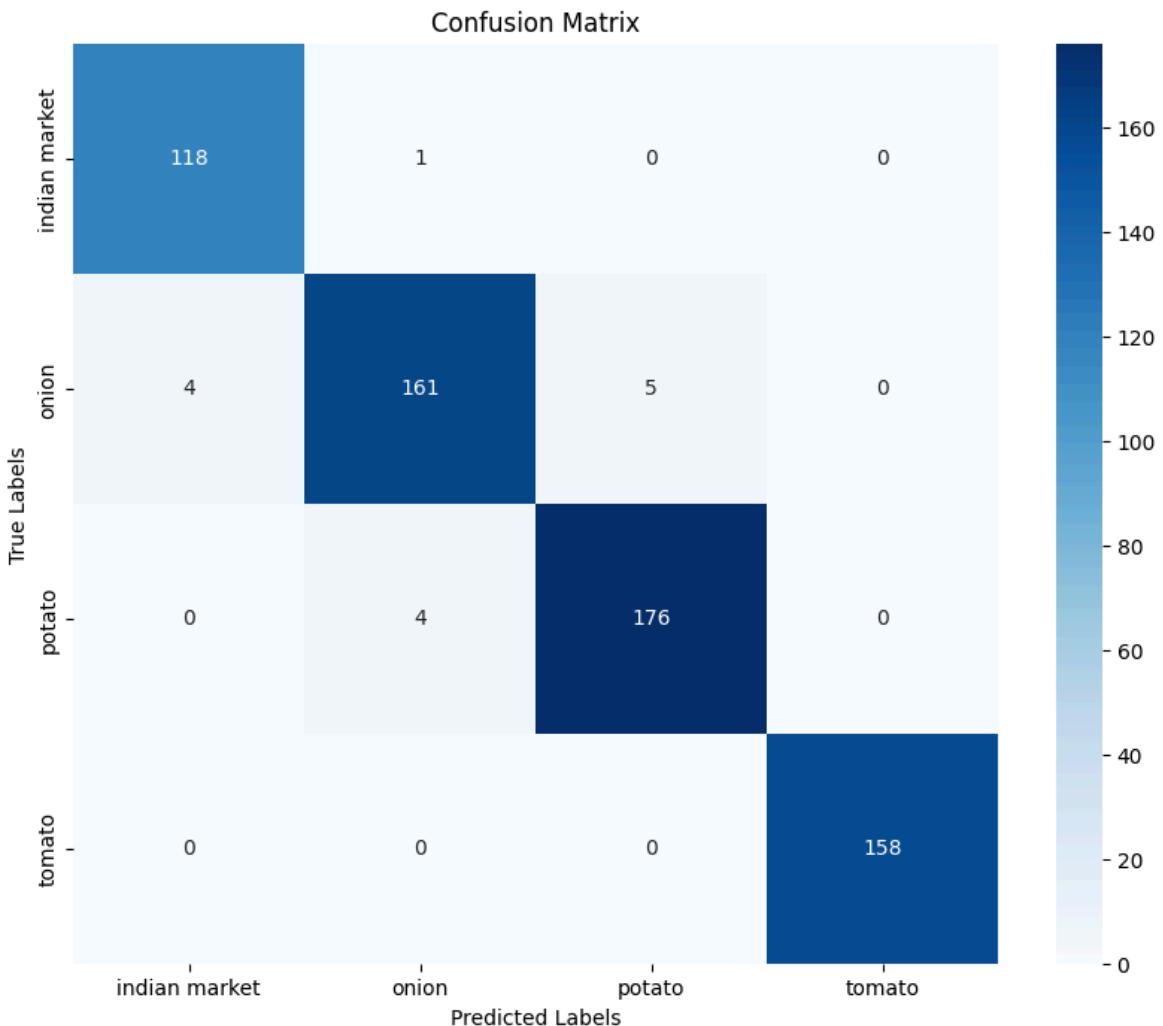
In [ ]: `plot_training_history(after_ft_inc_hist)`

History keys: dict\_keys(['accuracy', 'loss', 'val\_accuracy', 'val\_loss'])



```
In [ ]: class_names = list(val_data.class_indices.keys())
plot_confusion_matrix(inception_model, val_data, class_names)
```

20/20 ━━━━━━ 21s 642ms/step



### Classification Report:

	precision	recall	f1-score	support
indian market	0.97	0.99	0.98	119
onion	0.97	0.95	0.96	170
potato	0.97	0.98	0.98	180
tomato	1.00	1.00	1.00	158
accuracy			0.98	627
macro avg	0.98	0.98	0.98	627
weighted avg	0.98	0.98	0.98	627

```
In [ ]: inception_model.save("/content/ninjacart_models/model9_Inception_GAP_DP_FineTune")
```

```
In [ ]: inception_model_finetune = tf.keras.models.load_model("/content/ninjacart_models/model9_Inception_GAP_DP_FineTune")
loss, acc = inception_model_finetune.evaluate(test_data, verbose=2)
print('After FineTuning')
print("Restored model, accuracy: {:.2f}%".format(100 * acc))
```

11/11 - 14s - 1s/step - accuracy: 0.9316 - loss: 0.2883

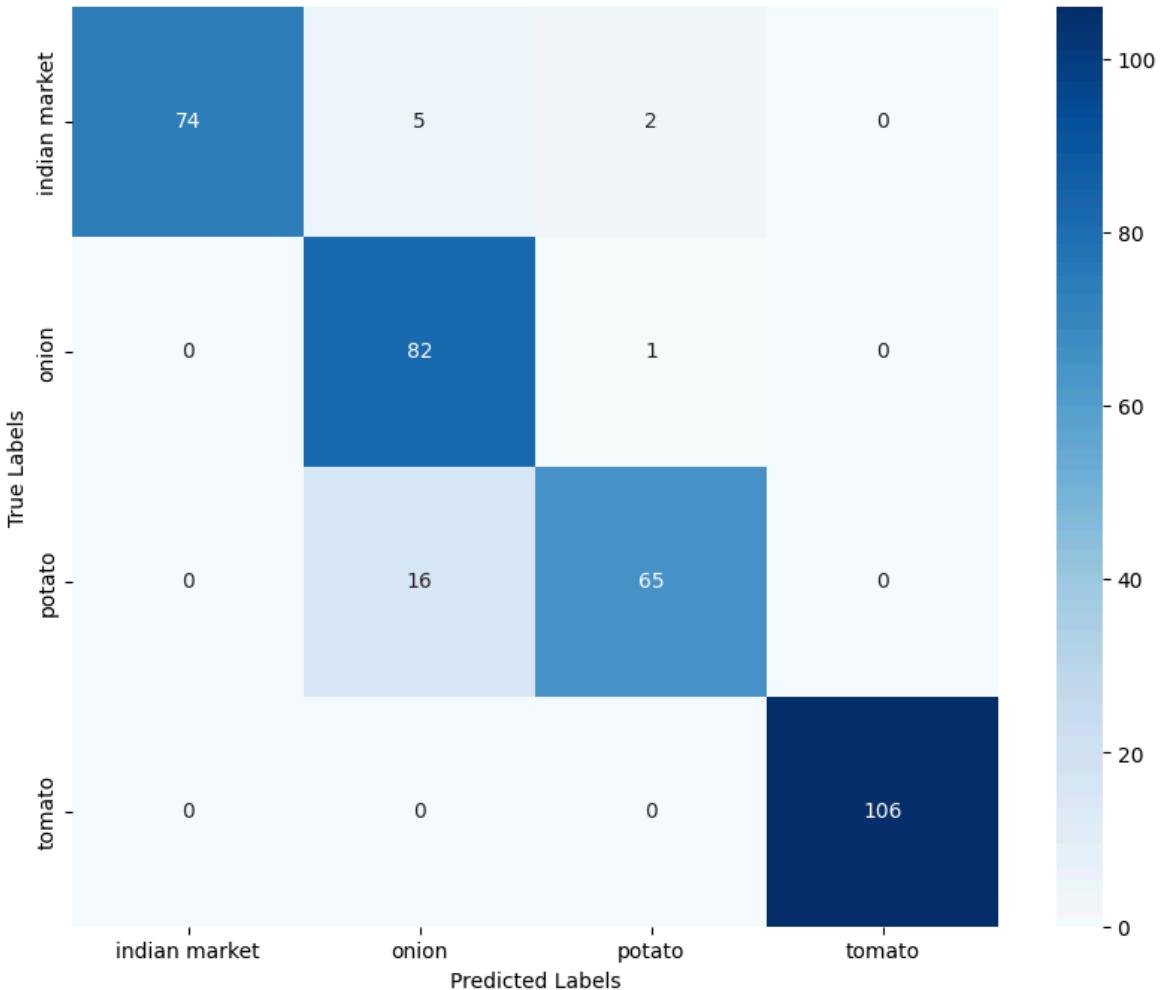
After FineTuning

Restored model, accuracy: 93.16%

```
In [ ]: class_names = list(test_data.class_indices.keys())
plot_confusion_matrix(inception_model, test_data, class_names)
```

11/11 ━━━━━━━━ 5s 509ms/step

Confusion Matrix



Classification Report:

	precision	recall	f1-score	support
indian market	1.00	0.91	0.95	81
onion	0.80	0.99	0.88	83
potato	0.96	0.80	0.87	81
tomato	1.00	1.00	1.00	106
accuracy			0.93	351
macro avg	0.94	0.93	0.93	351
weighted avg	0.94	0.93	0.93	351

## Resnet50

```
In [ ]: datagen = ImageDataGenerator(preprocessing_function=resnet50.preprocess_input)

train_data = datagen.flow_from_directory('/content/ninjacart_data/train', target_size=(224, 224))
val_data = datagen.flow_from_directory('/content/ninjacart_data/validation', target_size=(224, 224))
test_data = datagen.flow_from_directory('/content/ninjacart_data/test', target_size=(224, 224))

Found 2502 images belonging to 4 classes.
Found 627 images belonging to 4 classes.
Found 351 images belonging to 4 classes.
```

## Without finetuning

```
In [ ]: num_classes = 4

# Step 1: Load base model (without top layer)
base_model = ResNet50(weights='imagenet', include_top=False, input_shape=(224, 224, 3))

# Step 2: Freeze all layers (for feature extraction)
for layer in base_model.layers:
    layer.trainable = False

# Step 3: Add custom top layers
x = base_model.output
x = tf.keras.layers.GlobalAveragePooling2D()(x)
x = tf.keras.layers.Dropout(0.5)(x) # optional regularization
x = tf.keras.layers.Dense(512, activation='relu')(x)
predictions = tf.keras.layers.Dense(num_classes, activation='softmax')(x) # num_classes = 4

# Final model
resnet50_model = Model(inputs=base_model.input, outputs=predictions)

# Step 4: Compile the model
resnet50_model.compile(optimizer=Adam(learning_rate=0.001), loss='categorical_crossentropy')

# Step 5: Train only the top layers
before_fine_tuning = resnet50_model.fit(train_data, epochs=10, validation_data=val_data)
```

Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/resnet/resnet50\_weights\_tf\_dim\_ordering\_tf\_kernels\_notop.h5  
94765736/94765736 ————— 0s 0us/step

```
/usr/local/lib/python3.11/dist-packages/keras/src/trainers/data_adapters/py_data  
set_adapter.py:121: UserWarning: Your `PyDataset` class should call `super().__init__(  
**kwargs)` in its constructor. `**kwargs` can include `workers`, `use_multipro  
cessing`, `max_queue_size`. Do not pass these arguments to `fit()`, as they will  
be ignored.  
    self._warn_if_super_not_called()  
Epoch 1/10  
79/79 ━━━━━━━━━━ 43s 369ms/step - accuracy: 0.8178 - loss: 0.5439 - val  
_accuracy: 0.9585 - val_loss: 0.1233  
Epoch 2/10  
79/79 ━━━━━━━━━━ 16s 205ms/step - accuracy: 0.9680 - loss: 0.0834 - val  
_accuracy: 0.9649 - val_loss: 0.1008  
Epoch 3/10  
79/79 ━━━━━━━━━━ 16s 203ms/step - accuracy: 0.9704 - loss: 0.0798 - val  
_accuracy: 0.9697 - val_loss: 0.0889  
Epoch 4/10  
79/79 ━━━━━━━━━━ 16s 206ms/step - accuracy: 0.9850 - loss: 0.0401 - val  
_accuracy: 0.9681 - val_loss: 0.0823  
Epoch 5/10  
79/79 ━━━━━━━━━━ 18s 233ms/step - accuracy: 0.9826 - loss: 0.0475 - val  
_accuracy: 0.9633 - val_loss: 0.0973  
Epoch 6/10  
79/79 ━━━━━━━━━━ 16s 207ms/step - accuracy: 0.9881 - loss: 0.0345 - val  
_accuracy: 0.9777 - val_loss: 0.0612  
Epoch 7/10  
79/79 ━━━━━━━━━━ 17s 212ms/step - accuracy: 0.9900 - loss: 0.0276 - val  
_accuracy: 0.9729 - val_loss: 0.0968  
Epoch 8/10  
79/79 ━━━━━━━━━━ 16s 204ms/step - accuracy: 0.9827 - loss: 0.0320 - val  
_accuracy: 0.9681 - val_loss: 0.0691  
Epoch 9/10  
79/79 ━━━━━━━━━━ 18s 231ms/step - accuracy: 0.9891 - loss: 0.0312 - val  
_accuracy: 0.9761 - val_loss: 0.0820  
Epoch 10/10  
79/79 ━━━━━━━━━━ 16s 202ms/step - accuracy: 0.9901 - loss: 0.0289 - val  
_accuracy: 0.9697 - val_loss: 0.0651
```

In [ ]: `resnet50_model.summary()`

Model: "functional\_12"

Layer (type)	Output Shape	Param #	Connected to
input_layer_12 (InputLayer)	(None, 224, 224, 3)	0	-
conv1_pad (ZeroPadding2D)	(None, 230, 230, 3)	0	input_layer_12[0]
conv1_conv (Conv2D)	(None, 112, 112, 64)	9,472	conv1_pad[0][0]
conv1_bn (BatchNormalizatio...)	(None, 112, 112, 64)	256	conv1_conv[0][0]
conv1_relu (Activation)	(None, 112, 112, 64)	0	conv1_bn[0][0]
pool1_pad (ZeroPadding2D)	(None, 114, 114, 64)	0	conv1_relu[0][0]
pool1_pool (MaxPooling2D)	(None, 56, 56, 64)	0	pool1_pad[0][0]
conv2_block1_1_conv (Conv2D)	(None, 56, 56, 64)	4,160	pool1_pool[0][0]
conv2_block1_1_bn (BatchNormalizatio...)	(None, 56, 56, 64)	256	conv2_block1_1_c...
conv2_block1_1_relu (Activation)	(None, 56, 56, 64)	0	conv2_block1_1_b...
conv2_block1_2_conv (Conv2D)	(None, 56, 56, 64)	36,928	conv2_block1_1_r...
conv2_block1_2_bn (BatchNormalizatio...)	(None, 56, 56, 64)	256	conv2_block1_2_c...
conv2_block1_2_relu (Activation)	(None, 56, 56, 64)	0	conv2_block1_2_b...
conv2_block1_0_conv (Conv2D)	(None, 56, 56, 256)	16,640	pool1_pool[0][0]
conv2_block1_3_conv (Conv2D)	(None, 56, 56, 256)	16,640	conv2_block1_2_r...
conv2_block1_0_bn (BatchNormalizatio...)	(None, 56, 56, 256)	1,024	conv2_block1_0_c...
conv2_block1_3_bn (BatchNormalizatio...)	(None, 56, 56, 256)	1,024	conv2_block1_3_c...
conv2_block1_add (Add)	(None, 56, 56, 256)	0	conv2_block1_0_b... conv2_block1_3_b...
conv2_block1_out (Activation)	(None, 56, 56, 256)	0	conv2_block1_add... conv2_block1_out...

conv2_block2_1_conv (Conv2D)	(None, 56, 56, 64)	16,448	conv2_block1_out...
conv2_block2_1_bn (BatchNormalizatio...)	(None, 56, 56, 64)	256	conv2_block2_1_c...
conv2_block2_1_relu (Activation)	(None, 56, 56, 64)	0	conv2_block2_1_b...
conv2_block2_2_conv (Conv2D)	(None, 56, 56, 64)	36,928	conv2_block2_1_r...
conv2_block2_2_bn (BatchNormalizatio...)	(None, 56, 56, 64)	256	conv2_block2_2_c...
conv2_block2_2_relu (Activation)	(None, 56, 56, 64)	0	conv2_block2_2_b...
conv2_block2_3_conv (Conv2D)	(None, 56, 56, 256)	16,640	conv2_block2_2_r...
conv2_block2_3_bn (BatchNormalizatio...)	(None, 56, 56, 256)	1,024	conv2_block2_3_c...
conv2_block2_add (Add)	(None, 56, 56, 256)	0	conv2_block1_out... conv2_block2_3_b...
conv2_block2_out (Activation)	(None, 56, 56, 256)	0	conv2_block2_add...
conv2_block3_1_conv (Conv2D)	(None, 56, 56, 64)	16,448	conv2_block2_out...
conv2_block3_1_bn (BatchNormalizatio...)	(None, 56, 56, 64)	256	conv2_block3_1_c...
conv2_block3_1_relu (Activation)	(None, 56, 56, 64)	0	conv2_block3_1_b...
conv2_block3_2_conv (Conv2D)	(None, 56, 56, 64)	36,928	conv2_block3_1_r...
conv2_block3_2_bn (BatchNormalizatio...)	(None, 56, 56, 64)	256	conv2_block3_2_c...
conv2_block3_2_relu (Activation)	(None, 56, 56, 64)	0	conv2_block3_2_b...
conv2_block3_3_conv (Conv2D)	(None, 56, 56, 256)	16,640	conv2_block3_2_r...
conv2_block3_3_bn (BatchNormalizatio...)	(None, 56, 56, 256)	1,024	conv2_block3_3_c...
conv2_block3_add (Add)	(None, 56, 56, 256)	0	conv2_block2_out... conv2_block3_3_b...
conv2_block3_out (Activation)	(None, 56, 56, 256)	0	conv2_block3_add...

conv3_block1_1_conv (Conv2D)	(None, 28, 28, 128)	32,896	conv2_block3_out...
conv3_block1_1_bn (BatchNormalizatio...)	(None, 28, 28, 128)	512	conv3_block1_1_c...
conv3_block1_1_relu (Activation)	(None, 28, 28, 128)	0	conv3_block1_1_b...
conv3_block1_2_conv (Conv2D)	(None, 28, 28, 128)	147,584	conv3_block1_1_r...
conv3_block1_2_bn (BatchNormalizatio...)	(None, 28, 28, 128)	512	conv3_block1_2_c...
conv3_block1_2_relu (Activation)	(None, 28, 28, 128)	0	conv3_block1_2_b...
conv3_block1_0_conv (Conv2D)	(None, 28, 28, 512)	131,584	conv2_block3_out...
conv3_block1_3_conv (Conv2D)	(None, 28, 28, 512)	66,048	conv3_block1_2_r...
conv3_block1_0_bn (BatchNormalizatio...)	(None, 28, 28, 512)	2,048	conv3_block1_0_c...
conv3_block1_3_bn (BatchNormalizatio...)	(None, 28, 28, 512)	2,048	conv3_block1_3_c...
conv3_block1_add (Add)	(None, 28, 28, 512)	0	conv3_block1_0_b... conv3_block1_3_b...
conv3_block1_out (Activation)	(None, 28, 28, 512)	0	conv3_block1_add...
conv3_block2_1_conv (Conv2D)	(None, 28, 28, 128)	65,664	conv3_block1_out...
conv3_block2_1_bn (BatchNormalizatio...)	(None, 28, 28, 128)	512	conv3_block2_1_c...
conv3_block2_1_relu (Activation)	(None, 28, 28, 128)	0	conv3_block2_1_b...
conv3_block2_2_conv (Conv2D)	(None, 28, 28, 128)	147,584	conv3_block2_1_r...
conv3_block2_2_bn (BatchNormalizatio...)	(None, 28, 28, 128)	512	conv3_block2_2_c...
conv3_block2_2_relu (Activation)	(None, 28, 28, 128)	0	conv3_block2_2_b...
conv3_block2_3_conv (Conv2D)	(None, 28, 28, 512)	66,048	conv3_block2_2_r...
conv3_block2_3_bn (BatchNormalizatio...)	(None, 28, 28, 512)	2,048	conv3_block2_3_c...

conv3_block2_add (Add)	(None, 28, 28, 512)	0	conv3_block1_out... conv3_block2_3_b...
conv3_block2_out (Activation)	(None, 28, 28, 512)	0	conv3_block2_add...
conv3_block3_1_conv (Conv2D)	(None, 28, 28, 128)	65,664	conv3_block2_out...
conv3_block3_1_bn (BatchNormalizatio...)	(None, 28, 28, 128)	512	conv3_block3_1_c...
conv3_block3_1_relu (Activation)	(None, 28, 28, 128)	0	conv3_block3_1_b...
conv3_block3_2_conv (Conv2D)	(None, 28, 28, 128)	147,584	conv3_block3_1_r...
conv3_block3_2_bn (BatchNormalizatio...)	(None, 28, 28, 128)	512	conv3_block3_2_c...
conv3_block3_2_relu (Activation)	(None, 28, 28, 128)	0	conv3_block3_2_b...
conv3_block3_3_conv (Conv2D)	(None, 28, 28, 512)	66,048	conv3_block3_2_r...
conv3_block3_3_bn (BatchNormalizatio...)	(None, 28, 28, 512)	2,048	conv3_block3_3_c...
conv3_block3_add (Add)	(None, 28, 28, 512)	0	conv3_block2_out... conv3_block3_3_b...
conv3_block3_out (Activation)	(None, 28, 28, 512)	0	conv3_block3_add...
conv3_block4_1_conv (Conv2D)	(None, 28, 28, 128)	65,664	conv3_block3_out...
conv3_block4_1_bn (BatchNormalizatio...)	(None, 28, 28, 128)	512	conv3_block4_1_c...
conv3_block4_1_relu (Activation)	(None, 28, 28, 128)	0	conv3_block4_1_b...
conv3_block4_2_conv (Conv2D)	(None, 28, 28, 128)	147,584	conv3_block4_1_r...
conv3_block4_2_bn (BatchNormalizatio...)	(None, 28, 28, 128)	512	conv3_block4_2_c...
conv3_block4_2_relu (Activation)	(None, 28, 28, 128)	0	conv3_block4_2_b...
conv3_block4_3_conv (Conv2D)	(None, 28, 28, 512)	66,048	conv3_block4_2_r...
conv3_block4_3_bn (BatchNormalizatio...)	(None, 28, 28, 512)	2,048	conv3_block4_3_c...

conv3_block4_add (Add)	(None, 28, 28, 512)	0	conv3_block3_out... conv3_block4_3_b...
conv3_block4_out (Activation)	(None, 28, 28, 512)	0	conv3_block4_add...
conv4_block1_1_conv (Conv2D)	(None, 14, 14, 256)	131,328	conv3_block4_out...
conv4_block1_1_bn (BatchNormalizatio...)	(None, 14, 14, 256)	1,024	conv4_block1_1_c...
conv4_block1_1_relu (Activation)	(None, 14, 14, 256)	0	conv4_block1_1_b...
conv4_block1_2_conv (Conv2D)	(None, 14, 14, 256)	590,080	conv4_block1_1_r...
conv4_block1_2_bn (BatchNormalizatio...)	(None, 14, 14, 256)	1,024	conv4_block1_2_c...
conv4_block1_2_relu (Activation)	(None, 14, 14, 256)	0	conv4_block1_2_b...
conv4_block1_0_conv (Conv2D)	(None, 14, 14, 1024)	525,312	conv3_block4_out...
conv4_block1_3_conv (Conv2D)	(None, 14, 14, 1024)	263,168	conv4_block1_2_r...
conv4_block1_0_bn (BatchNormalizatio...)	(None, 14, 14, 1024)	4,096	conv4_block1_0_c...
conv4_block1_3_bn (BatchNormalizatio...)	(None, 14, 14, 1024)	4,096	conv4_block1_3_c...
conv4_block1_add (Add)	(None, 14, 14, 1024)	0	conv4_block1_0_b... conv4_block1_3_b...
conv4_block1_out (Activation)	(None, 14, 14, 1024)	0	conv4_block1_add...
conv4_block2_1_conv (Conv2D)	(None, 14, 14, 256)	262,400	conv4_block1_out...
conv4_block2_1_bn (BatchNormalizatio...)	(None, 14, 14, 256)	1,024	conv4_block2_1_c...
conv4_block2_1_relu (Activation)	(None, 14, 14, 256)	0	conv4_block2_1_b...
conv4_block2_2_conv (Conv2D)	(None, 14, 14, 256)	590,080	conv4_block2_1_r...
conv4_block2_2_bn (BatchNormalizatio...)	(None, 14, 14, 256)	1,024	conv4_block2_2_c...
conv4_block2_2_relu (Activation)	(None, 14, 14, 256)	0	conv4_block2_2_b...

conv4_block2_3_conv (Conv2D)	(None, 14, 14, 1024)	263,168	conv4_block2_2_r...
conv4_block2_3_bn (BatchNormalizatio...)	(None, 14, 14, 1024)	4,096	conv4_block2_3_c...
conv4_block2_add (Add)	(None, 14, 14, 1024)	0	conv4_block1_out... conv4_block2_3_b...
conv4_block2_out (Activation)	(None, 14, 14, 1024)	0	conv4_block2_add...
conv4_block3_1_conv (Conv2D)	(None, 14, 14, 256)	262,400	conv4_block2_out...
conv4_block3_1_bn (BatchNormalizatio...)	(None, 14, 14, 256)	1,024	conv4_block3_1_c...
conv4_block3_1_relu (Activation)	(None, 14, 14, 256)	0	conv4_block3_1_b...
conv4_block3_2_conv (Conv2D)	(None, 14, 14, 256)	590,080	conv4_block3_1_r...
conv4_block3_2_bn (BatchNormalizatio...)	(None, 14, 14, 256)	1,024	conv4_block3_2_c...
conv4_block3_2_relu (Activation)	(None, 14, 14, 256)	0	conv4_block3_2_b...
conv4_block3_3_conv (Conv2D)	(None, 14, 14, 1024)	263,168	conv4_block3_2_r...
conv4_block3_3_bn (BatchNormalizatio...)	(None, 14, 14, 1024)	4,096	conv4_block3_3_c...
conv4_block3_add (Add)	(None, 14, 14, 1024)	0	conv4_block2_out... conv4_block3_3_b...
conv4_block3_out (Activation)	(None, 14, 14, 1024)	0	conv4_block3_add...
conv4_block4_1_conv (Conv2D)	(None, 14, 14, 256)	262,400	conv4_block3_out...
conv4_block4_1_bn (BatchNormalizatio...)	(None, 14, 14, 256)	1,024	conv4_block4_1_c...
conv4_block4_1_relu (Activation)	(None, 14, 14, 256)	0	conv4_block4_1_b...
conv4_block4_2_conv (Conv2D)	(None, 14, 14, 256)	590,080	conv4_block4_1_r...
conv4_block4_2_bn (BatchNormalizatio...)	(None, 14, 14, 256)	1,024	conv4_block4_2_c...
conv4_block4_2_relu (Activation)	(None, 14, 14, 256)	0	conv4_block4_2_b...

conv4_block4_3_conv (Conv2D)	(None, 14, 14, 1024)	263,168	conv4_block4_2_r...
conv4_block4_3_bn (BatchNormalizatio...)	(None, 14, 14, 1024)	4,096	conv4_block4_3_c...
conv4_block4_add (Add)	(None, 14, 14, 1024)	0	conv4_block3_out... conv4_block4_3_b...
conv4_block4_out (Activation)	(None, 14, 14, 1024)	0	conv4_block4_add...
conv4_block5_1_conv (Conv2D)	(None, 14, 14, 256)	262,400	conv4_block4_out...
conv4_block5_1_bn (BatchNormalizatio...)	(None, 14, 14, 256)	1,024	conv4_block5_1_c...
conv4_block5_1_relu (Activation)	(None, 14, 14, 256)	0	conv4_block5_1_b...
conv4_block5_2_conv (Conv2D)	(None, 14, 14, 256)	590,080	conv4_block5_1_r...
conv4_block5_2_bn (BatchNormalizatio...)	(None, 14, 14, 256)	1,024	conv4_block5_2_c...
conv4_block5_2_relu (Activation)	(None, 14, 14, 256)	0	conv4_block5_2_b...
conv4_block5_3_conv (Conv2D)	(None, 14, 14, 1024)	263,168	conv4_block5_2_r...
conv4_block5_3_bn (BatchNormalizatio...)	(None, 14, 14, 1024)	4,096	conv4_block5_3_c...
conv4_block5_add (Add)	(None, 14, 14, 1024)	0	conv4_block4_out... conv4_block5_3_b...
conv4_block5_out (Activation)	(None, 14, 14, 1024)	0	conv4_block5_add...
conv4_block6_1_conv (Conv2D)	(None, 14, 14, 256)	262,400	conv4_block5_out...
conv4_block6_1_bn (BatchNormalizatio...)	(None, 14, 14, 256)	1,024	conv4_block6_1_c...
conv4_block6_1_relu (Activation)	(None, 14, 14, 256)	0	conv4_block6_1_b...
conv4_block6_2_conv (Conv2D)	(None, 14, 14, 256)	590,080	conv4_block6_1_r...
conv4_block6_2_bn (BatchNormalizatio...)	(None, 14, 14, 256)	1,024	conv4_block6_2_c...
conv4_block6_2_relu (Activation)	(None, 14, 14, 256)	0	conv4_block6_2_b...

conv4_block6_3_conv (Conv2D)	(None, 14, 14, 1024)	263,168	conv4_block6_2_r...
conv4_block6_3_bn (BatchNormalizatio...)	(None, 14, 14, 1024)	4,096	conv4_block6_3_c...
conv4_block6_add (Add)	(None, 14, 14, 1024)	0	conv4_block5_out... conv4_block6_3_b...
conv4_block6_out (Activation)	(None, 14, 14, 1024)	0	conv4_block6_add...
conv5_block1_1_conv (Conv2D)	(None, 7, 7, 512)	524,800	conv4_block6_out...
conv5_block1_1_bn (BatchNormalizatio...)	(None, 7, 7, 512)	2,048	conv5_block1_1_c...
conv5_block1_1_relu (Activation)	(None, 7, 7, 512)	0	conv5_block1_1_b...
conv5_block1_2_conv (Conv2D)	(None, 7, 7, 512)	2,359,808	conv5_block1_1_r...
conv5_block1_2_bn (BatchNormalizatio...)	(None, 7, 7, 512)	2,048	conv5_block1_2_c...
conv5_block1_2_relu (Activation)	(None, 7, 7, 512)	0	conv5_block1_2_b...
conv5_block1_0_conv (Conv2D)	(None, 7, 7, 2048)	2,099,200	conv4_block6_out...
conv5_block1_3_conv (Conv2D)	(None, 7, 7, 2048)	1,050,624	conv5_block1_2_r...
conv5_block1_0_bn (BatchNormalizatio...)	(None, 7, 7, 2048)	8,192	conv5_block1_0_c...
conv5_block1_3_bn (BatchNormalizatio...)	(None, 7, 7, 2048)	8,192	conv5_block1_3_c...
conv5_block1_add (Add)	(None, 7, 7, 2048)	0	conv5_block1_0_b... conv5_block1_3_b...
conv5_block1_out (Activation)	(None, 7, 7, 2048)	0	conv5_block1_add...
conv5_block2_1_conv (Conv2D)	(None, 7, 7, 512)	1,049,088	conv5_block1_out...
conv5_block2_1_bn (BatchNormalizatio...)	(None, 7, 7, 512)	2,048	conv5_block2_1_c...
conv5_block2_1_relu (Activation)	(None, 7, 7, 512)	0	conv5_block2_1_b...
conv5_block2_2_conv (Conv2D)	(None, 7, 7, 512)	2,359,808	conv5_block2_1_r...

conv5_block2_2_bn (BatchNormalizatio...)	(None, 7, 7, 512)	2,048	conv5_block2_2_c...
conv5_block2_2_relu (Activation)	(None, 7, 7, 512)	0	conv5_block2_2_b...
conv5_block2_3_conv (Conv2D)	(None, 7, 7, 2048)	1,050,624	conv5_block2_2_r...
conv5_block2_3_bn (BatchNormalizatio...)	(None, 7, 7, 2048)	8,192	conv5_block2_3_c...
conv5_block2_add (Add)	(None, 7, 7, 2048)	0	conv5_block1_out... conv5_block2_3_b...
conv5_block2_out (Activation)	(None, 7, 7, 2048)	0	conv5_block2_add...
conv5_block3_1_conv (Conv2D)	(None, 7, 7, 512)	1,049,088	conv5_block2_out...
conv5_block3_1_bn (BatchNormalizatio...)	(None, 7, 7, 512)	2,048	conv5_block3_1_c...
conv5_block3_1_relu (Activation)	(None, 7, 7, 512)	0	conv5_block3_1_b...
conv5_block3_2_conv (Conv2D)	(None, 7, 7, 512)	2,359,808	conv5_block3_1_r...
conv5_block3_2_bn (BatchNormalizatio...)	(None, 7, 7, 512)	2,048	conv5_block3_2_c...
conv5_block3_2_relu (Activation)	(None, 7, 7, 512)	0	conv5_block3_2_b...
conv5_block3_3_conv (Conv2D)	(None, 7, 7, 2048)	1,050,624	conv5_block3_2_r...
conv5_block3_3_bn (BatchNormalizatio...)	(None, 7, 7, 2048)	8,192	conv5_block3_3_c...
conv5_block3_add (Add)	(None, 7, 7, 2048)	0	conv5_block2_out... conv5_block3_3_b...
conv5_block3_out (Activation)	(None, 7, 7, 2048)	0	conv5_block3_add...
global_average_poo... (GlobalAveragePool...)	(None, 2048)	0	conv5_block3_out...
dropout_7 (Dropout)	(None, 2048)	0	global_average_p...
dense_18 (Dense)	(None, 512)	1,049,088	dropout_7[0][0]
dense_19 (Dense)	(None, 4)	2,052	dense_18[0][0]

Total params: 26,741,134 (102.01 MB)

```
Trainable params: 1,051,140 (4.01 MB)
Non-trainable params: 23,587,712 (89.98 MB)
Optimizer params: 2,102,282 (8.02 MB)
```

```
In [ ]: create_directory_if_not_exists('/content/ninjacart_models/model10_Resent_GAP_DP_
resnet50_model.save("/content/ninjacart_models/model10_Resent_GAP_DP_FineTune/mo
```

```
Created directory: /content/ninjacart_models/model10_Resent_GAP_DP_FineTune
```

```
In [ ]: resnet50_model_finetune = tf.keras.models.load_model("/content/ninjacart_models/
loss, acc = resnet50_model_finetune.evaluate(test_data, verbose=2)
print('Before FineTuning')
print("Restored model, accuracy: {:.2f}%".format(100 * acc))
```

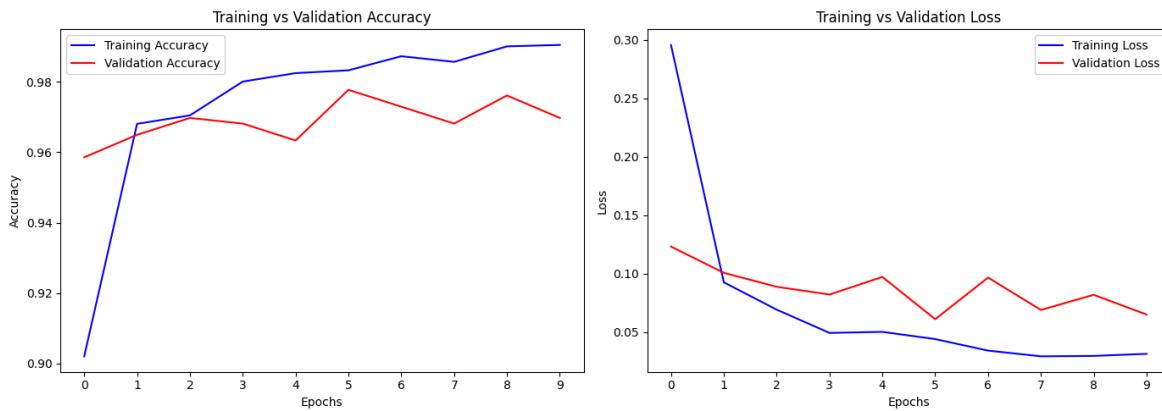
```
11/11 - 12s - 1s/step - accuracy: 0.9288 - loss: 0.3587
```

```
Before FineTuning
```

```
Restored model, accuracy: 92.88%
```

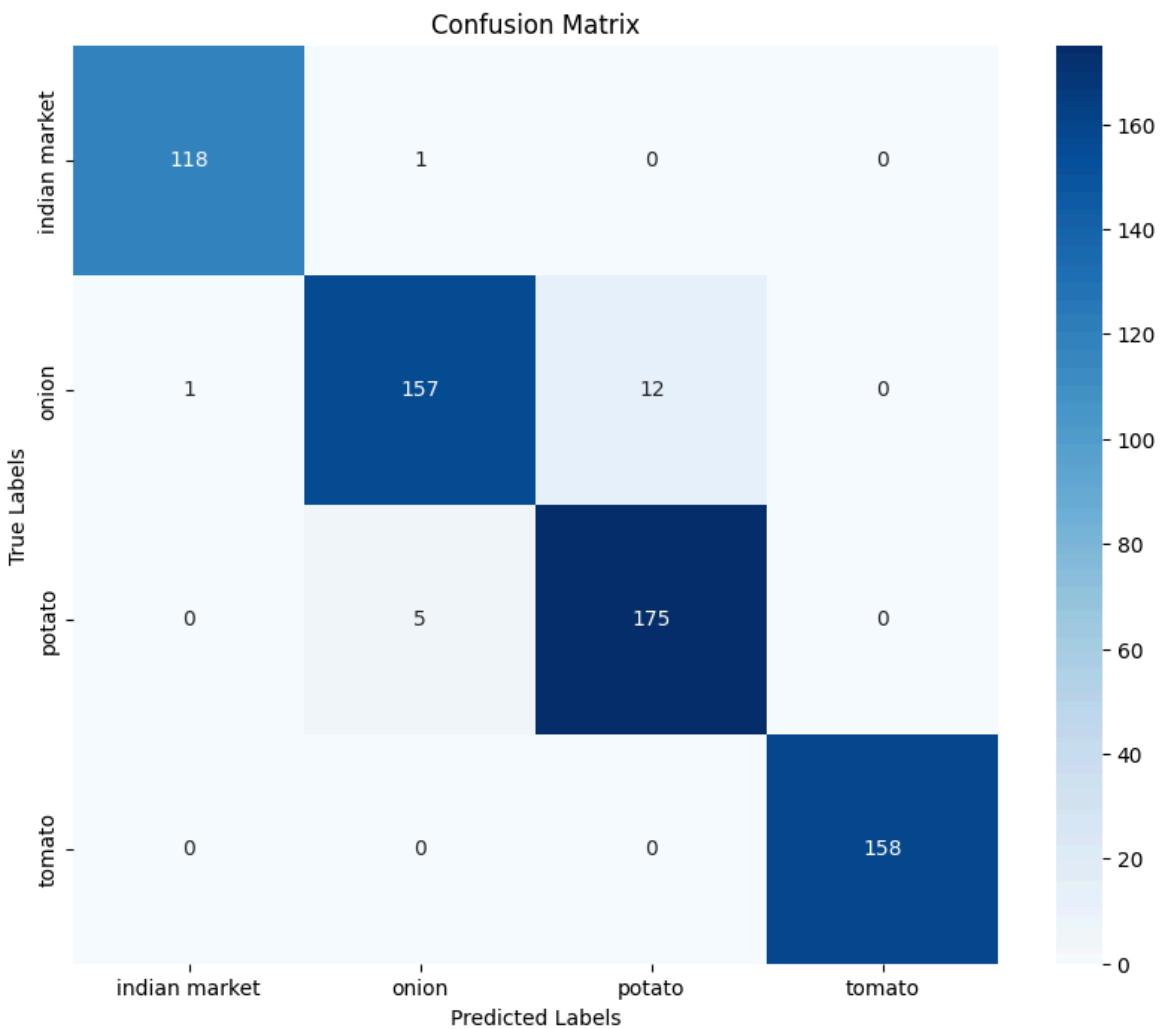
```
In [ ]: plot_training_history(before_fine_tuning)
```

```
History keys: dict_keys(['accuracy', 'loss', 'val_accuracy', 'val_loss'])
```



```
In [ ]: class_names = list(val_data.class_indices.keys())
plot_confusion_matrix(resnet50_model, val_data, class_names)
```

```
20/20 ━━━━━━━━ 12s 399ms/step
```



Classification Report:

	precision	recall	f1-score	support
indian market	0.99	0.99	0.99	119
onion	0.96	0.92	0.94	170
potato	0.94	0.97	0.95	180
tomato	1.00	1.00	1.00	158
accuracy			0.97	627
macro avg	0.97	0.97	0.97	627
weighted avg	0.97	0.97	0.97	627

```
In [ ]: inception_model_finetune = tf.keras.models.load_model("/content/ninjacart_models")
loss, acc = inception_model_finetune.evaluate(test_data, verbose=2)
print('Before FineTuning')
print("Restored model, accuracy: {:.5f}%".format(100 * acc))
```

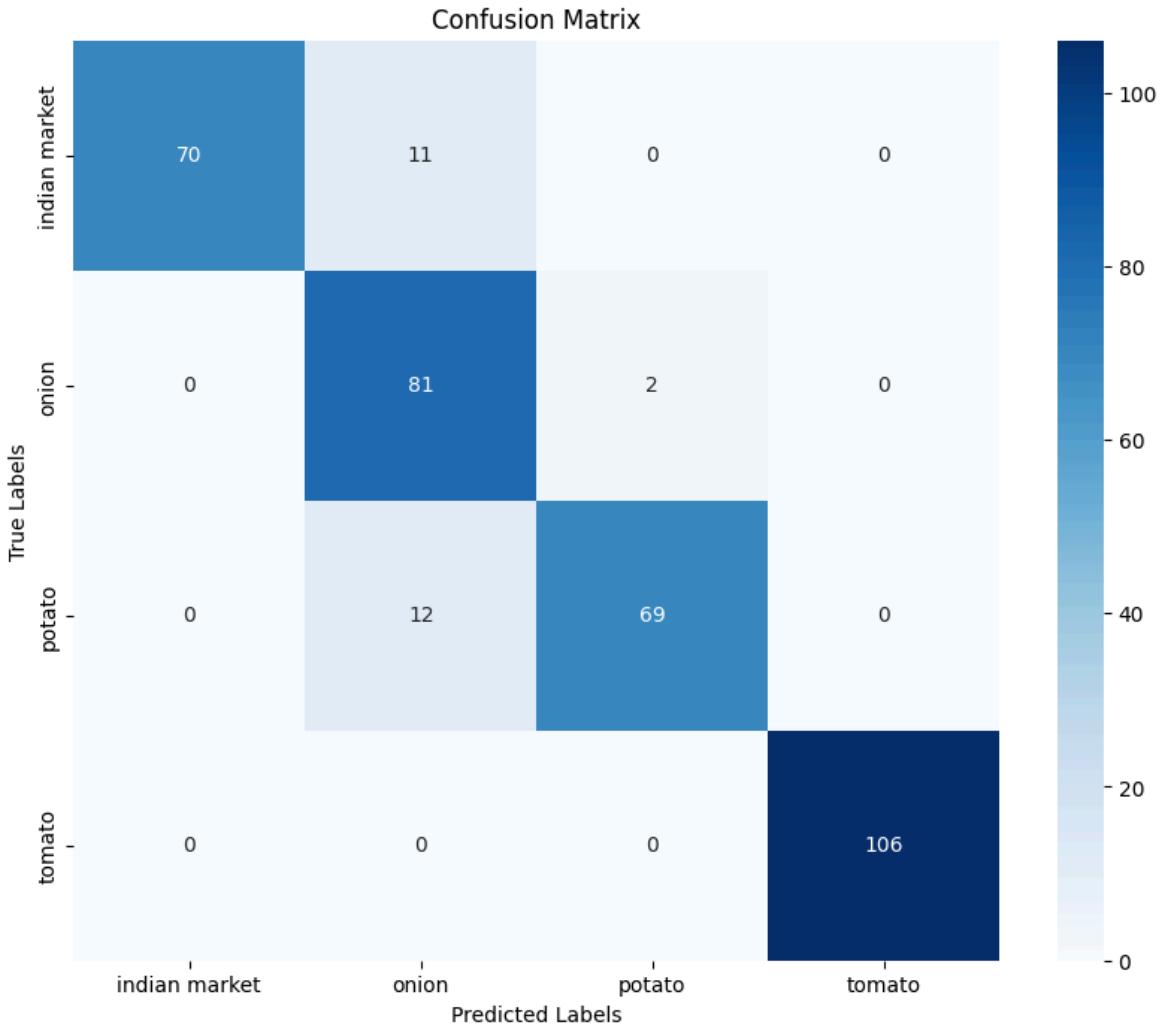
11/11 - 9s - 830ms/step - accuracy: 0.9288 - loss: 0.3587

Before FineTuning

Restored model, accuracy: 92.88%

```
In [ ]: class_names = list(test_data.class_indices.keys())
plot_confusion_matrix(resnet50_model, test_data, class_names)
```

11/11 ━━━━━━━━ 3s 322ms/step



Classification Report:

	precision	recall	f1-score	support
indian market	1.00	0.86	0.93	81
onion	0.78	0.98	0.87	83
potato	0.97	0.85	0.91	81
tomato	1.00	1.00	1.00	106
accuracy			0.93	351
macro avg	0.94	0.92	0.93	351
weighted avg	0.94	0.93	0.93	351

## After Finetuning

```
In [ ]: # Step 6: Unfreeze Last few Layers of ResNet for fine-tuning
for layer in base_model.layers[-30:]: # last 30 layers
    if not isinstance(layer, tf.keras.layers.BatchNormalization):
        layer.trainable = True

# Re-compile with a smaller Learning rate
resnet50_model.compile(optimizer=Adam(learning_rate=1e-5), loss='categorical_crossentropy')

# Step 7: Continue training (fine-tuning)
after_fine_tuning = resnet50_model.fit(train_data, epochs=10, validation_data=validation_data)
```

```
Epoch 1/10
79/79 48s 413ms/step - accuracy: 0.9917 - loss: 0.0250 - val
_accuracy: 0.9713 - val_loss: 0.0784
Epoch 2/10
79/79 20s 216ms/step - accuracy: 0.9970 - loss: 0.0060 - val
_accuracy: 0.9713 - val_loss: 0.0806
Epoch 3/10
79/79 18s 224ms/step - accuracy: 0.9979 - loss: 0.0048 - val
_accuracy: 0.9777 - val_loss: 0.0533
Epoch 4/10
79/79 20s 211ms/step - accuracy: 0.9971 - loss: 0.0052 - val
_accuracy: 0.9713 - val_loss: 0.0766
Epoch 5/10
79/79 17s 215ms/step - accuracy: 0.9972 - loss: 0.0054 - val
_accuracy: 0.9777 - val_loss: 0.0554
Epoch 6/10
79/79 17s 214ms/step - accuracy: 0.9974 - loss: 0.0085 - val
_accuracy: 0.9745 - val_loss: 0.0667
Epoch 7/10
79/79 18s 225ms/step - accuracy: 0.9998 - loss: 0.0022 - val
_accuracy: 0.9729 - val_loss: 0.0812
Epoch 8/10
79/79 17s 215ms/step - accuracy: 0.9991 - loss: 0.0059 - val
_accuracy: 0.9729 - val_loss: 0.0721
Epoch 9/10
79/79 17s 219ms/step - accuracy: 0.9997 - loss: 0.0019 - val
_accuracy: 0.9745 - val_loss: 0.0844
Epoch 10/10
79/79 17s 218ms/step - accuracy: 0.9986 - loss: 0.0026 - val
_accuracy: 0.9809 - val_loss: 0.0676
```

```
In [ ]: resnet50_model.summary()
```

```
Model: "functional_12"
```

Layer (type)	Output Shape	Param #	Connected to
input_layer_12 (InputLayer)	(None, 224, 224, 3)	0	-
conv1_pad (ZeroPadding2D)	(None, 230, 230, 3)	0	input_layer_12[0]
conv1_conv (Conv2D)	(None, 112, 112, 64)	9,472	conv1_pad[0][0]
conv1_bn (BatchNormalizatio...)	(None, 112, 112, 64)	256	conv1_conv[0][0]
conv1_relu (Activation)	(None, 112, 112, 64)	0	conv1_bn[0][0]
pool1_pad (ZeroPadding2D)	(None, 114, 114, 64)	0	conv1_relu[0][0]
pool1_pool (MaxPooling2D)	(None, 56, 56, 64)	0	pool1_pad[0][0]
conv2_block1_1_conv (Conv2D)	(None, 56, 56, 64)	4,160	pool1_pool[0][0]
conv2_block1_1_bn (BatchNormalizatio...)	(None, 56, 56, 64)	256	conv2_block1_1_c...
conv2_block1_1_relu (Activation)	(None, 56, 56, 64)	0	conv2_block1_1_b...
conv2_block1_2_conv (Conv2D)	(None, 56, 56, 64)	36,928	conv2_block1_1_r...
conv2_block1_2_bn (BatchNormalizatio...)	(None, 56, 56, 64)	256	conv2_block1_2_c...
conv2_block1_2_relu (Activation)	(None, 56, 56, 64)	0	conv2_block1_2_b...
conv2_block1_0_conv (Conv2D)	(None, 56, 56, 256)	16,640	pool1_pool[0][0]
conv2_block1_3_conv (Conv2D)	(None, 56, 56, 256)	16,640	conv2_block1_2_r...
conv2_block1_0_bn (BatchNormalizatio...)	(None, 56, 56, 256)	1,024	conv2_block1_0_c...
conv2_block1_3_bn (BatchNormalizatio...)	(None, 56, 56, 256)	1,024	conv2_block1_3_c...
conv2_block1_add (Add)	(None, 56, 56, 256)	0	conv2_block1_0_b... conv2_block1_3_b...
conv2_block1_out (Activation)	(None, 56, 56, 256)	0	conv2_block1_add... conv2_block1_out...

conv2_block2_1_conv (Conv2D)	(None, 56, 56, 64)	16,448	conv2_block1_out...
conv2_block2_1_bn (BatchNormalizatio...)	(None, 56, 56, 64)	256	conv2_block2_1_c...
conv2_block2_1_relu (Activation)	(None, 56, 56, 64)	0	conv2_block2_1_b...
conv2_block2_2_conv (Conv2D)	(None, 56, 56, 64)	36,928	conv2_block2_1_r...
conv2_block2_2_bn (BatchNormalizatio...)	(None, 56, 56, 64)	256	conv2_block2_2_c...
conv2_block2_2_relu (Activation)	(None, 56, 56, 64)	0	conv2_block2_2_b...
conv2_block2_3_conv (Conv2D)	(None, 56, 56, 256)	16,640	conv2_block2_2_r...
conv2_block2_3_bn (BatchNormalizatio...)	(None, 56, 56, 256)	1,024	conv2_block2_3_c...
conv2_block2_add (Add)	(None, 56, 56, 256)	0	conv2_block1_out... conv2_block2_3_b...
conv2_block2_out (Activation)	(None, 56, 56, 256)	0	conv2_block2_add...
conv2_block3_1_conv (Conv2D)	(None, 56, 56, 64)	16,448	conv2_block2_out...
conv2_block3_1_bn (BatchNormalizatio...)	(None, 56, 56, 64)	256	conv2_block3_1_c...
conv2_block3_1_relu (Activation)	(None, 56, 56, 64)	0	conv2_block3_1_b...
conv2_block3_2_conv (Conv2D)	(None, 56, 56, 64)	36,928	conv2_block3_1_r...
conv2_block3_2_bn (BatchNormalizatio...)	(None, 56, 56, 64)	256	conv2_block3_2_c...
conv2_block3_2_relu (Activation)	(None, 56, 56, 64)	0	conv2_block3_2_b...
conv2_block3_3_conv (Conv2D)	(None, 56, 56, 256)	16,640	conv2_block3_2_r...
conv2_block3_3_bn (BatchNormalizatio...)	(None, 56, 56, 256)	1,024	conv2_block3_3_c...
conv2_block3_add (Add)	(None, 56, 56, 256)	0	conv2_block2_out... conv2_block3_3_b...
conv2_block3_out (Activation)	(None, 56, 56, 256)	0	conv2_block3_add...

conv3_block1_1_conv (Conv2D)	(None, 28, 28, 128)	32,896	conv2_block3_out...
conv3_block1_1_bn (BatchNormalizatio...)	(None, 28, 28, 128)	512	conv3_block1_1_c...
conv3_block1_1_relu (Activation)	(None, 28, 28, 128)	0	conv3_block1_1_b...
conv3_block1_2_conv (Conv2D)	(None, 28, 28, 128)	147,584	conv3_block1_1_r...
conv3_block1_2_bn (BatchNormalizatio...)	(None, 28, 28, 128)	512	conv3_block1_2_c...
conv3_block1_2_relu (Activation)	(None, 28, 28, 128)	0	conv3_block1_2_b...
conv3_block1_0_conv (Conv2D)	(None, 28, 28, 512)	131,584	conv2_block3_out...
conv3_block1_3_conv (Conv2D)	(None, 28, 28, 512)	66,048	conv3_block1_2_r...
conv3_block1_0_bn (BatchNormalizatio...)	(None, 28, 28, 512)	2,048	conv3_block1_0_c...
conv3_block1_3_bn (BatchNormalizatio...)	(None, 28, 28, 512)	2,048	conv3_block1_3_c...
conv3_block1_add (Add)	(None, 28, 28, 512)	0	conv3_block1_0_b... conv3_block1_3_b...
conv3_block1_out (Activation)	(None, 28, 28, 512)	0	conv3_block1_add...
conv3_block2_1_conv (Conv2D)	(None, 28, 28, 128)	65,664	conv3_block1_out...
conv3_block2_1_bn (BatchNormalizatio...)	(None, 28, 28, 128)	512	conv3_block2_1_c...
conv3_block2_1_relu (Activation)	(None, 28, 28, 128)	0	conv3_block2_1_b...
conv3_block2_2_conv (Conv2D)	(None, 28, 28, 128)	147,584	conv3_block2_1_r...
conv3_block2_2_bn (BatchNormalizatio...)	(None, 28, 28, 128)	512	conv3_block2_2_c...
conv3_block2_2_relu (Activation)	(None, 28, 28, 128)	0	conv3_block2_2_b...
conv3_block2_3_conv (Conv2D)	(None, 28, 28, 512)	66,048	conv3_block2_2_r...
conv3_block2_3_bn (BatchNormalizatio...)	(None, 28, 28, 512)	2,048	conv3_block2_3_c...

conv3_block2_add (Add)	(None, 28, 28, 512)	0	conv3_block1_out... conv3_block2_3_b...
conv3_block2_out (Activation)	(None, 28, 28, 512)	0	conv3_block2_add...
conv3_block3_1_conv (Conv2D)	(None, 28, 28, 128)	65,664	conv3_block2_out...
conv3_block3_1_bn (BatchNormalizatio...)	(None, 28, 28, 128)	512	conv3_block3_1_c...
conv3_block3_1_relu (Activation)	(None, 28, 28, 128)	0	conv3_block3_1_b...
conv3_block3_2_conv (Conv2D)	(None, 28, 28, 128)	147,584	conv3_block3_1_r...
conv3_block3_2_bn (BatchNormalizatio...)	(None, 28, 28, 128)	512	conv3_block3_2_c...
conv3_block3_2_relu (Activation)	(None, 28, 28, 128)	0	conv3_block3_2_b...
conv3_block3_3_conv (Conv2D)	(None, 28, 28, 512)	66,048	conv3_block3_2_r...
conv3_block3_3_bn (BatchNormalizatio...)	(None, 28, 28, 512)	2,048	conv3_block3_3_c...
conv3_block3_add (Add)	(None, 28, 28, 512)	0	conv3_block2_out... conv3_block3_3_b...
conv3_block3_out (Activation)	(None, 28, 28, 512)	0	conv3_block3_add...
conv3_block4_1_conv (Conv2D)	(None, 28, 28, 128)	65,664	conv3_block3_out...
conv3_block4_1_bn (BatchNormalizatio...)	(None, 28, 28, 128)	512	conv3_block4_1_c...
conv3_block4_1_relu (Activation)	(None, 28, 28, 128)	0	conv3_block4_1_b...
conv3_block4_2_conv (Conv2D)	(None, 28, 28, 128)	147,584	conv3_block4_1_r...
conv3_block4_2_bn (BatchNormalizatio...)	(None, 28, 28, 128)	512	conv3_block4_2_c...
conv3_block4_2_relu (Activation)	(None, 28, 28, 128)	0	conv3_block4_2_b...
conv3_block4_3_conv (Conv2D)	(None, 28, 28, 512)	66,048	conv3_block4_2_r...
conv3_block4_3_bn (BatchNormalizatio...)	(None, 28, 28, 512)	2,048	conv3_block4_3_c...

conv3_block4_add (Add)	(None, 28, 28, 512)	0	conv3_block3_out... conv3_block4_3_b...
conv3_block4_out (Activation)	(None, 28, 28, 512)	0	conv3_block4_add...
conv4_block1_1_conv (Conv2D)	(None, 14, 14, 256)	131,328	conv3_block4_out...
conv4_block1_1_bn (BatchNormalizatio...)	(None, 14, 14, 256)	1,024	conv4_block1_1_c...
conv4_block1_1_relu (Activation)	(None, 14, 14, 256)	0	conv4_block1_1_b...
conv4_block1_2_conv (Conv2D)	(None, 14, 14, 256)	590,080	conv4_block1_1_r...
conv4_block1_2_bn (BatchNormalizatio...)	(None, 14, 14, 256)	1,024	conv4_block1_2_c...
conv4_block1_2_relu (Activation)	(None, 14, 14, 256)	0	conv4_block1_2_b...
conv4_block1_0_conv (Conv2D)	(None, 14, 14, 1024)	525,312	conv3_block4_out...
conv4_block1_3_conv (Conv2D)	(None, 14, 14, 1024)	263,168	conv4_block1_2_r...
conv4_block1_0_bn (BatchNormalizatio...)	(None, 14, 14, 1024)	4,096	conv4_block1_0_c...
conv4_block1_3_bn (BatchNormalizatio...)	(None, 14, 14, 1024)	4,096	conv4_block1_3_c...
conv4_block1_add (Add)	(None, 14, 14, 1024)	0	conv4_block1_0_b... conv4_block1_3_b...
conv4_block1_out (Activation)	(None, 14, 14, 1024)	0	conv4_block1_add...
conv4_block2_1_conv (Conv2D)	(None, 14, 14, 256)	262,400	conv4_block1_out...
conv4_block2_1_bn (BatchNormalizatio...)	(None, 14, 14, 256)	1,024	conv4_block2_1_c...
conv4_block2_1_relu (Activation)	(None, 14, 14, 256)	0	conv4_block2_1_b...
conv4_block2_2_conv (Conv2D)	(None, 14, 14, 256)	590,080	conv4_block2_1_r...
conv4_block2_2_bn (BatchNormalizatio...)	(None, 14, 14, 256)	1,024	conv4_block2_2_c...
conv4_block2_2_relu (Activation)	(None, 14, 14, 256)	0	conv4_block2_2_b...

conv4_block2_3_conv (Conv2D)	(None, 14, 14, 1024)	263,168	conv4_block2_2_r...
conv4_block2_3_bn (BatchNormalizatio...)	(None, 14, 14, 1024)	4,096	conv4_block2_3_c...
conv4_block2_add (Add)	(None, 14, 14, 1024)	0	conv4_block1_out... conv4_block2_3_b...
conv4_block2_out (Activation)	(None, 14, 14, 1024)	0	conv4_block2_add...
conv4_block3_1_conv (Conv2D)	(None, 14, 14, 256)	262,400	conv4_block2_out...
conv4_block3_1_bn (BatchNormalizatio...)	(None, 14, 14, 256)	1,024	conv4_block3_1_c...
conv4_block3_1_relu (Activation)	(None, 14, 14, 256)	0	conv4_block3_1_b...
conv4_block3_2_conv (Conv2D)	(None, 14, 14, 256)	590,080	conv4_block3_1_r...
conv4_block3_2_bn (BatchNormalizatio...)	(None, 14, 14, 256)	1,024	conv4_block3_2_c...
conv4_block3_2_relu (Activation)	(None, 14, 14, 256)	0	conv4_block3_2_b...
conv4_block3_3_conv (Conv2D)	(None, 14, 14, 1024)	263,168	conv4_block3_2_r...
conv4_block3_3_bn (BatchNormalizatio...)	(None, 14, 14, 1024)	4,096	conv4_block3_3_c...
conv4_block3_add (Add)	(None, 14, 14, 1024)	0	conv4_block2_out... conv4_block3_3_b...
conv4_block3_out (Activation)	(None, 14, 14, 1024)	0	conv4_block3_add...
conv4_block4_1_conv (Conv2D)	(None, 14, 14, 256)	262,400	conv4_block3_out...
conv4_block4_1_bn (BatchNormalizatio...)	(None, 14, 14, 256)	1,024	conv4_block4_1_c...
conv4_block4_1_relu (Activation)	(None, 14, 14, 256)	0	conv4_block4_1_b...
conv4_block4_2_conv (Conv2D)	(None, 14, 14, 256)	590,080	conv4_block4_1_r...
conv4_block4_2_bn (BatchNormalizatio...)	(None, 14, 14, 256)	1,024	conv4_block4_2_c...
conv4_block4_2_relu (Activation)	(None, 14, 14, 256)	0	conv4_block4_2_b...

conv4_block4_3_conv (Conv2D)	(None, 14, 14, 1024)	263,168	conv4_block4_2_r...
conv4_block4_3_bn (BatchNormalizatio...)	(None, 14, 14, 1024)	4,096	conv4_block4_3_c...
conv4_block4_add (Add)	(None, 14, 14, 1024)	0	conv4_block3_out... conv4_block4_3_b...
conv4_block4_out (Activation)	(None, 14, 14, 1024)	0	conv4_block4_add...
conv4_block5_1_conv (Conv2D)	(None, 14, 14, 256)	262,400	conv4_block4_out...
conv4_block5_1_bn (BatchNormalizatio...)	(None, 14, 14, 256)	1,024	conv4_block5_1_c...
conv4_block5_1_relu (Activation)	(None, 14, 14, 256)	0	conv4_block5_1_b...
conv4_block5_2_conv (Conv2D)	(None, 14, 14, 256)	590,080	conv4_block5_1_r...
conv4_block5_2_bn (BatchNormalizatio...)	(None, 14, 14, 256)	1,024	conv4_block5_2_c...
conv4_block5_2_relu (Activation)	(None, 14, 14, 256)	0	conv4_block5_2_b...
conv4_block5_3_conv (Conv2D)	(None, 14, 14, 1024)	263,168	conv4_block5_2_r...
conv4_block5_3_bn (BatchNormalizatio...)	(None, 14, 14, 1024)	4,096	conv4_block5_3_c...
conv4_block5_add (Add)	(None, 14, 14, 1024)	0	conv4_block4_out... conv4_block5_3_b...
conv4_block5_out (Activation)	(None, 14, 14, 1024)	0	conv4_block5_add...
conv4_block6_1_conv (Conv2D)	(None, 14, 14, 256)	262,400	conv4_block5_out...
conv4_block6_1_bn (BatchNormalizatio...)	(None, 14, 14, 256)	1,024	conv4_block6_1_c...
conv4_block6_1_relu (Activation)	(None, 14, 14, 256)	0	conv4_block6_1_b...
conv4_block6_2_conv (Conv2D)	(None, 14, 14, 256)	590,080	conv4_block6_1_r...
conv4_block6_2_bn (BatchNormalizatio...)	(None, 14, 14, 256)	1,024	conv4_block6_2_c...
conv4_block6_2_relu (Activation)	(None, 14, 14, 256)	0	conv4_block6_2_b...

conv4_block6_3_conv (Conv2D)	(None, 14, 14, 1024)	263,168	conv4_block6_2_r...
conv4_block6_3_bn (BatchNormalizatio...)	(None, 14, 14, 1024)	4,096	conv4_block6_3_c...
conv4_block6_add (Add)	(None, 14, 14, 1024)	0	conv4_block5_out... conv4_block6_3_b...
conv4_block6_out (Activation)	(None, 14, 14, 1024)	0	conv4_block6_add...
conv5_block1_1_conv (Conv2D)	(None, 7, 7, 512)	524,800	conv4_block6_out...
conv5_block1_1_bn (BatchNormalizatio...)	(None, 7, 7, 512)	2,048	conv5_block1_1_c...
conv5_block1_1_relu (Activation)	(None, 7, 7, 512)	0	conv5_block1_1_b...
conv5_block1_2_conv (Conv2D)	(None, 7, 7, 512)	2,359,808	conv5_block1_1_r...
conv5_block1_2_bn (BatchNormalizatio...)	(None, 7, 7, 512)	2,048	conv5_block1_2_c...
conv5_block1_2_relu (Activation)	(None, 7, 7, 512)	0	conv5_block1_2_b...
conv5_block1_0_conv (Conv2D)	(None, 7, 7, 2048)	2,099,200	conv4_block6_out...
conv5_block1_3_conv (Conv2D)	(None, 7, 7, 2048)	1,050,624	conv5_block1_2_r...
conv5_block1_0_bn (BatchNormalizatio...)	(None, 7, 7, 2048)	8,192	conv5_block1_0_c...
conv5_block1_3_bn (BatchNormalizatio...)	(None, 7, 7, 2048)	8,192	conv5_block1_3_c...
conv5_block1_add (Add)	(None, 7, 7, 2048)	0	conv5_block1_0_b... conv5_block1_3_b...
conv5_block1_out (Activation)	(None, 7, 7, 2048)	0	conv5_block1_add...
conv5_block2_1_conv (Conv2D)	(None, 7, 7, 512)	1,049,088	conv5_block1_out...
conv5_block2_1_bn (BatchNormalizatio...)	(None, 7, 7, 512)	2,048	conv5_block2_1_c...
conv5_block2_1_relu (Activation)	(None, 7, 7, 512)	0	conv5_block2_1_b...
conv5_block2_2_conv (Conv2D)	(None, 7, 7, 512)	2,359,808	conv5_block2_1_r...

conv5_block2_2_bn (BatchNormalizatio...)	(None, 7, 7, 512)	2,048	conv5_block2_2_c...
conv5_block2_2_relu (Activation)	(None, 7, 7, 512)	0	conv5_block2_2_b...
conv5_block2_3_conv (Conv2D)	(None, 7, 7, 2048)	1,050,624	conv5_block2_2_r...
conv5_block2_3_bn (BatchNormalizatio...)	(None, 7, 7, 2048)	8,192	conv5_block2_3_c...
conv5_block2_add (Add)	(None, 7, 7, 2048)	0	conv5_block1_out... conv5_block2_3_b...
conv5_block2_out (Activation)	(None, 7, 7, 2048)	0	conv5_block2_add...
conv5_block3_1_conv (Conv2D)	(None, 7, 7, 512)	1,049,088	conv5_block2_out...
conv5_block3_1_bn (BatchNormalizatio...)	(None, 7, 7, 512)	2,048	conv5_block3_1_c...
conv5_block3_1_relu (Activation)	(None, 7, 7, 512)	0	conv5_block3_1_b...
conv5_block3_2_conv (Conv2D)	(None, 7, 7, 512)	2,359,808	conv5_block3_1_r...
conv5_block3_2_bn (BatchNormalizatio...)	(None, 7, 7, 512)	2,048	conv5_block3_2_c...
conv5_block3_2_relu (Activation)	(None, 7, 7, 512)	0	conv5_block3_2_b...
conv5_block3_3_conv (Conv2D)	(None, 7, 7, 2048)	1,050,624	conv5_block3_2_r...
conv5_block3_3_bn (BatchNormalizatio...)	(None, 7, 7, 2048)	8,192	conv5_block3_3_c...
conv5_block3_add (Add)	(None, 7, 7, 2048)	0	conv5_block2_out... conv5_block3_3_b...
conv5_block3_out (Activation)	(None, 7, 7, 2048)	0	conv5_block3_add...
global_average_poo... (GlobalAveragePool...)	(None, 2048)	0	conv5_block3_out...
dropout_7 (Dropout)	(None, 2048)	0	global_average_p...
dense_18 (Dense)	(None, 512)	1,049,088	dropout_7[0][0]
dense_19 (Dense)	(None, 4)	2,052	dense_18[0][0]

Total params: 55,598,478 (212.09 MB)

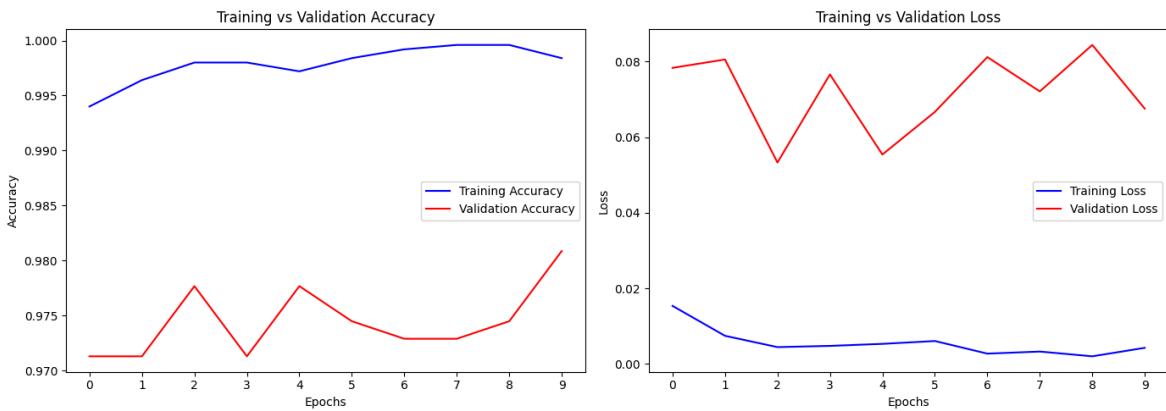
```
Trainable params: 15,479,812 (59.05 MB)
Non-trainable params: 9,159,040 (34.94 MB)
Optimizer params: 30,959,626 (118.10 MB)
```

```
In [ ]: create_directory_if_not_exists('/content/ninjacart_models/model10_Resent_GAP_DP_
resnet50_model.save("/content/ninjacart_models/model10_Resent_GAP_DP_FineTune/mo
```

```
Directory already exists: /content/ninjacart_models/model10_Resent_GAP_DP_FineTun
e
```

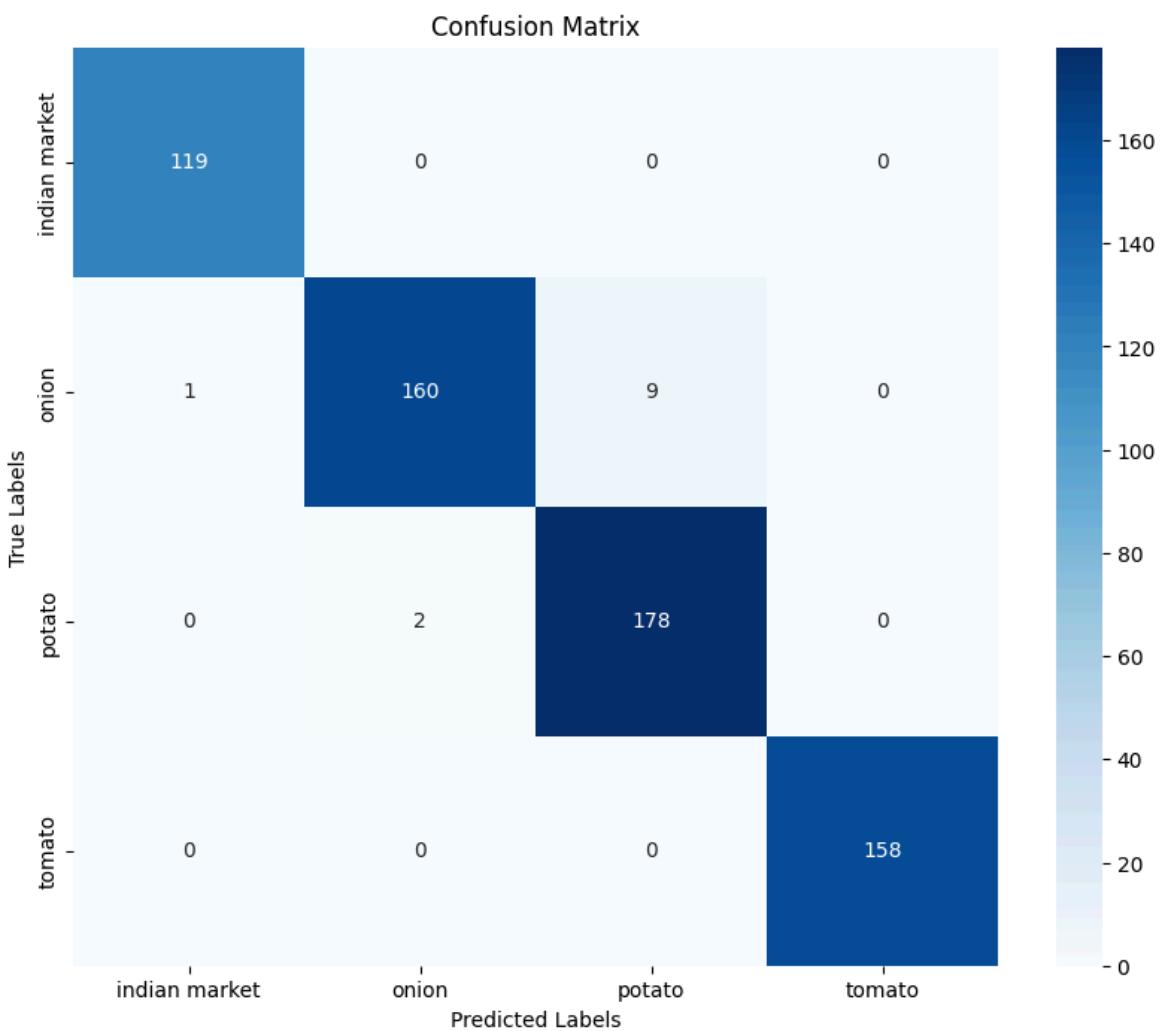
```
In [ ]: plot_training_history(after_fine_tuning)
```

```
History keys: dict_keys(['accuracy', 'loss', 'val_accuracy', 'val_loss'])
```



```
In [ ]: class_names = list(val_data.class_indices.keys())
plot_confusion_matrix(resnet50_model, val_data, class_names)
```

```
20/20 ————— 10s 327ms/step
```



Classification Report:

	precision	recall	f1-score	support
indian market	0.99	1.00	1.00	119
onion	0.99	0.94	0.96	170
potato	0.95	0.99	0.97	180
tomato	1.00	1.00	1.00	158
accuracy			0.98	627
macro avg	0.98	0.98	0.98	627
weighted avg	0.98	0.98	0.98	627

```
In [ ]: resnet50_model_finetune = tf.keras.models.load_model("/content/ninjacart_models/loss, acc = resnet50_model_finetune.evaluate(test_data, verbose=2)
print('After FineTuning')
print("Restored model, accuracy: {:.5f}%".format(100 * acc))
```

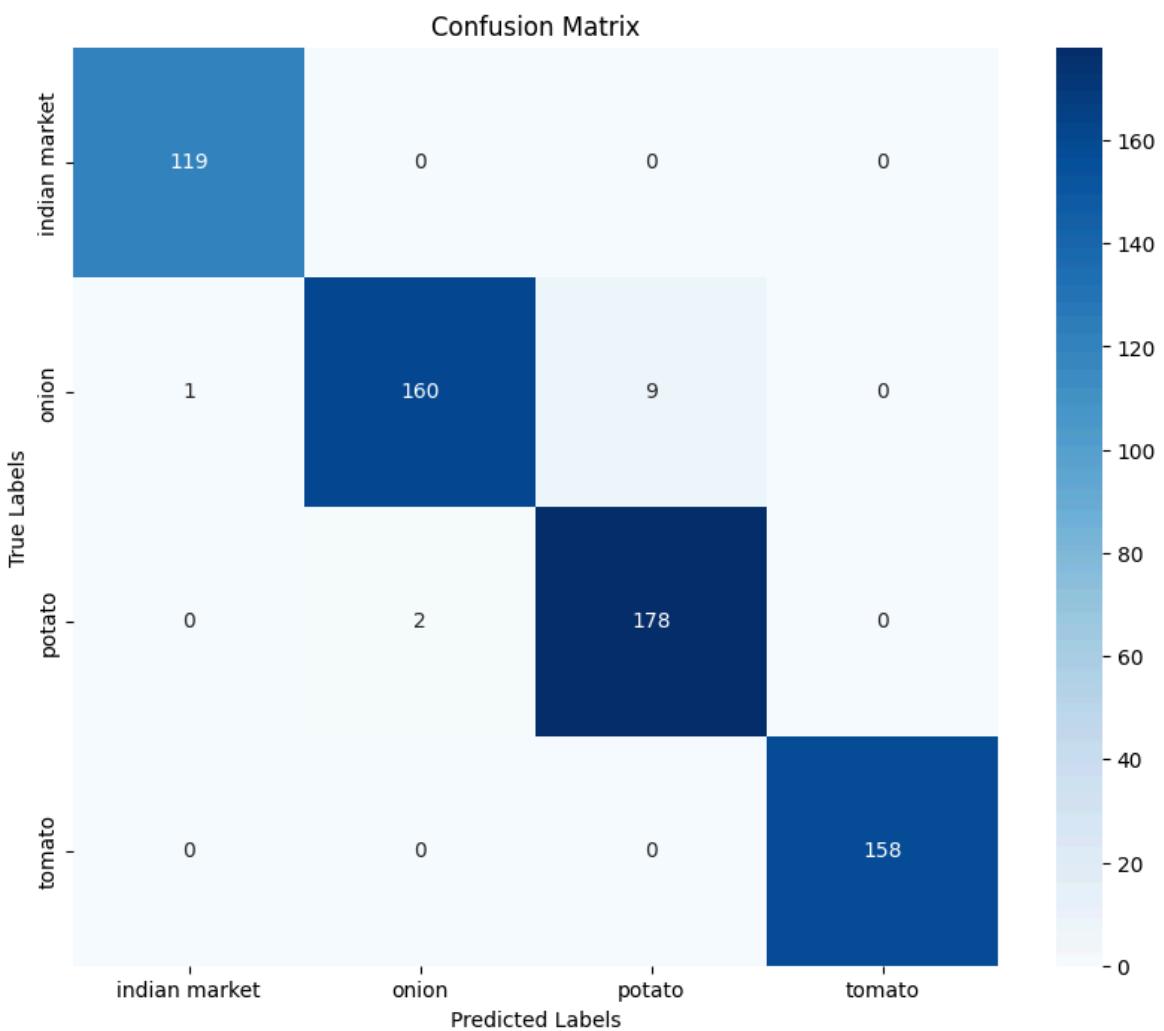
11/11 - 9s - 810ms/step - accuracy: 0.9459 - loss: 0.3997

After FineTuning

Restored model, accuracy: 94.59%

```
In [ ]: class_names = list(val_data.class_indices.keys())
plot_confusion_matrix(resnet50_model, val_data, class_names)
```

20/20 ————— 4s 182ms/step



Classification Report:

	precision	recall	f1-score	support
indian market	0.99	1.00	1.00	119
onion	0.99	0.94	0.96	170
potato	0.95	0.99	0.97	180
tomato	1.00	1.00	1.00	158
accuracy			0.98	627
macro avg	0.98	0.98	0.98	627
weighted avg	0.98	0.98	0.98	627

## MobileNet

### Before FineTuning

```
In [ ]: IMAGE_SIZE = (224, 224)
BATCH_SIZE = 32
NUM_CLASSES = 4
EPOCHS = 10

# 1. Data Preparation
train_gen = ImageDataGenerator(
    rescale=1./255,
```

```

        rotation_range=20,
        zoom_range=0.2,
        horizontal_flip=True,
    )

test_gen = ImageDataGenerator(rescale=1./255)
val_gen = ImageDataGenerator(rescale=1./255)

train_generator = train_gen.flow_from_directory(
    '/content/ninjacart_data/train',
    target_size=IMAGE_SIZE,
    batch_size=BATCH_SIZE,
    class_mode='categorical'
)

val_generator = val_gen.flow_from_directory(
    '/content/ninjacart_data/validation',
    target_size=IMAGE_SIZE,
    batch_size=BATCH_SIZE,
    class_mode='categorical',
    shuffle=False
)

test_generator = test_gen.flow_from_directory(
    '/content/ninjacart_data/test',
    target_size=IMAGE_SIZE,
    batch_size=BATCH_SIZE,
    class_mode='categorical',
    shuffle=False
)

# 2. Load Pretrained MobileNetV2
base_model = MobileNetV2(weights='imagenet', include_top=False, input_shape=(224
base_model.trainable = False # Freeze base layers initially

# 3. Add custom head
x = base_model.output
x = tf.keras.layers.GlobalAveragePooling2D()(x)
x = tf.keras.layers.Dropout(0.5)(x)
x = tf.keras.layers.Dense(256, activation='relu')(x)
output = tf.keras.layers.Dense(NUM_CLASSES, activation='softmax')(x)

model = Model(inputs=base_model.input, outputs=output)

# 4. Compile and train top Layers
model.compile(optimizer=Adam(1e-4), loss='categorical_crossentropy', metrics=['a
hist_before_ft =model.fit(train_generator, validation_data=val_generator, epochs

```

Found 2502 images belonging to 4 classes.  
 Found 627 images belonging to 4 classes.  
 Found 351 images belonging to 4 classes.  
 Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/mobilenet\_v2/mobilenet\_v2\_weights\_tf\_dim\_ordering\_tf\_kernels\_1.0\_224\_no\_top.h5  
 9406464/9406464 ————— 0s 0us/step

```

/usr/local/lib/python3.11/dist-packages/keras/src/trainers/data_adapters/py_datasets_adapter.py:121: UserWarning: Your `PyDataset` class should call `super().__init__(**kwargs)` in its constructor. `**kwargs` can include `workers`, `use_multiprocessing`, `max_queue_size`. Do not pass these arguments to `fit()`, as they will be ignored.
    self._warn_if_super_not_called()

```

```
Epoch 1/10
79/79 62s 636ms/step - accuracy: 0.5426 - loss: 1.1460 - val
_accuracy: 0.8788 - val_loss: 0.3333
Epoch 2/10
79/79 40s 511ms/step - accuracy: 0.8636 - loss: 0.4008 - val
_accuracy: 0.9123 - val_loss: 0.2318
Epoch 3/10
79/79 41s 517ms/step - accuracy: 0.9000 - loss: 0.2687 - val
_accuracy: 0.9346 - val_loss: 0.1889
Epoch 4/10
79/79 40s 513ms/step - accuracy: 0.9110 - loss: 0.2441 - val
_accuracy: 0.9362 - val_loss: 0.1736
Epoch 5/10
79/79 41s 515ms/step - accuracy: 0.9375 - loss: 0.1815 - val
_accuracy: 0.9410 - val_loss: 0.1649
Epoch 6/10
79/79 41s 522ms/step - accuracy: 0.9417 - loss: 0.1677 - val
_accuracy: 0.9506 - val_loss: 0.1501
Epoch 7/10
79/79 41s 523ms/step - accuracy: 0.9538 - loss: 0.1372 - val
_accuracy: 0.9506 - val_loss: 0.1430
Epoch 8/10
79/79 41s 524ms/step - accuracy: 0.9565 - loss: 0.1242 - val
_accuracy: 0.9522 - val_loss: 0.1379
Epoch 9/10
79/79 81s 512ms/step - accuracy: 0.9539 - loss: 0.1132 - val
_accuracy: 0.9553 - val_loss: 0.1317
Epoch 10/10
79/79 43s 539ms/step - accuracy: 0.9446 - loss: 0.1448 - val
_accuracy: 0.9537 - val_loss: 0.1324
```

In [ ]: `model.summary()`

Model: "functional\_13"

Layer (type)	Output Shape	Param #	Connected to
input_layer_13 (InputLayer)	(None, 224, 224, 3)	0	-
Conv1 (Conv2D)	(None, 112, 112, 32)	864	input_layer_13[0...]
bn_Conv1 (BatchNormalizatio...)	(None, 112, 112, 32)	128	Conv1[0][0]
Conv1_relu (ReLU)	(None, 112, 112, 32)	0	bn_Conv1[0][0]
expanded_conv_dept... (DepthwiseConv2D)	(None, 112, 112, 32)	288	Conv1_relu[0][0]
expanded_conv_dept... (BatchNormalizatio...)	(None, 112, 112, 32)	128	expanded_conv_de...
expanded_conv_dept... (ReLU)	(None, 112, 112, 32)	0	expanded_conv_de...
expanded_conv_proj... (Conv2D)	(None, 112, 112, 16)	512	expanded_conv_de...
expanded_conv_proj... (BatchNormalizatio...)	(None, 112, 112, 16)	64	expanded_conv_pr...
block_1_expand (Conv2D)	(None, 112, 112, 96)	1,536	expanded_conv_pr...
block_1_expand_BN (BatchNormalizatio...)	(None, 112, 112, 96)	384	block_1_expand[0...]
block_1_expand_relu (ReLU)	(None, 112, 112, 96)	0	block_1_expand_B...
block_1_pad (ZeroPadding2D)	(None, 113, 113, 96)	0	block_1_expand_r...
block_1_depthwise (DepthwiseConv2D)	(None, 56, 56, 96)	864	block_1_pad[0][0]
block_1_depthwise_... (BatchNormalizatio...)	(None, 56, 56, 96)	384	block_1_depthwis...
block_1_depthwise_... (ReLU)	(None, 56, 56, 96)	0	block_1_depthwis...
block_1_project (Conv2D)	(None, 56, 56, 24)	2,304	block_1_depthwis...
block_1_project_BN (BatchNormalizatio...)	(None, 56, 56, 24)	96	block_1_project[...]
block_2_expand (Conv2D)	(None, 56, 56, 144)	3,456	block_1_project[...]

block_2_expand_BN (BatchNormalizatio...)	(None, 56, 56, 144)	576	block_2_expand[0...]
block_2_expand_relu (ReLU)	(None, 56, 56, 144)	0	block_2_expand_B...
block_2_depthwise (DepthwiseConv2D)	(None, 56, 56, 144)	1,296	block_2_expand_r...
block_2_depthwise_... (BatchNormalizatio...)	(None, 56, 56, 144)	576	block_2_depthwis...
block_2_depthwise_... (ReLU)	(None, 56, 56, 144)	0	block_2_depthwis...
block_2_project (Conv2D)	(None, 56, 56, 24)	3,456	block_2_depthwis...
block_2_project_BN (BatchNormalizatio...)	(None, 56, 56, 24)	96	block_2_project[...]
block_2_add (Add)	(None, 56, 56, 24)	0	block_1_project_... block_2_project_...
block_3_expand (Conv2D)	(None, 56, 56, 144)	3,456	block_2_add[0][0]
block_3_expand_BN (BatchNormalizatio...)	(None, 56, 56, 144)	576	block_3_expand[0...]
block_3_expand_relu (ReLU)	(None, 56, 56, 144)	0	block_3_expand_B...
block_3_pad (ZeroPadding2D)	(None, 57, 57, 144)	0	block_3_expand_r...
block_3_depthwise (DepthwiseConv2D)	(None, 28, 28, 144)	1,296	block_3_pad[0][0]
block_3_depthwise_... (BatchNormalizatio...)	(None, 28, 28, 144)	576	block_3_depthwis...
block_3_depthwise_... (ReLU)	(None, 28, 28, 144)	0	block_3_depthwis...
block_3_project (Conv2D)	(None, 28, 28, 32)	4,608	block_3_depthwis...
block_3_project_BN (BatchNormalizatio...)	(None, 28, 28, 32)	128	block_3_project[...]
block_4_expand (Conv2D)	(None, 28, 28, 192)	6,144	block_3_project_...
block_4_expand_BN (BatchNormalizatio...)	(None, 28, 28, 192)	768	block_4_expand[0...]
block_4_expand_relu (ReLU)	(None, 28, 28, 192)	0	block_4_expand_B...

block_4_depthwise (DepthwiseConv2D)	(None, 28, 28, 192)	1,728	block_4_expand_r...
block_4_depthwise_BN (BatchNormalization)	(None, 28, 28, 192)	768	block_4_depthwis...
block_4_depthwise_relu (ReLU)	(None, 28, 28, 192)	0	block_4_depthwis...
block_4_project (Conv2D)	(None, 28, 28, 32)	6,144	block_4_depthwis...
block_4_project_BN (BatchNormalization)	(None, 28, 28, 32)	128	block_4_project[...
block_4_add (Add)	(None, 28, 28, 32)	0	block_3_project... block_4_project[...]
block_5_expand (Conv2D)	(None, 28, 28, 192)	6,144	block_4_add[0][0]
block_5_expand_BN (BatchNormalization)	(None, 28, 28, 192)	768	block_5_expand[0...]
block_5_expand_relu (ReLU)	(None, 28, 28, 192)	0	block_5_expand_B...
block_5_depthwise (DepthwiseConv2D)	(None, 28, 28, 192)	1,728	block_5_expand_r...
block_5_depthwise_BN (BatchNormalization)	(None, 28, 28, 192)	768	block_5_depthwis...
block_5_depthwise_relu (ReLU)	(None, 28, 28, 192)	0	block_5_depthwis...
block_5_project (Conv2D)	(None, 28, 28, 32)	6,144	block_5_depthwis...
block_5_project_BN (BatchNormalization)	(None, 28, 28, 32)	128	block_5_project[...]
block_5_add (Add)	(None, 28, 28, 32)	0	block_4_add[0][0] block_5_project[...]
block_6_expand (Conv2D)	(None, 28, 28, 192)	6,144	block_5_add[0][0]
block_6_expand_BN (BatchNormalization)	(None, 28, 28, 192)	768	block_6_expand[0...]
block_6_expand_relu (ReLU)	(None, 28, 28, 192)	0	block_6_expand_B...
block_6_pad (ZeroPadding2D)	(None, 29, 29, 192)	0	block_6_expand_r...
block_6_depthwise (DepthwiseConv2D)	(None, 14, 14, 192)	1,728	block_6_pad[0][0]

block_6_depthwise_(BatchNormalizatio... (...)	(None, 14, 14, 192)	768	block_6_depthwis...
block_6_depthwise_(ReLU) (...)	(None, 14, 14, 192)	0	block_6_depthwis...
block_6_project (Conv2D) (...)	(None, 14, 14, 64)	12,288	block_6_depthwis...
block_6_project_BN (BatchNormalizatio... (...)	(None, 14, 14, 64)	256	block_6_project[...]
block_7_expand (Conv2D) (...)	(None, 14, 14, 384)	24,576	block_6_project[...]
block_7_expand_BN (BatchNormalizatio... (...)	(None, 14, 14, 384)	1,536	block_7_expand[0...]
block_7_expand_relu (ReLU) (...)	(None, 14, 14, 384)	0	block_7_expand_B...
block_7_depthwise_(DepthwiseConv2D) (...)	(None, 14, 14, 384)	3,456	block_7_expand_r...
block_7_depthwise_(BatchNormalizatio... (...)	(None, 14, 14, 384)	1,536	block_7_depthwis...
block_7_depthwise_(ReLU) (...)	(None, 14, 14, 384)	0	block_7_depthwis...
block_7_project (Conv2D) (...)	(None, 14, 14, 64)	24,576	block_7_depthwis...
block_7_project_BN (BatchNormalizatio... (...)	(None, 14, 14, 64)	256	block_7_project[...]
block_7_add (Add) (...)	(None, 14, 14, 64)	0	block_6_project[...] block_7_project[...]
block_8_expand (Conv2D) (...)	(None, 14, 14, 384)	24,576	block_7_add[0][0]
block_8_expand_BN (BatchNormalizatio... (...)	(None, 14, 14, 384)	1,536	block_8_expand[0...]
block_8_expand_relu (ReLU) (...)	(None, 14, 14, 384)	0	block_8_expand_B...
block_8_depthwise_(DepthwiseConv2D) (...)	(None, 14, 14, 384)	3,456	block_8_expand_r...
block_8_depthwise_(BatchNormalizatio... (...)	(None, 14, 14, 384)	1,536	block_8_depthwis...
block_8_depthwise_(ReLU) (...)	(None, 14, 14, 384)	0	block_8_depthwis...
block_8_project (Conv2D) (...)	(None, 14, 14, 64)	24,576	block_8_depthwis...

block_8_project_BN (BatchNormalizatio...)	(None, 14, 14, 64)	256	block_8_project[...]
block_8_add (Add)	(None, 14, 14, 64)	0	block_7_add[0][0] + block_8_project[...]
block_9_expand (Conv2D)	(None, 14, 14, 384)	24,576	block_8_add[0][0]
block_9_expand_BN (BatchNormalizatio...)	(None, 14, 14, 384)	1,536	block_9_expand[0][0]
block_9_expand_relu (ReLU)	(None, 14, 14, 384)	0	block_9_expand_BN[0][0]
block_9_depthwise (DepthwiseConv2D)	(None, 14, 14, 384)	3,456	block_9_expand_relu[0][0]
block_9_depthwise_BN (BatchNormalizatio...)	(None, 14, 14, 384)	1,536	block_9_depthwise[0][0]
block_9_depthwise_ReLU (ReLU)	(None, 14, 14, 384)	0	block_9_depthwise_BN[0][0]
block_9_project (Conv2D)	(None, 14, 14, 64)	24,576	block_9_depthwise_ReLU[0][0]
block_9_project_BN (BatchNormalizatio...)	(None, 14, 14, 64)	256	block_9_project[0][0]
block_9_add (Add)	(None, 14, 14, 64)	0	block_8_add[0][0] + block_9_project[0][0]
block_10_expand (Conv2D)	(None, 14, 14, 384)	24,576	block_9_add[0][0]
block_10_expand_BN (BatchNormalizatio...)	(None, 14, 14, 384)	1,536	block_10_expand[0][0]
block_10_expand_relu (ReLU)	(None, 14, 14, 384)	0	block_10_expand_BN[0][0]
block_10_depthwise (DepthwiseConv2D)	(None, 14, 14, 384)	3,456	block_10_expand_relu[0][0]
block_10_depthwise_BN (BatchNormalizatio...)	(None, 14, 14, 384)	1,536	block_10_depthwise[0][0]
block_10_depthwise_ReLU (ReLU)	(None, 14, 14, 384)	0	block_10_depthwise_BN[0][0]
block_10_project (Conv2D)	(None, 14, 14, 96)	36,864	block_10_depthwise_ReLU[0][0]
block_10_project_BN (BatchNormalizatio...)	(None, 14, 14, 96)	384	block_10_project[0][0]
block_11_expand (Conv2D)	(None, 14, 14, 576)	55,296	block_10_project_BN[0][0]

block_11_expand_BN (BatchNormalizatio...)	(None, 14, 14, 576)	2,304	block_11_expand[...]
block_11_expand_re... (ReLU)	(None, 14, 14, 576)	0	block_11_expand [...]
block_11_depthwise (DepthwiseConv2D)	(None, 14, 14, 576)	5,184	block_11_expand [...]
block_11_depthwise... (BatchNormalizatio...)	(None, 14, 14, 576)	2,304	block_11_depthwi...
block_11_depthwise... (ReLU)	(None, 14, 14, 576)	0	block_11_depthwi...
block_11_project (Conv2D)	(None, 14, 14, 96)	55,296	block_11_depthwi...
block_11_project_BN (BatchNormalizatio...)	(None, 14, 14, 96)	384	block_11_project...
block_11_add (Add)	(None, 14, 14, 96)	0	block_10_project... block_11_project...
block_12_expand (Conv2D)	(None, 14, 14, 576)	55,296	block_11_add[0][...]
block_12_expand_BN (BatchNormalizatio...)	(None, 14, 14, 576)	2,304	block_12_expand[...]
block_12_expand_re... (ReLU)	(None, 14, 14, 576)	0	block_12_expand [...]
block_12_depthwise (DepthwiseConv2D)	(None, 14, 14, 576)	5,184	block_12_expand [...]
block_12_depthwise... (BatchNormalizatio...)	(None, 14, 14, 576)	2,304	block_12_depthwi...
block_12_depthwise... (ReLU)	(None, 14, 14, 576)	0	block_12_depthwi...
block_12_project (Conv2D)	(None, 14, 14, 96)	55,296	block_12_depthwi...
block_12_project_BN (BatchNormalizatio...)	(None, 14, 14, 96)	384	block_12_project...
block_12_add (Add)	(None, 14, 14, 96)	0	block_11_add[0][...] block_12_project...
block_13_expand (Conv2D)	(None, 14, 14, 576)	55,296	block_12_add[0][...]
block_13_expand_BN (BatchNormalizatio...)	(None, 14, 14, 576)	2,304	block_13_expand[...]
block_13_expand_re... (ReLU)	(None, 14, 14, 576)	0	block_13_expand [...]

block_13_pad (ZeroPadding2D)	(None, 15, 15, 576)	0	block_13_expand_...
block_13_depthwise (DepthwiseConv2D)	(None, 7, 7, 576)	5,184	block_13_pad[0][...]
block_13_depthwise... (BatchNormalizatio...)	(None, 7, 7, 576)	2,304	block_13_depthwi...
block_13_depthwise... (ReLU)	(None, 7, 7, 576)	0	block_13_depthwi...
block_13_project (Conv2D)	(None, 7, 7, 160)	92,160	block_13_depthwi...
block_13_project_BN (BatchNormalizatio...)	(None, 7, 7, 160)	640	block_13_project...
block_14_expand (Conv2D)	(None, 7, 7, 960)	153,600	block_13_project...
block_14_expand_BN (BatchNormalizatio...)	(None, 7, 7, 960)	3,840	block_14_expand[...]
block_14_expand_re... (ReLU)	(None, 7, 7, 960)	0	block_14_expand_...
block_14_depthwise (DepthwiseConv2D)	(None, 7, 7, 960)	8,640	block_14_expand_...
block_14_depthwise... (BatchNormalizatio...)	(None, 7, 7, 960)	3,840	block_14_depthwi...
block_14_depthwise... (ReLU)	(None, 7, 7, 960)	0	block_14_depthwi...
block_14_project (Conv2D)	(None, 7, 7, 160)	153,600	block_14_depthwi...
block_14_project_BN (BatchNormalizatio...)	(None, 7, 7, 160)	640	block_14_project...
block_14_add (Add)	(None, 7, 7, 160)	0	block_13_project... block_14_project...
block_15_expand (Conv2D)	(None, 7, 7, 960)	153,600	block_14_add[0][...]
block_15_expand_BN (BatchNormalizatio...)	(None, 7, 7, 960)	3,840	block_15_expand[...]
block_15_expand_re... (ReLU)	(None, 7, 7, 960)	0	block_15_expand_...
block_15_depthwise (DepthwiseConv2D)	(None, 7, 7, 960)	8,640	block_15_expand_...
block_15_depthwise... (BatchNormalizatio...)	(None, 7, 7, 960)	3,840	block_15_depthwi...

block_15_depthwise... (ReLU)	(None, 7, 7, 960)	0	block_15_depthwi...
block_15_project (Conv2D)	(None, 7, 7, 160)	153,600	block_15_depthwi...
block_15_project_BN (BatchNormalizatio...	(None, 7, 7, 160)	640	block_15_project...
block_15_add (Add)	(None, 7, 7, 160)	0	block_14_add[0][...] block_15_project...
block_16_expand (Conv2D)	(None, 7, 7, 960)	153,600	block_15_add[0][...]
block_16_expand_BN (BatchNormalizatio...	(None, 7, 7, 960)	3,840	block_16_expand[ ...]
block_16_expand_re... (ReLU)	(None, 7, 7, 960)	0	block_16_expand[...]
block_16_depthwise (DepthwiseConv2D)	(None, 7, 7, 960)	8,640	block_16_expand[...]
block_16_depthwise... (BatchNormalizatio...	(None, 7, 7, 960)	3,840	block_16_depthwi...
block_16_depthwise... (ReLU)	(None, 7, 7, 960)	0	block_16_depthwi...
block_16_project (Conv2D)	(None, 7, 7, 320)	307,200	block_16_depthwi...
block_16_project_BN (BatchNormalizatio...	(None, 7, 7, 320)	1,280	block_16_project...
Conv_1 (Conv2D)	(None, 7, 7, 1280)	409,600	block_16_project...
Conv_1_bn (BatchNormalizatio...	(None, 7, 7, 1280)	5,120	Conv_1[0][0]
out_relu (ReLU)	(None, 7, 7, 1280)	0	Conv_1_bn[0][0]
global_average_poo... (GlobalAveragePool...	(None, 1280)	0	out_relu[0][0]
dropout_8 (Dropout)	(None, 1280)	0	global_average_p...
dense_20 (Dense)	(None, 256)	327,936	dropout_8[0][0]
dense_21 (Dense)	(None, 4)	1,028	dense_20[0][0]

Total params: 3,244,878 (12.38 MB)

Trainable params: 328,964 (1.25 MB)

Non-trainable params: 2,257,984 (8.61 MB)

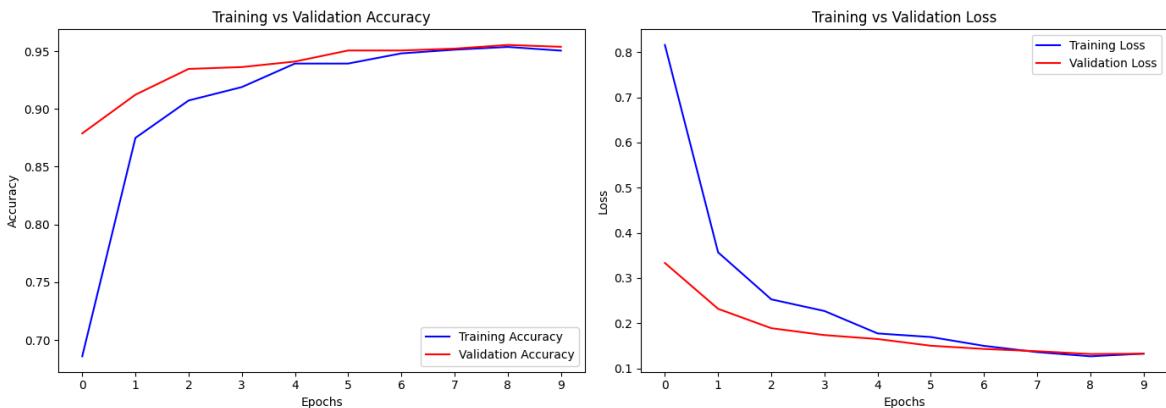
```
Optimizer params: 657,930 (2.51 MB)
```

```
In [ ]: create_directory_if_not_exists('/content/ninjacart_models/model11_mobilenet_GAP_resnet50_model.save("/content/ninjacart_models/model11_mobilenet_GAP_DP_FineTune
```

```
Created directory: /content/ninjacart_models/model11_mobilenet_GAP_DP_FineTune
```

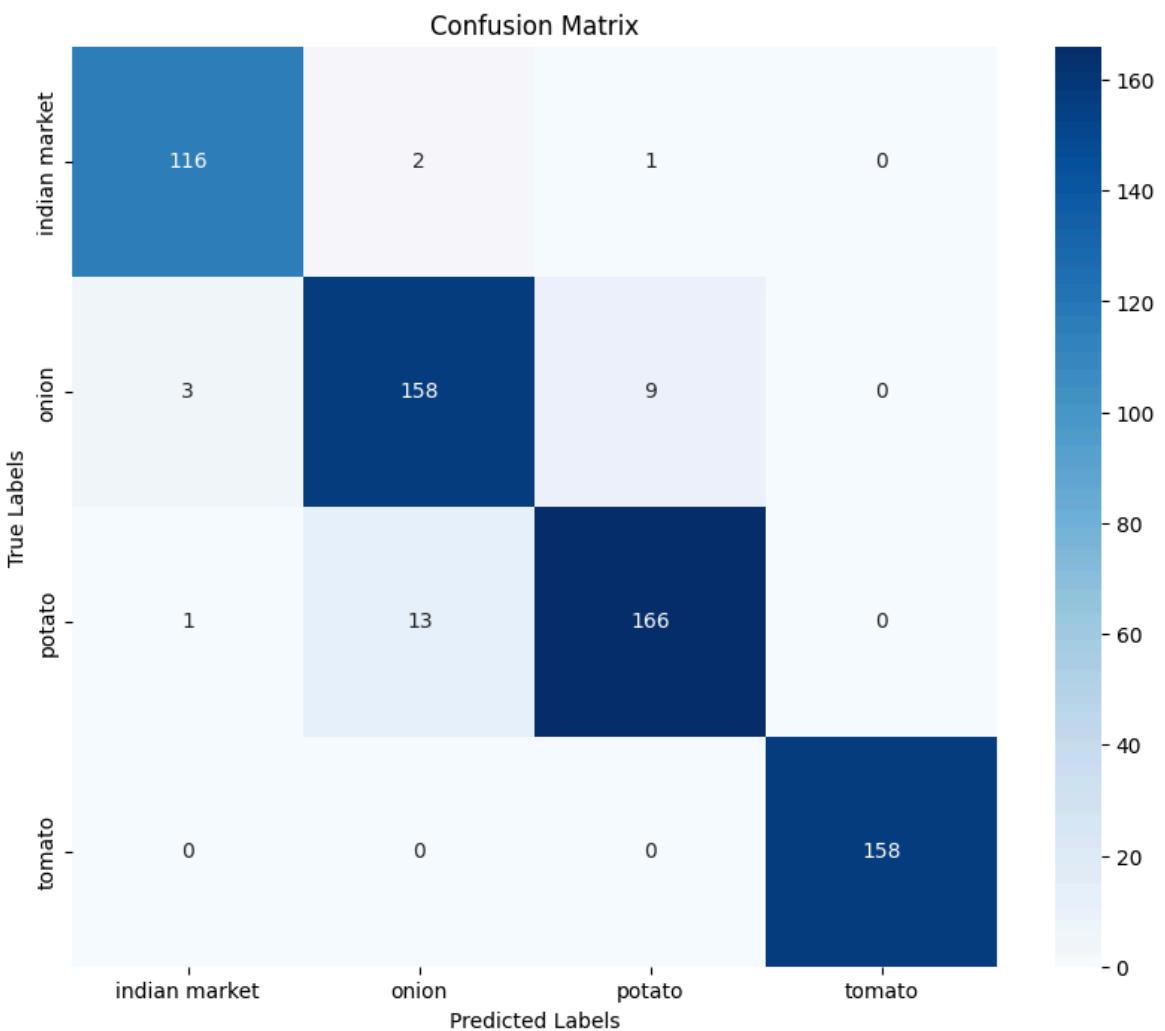
```
In [ ]: plot_training_history(hist_before_ft)
```

```
History keys: dict_keys(['accuracy', 'loss', 'val_accuracy', 'val_loss'])
```



```
In [ ]: class_names = list(val_generator.class_indices.keys())
plot_confusion_matrix(model, val_generator, class_names)
```

```
20/20 ━━━━━━━━ 10s 334ms/step
```



### Classification Report:

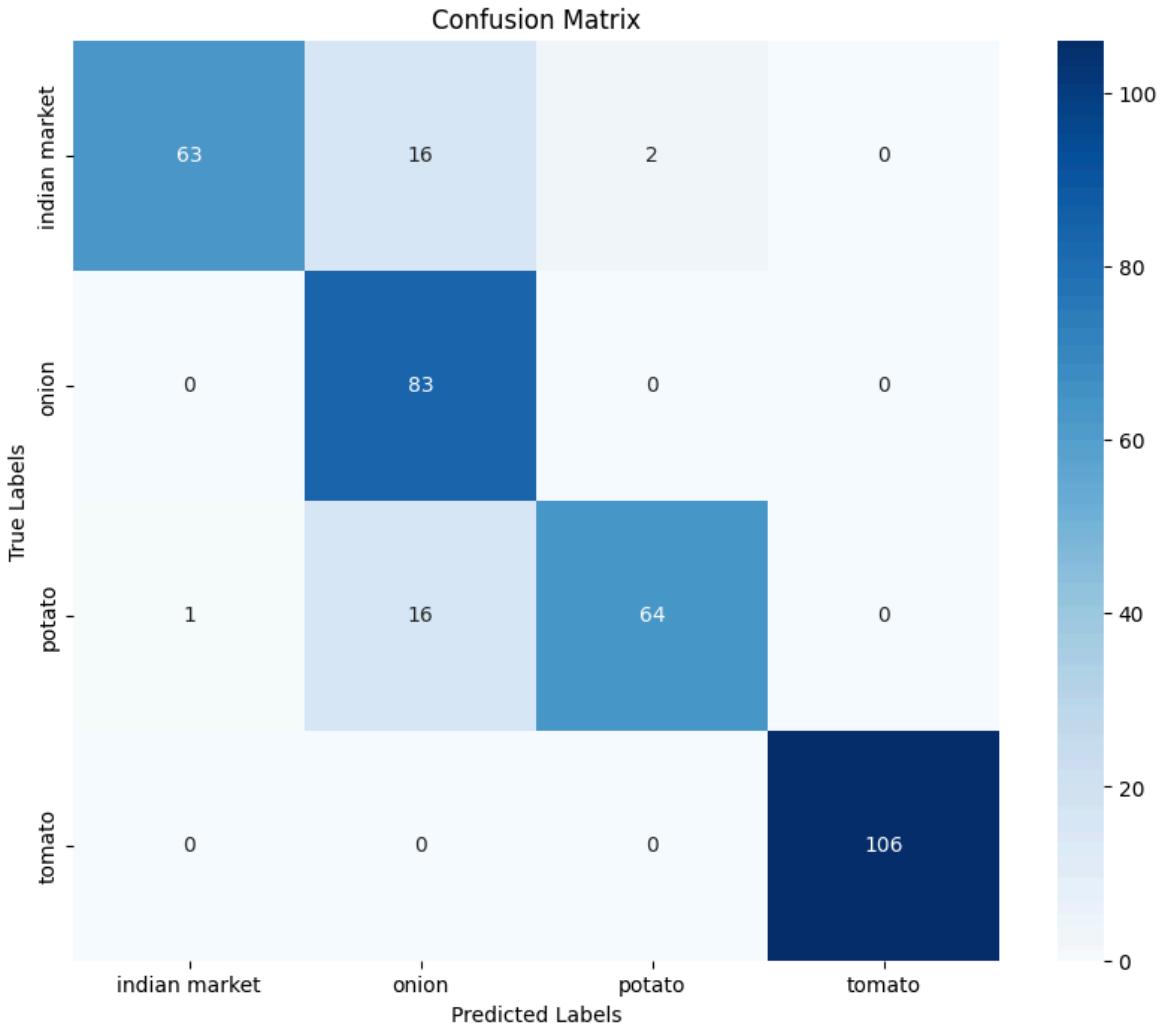
	precision	recall	f1-score	support
indian market	0.97	0.97	0.97	119
onion	0.91	0.93	0.92	170
potato	0.94	0.92	0.93	180
tomato	1.00	1.00	1.00	158
accuracy			0.95	627
macro avg	0.96	0.96	0.96	627
weighted avg	0.95	0.95	0.95	627

```
In [ ]: mobilenetv2_model_finetune = tf.keras.models.load_model("/content/ninjacart_mode
loss, acc = mobilenetv2_model_finetune.evaluate(test_generator, verbose=2)
print('After FineTuning')
print("Restored model, accuracy: {:.5.2f}%".format(100 * acc))
```

```
/usr/local/lib/python3.11/dist-packages/keras/src/trainers/data_adapters/py_datal
et_adapter.py:121: UserWarning: Your `PyDataset` class should call `super().__init__(**kwargs)` in its constructor. `**kwargs` can include `workers`, `use_multipro
cessing`, `max_queue_size`. Do not pass these arguments to `fit()`, as they will
be ignored.
    self._warn_if_super_not_called()
11/11 - 9s - 781ms/step - accuracy: 0.2308 - loss: 9.6986
After FineTuning
Restored model, accuracy: 23.08%
```

```
In [ ]: class_names = list(val_generator.class_indices.keys())
plot_confusion_matrix(model, test_generator, class_names)
```

```
11/11 ━━━━━━━━━━ 6s 568ms/step
```



Classification Report:

	precision	recall	f1-score	support
indian market	0.98	0.78	0.87	81
onion	0.72	1.00	0.84	83
potato	0.97	0.79	0.87	81
tomato	1.00	1.00	1.00	106
accuracy			0.90	351
macro avg	0.92	0.89	0.89	351
weighted avg	0.92	0.90	0.90	351

## After Finetuning

```
In [ ]: # 5. Unfreeze some deeper Layers for fine-tuning
base_model.trainable = True
for layer in base_model.layers[:100]:
    layer.trainable = False

model.compile(optimizer=Adam(1e-5), loss='categorical_crossentropy', metrics=['a
history = model.fit(train_generator, validation_data=val_generator, epochs=EPOCH
```

```
Epoch 1/10
79/79 83s 708ms/step - accuracy: 0.9698 - loss: 0.0861 - val
_accuracy: 0.9601 - val_loss: 0.1119
Epoch 2/10
79/79 56s 492ms/step - accuracy: 0.9687 - loss: 0.0919 - val
_accuracy: 0.9601 - val_loss: 0.1110
Epoch 3/10
79/79 39s 497ms/step - accuracy: 0.9790 - loss: 0.0671 - val
_accuracy: 0.9601 - val_loss: 0.1106
Epoch 4/10
79/79 39s 489ms/step - accuracy: 0.9654 - loss: 0.0738 - val
_accuracy: 0.9617 - val_loss: 0.1093
Epoch 5/10
79/79 39s 490ms/step - accuracy: 0.9810 - loss: 0.0590 - val
_accuracy: 0.9617 - val_loss: 0.1109
Epoch 6/10
79/79 39s 490ms/step - accuracy: 0.9818 - loss: 0.0557 - val
_accuracy: 0.9633 - val_loss: 0.1093
Epoch 7/10
79/79 38s 487ms/step - accuracy: 0.9840 - loss: 0.0536 - val
_accuracy: 0.9569 - val_loss: 0.1091
Epoch 8/10
79/79 38s 474ms/step - accuracy: 0.9822 - loss: 0.0575 - val
_accuracy: 0.9601 - val_loss: 0.1079
Epoch 9/10
79/79 38s 483ms/step - accuracy: 0.9848 - loss: 0.0489 - val
_accuracy: 0.9617 - val_loss: 0.1132
Epoch 10/10
79/79 38s 487ms/step - accuracy: 0.9906 - loss: 0.0386 - val
_accuracy: 0.9617 - val_loss: 0.1106
```

In [ ]: `model.summary()`

Model: "functional\_13"

Layer (type)	Output Shape	Param #	Connected to
input_layer_13 (InputLayer)	(None, 224, 224, 3)	0	-
Conv1 (Conv2D)	(None, 112, 112, 32)	864	input_layer_13[0...]
bn_Conv1 (BatchNormalizatio...)	(None, 112, 112, 32)	128	Conv1[0][0]
Conv1_relu (ReLU)	(None, 112, 112, 32)	0	bn_Conv1[0][0]
expanded_conv_dept... (DepthwiseConv2D)	(None, 112, 112, 32)	288	Conv1_relu[0][0]
expanded_conv_dept... (BatchNormalizatio...)	(None, 112, 112, 32)	128	expanded_conv_de...
expanded_conv_dept... (ReLU)	(None, 112, 112, 32)	0	expanded_conv_de...
expanded_conv_proj... (Conv2D)	(None, 112, 112, 16)	512	expanded_conv_de...
expanded_conv_proj... (BatchNormalizatio...)	(None, 112, 112, 16)	64	expanded_conv_pr...
block_1_expand (Conv2D)	(None, 112, 112, 96)	1,536	expanded_conv_pr...
block_1_expand_BN (BatchNormalizatio...)	(None, 112, 112, 96)	384	block_1_expand[0...]
block_1_expand_relu (ReLU)	(None, 112, 112, 96)	0	block_1_expand_B...
block_1_pad (ZeroPadding2D)	(None, 113, 113, 96)	0	block_1_expand_r...
block_1_depthwise (DepthwiseConv2D)	(None, 56, 56, 96)	864	block_1_pad[0][0]
block_1_depthwise_... (BatchNormalizatio...)	(None, 56, 56, 96)	384	block_1_depthwis...
block_1_depthwise_... (ReLU)	(None, 56, 56, 96)	0	block_1_depthwis...
block_1_project (Conv2D)	(None, 56, 56, 24)	2,304	block_1_depthwis...
block_1_project_BN (BatchNormalizatio...)	(None, 56, 56, 24)	96	block_1_project[...]
block_2_expand (Conv2D)	(None, 56, 56, 144)	3,456	block_1_project[...]

block_2_expand_BN (BatchNormalizatio...)	(None, 56, 56, 144)	576	block_2_expand[0...]
block_2_expand_relu (ReLU)	(None, 56, 56, 144)	0	block_2_expand_B...
block_2_depthwise (DepthwiseConv2D)	(None, 56, 56, 144)	1,296	block_2_expand_r...
block_2_depthwise_... (BatchNormalizatio...)	(None, 56, 56, 144)	576	block_2_depthwis...
block_2_depthwise_... (ReLU)	(None, 56, 56, 144)	0	block_2_depthwis...
block_2_project (Conv2D)	(None, 56, 56, 24)	3,456	block_2_depthwis...
block_2_project_BN (BatchNormalizatio...)	(None, 56, 56, 24)	96	block_2_project[...]
block_2_add (Add)	(None, 56, 56, 24)	0	block_1_project_... block_2_project_...
block_3_expand (Conv2D)	(None, 56, 56, 144)	3,456	block_2_add[0][0]
block_3_expand_BN (BatchNormalizatio...)	(None, 56, 56, 144)	576	block_3_expand[0...]
block_3_expand_relu (ReLU)	(None, 56, 56, 144)	0	block_3_expand_B...
block_3_pad (ZeroPadding2D)	(None, 57, 57, 144)	0	block_3_expand_r...
block_3_depthwise (DepthwiseConv2D)	(None, 28, 28, 144)	1,296	block_3_pad[0][0]
block_3_depthwise_... (BatchNormalizatio...)	(None, 28, 28, 144)	576	block_3_depthwis...
block_3_depthwise_... (ReLU)	(None, 28, 28, 144)	0	block_3_depthwis...
block_3_project (Conv2D)	(None, 28, 28, 32)	4,608	block_3_depthwis...
block_3_project_BN (BatchNormalizatio...)	(None, 28, 28, 32)	128	block_3_project[...]
block_4_expand (Conv2D)	(None, 28, 28, 192)	6,144	block_3_project_...
block_4_expand_BN (BatchNormalizatio...)	(None, 28, 28, 192)	768	block_4_expand[0...]
block_4_expand_relu (ReLU)	(None, 28, 28, 192)	0	block_4_expand_B...

block_4_depthwise (DepthwiseConv2D)	(None, 28, 28, 192)	1,728	block_4_expand_r...
block_4_depthwise_BN (BatchNormalization)	(None, 28, 28, 192)	768	block_4_depthwis...
block_4_depthwise_relu (ReLU)	(None, 28, 28, 192)	0	block_4_depthwis...
block_4_project (Conv2D)	(None, 28, 28, 32)	6,144	block_4_depthwis...
block_4_project_BN (BatchNormalization)	(None, 28, 28, 32)	128	block_4_project[...
block_4_add (Add)	(None, 28, 28, 32)	0	block_3_project... block_4_project[...]
block_5_expand (Conv2D)	(None, 28, 28, 192)	6,144	block_4_add[0][0]
block_5_expand_BN (BatchNormalization)	(None, 28, 28, 192)	768	block_5_expand[0...]
block_5_expand_relu (ReLU)	(None, 28, 28, 192)	0	block_5_expand_B...
block_5_depthwise (DepthwiseConv2D)	(None, 28, 28, 192)	1,728	block_5_expand_r...
block_5_depthwise_BN (BatchNormalization)	(None, 28, 28, 192)	768	block_5_depthwis...
block_5_depthwise_relu (ReLU)	(None, 28, 28, 192)	0	block_5_depthwis...
block_5_project (Conv2D)	(None, 28, 28, 32)	6,144	block_5_depthwis...
block_5_project_BN (BatchNormalization)	(None, 28, 28, 32)	128	block_5_project[...]
block_5_add (Add)	(None, 28, 28, 32)	0	block_4_add[0][0] block_5_project[...]
block_6_expand (Conv2D)	(None, 28, 28, 192)	6,144	block_5_add[0][0]
block_6_expand_BN (BatchNormalization)	(None, 28, 28, 192)	768	block_6_expand[0...]
block_6_expand_relu (ReLU)	(None, 28, 28, 192)	0	block_6_expand_B...
block_6_pad (ZeroPadding2D)	(None, 29, 29, 192)	0	block_6_expand_r...
block_6_depthwise (DepthwiseConv2D)	(None, 14, 14, 192)	1,728	block_6_pad[0][0]

block_6_depthwise_(BatchNormalizatio... (...)	(None, 14, 14, 192)	768	block_6_depthwis...
block_6_depthwise_(ReLU) (...)	(None, 14, 14, 192)	0	block_6_depthwis...
block_6_project (Conv2D) (...)	(None, 14, 14, 64)	12,288	block_6_depthwis...
block_6_project_BN (BatchNormalizatio... (...)	(None, 14, 14, 64)	256	block_6_project[...]
block_7_expand (Conv2D) (...)	(None, 14, 14, 384)	24,576	block_6_project[...]
block_7_expand_BN (BatchNormalizatio... (...)	(None, 14, 14, 384)	1,536	block_7_expand[0...]
block_7_expand_relu (ReLU) (...)	(None, 14, 14, 384)	0	block_7_expand_B...
block_7_depthwise_(DepthwiseConv2D) (...)	(None, 14, 14, 384)	3,456	block_7_expand_r...
block_7_depthwise_(BatchNormalizatio... (...)	(None, 14, 14, 384)	1,536	block_7_depthwis...
block_7_depthwise_(ReLU) (...)	(None, 14, 14, 384)	0	block_7_depthwis...
block_7_project (Conv2D) (...)	(None, 14, 14, 64)	24,576	block_7_depthwis...
block_7_project_BN (BatchNormalizatio... (...)	(None, 14, 14, 64)	256	block_7_project[...]
block_7_add (Add) (...)	(None, 14, 14, 64)	0	block_6_project[...] block_7_project[...]
block_8_expand (Conv2D) (...)	(None, 14, 14, 384)	24,576	block_7_add[0][0]
block_8_expand_BN (BatchNormalizatio... (...)	(None, 14, 14, 384)	1,536	block_8_expand[0...]
block_8_expand_relu (ReLU) (...)	(None, 14, 14, 384)	0	block_8_expand_B...
block_8_depthwise_(DepthwiseConv2D) (...)	(None, 14, 14, 384)	3,456	block_8_expand_r...
block_8_depthwise_(BatchNormalizatio... (...)	(None, 14, 14, 384)	1,536	block_8_depthwis...
block_8_depthwise_(ReLU) (...)	(None, 14, 14, 384)	0	block_8_depthwis...
block_8_project (Conv2D) (...)	(None, 14, 14, 64)	24,576	block_8_depthwis...

block_8_project_BN (BatchNormalizatio...)	(None, 14, 14, 64)	256	block_8_project[...]
block_8_add (Add)	(None, 14, 14, 64)	0	block_7_add[0][0] + block_8_project[...]
block_9_expand (Conv2D)	(None, 14, 14, 384)	24,576	block_8_add[0][0]
block_9_expand_BN (BatchNormalizatio...)	(None, 14, 14, 384)	1,536	block_9_expand[0][0]
block_9_expand_relu (ReLU)	(None, 14, 14, 384)	0	block_9_expand_BN[0][0]
block_9_depthwise (DepthwiseConv2D)	(None, 14, 14, 384)	3,456	block_9_expand_relu[0][0]
block_9_depthwise_BN (BatchNormalizatio...)	(None, 14, 14, 384)	1,536	block_9_depthwise[0][0]
block_9_depthwise_ReLU (ReLU)	(None, 14, 14, 384)	0	block_9_depthwise_BN[0][0]
block_9_project (Conv2D)	(None, 14, 14, 64)	24,576	block_9_depthwise_ReLU[0][0]
block_9_project_BN (BatchNormalizatio...)	(None, 14, 14, 64)	256	block_9_project[0][0]
block_9_add (Add)	(None, 14, 14, 64)	0	block_8_add[0][0] + block_9_project[0][0]
block_10_expand (Conv2D)	(None, 14, 14, 384)	24,576	block_9_add[0][0]
block_10_expand_BN (BatchNormalizatio...)	(None, 14, 14, 384)	1,536	block_10_expand[0][0]
block_10_expand_relu (ReLU)	(None, 14, 14, 384)	0	block_10_expand_BN[0][0]
block_10_depthwise (DepthwiseConv2D)	(None, 14, 14, 384)	3,456	block_10_expand_relu[0][0]
block_10_depthwise_BN (BatchNormalizatio...)	(None, 14, 14, 384)	1,536	block_10_depthwise[0][0]
block_10_depthwise_ReLU (ReLU)	(None, 14, 14, 384)	0	block_10_depthwise_BN[0][0]
block_10_project (Conv2D)	(None, 14, 14, 96)	36,864	block_10_depthwise_ReLU[0][0]
block_10_project_BN (BatchNormalizatio...)	(None, 14, 14, 96)	384	block_10_project[0][0]
block_11_expand (Conv2D)	(None, 14, 14, 576)	55,296	block_10_project_BN[0][0]

block_11_expand_BN (BatchNormalizatio...)	(None, 14, 14, 576)	2,304	block_11_expand[...]
block_11_expand_re... (ReLU)	(None, 14, 14, 576)	0	block_11_expand [...]
block_11_depthwise (DepthwiseConv2D)	(None, 14, 14, 576)	5,184	block_11_expand [...]
block_11_depthwise... (BatchNormalizatio...)	(None, 14, 14, 576)	2,304	block_11_depthwi...
block_11_depthwise... (ReLU)	(None, 14, 14, 576)	0	block_11_depthwi...
block_11_project (Conv2D)	(None, 14, 14, 96)	55,296	block_11_depthwi...
block_11_project_BN (BatchNormalizatio...)	(None, 14, 14, 96)	384	block_11_project...
block_11_add (Add)	(None, 14, 14, 96)	0	block_10_project... block_11_project...
block_12_expand (Conv2D)	(None, 14, 14, 576)	55,296	block_11_add[0][...]
block_12_expand_BN (BatchNormalizatio...)	(None, 14, 14, 576)	2,304	block_12_expand[...]
block_12_expand_re... (ReLU)	(None, 14, 14, 576)	0	block_12_expand [...]
block_12_depthwise (DepthwiseConv2D)	(None, 14, 14, 576)	5,184	block_12_expand [...]
block_12_depthwise... (BatchNormalizatio...)	(None, 14, 14, 576)	2,304	block_12_depthwi...
block_12_depthwise... (ReLU)	(None, 14, 14, 576)	0	block_12_depthwi...
block_12_project (Conv2D)	(None, 14, 14, 96)	55,296	block_12_depthwi...
block_12_project_BN (BatchNormalizatio...)	(None, 14, 14, 96)	384	block_12_project...
block_12_add (Add)	(None, 14, 14, 96)	0	block_11_add[0][...] block_12_project...
block_13_expand (Conv2D)	(None, 14, 14, 576)	55,296	block_12_add[0][...]
block_13_expand_BN (BatchNormalizatio...)	(None, 14, 14, 576)	2,304	block_13_expand[...]
block_13_expand_re... (ReLU)	(None, 14, 14, 576)	0	block_13_expand [...]

block_13_pad (ZeroPadding2D)	(None, 15, 15, 576)	0	block_13_expand_...
block_13_depthwise (DepthwiseConv2D)	(None, 7, 7, 576)	5,184	block_13_pad[0][...]
block_13_depthwise... (BatchNormalizatio...)	(None, 7, 7, 576)	2,304	block_13_depthwi...
block_13_depthwise... (ReLU)	(None, 7, 7, 576)	0	block_13_depthwi...
block_13_project (Conv2D)	(None, 7, 7, 160)	92,160	block_13_depthwi...
block_13_project_BN (BatchNormalizatio...)	(None, 7, 7, 160)	640	block_13_project...
block_14_expand (Conv2D)	(None, 7, 7, 960)	153,600	block_13_project...
block_14_expand_BN (BatchNormalizatio...)	(None, 7, 7, 960)	3,840	block_14_expand[...]
block_14_expand_re... (ReLU)	(None, 7, 7, 960)	0	block_14_expand_...
block_14_depthwise (DepthwiseConv2D)	(None, 7, 7, 960)	8,640	block_14_expand_...
block_14_depthwise... (BatchNormalizatio...)	(None, 7, 7, 960)	3,840	block_14_depthwi...
block_14_depthwise... (ReLU)	(None, 7, 7, 960)	0	block_14_depthwi...
block_14_project (Conv2D)	(None, 7, 7, 160)	153,600	block_14_depthwi...
block_14_project_BN (BatchNormalizatio...)	(None, 7, 7, 160)	640	block_14_project...
block_14_add (Add)	(None, 7, 7, 160)	0	block_13_project... block_14_project...
block_15_expand (Conv2D)	(None, 7, 7, 960)	153,600	block_14_add[0][...]
block_15_expand_BN (BatchNormalizatio...)	(None, 7, 7, 960)	3,840	block_15_expand[...]
block_15_expand_re... (ReLU)	(None, 7, 7, 960)	0	block_15_expand_...
block_15_depthwise (DepthwiseConv2D)	(None, 7, 7, 960)	8,640	block_15_expand_...
block_15_depthwise... (BatchNormalizatio...)	(None, 7, 7, 960)	3,840	block_15_depthwi...

block_15_depthwise... (ReLU)	(None, 7, 7, 960)	0	block_15_depthwi...
block_15_project (Conv2D)	(None, 7, 7, 160)	153,600	block_15_depthwi...
block_15_project_BN (BatchNormalizatio...	(None, 7, 7, 160)	640	block_15_project...
block_15_add (Add)	(None, 7, 7, 160)	0	block_14_add[0][...] block_15_project...
block_16_expand (Conv2D)	(None, 7, 7, 960)	153,600	block_15_add[0][...]
block_16_expand_BN (BatchNormalizatio...	(None, 7, 7, 960)	3,840	block_16_expand[ ...]
block_16_expand_re... (ReLU)	(None, 7, 7, 960)	0	block_16_expand[...]
block_16_depthwise (DepthwiseConv2D)	(None, 7, 7, 960)	8,640	block_16_expand[...]
block_16_depthwise... (BatchNormalizatio...	(None, 7, 7, 960)	3,840	block_16_depthwi...
block_16_depthwise... (ReLU)	(None, 7, 7, 960)	0	block_16_depthwi...
block_16_project (Conv2D)	(None, 7, 7, 320)	307,200	block_16_depthwi...
block_16_project_BN (BatchNormalizatio...	(None, 7, 7, 320)	1,280	block_16_project...
Conv_1 (Conv2D)	(None, 7, 7, 1280)	409,600	block_16_project...
Conv_1_bn (BatchNormalizatio...	(None, 7, 7, 1280)	5,120	Conv_1[0][0]
out_relu (ReLU)	(None, 7, 7, 1280)	0	Conv_1_bn[0][0]
global_average_poo... (GlobalAveragePool...	(None, 1280)	0	out_relu[0][0]
dropout_8 (Dropout)	(None, 1280)	0	global_average_p...
dense_20 (Dense)	(None, 256)	327,936	dropout_8[0][0]
dense_21 (Dense)	(None, 4)	1,028	dense_20[0][0]

Total params: 6,967,758 (26.58 MB)

Trainable params: 2,190,404 (8.36 MB)

Non-trainable params: 396,544 (1.51 MB)

```
Optimizer params: 4,380,810 (16.71 MB)
```

```
In [ ]: create_directory_if_not_exists('/content/ninjacart_models/model11_mobilenet_GAP_')
model.save("/content/ninjacart_models/model11_mobilenet_GAP_DP_FineTune/model11_
```

```
Directory already exists: /content/ninjacart_models/model11_mobilenet_GAP_DP_Fine
Tune
```

```
In [ ]: def plot_training_history(history):
    print("History keys:", history.history.keys()) # This will work now

    acc = history.history.get('accuracy', [])
    val_acc = history.history.get('val_accuracy', [])
    loss = history.history.get('loss', [])
    val_loss = history.history.get('val_loss', [])

    import matplotlib.pyplot as plt

    epochs = range(1, len(acc) + 1)

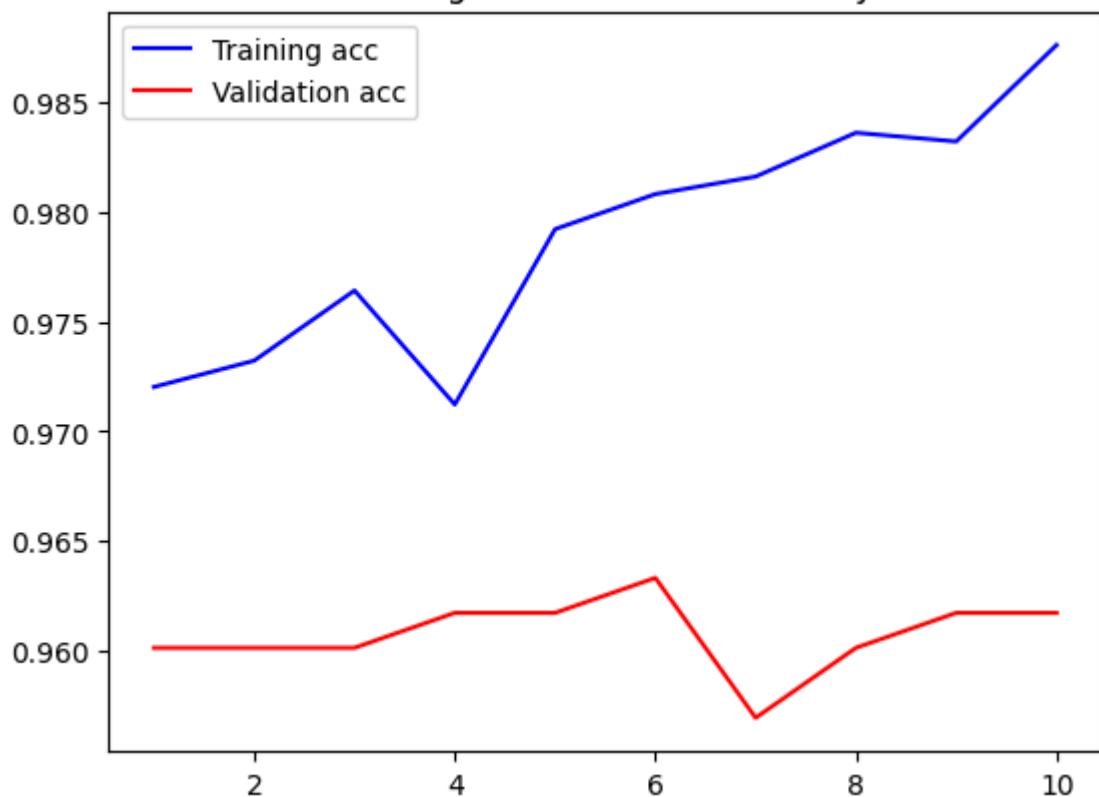
    plt.plot(epochs, acc, 'b', label='Training acc')
    plt.plot(epochs, val_acc, 'r', label='Validation acc')
    plt.title('Training and validation accuracy')
    plt.legend()
    plt.figure()

    plt.plot(epochs, loss, 'b', label='Training loss')
    plt.plot(epochs, val_loss, 'r', label='Validation loss')
    plt.title('Training and validation loss')
    plt.legend()
    plt.show()
```

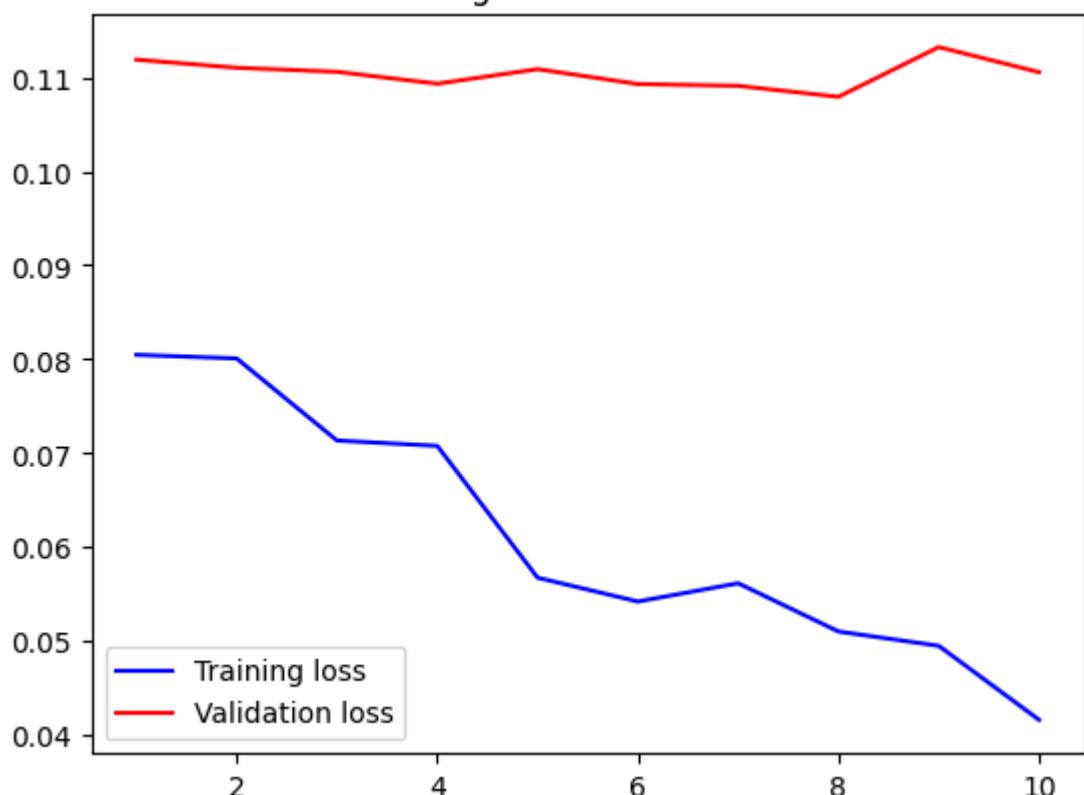
```
In [ ]: plot_training_history(history)
```

```
History keys: dict_keys(['accuracy', 'loss', 'val_accuracy', 'val_loss'])
```

### Training and validation accuracy

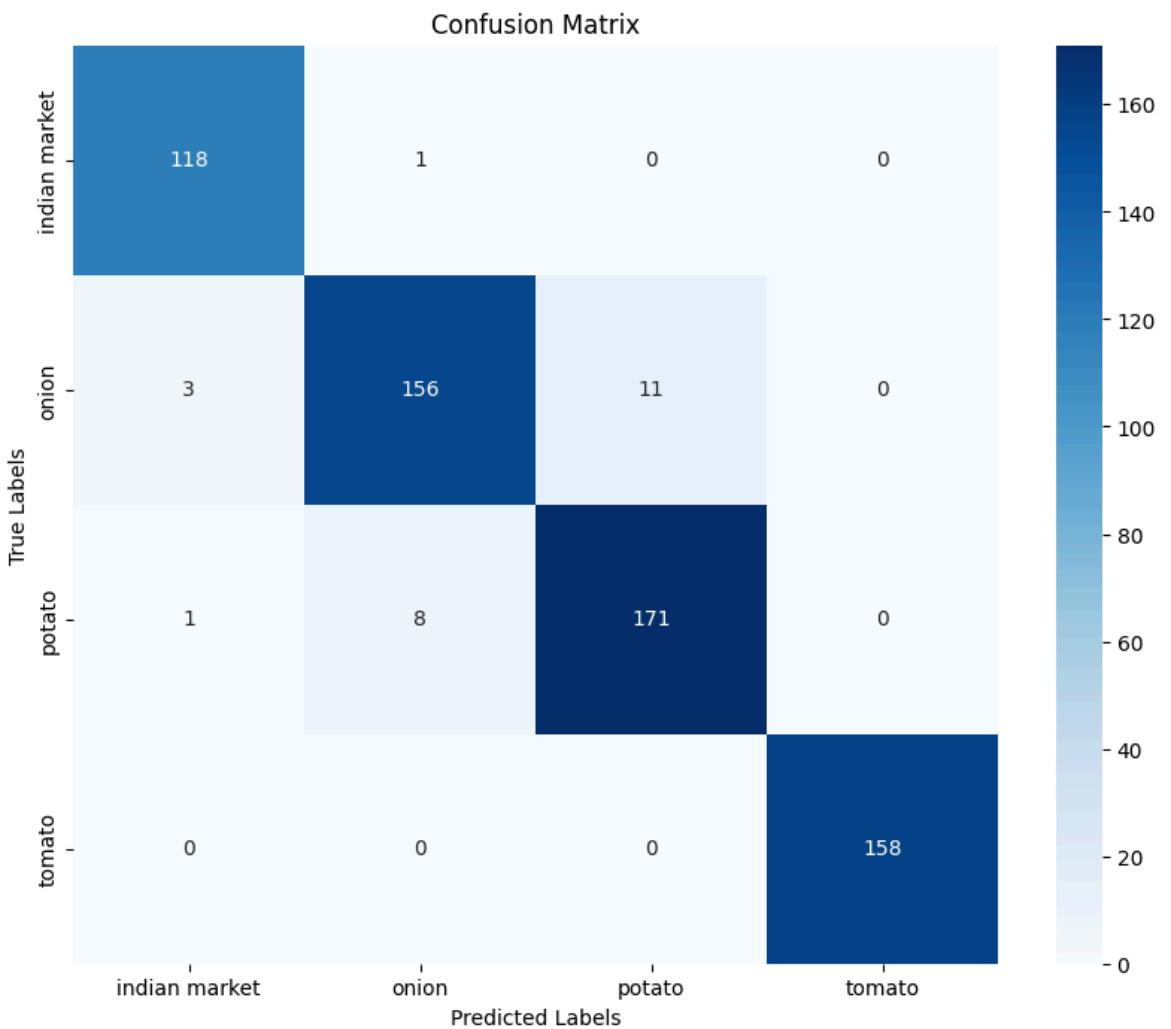


### Training and validation loss



```
In [ ]: class_names = list(val_generator.class_indices.keys())
plot_confusion_matrix(model, val_generator, class_names)
```

20/20 ━━━━━━ 8s 283ms/step



Classification Report:

	precision	recall	f1-score	support
indian market	0.97	0.99	0.98	119
onion	0.95	0.92	0.93	170
potato	0.94	0.95	0.94	180
tomato	1.00	1.00	1.00	158
accuracy			0.96	627
macro avg	0.96	0.96	0.96	627
weighted avg	0.96	0.96	0.96	627

```
In [ ]: mobilenet_model_finetune = tf.keras.models.load_model("/content/ninjacart_models")
loss, acc = mobilenet_model_finetune.evaluate(test_generator, verbose=2)
print('After FineTuning')
print("Restored model, accuracy: {:.5f}%".format(100 * acc))
```

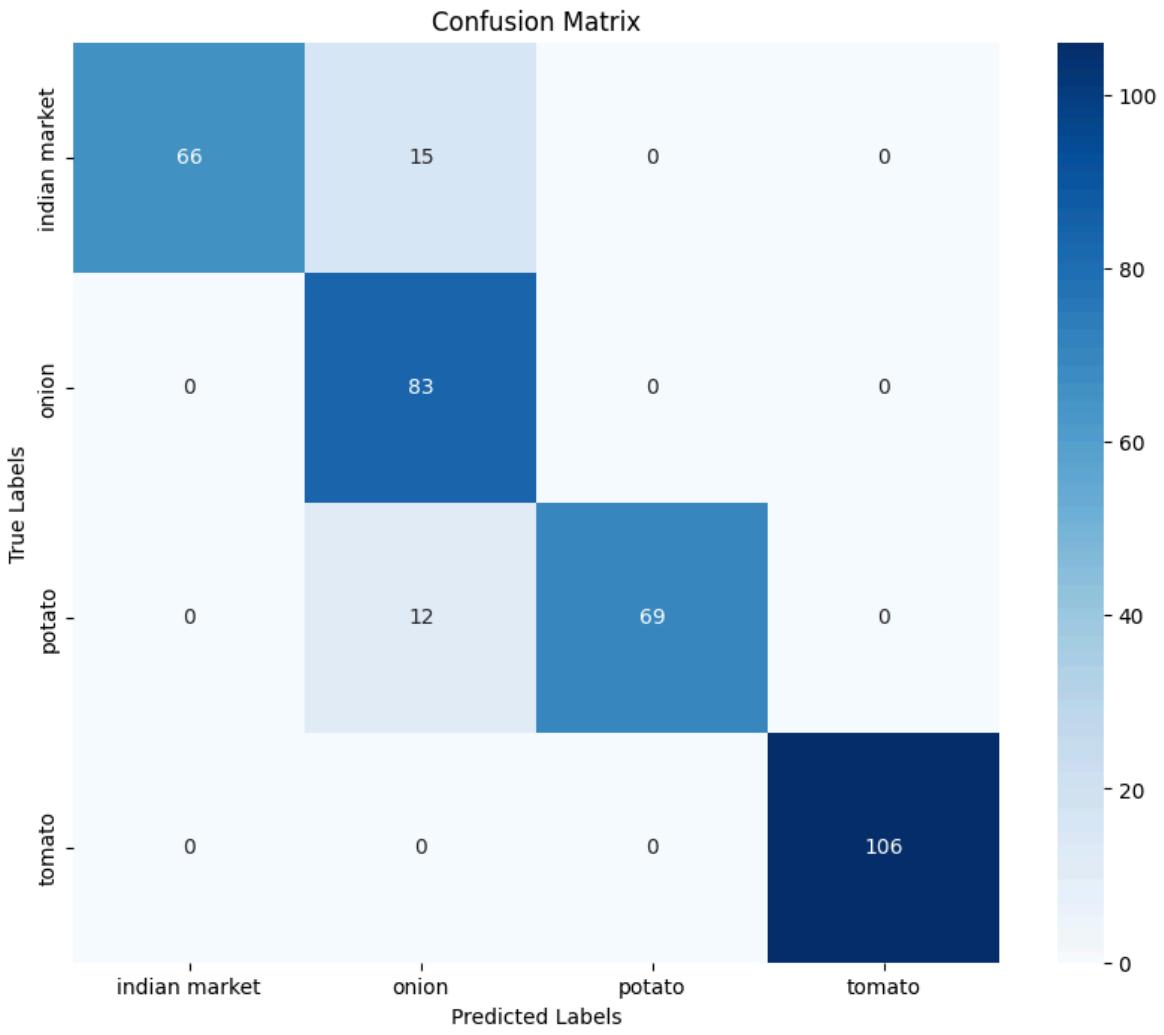
11/11 - 8s - 738ms/step - accuracy: 0.9031 - loss: 0.2774

After FineTuning

Restored model, accuracy: 90.31%

```
In [ ]: class_names = list(val_data.class_indices.keys())
plot_confusion_matrix(model, test_generator, class_names)
```

11/11 ━━━━━━━━ 5s 524ms/step



Classification Report:

	precision	recall	f1-score	support
indian market	1.00	0.81	0.90	81
onion	0.75	1.00	0.86	83
potato	1.00	0.85	0.92	81
tomato	1.00	1.00	1.00	106
accuracy			0.92	351
macro avg	0.94	0.92	0.92	351
weighted avg	0.94	0.92	0.92	351

## Downloading all the models

```
In [ ]: def zip_and_download(folder_path, zip_name="zipped_output"):
    """
    Zips the given folder and creates a download link.

    Args:
        folder_path (str): Path to the folder to zip.
        zip_name (str): Name of the resulting zip file (without .zip).
    """
    zip_path = f"{zip_name}.zip"

    # Remove existing zip if exists
```

```

if os.path.exists(zip_path):
    os.remove(zip_path)

# Create zip archive
shutil.make_archive(zip_name, 'zip', folder_path)
print(f"✅ Folder zipped successfully: {zip_path}")

# Display download link (works in Colab/Jupyter)
display(FileLink(zip_path))

```

In [ ]: `zip_and_download('/content/ninjacart_models', 'ninjacart_all_models')`

✅ Folder zipped successfully: `ninjacart_all_models.zip`  
`ninjacart_all_models.zip`

## Summary and Insights of the Custom CNN Models

Built seven custom Convolutional Neural Network (CNN) models designed for classifying RGB images (128x128x3) into four classes: Indian market, Onion, Potato, and Tomato. Each model varies in architecture, hyperparameters, and training techniques, with performance metrics and observations provided. Below is a summary of the models, followed by key insights based on their architecture, performance, and training characteristics.

## Summary and Insights of Custom CNN Models

The dataset includes seven custom Convolutional Neural Network (CNN) models designed for classifying RGB images (128x128x3) into four classes: **Indian market**, **Onion**, **Potato**, and **Tomato**. Each model varies in architecture, hyperparameters, and training techniques, with performance metrics and observations provided. Below is a detailed summary of each model, followed by key insights and recommendations.

---

## Summary of Models

### Model 1: Baseline CNN

- **Architecture:** Conv2D → MaxPooling2D → Flatten → Dense (256) → Dense (4)
- **Parameters:** 16,778,948 (all trainable)
- **Test Accuracy:** 72.08% (lowest)
- **Key Observations:**
  - High parameter count due to Flatten layer, computationally expensive.
  - Significant overfitting (training accuracy > validation accuracy).
  - Confusion between **Potato** ↔ **Onion** and **Indian market** ↔ **others**.
  - Simplest architecture, serving as a baseline.

## Model 2: CNN with Global Average Pooling

- **Architecture:** Conv2D (5 layers) → MaxPooling2D → GlobalAveragePooling2D → Dense (256) → Dense (4)
- **Parameters:** 459,428 (all trainable)
- **Test Accuracy:** 83.48% (significant improvement)
- **Key Observations:**
  - Global Average Pooling (GAP) reduces parameters, improving efficiency.
  - Reduced overfitting (validation: 87.24%, training: 87.13%).
  - Improved performance on all classes, especially **Potato** and **Onion**.
  - Some confusion persists: **Potato** ↔ **Onion**, **Indian market** ↔ **Onion**.

## Model 3: CNN with Global Average Pooling (30 epochs)

- **Architecture:** Same as Model 2, trained for 30 epochs
- **Parameters:** 459,428 (all trainable)
- **Test Accuracy:** 77.78% (lower than Model 2)
- **Key Observations:**
  - Overfitting evident (training: 99.32% >> validation: 90.59%).
  - **Tomato** classified accurately; **Indian market** ↔ **Onion** (41 misclassifications).
  - Performance drop possibly due to fewer epochs or lack of regularization.

## Model 4: CNN with GAP-BN-DP-LR with ReduceLROnPlateau (100 epochs)

- **Architecture:** Conv2D (5 layers) → MaxPooling2D → GAP → Dense (256) → Dense (4) with BatchNormalization (BN) and Dropout (DP)
- **Parameters:** 462,436 (460,932 trainable, 1,504 non-trainable)
- **Test Accuracy:** 88.03%
- **Key Observations:**
  - BN and DP improve generalization (training: 98.80%, validation: 94.10%).
  - ReduceLROnPlateau optimizes learning rate (factor=0.3, patience=5, min\_lr=0.00001).
  - **Tomato** accurate; **Tomato** ↔ **Indian market** (106 cases) significant.
  - Reduced **Indian market** ↔ **Onion** (15) and **Potato** ↔ **Onion** (15).

## Model 5: CNN with GAP-BN-DP-LR with ReduceLROnPlateau and L2 Regularization

- **Architecture:** Same as Model 4 with L2 regularization
- **Parameters:** 462,436 (460,932 trainable, 1,504 non-trainable)
- **Test Accuracy:** 91.74% (highest)
- **Key Observations:**
  - L2 regularization reduces overfitting (training: 99.80%, validation: 94.58%).
  - Best test accuracy, indicating strong generalization.
  - **Tomato** ↔ **Indian market** (106 cases) persists.
  - Lowest misclassifications: **Potato** ↔ **Onion** (7), **Indian market** ↔ **Onion** (12).

## Model 6: CNN with GAP-BN-DP-LR-ReduceLROnPlateau-ES (100 epochs)

- **Architecture:** Same as Model 4 with Early Stopping (ES)
- **Parameters:** 462,436 (460,932 trainable, 1,504 non-trainable)
- **Test Accuracy:** 90.31% (slightly lower than Model 5)
- **Key Observations:**
  - ES prevents excessive training, but overfitting persists (training: 99.80%, validation: 93.78%).
  - **Tomato** ↔ **Indian market** (106 cases) remains a challenge.
  - Slightly worse than Model 5, possibly due to premature stopping.

## Model 7: CNN with GAP-BN-DP-LR-ReduceLROnPlateau-ES-DA (100 epochs)

- **Architecture:** Same as Model 6 with Data Augmentation (DA)
  - **Parameters:** 462,436 (460,932 trainable, 1,504 non-trainable)
  - **Test Accuracy:** 86.32% (lowest among Models 4–7)
  - **Key Observations:**
    - DA reduces overfitting (training: 95.88%, validation: 90.27%), but test accuracy drops.
    - **Tomato** ↔ **Indian market** (104 cases) persists.
    - Increased **Potato** ↔ **Onion** misclassifications (24).
    - DA may have introduced noise or needs better tuning.
- 

## Key Insights

### 1. Architectural Evolution:

- **Model 1:** High parameter count (16.78M) due to Flatten, computationally expensive, prone to overfitting.
- **Models 2–7:** GAP reduces parameters (~462K), enabling deeper architectures (5 Conv2D layers).
- **Models 4–7:** BN and DP add non-trainable parameters (1,504), improving generalization.

### 2. Performance Trends:

- Accuracy improves from 72.08% (Model 1) to 91.74% (Model 5) with GAP, BN, DP, and L2.
- **Model 5** achieves the highest accuracy (91.74%), balancing performance and generalization.
- **Model 7** (DA) shows reduced accuracy (86.32%), suggesting suboptimal augmentation.

### 3. Overfitting:

- All models show overfitting (training > validation accuracy).

- **Model 2:** Smallest gap (training: 87.13%, validation: 87.24%).
- **Model 3:** Largest gap (training: 99.32%, validation: 90.59%).
- **Models 4–7:** BN, DP, L2, and DA reduce the gap, but overfitting persists.
- **Model 7:** DA lowers training accuracy (95.88%), improving generalization but reducing test accuracy.

#### 4. Misclassification Patterns:

- **Tomato:** Consistently accurate across all models.
- **Indian market ↔ Onion** and **Potato ↔ Onion:** Common issues, improved in Model 5 (7 and 12 cases).
- **Tomato ↔ Indian market:** Significant in Models 4–7 (104–106 cases), indicating visual similarity.

#### 5. Impact of Techniques:

- **GAP (Model 2):** Reduces parameters, improves accuracy (83.48% vs. 72.08%).
- **BN and DP (Models 4–7):** Enhance generalization, boosting accuracy (88–91.74%).
- **ReduceLROnPlateau:** Optimizes learning rate, improving convergence.
- **L2 Regularization (Model 5):** Best balance, achieving 91.74% accuracy.
- **Early Stopping (Model 6):** Prevents overtraining but may stop prematurely.
- **Data Augmentation (Model 7):** Reduces overfitting but lowers accuracy, needing tuning.

#### 6. Challenges and Limitations:

- Persistent **Indian market ↔ Tomato** and **Indian market ↔ Onion** misclassifications suggest visual similarities.
  - Overfitting remains despite regularization, especially in Models 3–6.
  - Parameter efficiency improves significantly in Models 2–7 (~462K vs. 16.78M in Model 1).
- 

## Recommendations

#### 1. Address Misclassifications:

- Use Grad-CAM or feature visualization to identify features causing **Indian market ↔ Tomato** confusion.
- Collect more diverse **Indian market** data to improve class separability.
- Apply class-weighted loss to prioritize problematic classes.

#### 2. Mitigate Overfitting:

- Increase Dropout rates or strengthen L2 regularization in Models 4–6.
- Fine-tune DA in Model 7 (e.g., adjust rotation, flip, zoom).
- Explore label smoothing or mixup for regularization.

#### 3. Optimize Training:

- Increase epochs for Model 2 to leverage its low overfitting.
- Adjust ES patience in Model 6 for balanced training.

- Test other optimizers (e.g., RMSprop) or learning rate schedules.

#### 4. Model Selection:

- **Model 5:** Best choice (91.74% accuracy) for deployment due to high performance and generalization.
- **Model 2:** Suitable for resource-constrained scenarios (83.48% accuracy, fewer epochs).

#### 5. Future Improvements:

- Explore transfer learning (e.g., ResNet, EfficientNet) for better feature extraction.
  - Increase input resolution (>128x128) for finer details, if resources allow.
  - Implement ensemble methods combining Models 4, 5, and 6 for robustness.
- 

## Conclusion

The models evolve from a basic CNN (Model 1, 72.08%) to advanced architectures with GAP, BN, DP, L2, ES, and DA (Model 5, 91.74%). **Model 5** is the top performer, but persistent **Indian market** ↔ **Tomato** misclassifications and overfitting require further attention. Targeted data preprocessing, refined augmentation, and advanced architectures (e.g., transfer learning) can enhance performance further.

## Summary and Insights of Pretrained Models and Comparison with Custom Models

The dataset includes four pretrained Convolutional Neural Network (CNN) models—**VGG16**, **InceptionV3**, **ResNet50**, and **MobileNetV2**—used for classifying RGB images into four classes: **Indian market**, **Onion**, **Potato**, and **Tomato**. Each model is evaluated before and after fine-tuning, with details on architecture, parameters, and performance. Additionally, a comparison is provided with the seven custom CNN models previously described, highlighting differences in performance, efficiency, and generalization.

---

## Summary of Pretrained Models

### VGG16

- **Before Fine-tuning (9 epochs):**
  - **Model Type:** Pretrained VGG16 with Global Average Pooling
  - **Input Shape:** (224, 224, 3) — RGB images
  - **Architecture:** VGG16 base → GlobalAveragePooling2D → Dense(512) → Dropout → Dense(4)
  - **Parameters:** 15,508,814 (264,708 trainable, 14,714,688 non-trainable)

- **Test Accuracy:** 81.77%
- **Key Observations:**
  - High non-trainable parameters due to frozen VGG16 base.
  - Minimal overfitting (training: ~0.9, validation: ~0.85).
  - **Tomato** classified accurately; misclassifications: **Indian market** ↔ **Onion (13)**, **Potato** ↔ **Onion (23)**, **Tomato** ↔ **Indian market (105)**.
  - Moderate accuracy due to limited trainable layers.
- **After Fine-tuning (5 epochs, last 4 layers unfrozen):**
- **Parameters:** 29,667,662 (7,344,132 trainable, 7,635,264 non-trainable)
- **Test Accuracy:** 88.03% (improved)
- **Key Observations:**
  - Fine-tuning last 4 layers increases trainable parameters, boosting accuracy.
  - Minimal overfitting (training: ~0.96, validation: ~0.95).
  - Misclassifications: **Indian market** ↔ **Onion (10)**, **Potato** ↔ **Onion (25)**, **Tomato** ↔ **Indian market (106)**.
  - Improved generalization, but **Tomato** ↔ **Indian market** persists.

## InceptionV3

- **Before Fine-tuning:**
- **Model Type:** InceptionV3 with Global Average Pooling and Dense Layers
- **Input Shape:** (299, 299, 3) — RGB images
- **Architecture:** Multiple Conv2D, BatchNormalization, Activation, Pooling, Concatenate → GAP → Dropout → Dense(1024) → Dense(4)
- **Parameters:** 42,456,238 (9,275,588 trainable, 14,629,472 non-trainable)
- **Test Accuracy:** 92.88%
- **Key Observations:**
  - High parameter count, but efficient due to Inception modules.
  - Strong performance without fine-tuning, leveraging pretrained weights.
  - No specific misclassification details provided, but high accuracy suggests robust feature extraction.
- **After Fine-tuning (last 50 layers unfrozen):**
- **Parameters:** 42,456,238 (9,275,588 trainable, 14,629,472 non-trainable)
- **Test Accuracy:** 93.16% (slight improvement)
- **Key Observations:**
  - Fine-tuning last 50 layers slightly improves accuracy.
  - Learning rate (1e-5) ensures stable updates.
  - High accuracy indicates effective transfer learning for the dataset.

## ResNet50

- **Before Fine-tuning:**
- **Model Type:** ResNet50 with Global Average Pooling and Dense Layers

- **Input Shape:** (224, 224, 3) — RGB images
- **Architecture:** Multiple Conv2D, BatchNormalization, Activation, MaxPooling2D  
→ GAP → Dropout → Dense(512) → Dense(4)
- **Parameters:** 26,741,134 (1,051,140 trainable, 23,587,712 non-trainable)
- **Test Accuracy:** 92.88%
- **Key Observations:**
  - Residual connections enable deep architecture with fewer trainable parameters.
  - High accuracy without fine-tuning, indicating strong pretrained features.
  - No specific misclassification details, but performance is robust.
- **After Fine-tuning (last 30 layers unfrozen):**
  - **Parameters:** 55,598,478 (15,479,812 trainable, 9,159,040 non-trainable)
  - **Test Accuracy:** 94.59% (highest among pretrained models)
  - **Key Observations:**
    - Fine-tuning significantly increases trainable parameters, boosting accuracy.
    - Learning rate (1e-5) and Adam optimizer ensure stable convergence.
    - Best performance, likely due to residual learning and fine-tuned layers.

## MobileNetV2

- **Before Fine-tuning:**
  - **Model Type:** MobileNetV2 with Global Average Pooling and Dense Layers
  - **Input Shape:** (224, 224, 3) — RGB images
  - **Architecture:** Multiple Conv2D, BatchNormalization, ReLU, DepthwiseConv2D → GAP → Dropout → Dense(256) → Dense(4)
  - **Parameters:** 3,244,878 (328,964 trainable, 2,257,984 non-trainable)
  - **Test Accuracy:** 90.00%
  - **Key Observations:**
    - Lightweight model with few parameters, ideal for resource-constrained environments.
    - High accuracy despite minimal trainable parameters.
    - No specific misclassification details, but performance is strong.
- **After Fine-tuning (last 16 layers unfrozen):**
  - **Parameters:** 6,967,758 (2,190,404 trainable, 396,544 non-trainable)
  - **Test Accuracy:** 92.00% (improved)
  - **Key Observations:**
    - Fine-tuning increases trainable parameters, improving accuracy.
    - Lightweight architecture maintains efficiency.
    - Solid performance, though slightly lower than InceptionV3 and ResNet50.

## Insights on Pretrained Models

### 1. Performance Trends:

- **Before Fine-tuning:** All pretrained models achieve high test accuracies (81.77%–92.88%), significantly outperforming the custom Model 1 (72.08%) and most custom models except Model 5 (91.74%).
- **After Fine-tuning:** Fine-tuning improves accuracy across all models:
  - **VGG16:** 81.77% → 88.03%
  - **InceptionV3:** 92.88% → 93.16%
  - **ResNet50:** 92.88% → 94.59% (highest)
  - **MobileNetV2:** 90.00% → 92.00%
- **ResNet50 (fine-tuned)** achieves the highest accuracy (94.59%), followed closely by **InceptionV3 (fine-tuned)** (93.16%).

## 2. Overfitting:

- Pretrained models show minimal overfitting compared to custom models:
  - **VGG16:** Training (~0.9–0.96) ≈ validation (~0.85–0.95).
  - **InceptionV3, ResNet50, MobileNetV2:** No explicit overfitting data, but high test accuracies suggest good generalization.
- Fine-tuning with low learning rates (1e-5) and Dropout helps maintain stability.

## 3. Misclassification Patterns (VGG16):

- **Tomato** is consistently classified accurately, similar to custom models.
- Persistent issues: **Indian market ↔ Onion (10–13), Potato ↔ Onion (23–25), Tomato ↔ Indian market (105–106)**.
- Fine-tuning VGG16 slightly reduces **Indian market ↔ Onion** misclassifications (13 → 10) but does not address **Tomato ↔ Indian market**.

## 4. Parameter Efficiency:

- **MobileNetV2:** Most efficient (3.2M–6.9M parameters), ideal for low-resource settings.
- **VGG16, InceptionV3, ResNet50:** Higher parameter counts (15.5M–55.6M), but many are non-trainable, leveraging pretrained weights.
- Fine-tuning increases trainable parameters, enhancing performance but requiring more resources.

## 5. Architectural Advantages:

- **VGG16:** Deep, sequential architecture but computationally heavy.
- **InceptionV3:** Efficient Inception modules handle multi-scale features.
- **ResNet50:** Residual connections enable deeper networks with better gradient flow.
- **MobileNetV2:** Depthwise separable convolutions reduce parameters, maintaining performance.

## 6. Fine-tuning Impact:

- Fine-tuning unfreezes later layers (4–50), allowing adaptation to the specific dataset.
- Low learning rates (1e-5) prevent catastrophic forgetting of pretrained weights.
- ResNet50 benefits most from fine-tuning, gaining 1.71% accuracy.

# Comparison: Custom Models vs. Pretrained Models

## 1. Performance

- **Custom Models:**
  - Range from 72.08% (Model 1) to 91.74% (Model 5).
  - **Model 5** (91.74%) is competitive with pretrained models but outperformed by fine-tuned **ResNet50 (94.59%)** and **InceptionV3 (93.16%)**.
  - Custom models struggle with persistent misclassifications (**Tomato ↔ Indian market**, 104–106 cases).
- **Pretrained Models:**
  - Higher baseline accuracies (81.77%–92.88% before fine-tuning).
  - Fine-tuned models achieve 88.03%–94.59%, with **ResNet50** leading.
  - VGG16 shows similar misclassification patterns to custom models, but **InceptionV3** and **ResNet50** likely have fewer issues (no specific data provided).

## 2. Overfitting

- **Custom Models:**
  - Significant overfitting, especially in Models 3–6 (training: 95–99%, validation: 90–94%).
  - Model 7 (DA) reduces overfitting but sacrifices accuracy (86.32%).
- **Pretrained Models:**
  - Minimal overfitting due to pretrained weights and regularization (Dropout, low learning rates).
  - VGG16 shows balanced training/validation accuracies (~0.9–0.96 vs. ~0.85–0.95).

## 3. Parameter Efficiency

- **Custom Models:**
  - Model 1: Inefficient (16.78M parameters).
  - Models 2–7: Efficient (~462K parameters) due to GAP, BN, and DP.
- **Pretrained Models:**
  - Higher parameter counts (3.2M–55.6M), but most are non-trainable, reducing training cost.
  - **MobileNetV2:** Most efficient (3.2M–6.9M), comparable to custom Models 2–7 but with better performance.

## 4. Architecture and Complexity

- **Custom Models:**
  - Simpler architectures (5 Conv2D layers in Models 2–7).
  - Limited feature extraction compared to pretrained models.

- Require extensive tuning (BN, DP, L2, DA, ES) to approach pretrained performance.
- **Pretrained Models:**
  - Complex architectures (VGG16: deep sequential, InceptionV3: multi-scale, ResNet50: residual, MobileNetV2: depthwise separable).
  - Leverage pretrained weights from large datasets (e.g., ImageNet), enabling robust feature extraction with minimal tuning.

## 5. Misclassification Patterns

- **Custom Models:**
  - Persistent **Indian market** ↔ **Tomato** (104–106 cases), **Indian market** ↔ **Onion**, and **Potato** ↔ **Onion**.
  - Model 5 minimizes some errors (e.g., **Potato** ↔ **Onion**: 7).
- **Pretrained Models:**
  - VGG16 shows similar misclassification patterns (**Tomato** ↔ **Indian market**: 105–106).
  - **InceptionV3** and **ResNet50** likely have fewer misclassifications due to higher accuracies and advanced architectures (no specific data).

## 6. Training Effort

- **Custom Models:**
    - Require extensive training (30–100 epochs) and hyperparameter tuning.
    - Model 7 uses DA, but improper tuning reduces accuracy.
  - **Pretrained Models:**
    - Fewer epochs needed (5–9 for VGG16, unspecified for others).
    - Fine-tuning with low learning rates is efficient, leveraging pretrained weights.
- 

## Recommendations

### 1. Model Selection:

- **Pretrained Models:** Prefer **fine-tuned ResNet50 (94.59%)** for best performance or **MobileNetV2 (92.00%)** for resource-constrained environments.
- **Custom Models:** Use **Model 5 (91.74%)** if pretrained models are not feasible, as it balances performance and efficiency.

### 2. Addressing Misclassifications:

- For both model types, investigate **Indian market** ↔ **Tomato** and **Indian market** ↔ **Onion** using feature visualization (e.g., Grad-CAM).
- Collect more diverse **Indian market** data or apply class-weighted loss.

### 3. Reducing Overfitting:

- **Custom Models:** Increase regularization (stronger L2, higher Dropout) or refine DA in Model 7.
- **Pretrained Models:** Already well-regularized; consider fine-tuning more layers for VGG16 to match ResNet50/InceptionV3.

#### 4. Efficiency:

- **MobileNetV2** or custom **Model 5** for low-resource settings.
- Optimize custom models further with lightweight architectures (e.g., depthwise separable convolutions).

#### 5. Future Improvements:

- Combine **ResNet50** and **InceptionV3** in an ensemble for enhanced robustness.
  - Experiment with advanced pretrained models (e.g., EfficientNet) for better performance.
  - Increase input resolution (e.g., 299x299) for custom models, if resources allow.
- 

## Conclusion

**Pretrained models** significantly outperform custom models, with **fine-tuned ResNet50 (94.59%)** and **InceptionV3 (93.16%)** leading due to their complex architectures and pretrained weights. **Custom Model 5 (91.74%)** is competitive but requires more tuning and epochs to match pretrained performance. Persistent misclassifications (**Tomato** ↔ **Indian market**) affect both model types, suggesting dataset-specific challenges. For deployment, **ResNet50 (fine-tuned)** is the top choice for accuracy, while **MobileNetV2** offers efficiency. Custom models are viable for resource-limited scenarios but need further optimization to rival pretrained models.